



# jsc::chunk\_evenly — Range adaptor for distributing work across tasks



Mateusz Zych, Ivo Kabadshow

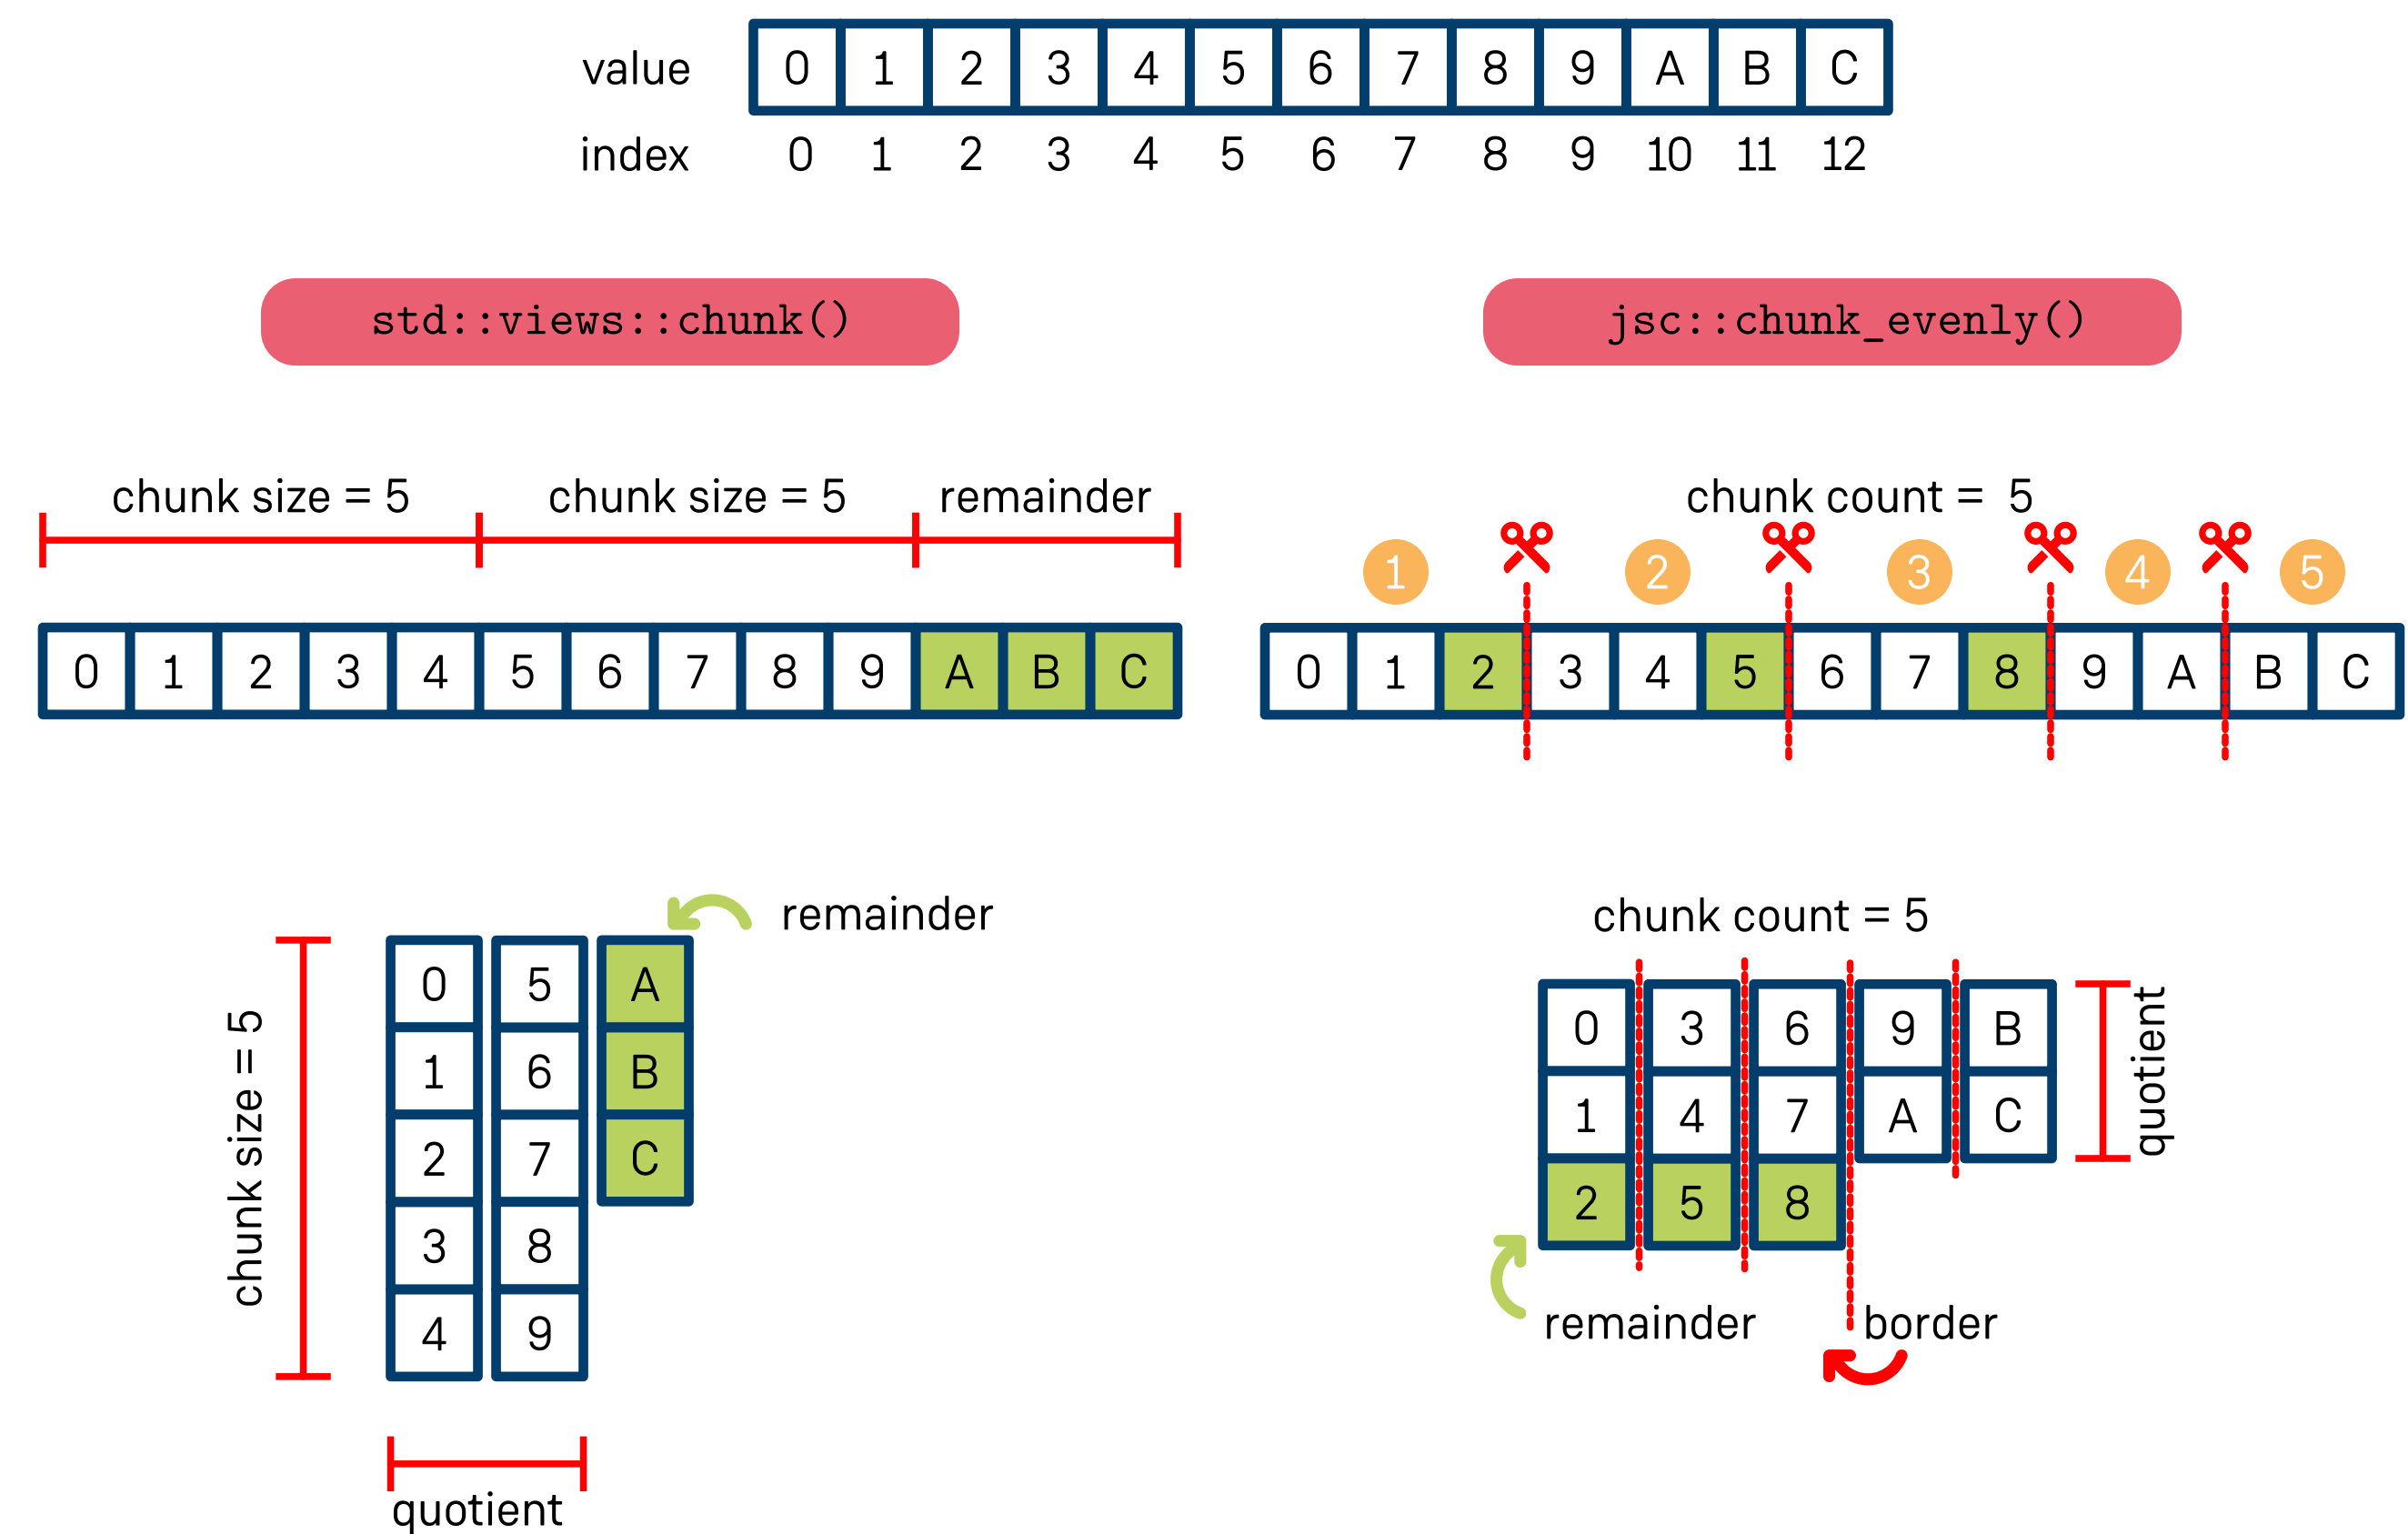
Jülich Supercomputing Centre, Research Centre Jülich, Germany

## ① How to chunk a sequence?

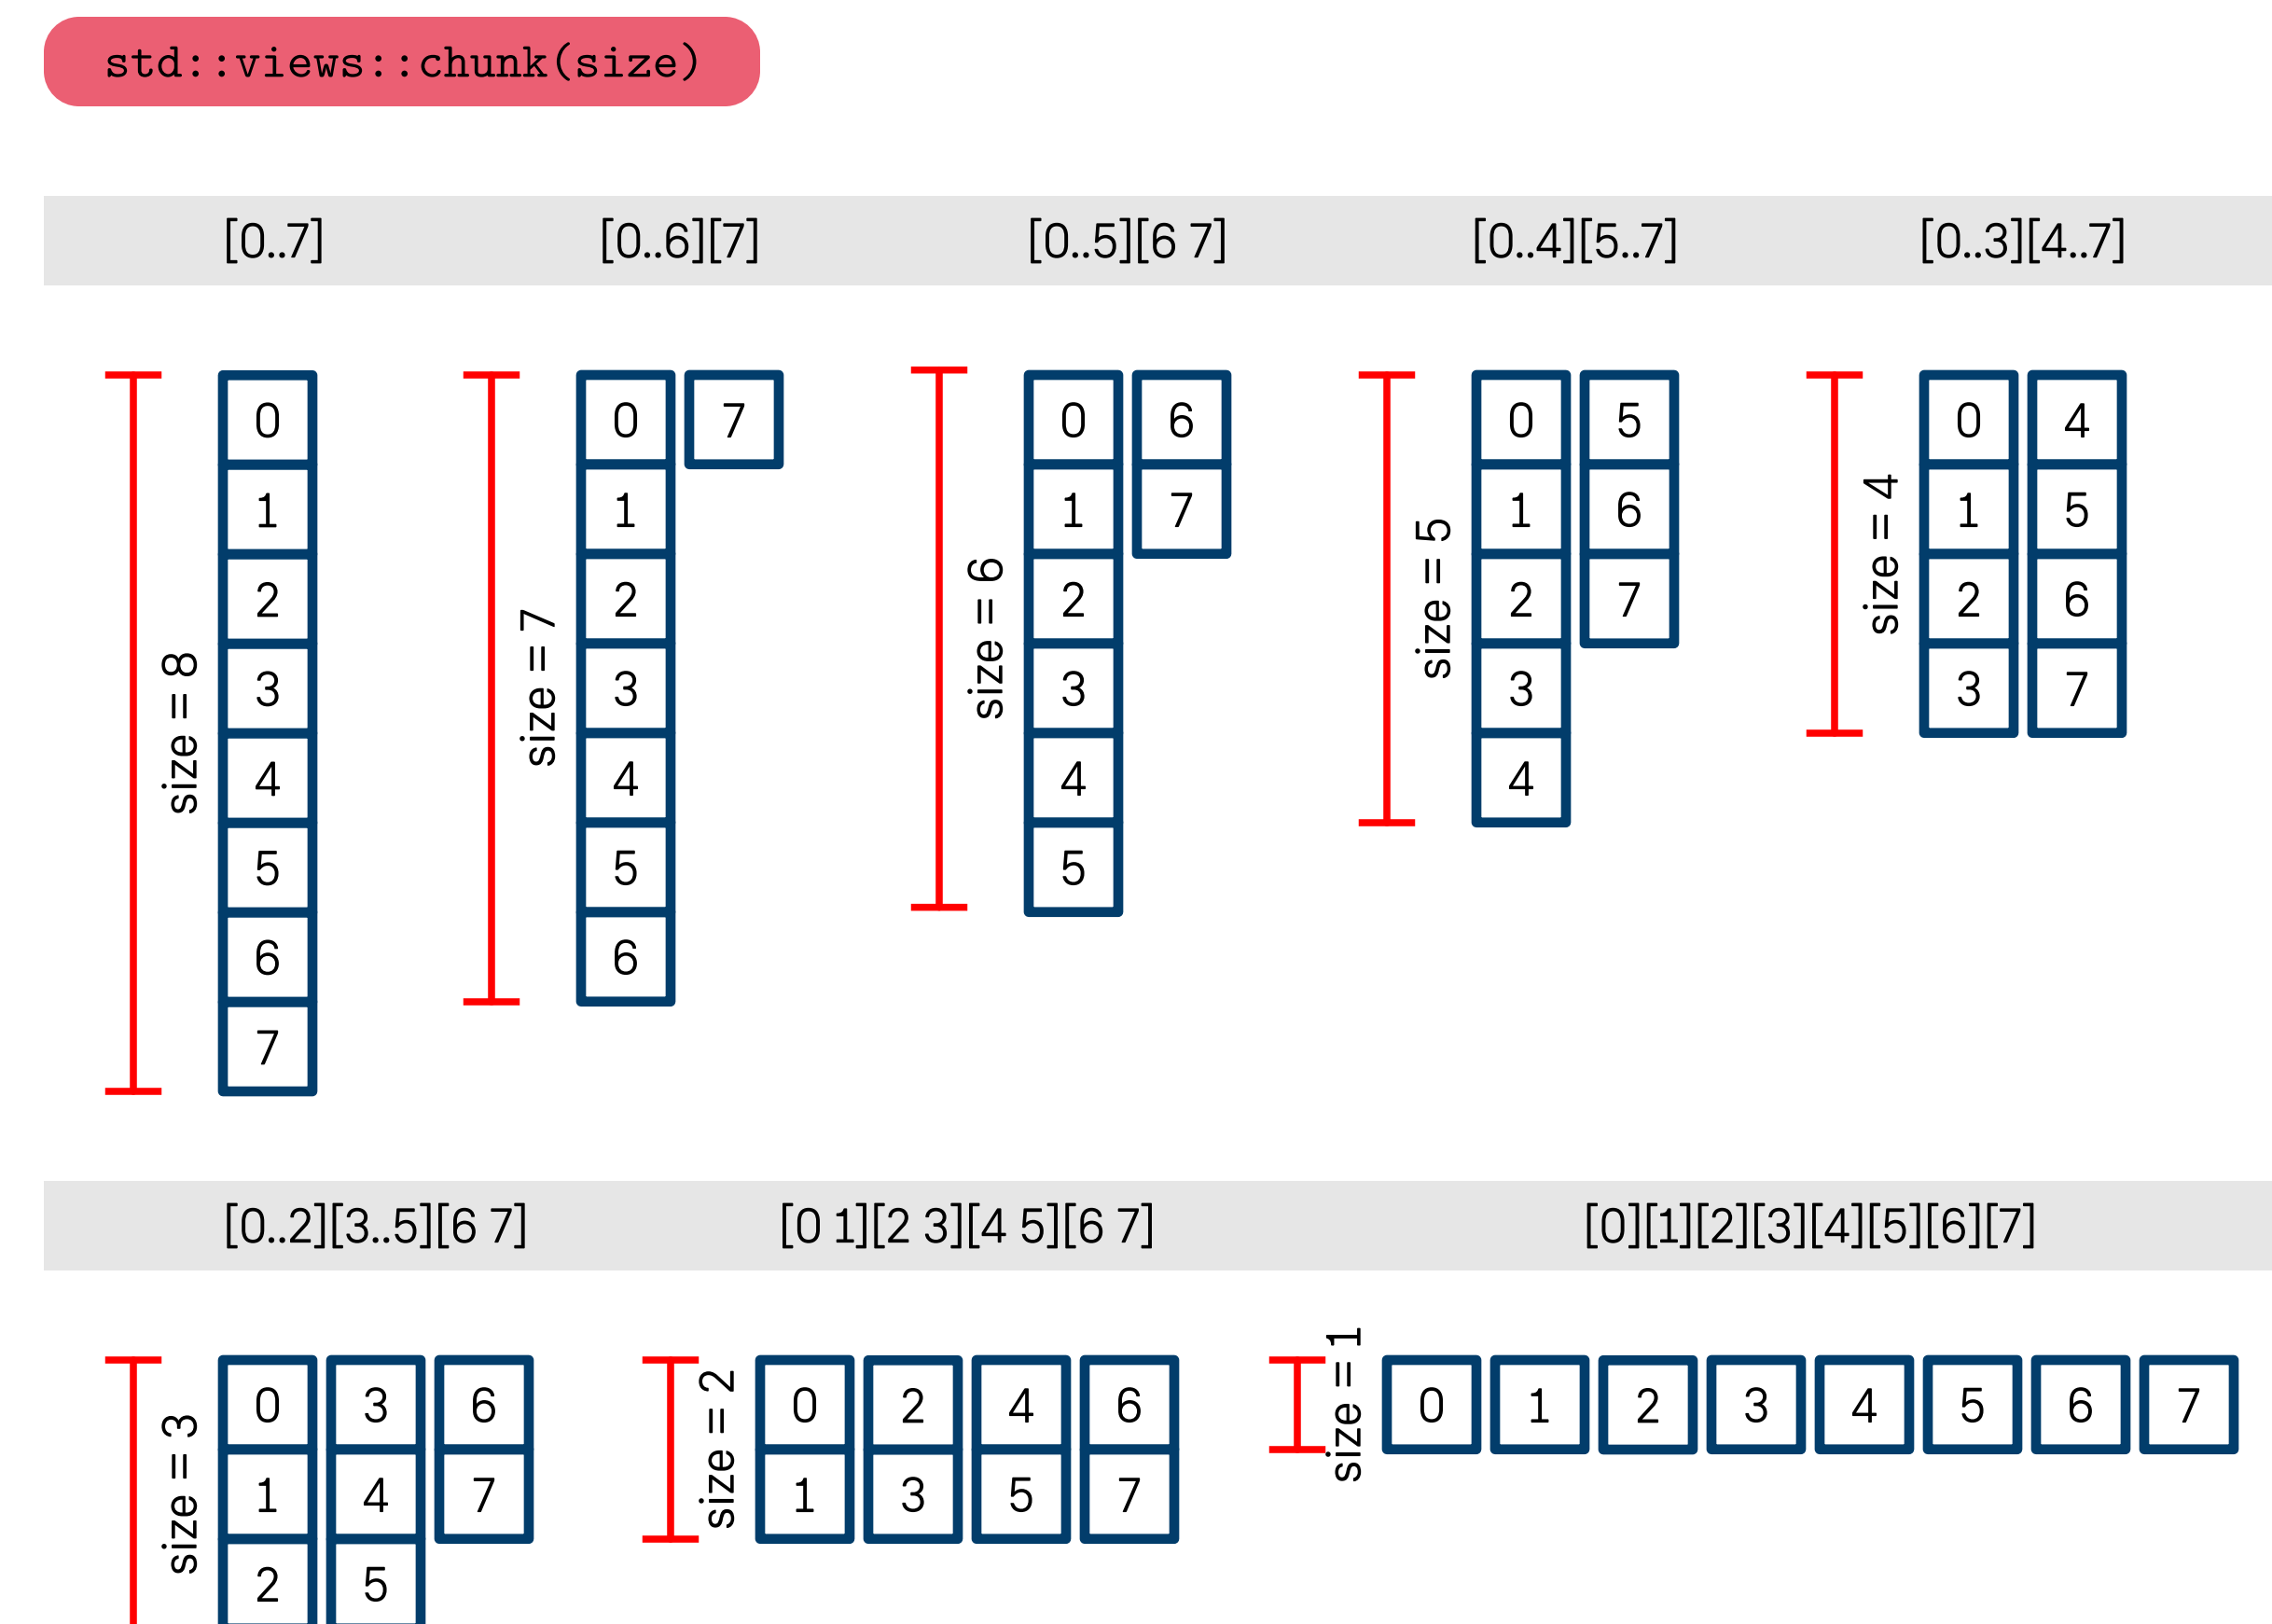
➤ Different range adaptors dividing a range into subranges available:

```
std::views::chunk(chunk_size)
std::views::chunk_by(predicate)
std::views::split(delimiter)
```

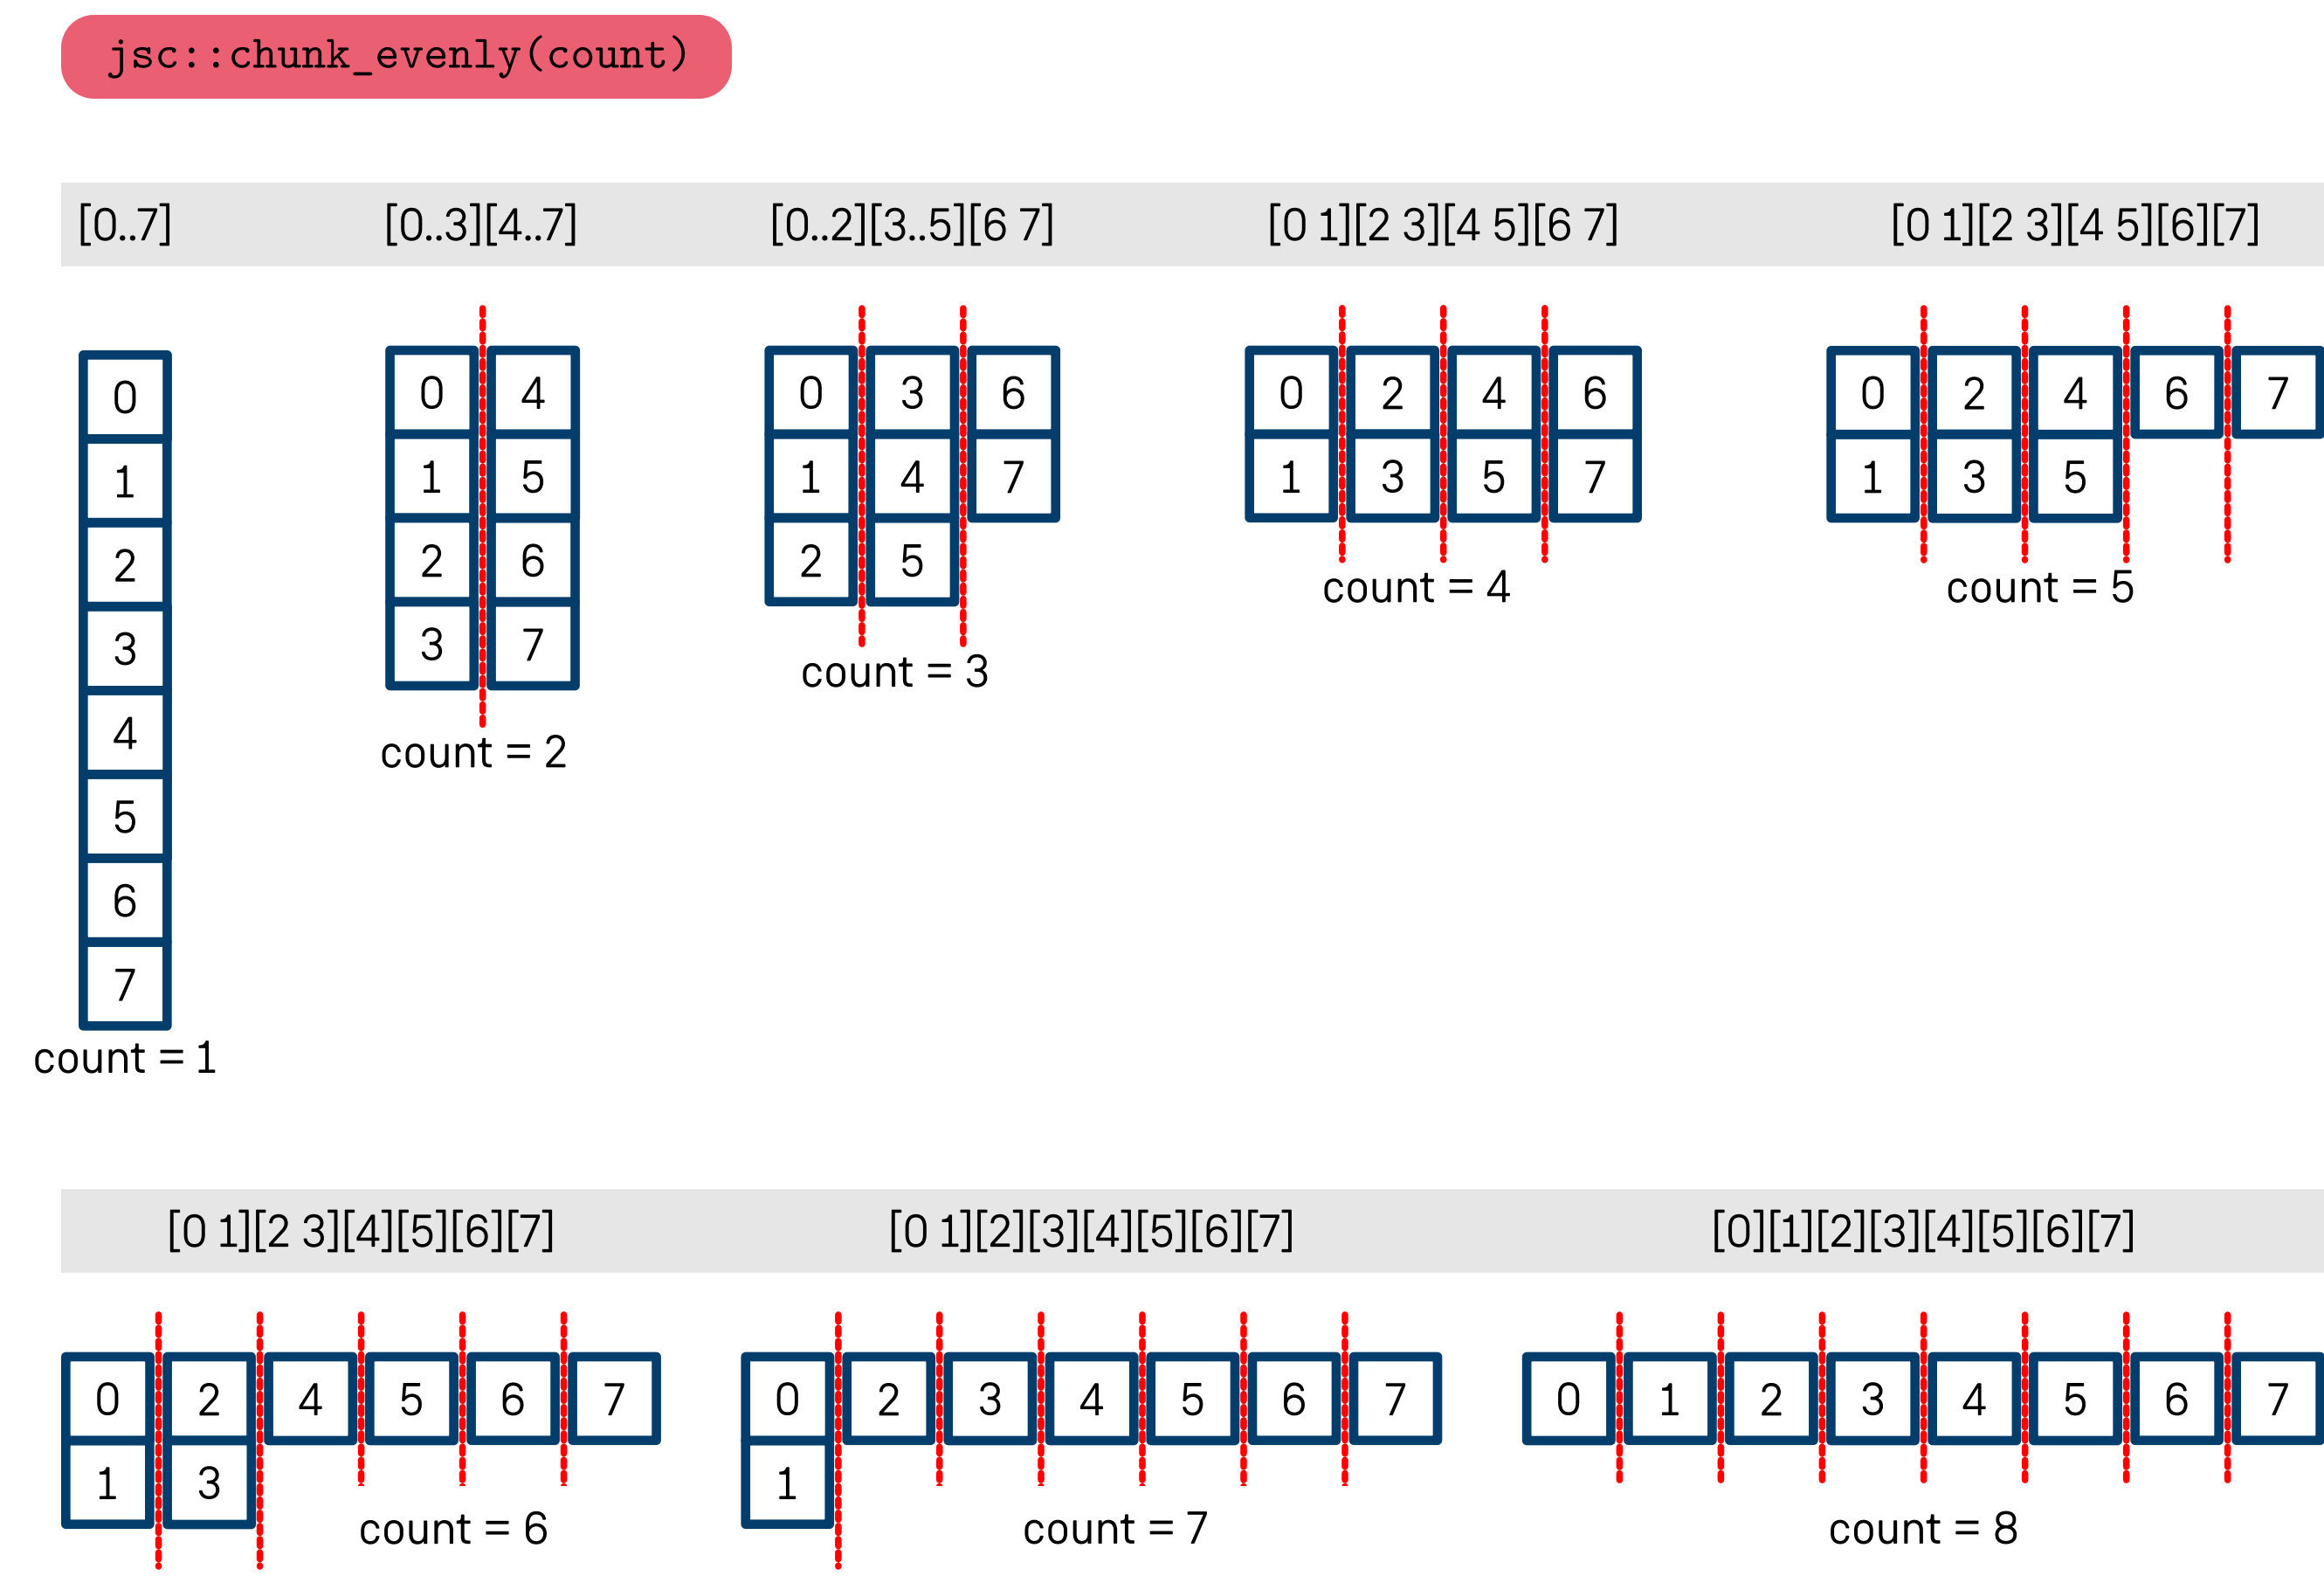
⚡ none is suitable to distribute work across tasks



## ② Chunk by size



## ③ Chunk by count



## ④ How to chunk & chunk\_evenly

```
#include <jsc/chunk_evenly.hpp>
#include <fmt/ranges.h>
#include <array>
#include <ranges>

auto main () -> int
{
    auto range = std::array
    { 0x0,0x1,0x2,0x3,0x4,0x5,0x6,0x7,0x8,0x9,0xA,0xB,0xC };

    fmt::println("{}:X", range | std::views::chunk(5));
    fmt::println("{}:X", range | jsc::chunk_evenly(5));
}
```

Output

## ⑤ Code gen comparison


```
#include <ranges>
#include <cstdlib>

#include <jsc/chunk_evenly.hpp>

auto func (std::ranges::subrange<
    std::ranges::iterator_t<
    std::ranges::iota_view<std::ptrdiff_t,
    std::ranges::iota_view<std::ptrdiff_t,
    std::ptrdiff_t>>> -> void;

auto std_views_chunk jsc_chunk_evenly (std::ptrdiff_t n,
    std::ptrdiff_t chunk_size chunk_count) -> void
{
    for (auto chunk : std::views::iota(std::ptrdiff_t { 0 }, n)
        | std::views::chunk(chunk_size) jsc::chunk_evenly(chunk_count))
    {
        func(chunk);
    }
}
```

## ⑥ Number of instructions & branches



## ⑦ Following zero-overhead principle

```
auto func (std::ranges::subrange<
    std::ranges::iterator_t<
    jsc::iota_view<std::ptrdiff_t,
    std::ptrdiff_t>>> -> void;

auto jsc_chunk_evenly (std::ptrdiff_t n,
    std::ptrdiff_t chunk_count) -> void
{
    for (auto chunk : jsc::iota(std::ptrdiff_t { 0 }, n)
        | jsc::chunk_evenly(chunk_count))
    {
        func(chunk);
    }
}
```

```
struct chunk { std::ptrdiff_t begin; std::ptrdiff_t end; };

auto func (chunk) -> void;

auto manual_loop (std::ptrdiff_t n, std::ptrdiff_t chunk_count) -> void
{
    [[assume(n >= 0)]];
    [[assume(chunk_count > 0)]];

    const auto quotient = n / chunk_count;
    const auto remainder = n % chunk_count;

    auto chunk_size = quotient + std::ptrdiff_t { remainder > 0 };
    auto chunk_begin = std::ptrdiff_t { 0 };

    for (auto chunk_index = std::ptrdiff_t { 0 }; chunk_index != chunk_count;)
    {
        func({ chunk_begin, chunk_begin + chunk_size });

        chunk_begin += chunk_size;
        ++chunk_index;
        chunk_size -= std::ptrdiff_t { chunk_index == remainder };
    }
}
```


## ⑧ ASM comparison

➤ GCC RISC-V 64-bit assembly

```
manual_loop(long, long):
    addi sp,sp,-64
    sd s4,16(sp)
    rem s4,a0,a1
    sd s2,32(sp)
    sd s0,48(sp)
    sd s1,40(sp)
    sd s3,24(sp)
    sd ra,56(sp)
    mv s3,a1
    li s1,0
    li s0,0
    div a0,a0,a1
    snez s2,s4
    add s2,s2,a0
.L2:
    mv a0,s0
    add s0,s0,s2
    mv a1,s0
    call func(chunk)
    addi s1,s1,1
    sub a5,s4,s1
    seqz a5,a5
    mv a4,s1
    sub s2,s2,a5
    bne s3,s1,.L2
    ld ra,56(sp)
    ld s0,48(sp)
    ld s1,40(sp)
    ld s2,32(sp)
    ld s3,24(sp)
    ld s4,16(sp)
    addi sp,sp,64
    jr ra

jsc_chunk_evenly(long, long):
    addi sp,sp,-64
    sd s4,16(sp)
    rem s4,a0,a1
    sd s0,48(sp)
    sd s1,40(sp)
    sd s3,24(sp)
    sd ra,56(sp)
    mv s3,a1
    li s1,1
    li s0,0
    div a0,a0,a1
    sgt a5,s4,zero
    add s2,a0,a5
.L2:
    mv a0,s0
    add s0,s0,s2
    mv a1,s0
    call func()
    sub a5,s1,s4
    seqz a5,a5
    mv a4,s1
    sub s2,s2,a5
    addi s1,s1,1
    bne s3,s4,.L2
    ld ra,56(sp)
    ld s0,48(sp)
    ld s1,40(sp)
    ld s2,32(sp)
    ld s3,24(sp)
    ld s4,16(sp)
    addi sp,sp,64
    jr ra
```

## ⑨ CP and LCD comparison



## ⑩ Design principles

- Idea inspired by `std::range::evenChunks()` from the D language
- Expensive computations (div, mod) are computed once per iterator
- Branch-less programming for crossing border between smaller and larger chunks
- Zero overhead

## ⑪ Improving std::execution example

```
stdexec::sender
auto async_inclusive_scan (stdexec::scheduler auto scheduler,
    std::span<const double> input_range,
    std::span<double> output_range,
    std::ptrdiff_t chunk_count)
{
    auto partials = std::vector<double>(chunk_count + 1);

    return stdexec::just(std::move(partials))
        | stdexec::continue_on(scheduler)
        | stdexec::bulk(chunk_count, [=] (std::ptrdiff_t chunk_index,
            std::vector<double>& partials)
        {
            auto input_chunks = input_range | jsc::chunk_evenly(chunk_count);
            auto output_chunks = output_range | jsc::chunk_evenly(chunk_count);

            auto input_chunk = input_chunks[chunk_index];
            auto output_chunk = output_chunks[chunk_index];

            std::inclusive_scan(input_chunk.begin(),
                input_chunk.end(), output_chunk.begin());

            partials[chunk_index + 1] = output_chunk.back();
        })
        | stdexec::then([] (std::vector<double>&& partials)
        {
            std::inclusive_scan(partials.begin(),
                partials.end(), partials.begin());

            return std::move(partials);
        })
        | stdexec::bulk(chunk_count, [=] (std::ptrdiff_t chunk_index,
            std::vector<double>&& partials)
        {
            auto output_chunks = output_range | jsc::chunk_evenly(chunk_count);

            for (auto& element : output_chunks[chunk_index])
            {
                element += partials[chunk_index];
            }
        })
        | stdexec::then([=] (std::vector<double>&& partials)
        {
            return output_range;
        }));
}
```

## ⑫ Discussion

➤ Original version chunks manually the input range using the same schema as `std::views::chunk()`, therefore will not always process `tile_count` tiles

① `std::size_t const tile_size = (input_range.size() + tile_count - 1) / tile_count;`  
② `input_range.size() == 8`  
`tile_count = 6`  
③ `tile_size=(8+6-1)/6 = 2`  
④ `input_range | std::views::chunk(tile_size)`

0 1 2 3 4 5 6 7 tile\_count != 6

- Also, fork & join model introduces unnecessary waiting
- If tasks are scheduled as early as possible, then distributing work evenly across tasks can improve performance

## ⑬ Future Directions

- Support random access in `jsc::chunk_evenly_view<>`
- Find more generic components in the JSC C++ library