



The banner features a dark blue background with a green geometric graphic in the top right corner containing a white plus sign and the number 24. The title "Back To Basics" is in white, and "Function Call Resolution" is in orange. The speaker's name "BEN SAKS" is in orange. The Cppcon logo is in the bottom left, and the date "2024" with the event dates "September 15 - 20" is in the bottom right.

Back To Basics
Function Call Resolution

BEN SAKS

 **Cppcon**
The C++ Conference

2024 | 
September 15 - 20

Saks & Associates

These notes are Copyright © 2024
by Ben Saks and Dan Saks
and distributed with their permission by:

Saks & Associates
393 Leander Dr.
Springfield, OH 45504-4906 USA
+1-412-521-4117 (voice)
service@saksandassociates.com
saksandassociates.com

About Ben Saks

Ben Saks is the chief engineer of Saks & Associates. He is the principal editor and presenter for much of Saks & Associates' training curriculum on the use of C and C++ in embedded systems.

Ben has represented Saks & Associates on the ISO C++ Standards committee as well as two of the committee's study groups:

- SG14 — low-latency
- SG20 — education

Ben has spoken at industry conferences, including *CppCon: The C++ Conference*, the *C++ and System Software Summit*, the *Embedded Systems Conference*, and *NDC Techtown*. At *CppCon*, he's the chair of the Embedded Track and a member of the program committee.

3

More About Ben Saks

Ben previously worked as a software engineer for Vorne Industries, where he used C++ and JavaScript to develop embedded systems that help improve manufacturing productivity in factories all over the world. He is a contributing author on multiple Vorne patents.

Ben earned a B.A. with Distinction in Computer Science from Carleton College.

4

About Dan Saks

Dan Saks is the president of Saks & Associates, which offers training and consulting in C and C++ and their use in developing embedded systems.

Dan wrote the “Programming Pointers” column for *embedded.com* online. He also wrote columns for numerous print publications (when such things existed) including *The C/C++ Users Journal*, *The C++ Report*, *Software Development*, and *Embedded Systems Design*. With Thomas Plum, he wrote *C++ Programming Guidelines*, which won a 1992 *Computer Language Magazine Productivity Award*.

Dan has taught C and C++ to thousands of programmers worldwide. He has delivered hundreds of lectures, including a few keynote addresses, at conferences such as the *ACCU (Association of C and C++ Users) Conference*, *CppCon: The C++ Conference*, *C++ World*, the *Embedded Systems Conference*, *Meeting Embedded*, *NDC Techtown*, and the *Software Development Conference*.

5

More About Dan Saks

Dan served as secretary of the ANSI and ISO C++ Standards committees and as a member of the ANSI C Standards committee. He also contributed to the *CERT Secure C Coding Standard* and the *CERT Secure C++ Coding Standard*.

Dan collaborated with Thomas Plum in writing and maintaining *Suite++™*, the *Plum Hall Validation Suite for C++*, which tests C++ compilers for conformance with the international standard. He was a Senior Software Engineer for Fischer and Porter (now ABB), where he designed languages and tools for distributed process control. He also worked as a programmer with Sperry Univac (now Unisys).

Dan earned an M.S.E. in Computer Science from the University of Pennsylvania, and a B.S. with Highest Honors in Mathematics/Information Science from Case Western Reserve University.

6

Introduction

- Some C++ language features allow you to create multiple functions with the same name, including:
 - name hiding
 - function overloading
 - function templates
- When a compiler encounters an expression like `f(x, y, z)`, it must consider each of these language features to determine which `f` function to call.

7

Introduction

- When used well, these facilities can produce interfaces that are flexible, easy to use correctly, and hard to use incorrectly.
- When used poorly, the same facilities can produce confusing interfaces that are hard to use correctly.
- To use these facilities effectively, you need a solid grasp of:
 - how they work individually, and
 - how they interact.

8

Introduction

- This session will examine each of these features in turn:
 - Function overloading and overload resolution
 - Name lookup
 - Default function arguments
 - Function templates
- For each feature, we'll explain first how it works on its own, and then how it interacts with the other features.
- We'll close with an example that combines several of these features.

9

Outline

- Function Overloading
- Name Lookup
- Default Function Arguments
- Function Templates
- Tying It All Together

10

Overloading

- A function is **overloaded** if there's another function declared:
 - with the same name,
 - in the same scope.
- For example, here's a group of overloaded functions name put:

```
int put(int c);  
int put(int c, FILE *f);  
int put(char const *s);  
int put(char const *s, FILE *f);
```

11

Overload Resolution

- When the compiler encounters a call to an overloaded function, it selects the function to call using **overload resolution**.
- The compiler matches:
 - the type(s) of the argument(s) in the call against
 - the type(s) of the parameter(s) in the function declarations.
- For example,

```
                // calls...  
put('a', stdout); // int put(int c, FILE *f);  
put("Hello\n");   // int put(char const *s);
```

12

Overload Resolution

- If the compiler can't find a function that accepts the specified arguments, it issues an error message.
- For example, this is an error:

```
put("n = ", n, stdout);    // compile error: too many arguments
```

- In this case, there's no `put` function that accepts three arguments.

13

Argument Conversions

- The type of each parameter need not be exactly the same as the type of its corresponding argument.
 - C++ compilers apply some conversions in an attempt to find a match.
- The Standard C `getchar` returns the character read as an `int`, not a `char`.
 - `getchar` uses the integer value `EOF` to represent end-of-file and I/O failures.
- `putchar` has a parameter of type `int` (not `char`) to match `getchar`'s convention.
- So do these overloaded functions...

14

Argument Conversions

```
int put(int c);           // write c to standard output
int put(int c, FILE *f);  // write c to a designated file *f
```

- Even though they use `int` as a parameter type, you can pass them a `char` as the argument type:

```
char c;
~~~
put(c);           // calls int put(int c);
put(c, stdout);   // calls int put(int c, FILE *f);
```

- In each case, the compiler promotes `char` to `int` to achieve the match.
 - That promotion may generate additional code that executes at run time.

15

Best Matches and Ambiguities

- In a group of overloaded functions, more than one function might be capable of satisfying a particular function call.
- For example, the first three of these functions are all capable of satisfying a call to `f(0)`:

```
int f(int i);           // f(0)?  maybe
long int f(long int li); // f(0)?  maybe
char *f(char *p);       // f(0)?  maybe
int f(double d, int i);  // f(0)?  definitely not
```

- Clearly, the fourth function can't satisfy the call:
 - `f(0)` passes only one argument, while the fourth function requires two.

16

Best Matches and Ambiguities

- We're left with three viable functions:

```
int f(int i);  
long int f(long int li);  
char *f(char *p);
```

- The call `f(0)` can pass `0`:
 - as an `int`,
 - as a `long int`, or even
 - as a `char *` whose value is the null pointer.

17

Best Matches and Ambiguities

- When confronted with a choice of functions:
 - Overload resolution uses a **ranking** of the conversions from the argument type into the parameter type.
 - The ranking helps determine which function, if any, is the **best match**.
- Thus, overload resolution depends on the exact type of each function call argument as well as the type of its corresponding parameter.
- In the case of calling `f(0)`, it depends on the exact type of argument `0`.
 - Obviously, `0` is an `int`.
 - But is it signed or unsigned?

18

Classifying Conversions

- The literal `0` is a plain `int`, which is **signed** by default.
- To find the best match for `f(0)`, consider the conversions...
- Calling `f(int)` requires no conversion at all.
 - It's an ***exact match***.
- Calling `f(long int)` requires converting `int` into "long int".
 - This is an ***integral conversion***.
- Calling `f(char *)` requires converting `int` into "pointer to char".
 - This is a ***pointer conversion***.

19

Exact Matches Always Win

- The compiler prefers "cheap" conversions.
- Exact matches are the cheapest of all.
- Thus, `f(0)` calls:

```
int f(int i);      // the winner for f(0)!
```
- It's the obvious choice.

20

Best Matches and Ambiguities

- Now, suppose we take away the exact match:

```
int f(int i);           // f(int)
long int f(long int li); // f(long)
char *f(char *pc);      // f(char *)
~~~
f(0);                  // best match?
```

21

Best Matches and Ambiguities

- This one might surprise you:

```
int f(int i);           // f(int)
long int f(long int li); // f(long)
char *f(char *pc);      // f(char *)
~~~
f(0);                  // compile error; why?
```

- C++ ranks all the standard conversions...

22

Standard Conversions by Rank

Conversion	Rank
no conversion	Exact Match "cheap"
array-to-pointer conversion	
qualification conversion	
etc.	
integral promotion	Promotion
floating-point promotion	
integral conversion	Conversion "expensive"
floating-point conversion	
pointer conversion	
boolean conversion	
etc.	

The choice:

$f(\theta) \Rightarrow f(\text{Long int})$?

$f(\theta) \Rightarrow f(\text{char *})$?

23

Standard Conversions by Rank

Conversion	Rank
no conversion	Exact Match "cheap"
array-to-pointer conversion	
qualification conversion	
etc.	
integral promotion	Promotion
floating-point promotion	
integral conversion	Conversion "expensive"
floating-point conversion	
pointer conversion	
boolean conversion	
etc.	

$f(\theta) \Rightarrow f(\text{Long int})$

$f(\theta) \Rightarrow f(\text{char *})$

24

Best Matches and Ambiguities

- In this case:
 - Calling `f(long int)` requires an *integral conversion*.
 - Calling `f(char *)` requires a *pointer conversion*.
- C++ considers these conversions to be equally “expensive”.

```
long int f(long int li);
char *f(char *p);
~~~~
f(0);           // compile error; a tie
```

- Thus, there’s no unique best match.

25

Best Matches and Ambiguities

- When the compiler can’t find at least one function that satisfies a particular call, it typically complains that there’s “*no matching function*”.
- When the compiler finds more than one function that satisfies the call, but no unique best match, it typically complains that the call is “*ambiguous*”.

```
long int f(long int li);
char *f(char *p);
~~~~
f();           // compile error: no matching function
f(0);          // compile error: ambiguous
```

26

Best Matches and Ambiguities

- The literal 0 is the only integer with an implicit conversion to a pointer type.
- The call f(x) is ambiguous only when x is the integer literal 0:

```
long int f(long int li);
char *f(char *p);
int n;
~~~
f(n);           // calls long int f(Long int)
f(1);           // calls long int f(Long int)
```

27

Functions with Multiple Arguments

- How does overload resolution select the best match among functions with multiple arguments?
- For example, suppose that we have the following functions:

```
void f(double x, double y, double z);
void f(double x, int y, double z);
```

- If you eliminate either function, the following call will compile:

```
f(1.1, 2, 3);           // could work with either function by itself
```

- But which is the best match when they're overloaded?

28

Functions with Multiple Arguments

- Here's the general rule.
- Suppose we have a group of overloaded functions F_1, F_2, \dots, F_n .
- For a given call, function F_i is a better match than F_j if:
 - for every argument A_k in the call, F_i 's conversion for A_k is no worse than F_j 's conversion for A_k , and
 - for at least one argument A in the call, F_i 's conversion for A is better than F_j 's conversion for A .

29

Converting Multiple Arguments

- Here, again, is our set of overloaded functions:


```
void f(double x, double y, double z);
void f(double x, int y, double z);
```
- The best match for the following call is the 2nd function:


```
f(1.1, 2, 3); // calls f(double, int, double)
```
- The conversions are the same for both functions on the 1st and 3rd arguments.
- The 2nd function is a better match on the 2nd (middle) argument.

30

Function Signatures

- Again, overloaded functions share the same name.
- Thus, it takes more than just a name to uniquely identify a particular C++ function.
- Thus, each function in a group of overloaded functions must have other properties that make it unique.
- The set of properties that uniquely identify a given function or function template is its ***signature***...

31

Function Signatures

- A function's ***signature*** is primarily the function's ***name*** and ***parameter type list***.
- The function's ***parameter type list*** is the sequence of types in the function's parameter list.
- For example, the declaration for one of standard operator new functions is:

```
void *operator new(std::size_t size, std::align_val_t alignment);
```
- Its parameter type list is:

```
(std::size_t, std::align_val_t)
```

32

Function Signatures

- The exact combination of additional properties that make up a function's signature depends on what "kind" of function it is.
- For example, a ***non-member function***'s signature also includes its enclosing namespace, if any.
- A class member function's signature also includes:
 - the class
 - *cv-qualifiers* (const and volatile), and
 - *ref-qualifiers* (& and &&).

33

Outline

- Function Overloading
- Name Lookup
- Default Function Arguments
- Function Templates
- Tying It All Together

34

Scope Regions and Name Lookup

- When the compiler encounters the declaration of a name, it stores that name and its attributes in a **symbol table**.
- When the compiler encounters a reference to a name, it looks up the name in the symbol table to find those attributes.

```
class widget {           // stores widget
public:
    string name() const; // looks up string; stores name
    ~~~
};

string s;                // looks up string; stores s
widget w;                // looks up widget; stores w
s = w.name();           // looks up s, w, and name
```

35

Name Lookup

- In C++, declarations can appear at:
 - **local scope**: inside a function declaration, including that function's parameter list or a block nested inside a function definition
 - A name declared at local scope is in scope to the end of the function declaration or block containing that name.
 - **class scope**: in the brace-enclosed body of a class definition
 - A name declared at class scope is in scope to the end of its class definition and within the parameter list and body of a member definition of the same class.
 - **namespace scope**: outside of any function, class, structure, or union, *whether global or in some other namespace*
 - A name declared at *namespace* scope is in scope to the end of its *namespace* definition, which for the global scope is the end of its translation unit.

36

Qualified vs. Unqualified Names

- A name appearing just to the right of a `::`, `.` or `->` is a *qualified name*.
- A name that's not qualified is an *unqualified name*.
- For example:

```

string::assign    // string is unqualified
                  // assign is qualified
other.Length()   // other is unqualified
                  // Length is qualified
words[i]->Length() // words is unqualified
                  // i is unqualified
                  // Length is qualified

```

37

Qualified Name Lookup

- Lookup for qualified names is different from lookup for unqualified names.
- Lookup for qualified names is fairly simple:
 - For `S::n`, where `S` is a namespace:
 - Look for `n` in `S`.
 - For `T::n`, where `T` is a class type, or
 - for `x.n`, where `x` is a `T` object, or
 - for `p->n`, where `p` is a pointer to `T`:
 - Look for `n` in `T`.
 - If `n` isn't in `T`, look for `n` in `T`'s base classes (if any) from the direct base to the most indirect base.
- Failure to find the name is a compile error.

38

Unqualified Name Lookup

- Name lookup for unqualified names in C++ is an extension of name lookup in C.
- In this example, `m` and `n` are unqualified names appearing in a non-member function at namespace scope:

```
namespace S {
    int m, n;
    void f(int n) {
        m = n;           // m refers to S::m; n refers to parameter
    }
}
```

39

Unqualified Name Lookup

- For an unqualified name `n` appearing in a non-member function `f` at namespace scope:
 - **Look in the local scope(s).**
 - That is, look for `n` in the scope of the block in which `n` appears.
 - Work outward to the scope of `f`.
 - **Look in the namespace scope(s).**
 - That is, look for `n` in the namespace scope(s) enclosing `f`.
 - Start in the namespace immediately enclosing `f`.
 - Work outward to the global scope.
 - **Stop as soon as you find any declaration for `n`.**
- Again, failure to find the name is a compile error.

40

Unqualified Name Lookup

- In this example, `m` and `n` are unqualified names appearing in function `f` that's a member of a class `T` within a namespace `S`:

```
namespace S {
    int m, n;
    class T {
        void f(int n);
    };
    void T::f(int n) {
        m = n;           // m refers to S::m; n refers to parameter
    }
}
```

41

Unqualified Name Lookup

- For an unqualified name `n` appearing in a function `f` that's a member of class `T`:
 - ***Look in `f`'s local scope(s).***
 - ***Look in the class scope(s).***
 - That is, look for `n` in the scope of `T`.
 - If `n` isn't in `T`, look for `n` in `T`'s base classes (if any) from the direct base to the most indirect base.
 - If the compiler finds `n` as a class member, it interprets `n` as `this->n`.
 - ***Look in the namespace scope(s).***
 - ***Again, stop as soon as you find any declaration for `n`.***
- As always, failure to find the name is a compile error.

42

Name Lookup Precedes Overload Resolution

- Name lookup and overload resolution are distinct compilation steps.
- When a name, say *f*, occurs as a function name in a call, name lookup seeks the first scope that declares at least one instance of *f*.
- If name lookup succeeds — it finds a scope containing at least one *f* — then all the *f*'s in that scope become the *candidates* for the call...

43

Name Lookup Precedes Overload Resolution

- After name lookup, overload resolution tries to identify the best matching function among the candidates.
- Overload resolution must find a unique best match.
- Otherwise, the call produces a compile error.
- If there's a better matching function in an outer scope...
 - name lookup won't find it, and
 - overload resolution won't consider it.

44

Name Lookup Precedes Overload Resolution

- For example, the call on (5) fails because name lookup finds only the `put` declarations at (3) and (4), not the declarations at (1) and (2):

```
int put(int c);                // (1) not found
int put(int c, FILE *f);       // (2) not found
namespace N {
    int put(char const *s);    // (3) found
    int put(char const *s, FILE *f); // (4) found
    void f(int c) {
        put(c);                // (5) error: no valid match
        ~~~
    }
}
```

- Neither the declaration at (3) nor (4) is a viable match, so the call fails.

45

Unqualified Name Lookup and Inheritance

- Similarly, `Derived::f` hides `Base::f` rather than overloading with it:

```
class Base {
public:
    void f(int n);
};

class Derived: public Base {
public:
    void f(double d);    // hides Base::f(int)
};

Derived dx;
dx.f(3);                // calls Derived::f(double)
```

46

Argument-Dependent Lookup

- Actually, it's not *quite* true that overloaded functions “must be in the same scope”.
- There's one more facet of unqualified name lookup to consider: ***argument-dependent lookup (ADL)***.
- ADL is specifically for unqualified function names in function calls.
- ADL adds this name lookup rule:
 - For each argument in the function call whose type is declared in a namespace, look in that namespace for the function name, as well as in other scopes searched by the usual name lookup.

47

Argument-Dependent Lookup

- Here's an example of what would happen if ADL weren't a part of C++:

```
namespace N {
    class T;
    void f(T &r);
}

N::T x;
~~~
f(x);          // compile error w/o ADL: never looks in N
N::f(x);       // OK w/o ADL: finds f in N
```

48

Argument-Dependent Lookup

- At first glance, this may not look so bad — the programmer simply needs to specify that `f` comes from namespace `N`.

`N::f(x);`

- However, suppose that:
 - namespace `N` is actually `std`,
 - class `T` is actually `string`, and
 - function `f` is actually `operator<<`.
- Altogether, those changes yield this simplified version of code from the standard `<string>` header...

49

Argument-Dependent Lookup

```
// <string>

#include <iosfwd>

namespace std {
    class string;
    ostream &operator<<(ostream &, string const &);
    ~~~
}
```

- Again, this is a simplified version of the `<string>` header.
- The following code uses this header...

50

Argument-Dependent Lookup

- Without ADL, the compiler would reject these calls to `operator<<` because it wouldn't look for `operator<<` in namespace `std`:

```
#include <iostream>
#include <string>

std::string s;
std::cout << s;           // compile error w/o ADL
operator<<(std::cout, s);  // compile error w/o ADL
```

- You'd need to call it using the function syntax and explicitly specify that `operator<<` is a member of `std`:

```
std::operator<<(std::cout, s);  // OK w/o ADL
```

51

Sutter's Interface Principle

- Sutter [2000] offers some advice for grouping code into namespaces based on what he calls the **Interface Principle**:

- For a class `X`, all functions, including free functions, that both "mention" `X` and are "supplied with" `X` are logically part of `X`, because they form part of the interface of `X`.

- For example, this function "mentions" `string` and is "supplied with" `string`:

```
ostream &operator<<(ostream &, string const &);
```

- Thus, it's part of `string`'s interface, even though it's not a member, and may not even be a friend.

52

Sutter's Interface Principle

- He then offers this advice:
 - ✓ *If you put a class into a namespace, be sure to put of its interface functions into the same namespace.*
- If you don't, you may be unhappy with how your code behaves.

53

ADL and Overloading

- In a function call, overload resolution also considers function declarations found by ADL.
- That is, function declarations found by ADL are added to the candidate set as if they had been found through unqualified name lookup.

54

Outline

- Function Overloading
- Name Lookup
- Default Function Arguments
- Function Templates
- Tying It All Together

55

Default Arguments

- In the Standard Library, `stdout` is an object that can be passed as a `FILE *` to represent the standard output stream.
- Here is a pair of overloaded functions we saw earlier:

```
int put(int c);           // write c to standard output
int put(int c, FILE *f);  // write c to a designated file *f
```

- You could declare it as a single function with a ***default argument*** for the second parameter:

```
int put(int c, FILE *f = stdout);
```

- With this change, the compiler produces slightly different code...

56

How Many Functions?

- When declared as a pair of overloaded function, each function has a distinct signature:
 - `int put(int c)`
 - is a function named “put” whose parameter type list is `(int)`.
 - `int put(int c, FILE *f)`
 - is a function named “put” whose parameter type list is `(int, FILE *)`.
- When declared as a single function with a default argument, it has the same signature as the second overloaded function:
 - `int put(int c, FILE *f = stdout)`
 - is a function named “put” whose parameter type list is `(int, FILE *)`.
- The default argument value is *not* part of the signature.

57

Default Arguments and Overload Resolution

- When measuring the “cost” of selecting an overloaded function, filling in a default argument is considered “free”.
- For example, suppose that we have the following two overloaded functions:


```
void g(double d);
void g(int x, int y = 1);
```
- Which is the best match for calling `g(0)`?
- Consider the conversions...

58

Default Arguments and Overload Resolution

- Calling `g(double)` requires converting `int` into `double`.
 - This is a ***floating-point conversion***.
- Calling `g(int, int = 1)` requires no conversion at all.
 - `0` is an ***exact match*** for the first parameter `int`.
 - Applying the default argument to the second `int` parameter isn't considered a "conversion" of any kind.
- Thus, `g(0)` calls:

```
void g(int x, int y = 1);      // the winner for g(0)!
```

59

Outline

- Function Overloading
- Name Lookup
- Default Function Arguments
- Function Templates
- Tying It All Together

60

Function Templates

- A *function template* is a generalization of an algorithm.
 - It's *not an actual function*.
- Rather, it's a single declaration that can generate declarations for similar, but distinct functions.
- Each generated function implements the algorithm for operands of different types.

61

Function Templates

- For example, the standard *<algorithm>* header provides a `max` function template that looks something like this:

```
template <typename T>
constexpr T const &max(T const &a, T const &b) {
    return (a > b) ? a : b;
}
```

62

Template Argument Deduction

- C++ often lets you omit the angle-bracketed template argument list from a call to a function template.
- That is, the call can use just the template name as the function name.

```
int x = 10;
int y = 20;
int z = max(x, y); // max, not max<int>
```

- In this case, the compiler performs **template argument deduction**:
 - It deduces the template argument(s) from the function call argument(s).

63

Template Argument Deduction

- Template argument deduction makes a function template look more like an unbounded set of overloaded functions, as in:

```
template <typename T>
constexpr T const &max(T const &a, T const &b);
```

- C++ can deduce the type argument for each of these calls:

```
int i, j;
float f, g;
~~~
int k = max(i, j);    // calls max<int>(i, j)
float h = max(f, g);  // calls max<float>(f, g)
```

64

Specified and Deduced Type Arguments

- When you call a function template, you can either specify the template arguments explicitly or let the compiler deduce them:

```
int i, j;
float f, g;
~~~
int k = max<int>(i, j);    // (1) type specified as int
float h = max(f, g);      // (2) type deduced as float
```

- However, these calls produce slightly different behavior...

65

Specified and Deduced Type Arguments

- When performing template argument deduction, the compiler doesn't consider most type conversions.
- For example, although there's normally a standard conversion from `int` to `double`, the compiler won't convert `x` into a `double` in this call:

```
int x = 1;
double y = 2.5;
double z = max(x, y);    // compile error: can't deduce T
```

- Instead, the compiler rejects the call.

66

Specified and Deduced Type Arguments

- However, if you explicitly call `max<double>`, the compiler will perform the implicit conversion:

```
int x = 1;
double y = 2.5;
double z = max<double>(x, y);    // OK: converts x to double
```

- Similarly, calling `max<int>` would implicitly convert `y` into an `int`:

```
int n = max<int>(x, y);          // OK: converts y to int
```

67

Overloading and Templates

- You can overload functions and function templates with each other.
- When an overloaded function and function template provide equally good matches, the compiler chooses the non-template function as the better match.
- Using this behavior, you can effectively customize a function template for specific argument types...

68

Overloading and Templates

- For example, you could write a `max` function for C-strings that uses `strcmp` to go along with the general `max` template:

```
template <typename T>
constexpr T const &max(T const &a, T const &b);

constexpr char const *max(char const *a, char const *b) {
    return strcmp(a, b) > 0 ? a : b;
}
```

69

Overloading and Templates

- This call invokes the `max` that specifically takes `char const *` objects:

```
char const N[] = "Nancy";
char const D[] = "Dan";
char const *p;
~~~
p = max(D, N); // calls non-template max
```

- In the very-unlikely event that you actually want to call the template version of `max` on `char const *` objects, you can do so explicitly:

```
p = max<char const *>(D, N);
```

70

Outline

- Function Overloading
- Name Lookup
- Default Function Arguments
- Function Templates
- Tying It All Together

71

The `std::swap` Two-Step

- The most common example that combines most of these features is probably the “`std::swap` Two-Step” [Niebler 2014].
- The “Two-Step” is a programming technique that allows for a ***customization point***.
 - It’s especially valuable for library authors writing templates that will be widely used, though it also has other uses.
- In the titular example, that customization point is a `swap` function...

72

The std::swap Two-Step

- Here's a possible implementation for the std::swap function template from *<algorithm>*:

```
template <typename T>
void std::swap(T &a, T &b) {
    T temp {std::move(a)};
    a = std::move(b);
    b = std::move(temp);
}
```

- Because this function template uses basic move semantics, it accepts a wide variety of types.
- However, a type-specific swap function can be more efficient...

73

The std::swap Two-Step

- For example, the Saks::string class from earlier might benefit from providing its own “custom” swap function:

```
namespace Saks {
    class string {
        ~~~
        friend void swap(string &a, string &b);
    private:
        std::size_t stored_length;
        char *actual_str;
    };
    void swap(string &a, string &b);
}
```

- You'll see why I've written swap as a friend function shortly.

74

The std::swap Two-Step

- This swap function is probably more efficient for Saks::strings, even taking move semantics into account:

```
void Saks::swap(string &a, string &b) {
    std::swap(a.stored_length, b.stored_length);
    std::swap(a.actual_str, b.actual_str);
}
```

- Unfortunately, it's still possible to use std::swap with Saks::strings by accident...

75

The std::swap Two-Step

- For example, here's another type Person with a custom swap function that has a Saks::string as a data member:

```
class Person {
public:
    ~~~
    void swap(Person &other);
    ~~~
private:
    Saks::string name;
    unsigned idnum;
};
```

76

The std::swap Two-Step

- If you didn't know that Saks::string had a custom swap function, you might write Person::swap like this:

```
void Person::swap(Person &other) {
    std::swap(name, other.name);    // oops, uses std::swap
    std::swap(idnum, other.idnum);
};
```

- Fortunately, because of how Saks::string provides its custom swap function, there's a simple technique that you can use to avoid this mistake...

77

The std::swap Two-Step

- If you instead write Person::swap like this, the compiler will find the custom swap function for Saks::string:

```
void Person::swap(Person &other) {
    using std::swap;
    swap(name, other.name);    // OK, uses Saks::swap
    swap(idnum, other.idnum);
};
```

- If Saks::string doesn't have a custom swap function, the compiler will automatically fall back to using std::swap.
- Although you can apply this technique easily, the effect is fairly complex...

78

How It Works

- If you call the function as `std::swap`, the compiler uses qualified name lookup and goes directly to namespace `std` to find the `swap` function:

```
void Person::swap(Person &other) {  
    std::swap(name, other.name);  
    std::swap(idnum, other.idnum);  
};
```

79

How It Works

- However, if you call the function as simply `swap`, the compiler searches for `swap` using *unqualified* name lookup:

```
void Person::swap(Person &other) {  
    using std::swap;  
    swap(name, other.name);  
    swap(idnum, other.idnum);  
};
```

- In other words, the compiler:
 - searches outward for declarations for `swap` from the point of the call, and
 - allows argument-dependent lookup (ADL).

80

How It Works — Name Lookup

- Because of the local using-declaration, the compiler immediately finds `std::swap` as if it were declared locally:

```
void Person::swap(Person &other) {
    using std::swap;
    swap(name, other.name);
    swap(idnum, other.idnum);
};
```

- Thus, the `std::swap` function template is a candidate for the calls to `swap`.
- However, the compiler also performs ADL...

81

How It Works — Name Lookup

- Because `name` is a `Saks::string` object, the compiler also finds the `swap` in namespace `Saks`:

```
namespace Saks {
    class string { ~~~ };
    void swap(string &a, string &b);
}
```

- Thus, `Saks::swap` is also a candidate for the calls to `swap`.
- Now, the compiler performs overload resolution to choose between `std::swap` and `Saks::swap`...

82

How It Works — Overload Resolution

- Both `std::swap<Saks::string>` and `Saks::swap` take two arguments of type `string &`.
- Thus, both functions provide exact matches for both arguments.
- However, `std::swap` is a function template, while `Saks::swap` is an ordinary function.
- Thus, the compiler selects `Saks::swap` as the best match.
- If `Saks::swap` hadn't existed, `std::swap` would've been the only candidate function, and thus would've been the best match by default.

83

Important Details

- Note that this technique works specifically because both `std::swap` and `Saks::swap` are non-member functions that take two arguments.
- If `Saks::string`'s `swap` function had been a member function, you would've needed to write the calls differently:

```
void Person::swap(Person &other) {  
    name.swap(other.name);  
    std::swap(idnum, other.idnum);  
};
```

- That being the case, `Person::swap` should probably be rewritten as a non-member function as well.

84

Important Details

- It's also important that `Saks::string`'s `swap` function is part of the `Saks` namespace.
- If the custom `swap` function was somewhere else, ADL wouldn't find it.
- Remember: the local using-declaration for `std::swap` ends the normal name lookup process immediately.
 - There's nowhere else that you could put the custom `swap` function where it would be selected as a candidate.

85

Takeaways

- Name lookup precedes overload resolution.
- To allow for ADL, remember Sutter's advice regarding interface design:
 - If you put a class into a namespace, be sure to put all helper functions and operators into the same namespace.
- When a function and a function template are equally good matches for a call, overload resolution favors the non-template.
- Be aware of customization points, and use the Two-Step technique to avoid accidentally calling the non-custom version of a function.

86

Bibliography

- Niebler [2014]. *Customization Point Design in C++11 and Beyond*.
<https://ericniebler.com/2014/10/21/customization-point-design-in-c11-and-beyond/>
- Sutter [2000]. Herb Sutter, *Exceptional C++: 47 Engineering Puzzles, Programming Problems, and Solutions*. Addison-Wesley.