

C++ Exceptions for Smaller Firmware

By: Khalil Estell (@kammce)

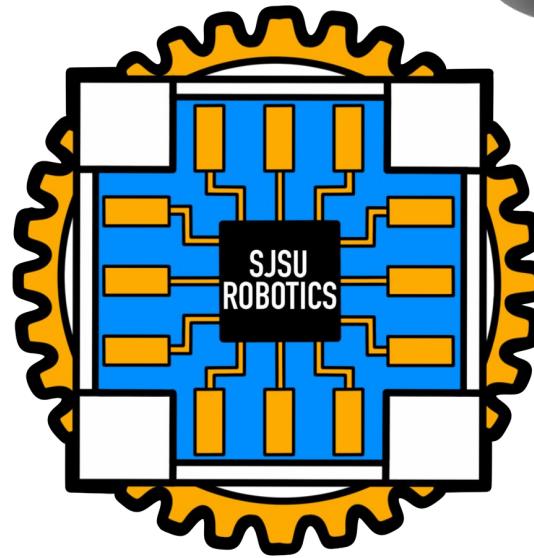
@kammce

Graduated 2017 from SJSU

Ex-Googler (~6y)

Volunteer @ SJSU

ISO C++ Committee Member



Motivation

A man in a dark long-sleeved shirt is standing in front of a whiteboard, pointing with a pen towards a diagram. He appears to be explaining the I2C bus architecture. The whiteboard shows a hierarchical tree diagram with 'I2C Master1' at the top, connected to 'SDA' and 'SCL' lines, which then branch down to multiple 'I2C Slave1' nodes. In the top right corner of the video frame, there are two black rounded rectangles containing white icons: a clock and a play button. In the bottom right corner, a timestamp displays '1:32:10'. Below the video, the title 'SJSU CMPE 146 I2C Lecture 10-03-17' is shown in large white text, followed by a three-dot menu icon. At the bottom left, the view count '283 views • 6 years ago' is displayed.

SJSU CMPE 146 I2C Lecture 10-03-17

283 views • 6 years ago

libhal / SJSU-Dev2 Public

<> Code Issues 129 Pull request

master ▾

7 Branches



FluffyFTW and kammce fix



February 15, 2024

Editors:

- Bjarne Stroustrup
- Herb Sutter



Microcontroller Constraints: No Operating System



~~fopen()~~



~~networking~~



~~std::printf()~~



~~std::chrono::steady_clock::now();~~



~~std::thread~~

Microcontroller Constraints: Limited Storage & Memory

< 1 MB = 33,725 devices

,779 | Smart Filtering ? Results remaining: 33725

Semiconductors > Embedded Processors & Controllers > Microcontrollers - MCU X

Program Memory Size	Data Bus Wid
576 kB	8 bit
640 kB	8 bit/16 bit
768 kB	16 bit
832 kB	16 bit/32 bit
1 MB	32 bit
1.024 MB	32 bit/16 bit
1024 kB	
1.06 MB	
1088 kB	
1.16 MB	
1.25 MB	
1.28 MB	
1.3 MB	
1.5 MB	

≤ ≥

Reset

>= 1 MB = 4,964 devices

,779 | Smart Filtering ? Results remaining: 4964

Semiconductors > Embedded Processors & Controllers > Microcontrollers - MCU X

Program Memory Size	Data Bus Wid
576 kB	8 bit
640 kB	8 bit/16 bit
768 kB	16 bit
832 kB	16 bit/32 bit
1 MB	32 bit
1.024 MB	32 bit/16 bit
1024 kB	
1.06 MB	
1088 kB	
1.16 MB	
1.25 MB	
1.28 MB	
1.3 MB	
1.5 MB	

≤ ≥

Reset

Microcontroller Constraints: Limited Storage & Memory

< 100kB = 31,098 devices

,779 | Smart Filtering ? Results remaining: 31098

Semiconductors	Embedded Processors & Controllers	Microcontrollers - MCU																																																	
<table border="1"> <thead> <tr> <th>Maximum Clock Frequency</th> <th>Number of I/Os</th> <th>Data RAM Size</th> </tr> </thead> <tbody> <tr> <td>625 kHz</td> <td>2 I/O</td> <td>66 kB</td> </tr> <tr> <td>1 MHz</td> <td>4 I/O</td> <td>68 kB</td> </tr> <tr> <td>1.2 MHz</td> <td>5 I/O</td> <td>69 kB</td> </tr> <tr> <td>1.3 MHz</td> <td>6 I/O</td> <td>76 kB</td> </tr> <tr> <td>2 MHz</td> <td>7 I/O</td> <td>80 kB</td> </tr> <tr> <td>2.1 MHz</td> <td>8 I/O</td> <td>82 kB</td> </tr> <tr> <td>3 MHz</td> <td>9 I/O</td> <td>92 kB</td> </tr> <tr> <td>4 MHz</td> <td>10 I/O</td> <td>94 kB</td> </tr> <tr> <td>4.2 MHz</td> <td>11 I/O</td> <td>96 kB</td> </tr> <tr> <td>5 MHz</td> <td>12 I/O</td> <td>96 kB, 172 kB</td> </tr> <tr> <td>8 MHz</td> <td>13 I/O</td> <td>96 kB, 216 kB</td> </tr> <tr> <td>8.2 MHz</td> <td>14 I/O</td> <td>98 kB</td> </tr> <tr> <td>8.4 MHz</td> <td>15 I/O</td> <td>100 kB</td> </tr> <tr> <td>10 MHz</td> <td>16 I/O</td> <td>104 kB</td> </tr> </tbody> </table>	Maximum Clock Frequency	Number of I/Os	Data RAM Size	625 kHz	2 I/O	66 kB	1 MHz	4 I/O	68 kB	1.2 MHz	5 I/O	69 kB	1.3 MHz	6 I/O	76 kB	2 MHz	7 I/O	80 kB	2.1 MHz	8 I/O	82 kB	3 MHz	9 I/O	92 kB	4 MHz	10 I/O	94 kB	4.2 MHz	11 I/O	96 kB	5 MHz	12 I/O	96 kB, 172 kB	8 MHz	13 I/O	96 kB, 216 kB	8.2 MHz	14 I/O	98 kB	8.4 MHz	15 I/O	100 kB	10 MHz	16 I/O	104 kB	<table border="1"> <thead> <tr> <th>≤</th> <th>≥</th> </tr> </thead> </table>	≤	≥	<table border="1"> <thead> <tr> <th>≤</th> <th>≥</th> </tr> </thead> </table>	≤	≥
Maximum Clock Frequency	Number of I/Os	Data RAM Size																																																	
625 kHz	2 I/O	66 kB																																																	
1 MHz	4 I/O	68 kB																																																	
1.2 MHz	5 I/O	69 kB																																																	
1.3 MHz	6 I/O	76 kB																																																	
2 MHz	7 I/O	80 kB																																																	
2.1 MHz	8 I/O	82 kB																																																	
3 MHz	9 I/O	92 kB																																																	
4 MHz	10 I/O	94 kB																																																	
4.2 MHz	11 I/O	96 kB																																																	
5 MHz	12 I/O	96 kB, 172 kB																																																	
8 MHz	13 I/O	96 kB, 216 kB																																																	
8.2 MHz	14 I/O	98 kB																																																	
8.4 MHz	15 I/O	100 kB																																																	
10 MHz	16 I/O	104 kB																																																	
≤	≥																																																		
≤	≥																																																		

≥ 100kB = 7095

,779 | Smart Filtering ? Results remaining: 7095

Semiconductors	Embedded Processors & Controllers	Microcontrollers - MCU																																																	
<table border="1"> <thead> <tr> <th>Maximum Clock Frequency</th> <th>Number of I/Os</th> <th>Data RAM Size</th> </tr> </thead> <tbody> <tr> <td>625 kHz</td> <td>2 I/O</td> <td>66 kB</td> </tr> <tr> <td>1 MHz</td> <td>4 I/O</td> <td>68 kB</td> </tr> <tr> <td>1.2 MHz</td> <td>5 I/O</td> <td>69 kB</td> </tr> <tr> <td>1.3 MHz</td> <td>6 I/O</td> <td>76 kB</td> </tr> <tr> <td>2 MHz</td> <td>7 I/O</td> <td>80 kB</td> </tr> <tr> <td>2.1 MHz</td> <td>8 I/O</td> <td>82 kB</td> </tr> <tr> <td>3 MHz</td> <td>9 I/O</td> <td>92 kB</td> </tr> <tr> <td>4 MHz</td> <td>10 I/O</td> <td>94 kB</td> </tr> <tr> <td>4.2 MHz</td> <td>11 I/O</td> <td>96 kB</td> </tr> <tr> <td>5 MHz</td> <td>12 I/O</td> <td>96 kB, 172 kB</td> </tr> <tr> <td>8 MHz</td> <td>13 I/O</td> <td>96 kB, 216 kB</td> </tr> <tr> <td>8.2 MHz</td> <td>14 I/O</td> <td>98 kB</td> </tr> <tr> <td>8.4 MHz</td> <td>15 I/O</td> <td>100 kB</td> </tr> <tr> <td>10 MHz</td> <td>16 I/O</td> <td>104 kB</td> </tr> </tbody> </table>	Maximum Clock Frequency	Number of I/Os	Data RAM Size	625 kHz	2 I/O	66 kB	1 MHz	4 I/O	68 kB	1.2 MHz	5 I/O	69 kB	1.3 MHz	6 I/O	76 kB	2 MHz	7 I/O	80 kB	2.1 MHz	8 I/O	82 kB	3 MHz	9 I/O	92 kB	4 MHz	10 I/O	94 kB	4.2 MHz	11 I/O	96 kB	5 MHz	12 I/O	96 kB, 172 kB	8 MHz	13 I/O	96 kB, 216 kB	8.2 MHz	14 I/O	98 kB	8.4 MHz	15 I/O	100 kB	10 MHz	16 I/O	104 kB	<table border="1"> <thead> <tr> <th>≤</th> <th>≥</th> </tr> </thead> </table>	≤	≥	<table border="1"> <thead> <tr> <th>≤</th> <th>≥</th> </tr> </thead> </table>	≤	≥
Maximum Clock Frequency	Number of I/Os	Data RAM Size																																																	
625 kHz	2 I/O	66 kB																																																	
1 MHz	4 I/O	68 kB																																																	
1.2 MHz	5 I/O	69 kB																																																	
1.3 MHz	6 I/O	76 kB																																																	
2 MHz	7 I/O	80 kB																																																	
2.1 MHz	8 I/O	82 kB																																																	
3 MHz	9 I/O	92 kB																																																	
4 MHz	10 I/O	94 kB																																																	
4.2 MHz	11 I/O	96 kB																																																	
5 MHz	12 I/O	96 kB, 172 kB																																																	
8 MHz	13 I/O	96 kB, 216 kB																																																	
8.2 MHz	14 I/O	98 kB																																																	
8.4 MHz	15 I/O	100 kB																																																	
10 MHz	16 I/O	104 kB																																																	
≤	≥																																																		
≤	≥																																																		

With so little memory how do
you prevent
running out of memory
or
memory fragmentation?

Typical Best Practice: Do not/avoid dynamic allocations

~~new~~

~~call~~ee~~~~

~~free~~

~~malloc~~

~~delete~~

~~realloc~~

Typical Best Practice: Do not/avoid dynamic allocations

~~new~~

~~std::string~~

~~calloc~~

~~std::unique_ptr~~

~~free~~

~~std::vector~~

~~malloc~~

~~delete~~

~~std::unordered_map~~

~~realloc~~

~~std::shared_ptr~~



Designed and maintained by
John Wellbelove

Embedded Template Library

A C++ template library for embedded applications
MIT licensed

 [Is the ETL free?](#)

[Home](#) [Overview](#) [Documentation](#) [Tutorials](#) [Links](#) [About](#) [Contact](#)

release date **december 2023** c++ **98/03/11/14/17** license **MIT** contributors **122**

 **msvc**   **gcc**   **clang**   **build**   **code quality** 

 [Arctic Code Vault](#)

Motivation

C++ is a great language to use for embedded applications and templates are a powerful aspect of it. The standard library can offer a great deal of well tested functionality, but there are some parts that do not fit well with deterministic behaviour and limited resource requirements. These limitations usually preclude the use of dynamically allocated memory which means that the STL containers are unusable.

What is needed is a template library where the user can declare the size, or maximum size of any object upfront. As most embedded compilers do not currently support the standard beyond C++ 03, it would also be nice to have access to some of the features introduced in the later library.

Lets not forget about

- `std::inplace_vector`: [p0843r8](#) voted into C++26 @ St Louis 2024 ISO C++ meeting

inplace_vector

A dynamically-resizable vector with fixed capacity and embedded storage

Document number: P0843R8.

Date: 2023-06-16.

Authors: Gonzalo Brito Gadeschi, Timur Doumler, Nevin Liber, David Sankel <dsankel_at_adobe.com>.

Reply to: Timur Doumler <papers_at_timur.audio>.

Audience: LEWG.

Table of Contents

- [inplace_vector](#)
- [Introduction](#)
- [Motivation and Scope](#)
- [Existing practice](#)
- [Design](#)

So why do firmware developers avoid C++ exceptions?

SPACE

Requires Dynamic Memory/Heap

Requires Malloc

Unbounded Memory Usage

Increases Binary Size

Requires whole
C++ STL

Run Time Type Info
(RTTI)

Exception Tables

Exception Code

TIME

Nondeterministic

Type Comparison (dynamic_cast)



Slow

Binary Search

Frame Unwinding

Frame Evaluation

Back to Class!

[← Files](#)

master ▾

SJSU-Dev2 / library / peripherals / lpc40xx /

[↑ Top](#)

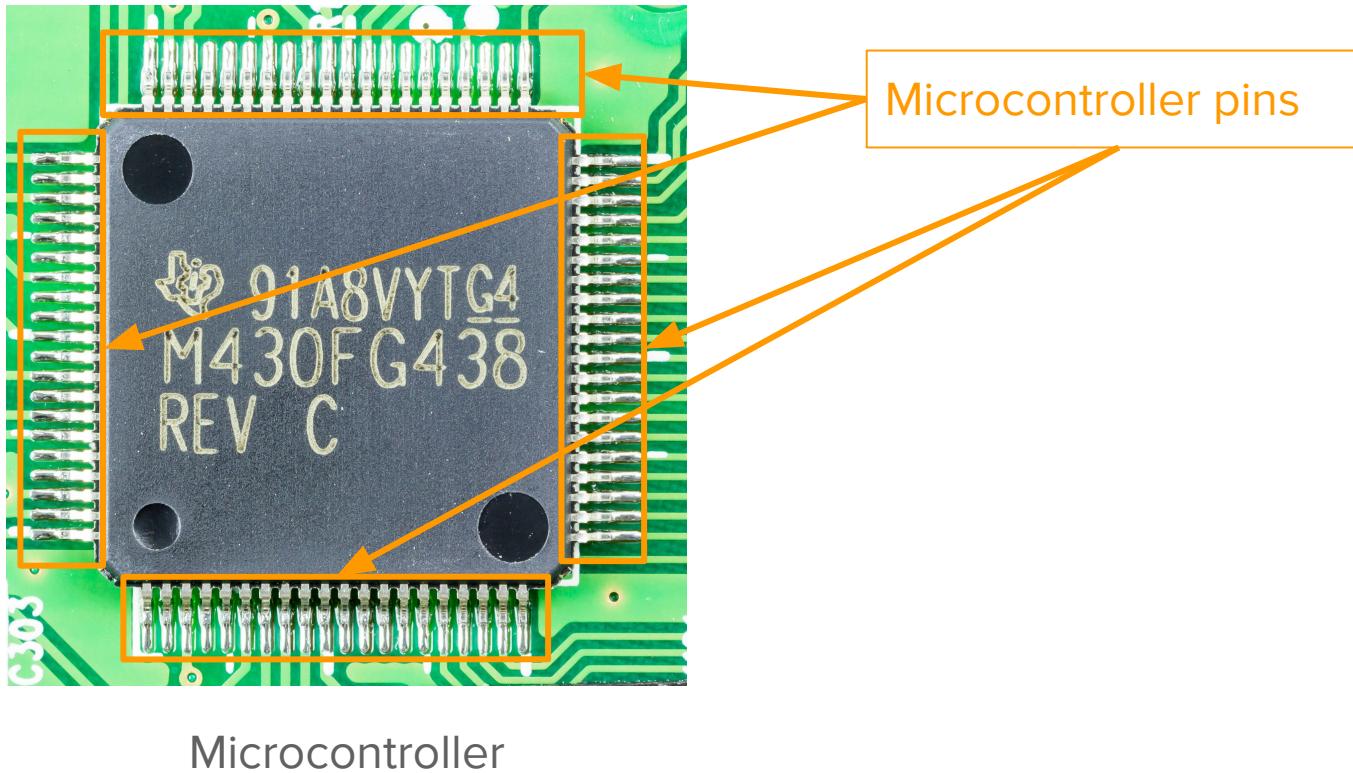
gpio.hpp	Reorganize SJSU-Dev2 Library structure	3 years ago
i2c.hpp	fixed lpx40xx i2c pin functions and pin	last year
pin.hpp	Reorganize SJSU-Dev2 Library structure	3 years ago
pulse_capture.hpp	Reorganize SJSU-Dev2 Library structure	3 years ago
pwm.hpp	Reorganize SJSU-Dev2 Library structure	3 years ago
spi.hpp	Reorganize SJSU-Dev2 Library structure	3 years ago
system_controller.hpp	Fix Issue with LPC40xx UART Baud Rat...	3 years ago
timer.hpp	Reorganize SJSU-Dev2 Library structure	3 years ago
uart.hpp	Add IsConnected implementation to ES...	2 years ago

Consider the output pin

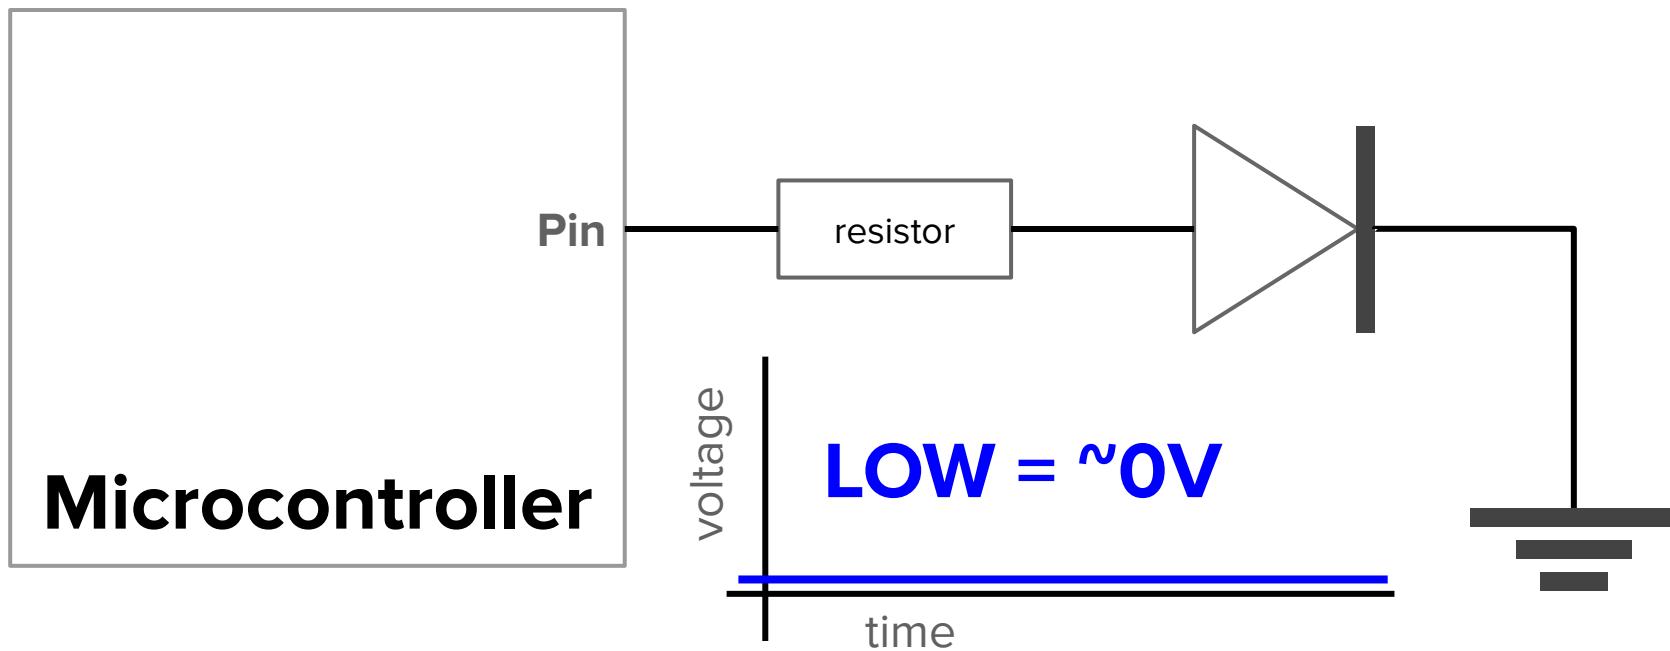


Microcontroller

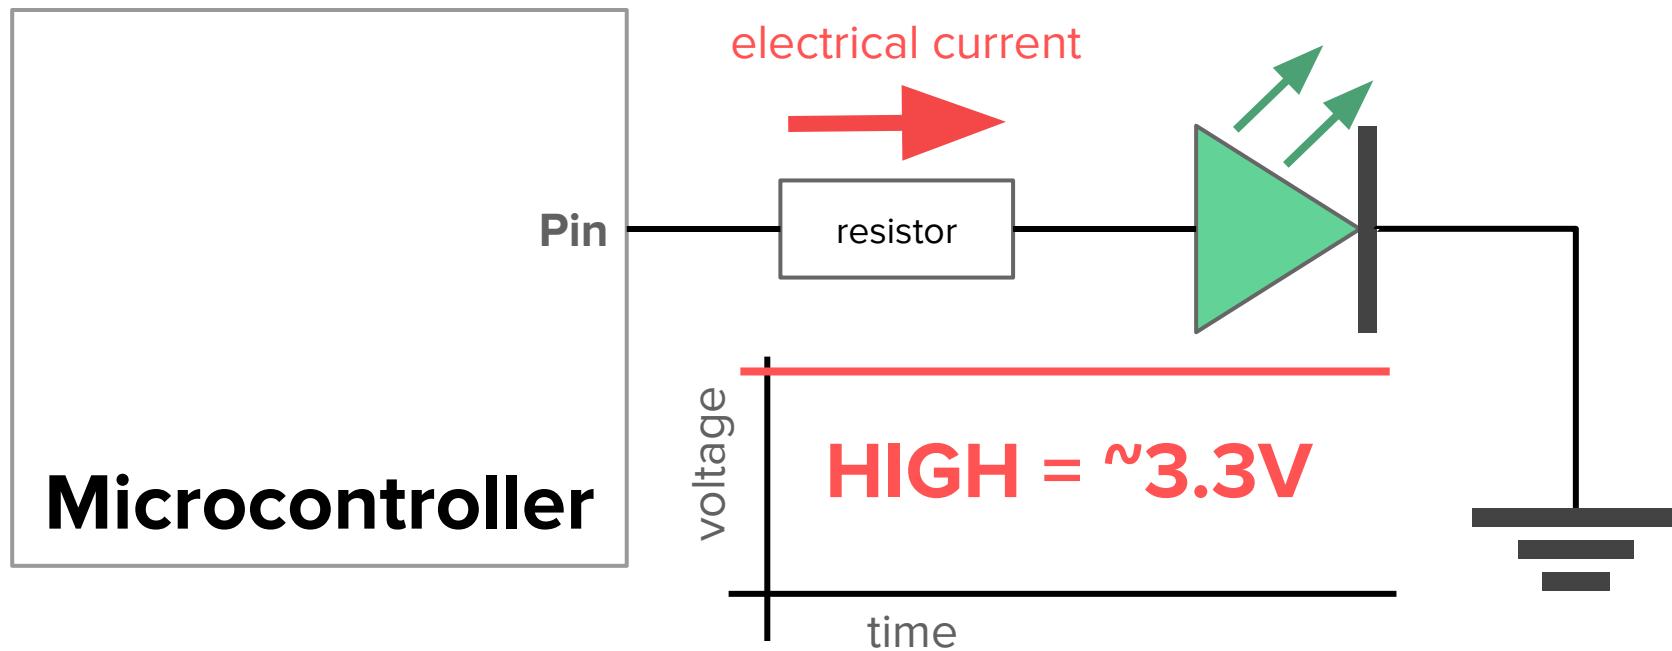
Consider the output pin



Consider the output pin



Consider the output pin



```
class output_pin

{
public:

    virtual ~output_pin() = default;

    virtual void level(bool p_high) = 0;

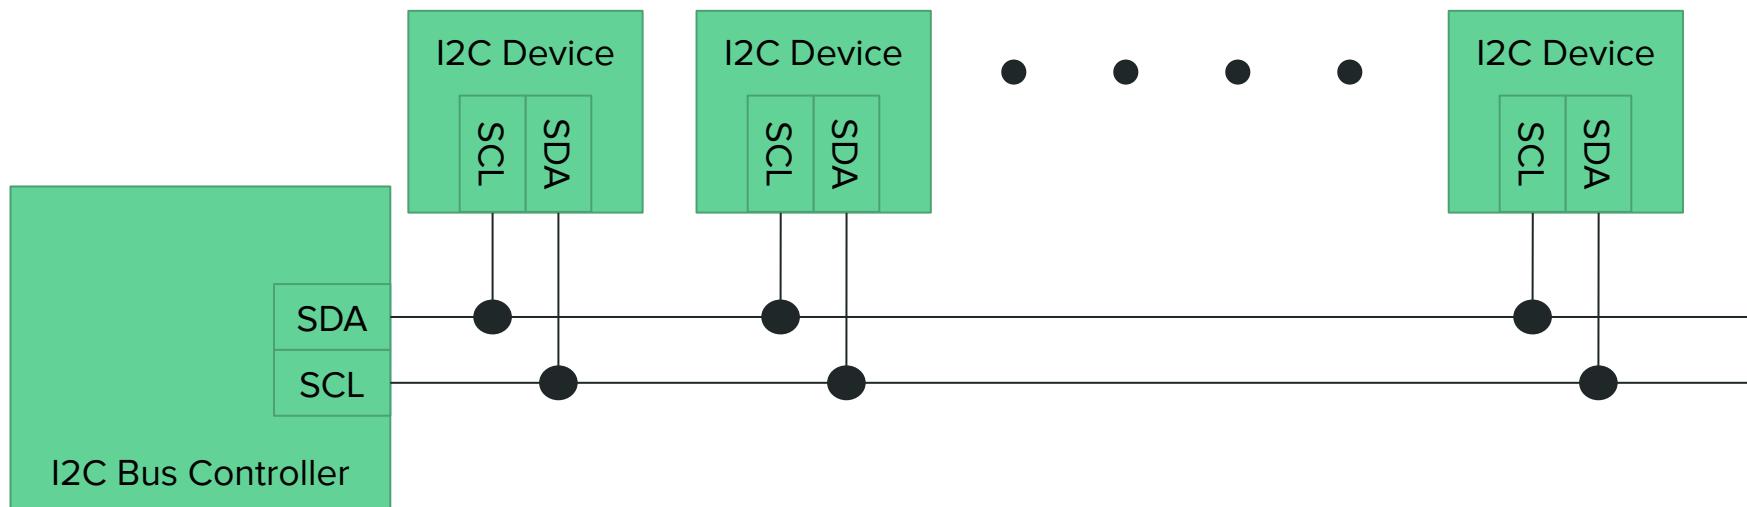
};
```

```
class output_pin  
{  
public:  
    virtual ~output_pin() = default;  
    virtual void level(bool p_high) = 0;  
};
```

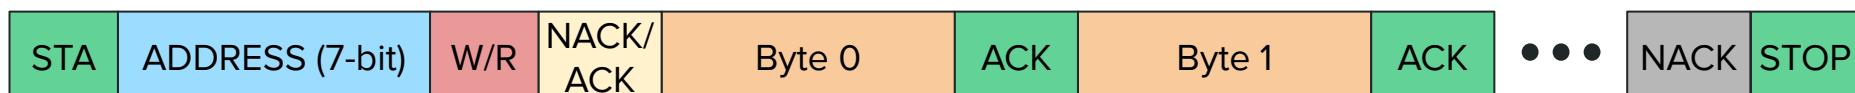


Is returning void okay?

A simplified Introduction to I2C



Wanna Learn More? i2c-bus.org



Possible I2C Errors



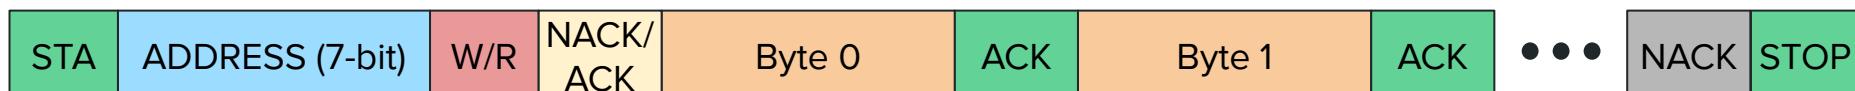
ADDR NACK

Device Not Present



IO ERROR

Unexpected Bus State



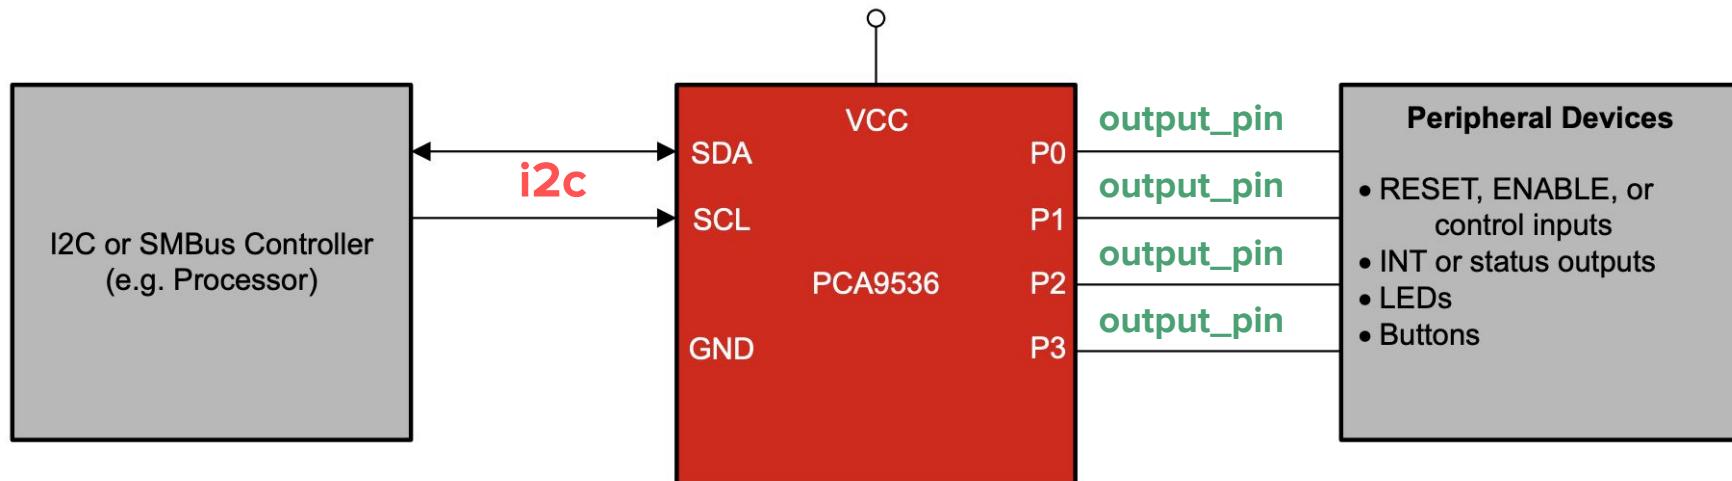
Let's consider an I²C I/O expander



PCA9536

SCPS125H – APRIL 2006 – REVISED MARCH 2022

PCA9536 Remote 4-Bit I²C and SMBus I/O Expander with Configuration Registers



Simplified Schematic

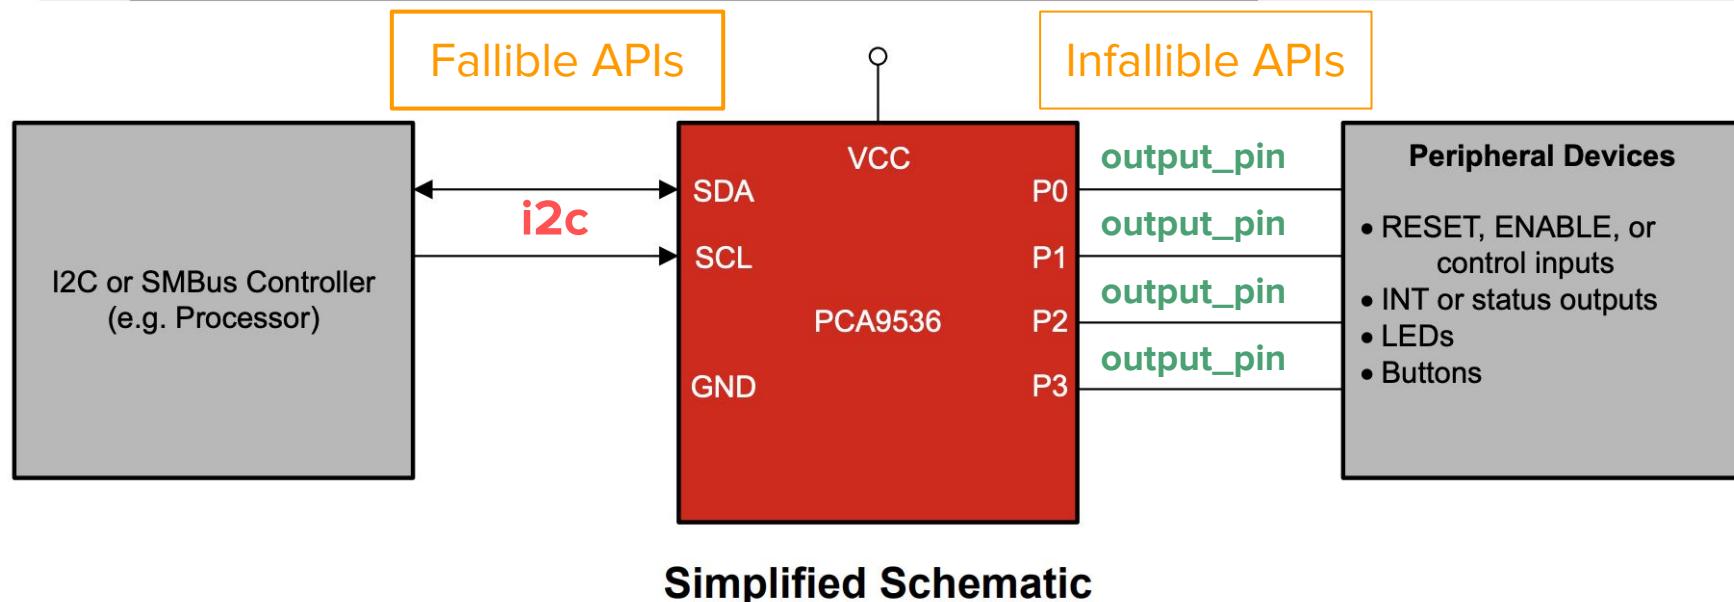
Let's consider an I²C I/O expander



PCA9536

SCPS125H – APRIL 2006 – REVISED MARCH 2022

PCA9536 Remote 4-Bit I²C and SMBus I/O Expander with Configuration Registers



Illustrating the problem: Return mismatch

```
/// @returns false if this failed!  
  
bool io_expander_level(std::uint8_t p_pin_number, bool p_high_voltage) {  
    // Implementation details...  
}  
  
void some_output_pin::level(bool p_high_voltage) {  
    auto success = io_expander_level(m_pin, p_high_voltage);  
    if (not success) {  
        // What do you do here ???  
    }  
}
```

Consider how output pin may be used?

```
class led {  
    led(output_pin& p_pin);  
};
```



not severe

```
class heating_element {  
    heating_element (output_pin& p_control);  
};
```

!! severe A red circular icon containing a white exclamation mark, indicating a severe warning.

```
class emergency_relay_cut_off {  
    emergency_relay_cut_off (output_pin&  
    p_control);  
};
```

```
class output_pin  
{  
public:  
    virtual ~output_pin() = default;  
    virtual status level(bool p_high) = 0;  
    virtual result<bool> level() = 0;  
};
```



std::expected<T, my_error_t>

std::expected<T, E>

```
#include <expected>

enum class my_error {
    precondition_violated, invalid_arguments, // etc ...
};

std::expected<int, my_error> get_version() {
    if (device_disconnected) {
        return std::unexpected(my_error::precondition_violated);
    }
    return 15;
}
```

The code if we used **exceptions**...

```
void toggle_led(output_pin& p_pin,
                 milliseconds p_transition_time) {
    p_pin.level(true);
    delay(p_transition_time / 2);
    p_pin.level(false);
    delay(p_transition_time / 2);
}
```

The code if we used **std::expected**

Notice the return type change?

```
status toggle_led(
    output_pin& p_pin,
    milliseconds p_transition_time) {

    if (auto status = p_pin.level(true);
        !status) {
        return std::unexpected(status.error());
    }

    delay(p_transition_time / 2);

    if (auto status = p_pin.level(false);
        !status) {
        return std::unexpected(status.error());
    }

    delay(p_transition_time / 2);
}
```

The code if we used **Error Propagator P2561R1 + std::expected...**

```
status toggle_led(
    output_pin& p_pin,
    milliseconds p_transition_time) {
    p_pin.level(true) ??
    delay(p_transition_time / 2);
    p_pin.level(false) ??
    delay(p_transition_time / 2);
}
```

Return type is still status

Settled on **std::expected** + Error Propagator Macro

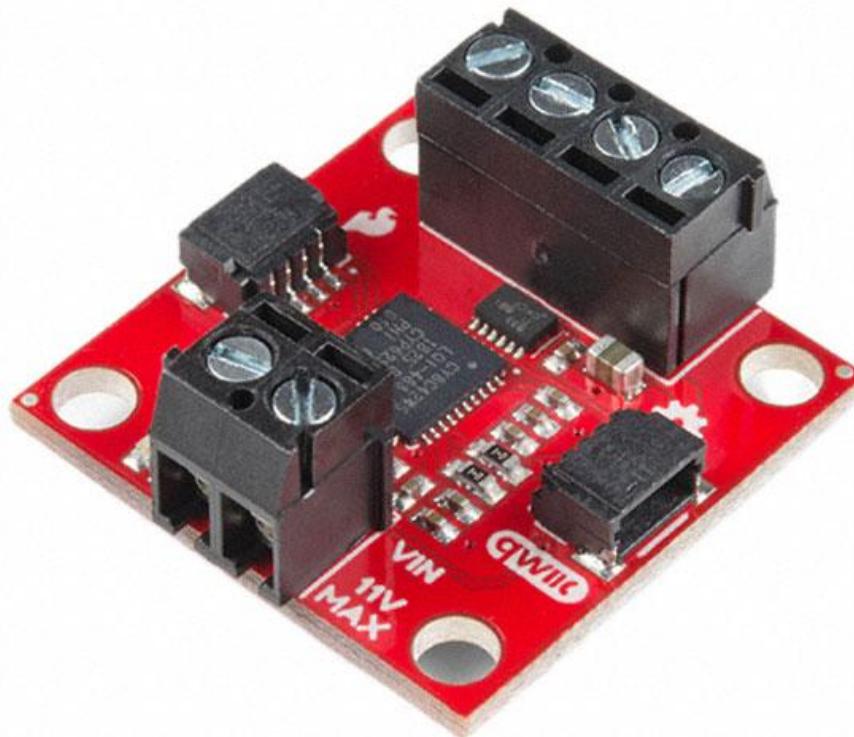
```
status toggle_led(output_pin& p_pin,
                   milliseconds p_delay) {
    SJ_CHECK(p_pin.level(true));
    delay(p_delay / 2);
    SJ_CHECK(p_pin.level(false));
    delay(p_delay / 2);
    return success();
}
```

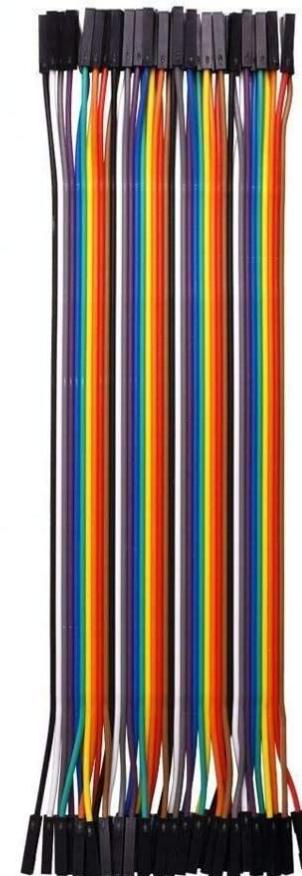


Which APIs need to return errors?

Almost all of them







So I moved **almost** all APIs to
use **std::expected...**

After **2.5 years**,
what did the students think?

IF/ELSE everywhere! Just (void) cast away the problems!

```
void blink_led(  
    output_pin& p_pin,  
    milliseconds p_transition_time) {  
    (void)p_pin.level(true);  
    delay(p_transition_time / 2);  
    (void)p_pin.level(false);  
    delay(p_transition_time / 2);  
}
```

Annoyed by the virality of error return types

```
class motion_controller {  
public:  
    motion_controller create(/* ... */);  
    status rotate(degrees_t p_degrees);  
    status move_to(coord_t p_x, coord_t p_y, coord_t p_z);  
    status configure_pid(std::uint16_t p_p,  
                          std::uint16_t p_i,  
                          std::uint16_t p_d);  
private:  
    status write_command(std::uint8_t p_register_address,  
                         std::uint32_t p_payload);  
    result<status_t> read_status(std::uint8_t  
                                  p_register_address);  
    status wait_for_eeprom_to_finish_loading();  
    /* ... */  
}
```

Everything returns status

MACROs were OK, but noisy

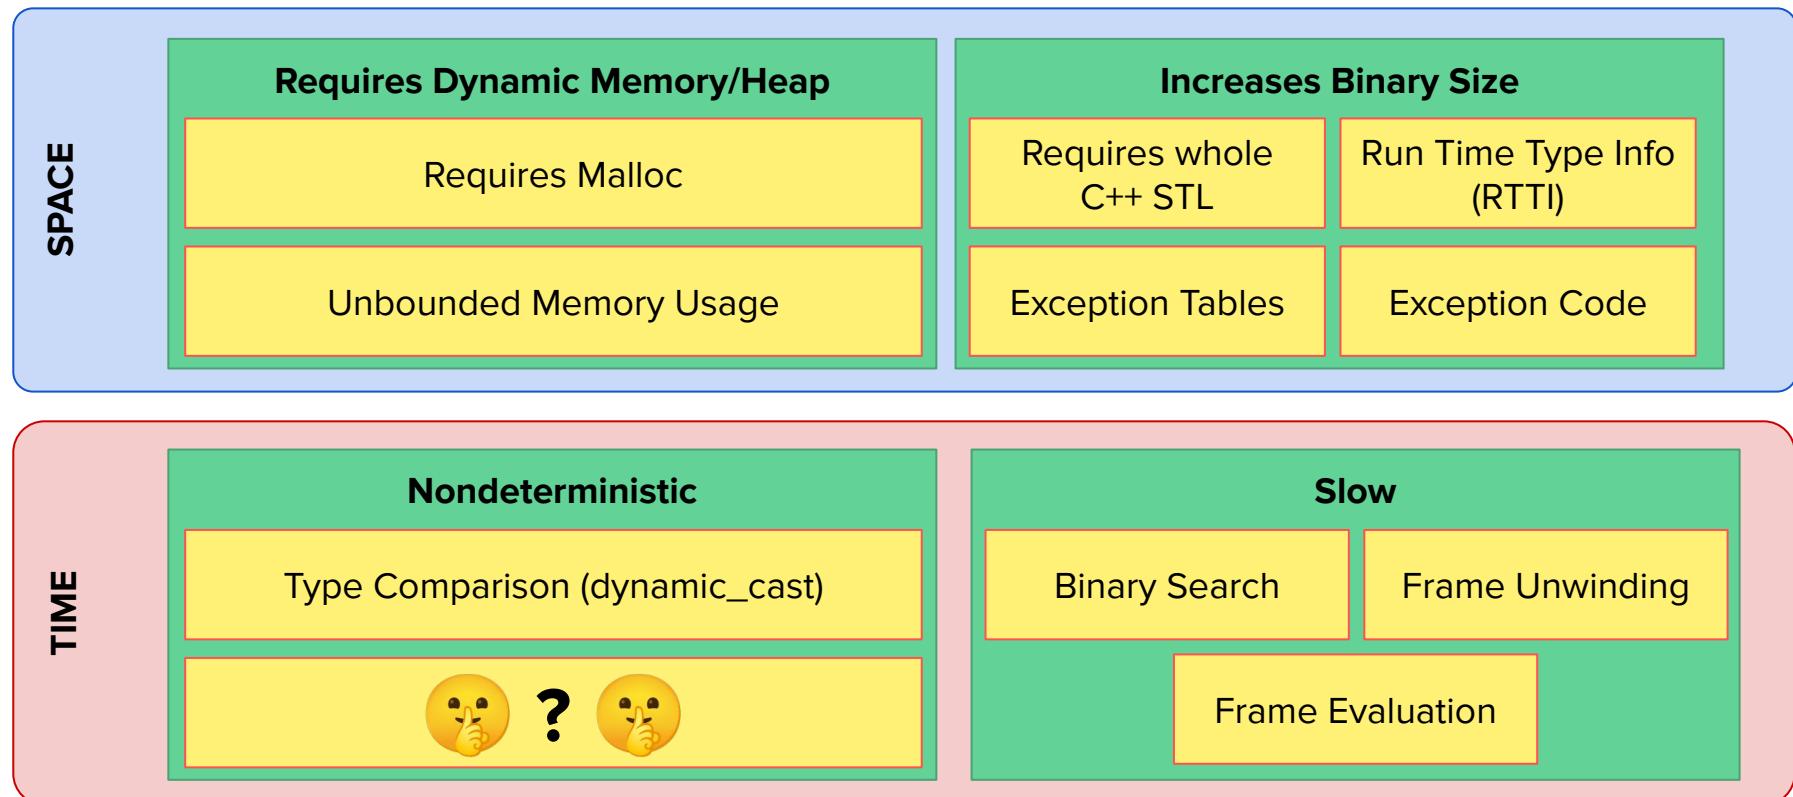
```
std::uint8_t const device_select =  
    SJ_CHECK(pin1.level()) << 1 | SJ_CHECK(pin1.level());  
auto const voltage = SJ_CHECK(v_sense[device_select].read());
```

MACROs were OK, but noisy

```
std::uint8_t const device_select =  
    SJ_CHECK(pin1.level()) << 1 | SJ_CHECK(pin1.level());  
auto const voltage = SJ_CHECK(v_sense[device_select].read());
```

So my students asked me,
why not use
C++ exceptions?

I verbally explained them this diagram...



Most accepted it...
Some asked why?

Why is it non-deterministic?

Why is it slow?

Why does it bloat the code?

Why does it need heap?

I would give an answer...
They'd accept...

Let's Make exceptions work on ARM!

Let's start simple

Arm GNU Toolchain GCC 11.3

```
arm-none-eabi-g++  
-o except.elf  
except.cpp  
-std=c++20  
-Os -g  
-fexceptions  
-frtti  
-mcpu=cortex-m4  
-mfloat-abi=hard -mthumb  
--specs=nosys.specs  
--specs=nano.specs  
--Wl,-T linker.ld
```

```
[[gnu::noinline]] int start()  
{  
    throw 5;  
}  
  
int main()  
{  
    volatile int return_code = 0;  
    try {  
        return_code = start();  
    } catch (...) {  
        return_code = -1;  
    }  
    return return_code;  
}
```

Barrier #1

Exceptions **disabled!**

☠ TERMINATED ☠
(any time the code throws)



Breaking Barrier #1

Enabling Exceptions in
ARM GCC

Step 1: Download & Build "ARM GNU Toolchain download"

Google search results for "arm gnu toolchain download". The top result is the official Arm Developer website.

Arm Developer
[https://developer.arm.com › downloads › arm-gnutoolchain... ›](https://developer.arm.com/downloads/arm-gnutoolchain)

Arm GNU Toolchain Downloads

Download the **Arm GNU Toolchain**, an open-source suite of tools for C, C++, and Assembly programming for the **Arm** architecture.

11.2-2022.02
 Arm GNU Toolchain releases are integrated and validated ...

12.2.Rel1
 Arm GNU Toolchain releases are integrated and validated ...

13.2.Rel1
 Arm GNU Toolchain releases package pre-built binaries of ...

Arm GNU Toolchain
 Free to download and use. Quality policy. Arm GNU Toolchain ...

12.3.Rel1
 Arm GNU Toolchain releases package pre-built binaries of ...

[More results from arm.com »](#)

≡ **arm** Developer

Source code

- [arm-gnutoolchain-src-snapshot-13.3.rel1.tar.xz](#)
- [arm-gnutoolchain-src-snapshot-13.3.rel1.tar.xz.asc](#)
- [arm-gnutoolchain-src-snapshot-13.3.rel1.tar.xz.sha256asc](#)
- [arm-gnutoolchain-src-snapshot-13.3.rel1-manifest.txt](#)

Linaro ABE example manifest files for Linux hosted cross toolchains

- [arm-gnutoolchain-arm-none-eabi-abe-manifest.txt](#)
- [arm-gnutoolchain-arm-none-eabi-nano-abe-manifest.txt](#)
- [arm-gnutoolchain-arm-none-linux-gnueabihf-abe-manifest.txt](#)
- [arm-gnutoolchain-aarch64-none-elf-abe-manifest.txt](#)
- [arm-gnutoolchain-aarch64-none-linux-gnu-abe-manifest.txt](#)
- [arm-gnutoolchain-aarch64_be-none-linux-gnu-abe-manifest.txt](#)
- [copy_nano_libraries.sh](#)

[Release Note](#) [EULA](#)

Step 2. Find and modify -fno-exceptions → -fexceptions

Step 3: Define low level C API: `_sbrk`

```
extern "C"

{
    std::array<std::byte, 1024> heap_memory {};
    std::span<std::byte> available_memory = heap_memory;

    void * _sbrk(int p_amount)
    {
        if (p_amount > available_memory.size())
        {
            return nullptr;
        }

        auto result      = available_memory.subspan(0, p_amount);
        available_memory = available_memory.subspan(p_amount);
        return result.data();
    }
}
```



IT WORKS!



but ...

Barrier #2

Large Binary Sizes

Size Comparison



Using Exceptions:

text	data	bss	dec	hex	filename
150008	2016	1328	89352	15d08	except.elf

Using a `std::expected<int, int>` type:

text	data	bss	dec	hex	filename
2680	1372	840	7516	1d5c	expected.elf

Baseline size type:

text	data	bss	dec	hex	filename
2656	1372	840	7516	1d5c	baseline.elf

+56x bigger
Uses ~29% of 512kB Flash!!



Breaking Barrier #2

Eliminating the bloat

Breaking Barrier #2: Searching the assembly

```
14496 0000a344 <d_growable_string_callback_adapter>:  
14497    a344: b5f0      push   {r4, r5, r6, r7, lr}  
14498    a346: 4614      mov    r4, r2  
14499    a348: 6852      ldr    r2, [r2, #4]  
14500    a34a: 68a5      ldr    r5, [r4, #8]  
14501    a34c: 1c4b      adds   r3, r1, #1  
14502    a34e: 4413      add    r3, r2  
14503    a350: 42ab      cmp    r3, r5  
14504    a352: b083      sub    sp, #12  
14505    a354: 460e      mov    r6, r1  
14506    a356: 4607      mov    r7, r0  
14507    a358: d811      bhi.n a37e <d_growable_string_callback_adapter+0x3a>  
14508    a35a: 68e5      ldr    r5, [r4, #12]  
14509    a35c: b96d      cbnz  r5, a37a <d_growable_string_callback_adapter  
+0x36>  
14510    a35e: 6863      ldr    r3, [r4, #4]  
14511    a360: 6820      ldr    r0, [r4, #0]  
14512    a362: 4632      mov    r2, r6  
14513    a364: 4418      add    r0, r3  
14514    a366: 4639      mov    r1, r7  
14515    a368: f008 f9a2 bl 126b0 <__aeabi_memcpy>  
14516    a36c: e9d4 3200 ldrd  r3, r2, [r4]  
14517    a370: 4433      add    r3, r6  
14518    a372: 549d      strb  r5, [r3, r2]  
14519    a374: 6863      ldr    r3, [r4, #4]  
14520    a376: 4433      add    r3, r6
```

Breaking Barrier #2: Searching the assembly

```
14778 0000a5c4 <d_append_num>:  
14779      a5c4: e92d 41f0    stmdb   sp!, {r4, r5, r6, r7, r8, lr}  
14780      a5c8: b088        sub sp, #32  
14781      a5ca: 460a        mov r2, r1  
14782      a5cc: 4604        mov r4, r0  
14783      a5ce: 491a        ldr r1, [pc, #104] ; (a638 <d_append_num+0x74>)  
14784      a5d0: a801        add r0, sp, #4  
14785      a5d2: f006 fd91    bl 110f8 <sprintf>  
14786      a5d6: a801        add r0, sp, #4  
14787      a5d8: f008 f880    bl 126dc <strlen>  
14788      a5dc: b340        cbz r0, a630 <d_append_num+0x6c>  
14789      a5de: ad01        add r5, sp, #4  
14790      a5e0: f8d4 1100    ldr.w  r1, [r4, #256] ; 0x100  
14791      a5e4: 182f        adds   r7, r5, r0  
14792      a5e6: f04f 0800    mov.w  r8, #0  
14793      a5ea: e009        b.n  a600 <d_append_num+0x3c>  
14794      a5ec: 460b        mov r3, r1  
14795      a5ee: 42bd        cmp r5, r7
```

Breaking Barrier #2: Find & eliminate the bloat...

```
13923 00009e14 <__gnu_cxx::__verbose_terminate_handler()>:  
13924 9e14: b570      push   {r4, r5, r6, lr}  
13925 9e16: 4b3b      ldr r3, [pc, #236] ; (9f04  
    <__gnu_cxx::__verbose_terminate_handler()>+0xf0>)  
13926 9e18: 781a      ldrb   r2, [r3, #0]  
13927 9e1a: b082      sub sp, #8  
13928 9e1c: 2a00      cmp r2, #0  
13929 9e1e: d141      bne.n 9ea4 <__gnu_cxx::__verbose_terminate_handler()  
    +0x90>  
13930 9e20: 2401      movs   r4, #1  
13931 9e22: 701c      strb   r4, [r3, #0]  
13932 9e24: f006 f8f0  bl 10008 <__cxa_current_exception_type>  
13933 9e28: 2800      cmp r0, #0  
13934 9e2a: d031      beq.n 9e90 <__gnu_cxx::__verbose_terminate_handler()  
    +0x7c>  
13935 9e2c: 6844      ldr r4, [r0, #4]  
13936 9e2e: 4d36      ldr r5, [pc, #216] ; (9f08  
    <__gnu_cxx::__verbose_terminate_handler()>+0xf4>)  
13937 9e30: 7823      ldrb   r3, [r4, #0]  
13938 9e32: 2b2a      cmp r3, #42 ; 0x2a  
13939 9e34: bf08      it eq  
13940 9e36: 3401      addeq r4, #1  
13941 9e38: 2200      movs   r2, #0  
13942 9e3a: f04f 30ff  mov.w r0, #4294967295 ; 0xffffffff  
13943 9e3e: 4611      mov r1, r2  
13944 9e40: ab01      add r3, sp, #4  
13945 9e42: 9001      str r0, [sp, #4]  
13946 9e44: 4620      mov r0, r4  
13947 9e46: f006 f849  bl fedc <__cxa_demangle>
```

Breaking Barrier #2: Within GCC

Output of x86-64 gcc (trunk) (Compiler #2) ✘ ×

A ▾ Wrap lines Select all

ASM generation compiler returned: 0

Execution build compiler returned: 0

Program returned: 139

terminate called after throwing an instance of 'my_error'

Program terminated with signal: SIGSEGV



So what can we do?

Breaking Barrier #2: with 5 lines of code

```
namespace __cxxabiv1 {  
    std::terminate_handler __terminate_handler = +[] () {  
        while (true) { continue; }  
    };  
}
```



Overrides GCC's weak symbol

Code Size Reduced!

text	data	bss	dec	hex filename
13632	120	660	14412	384c except.elf



-91% size reduction

Barrier #3

Uses **dynamic** allocation

Barrier #3: What is a throw expression in assembly?

```
0000800c <start()>:  
 800c: b508      push    {r3, lr}  
 800e: 2004      movs    r0, #4  
 8010: f000 f93c  bl     828c <__cxa_allocate_exception>  
 8014: 2305      movs    r3, #5  
 8016: 4902      ldr    r1, [pc, #8]  
 8018: 6003      str    r3, [r0, #0]  
 801a: 2200      movs    r2, #0  
 801c: f000 fd0a  bl     8a34 <__cxa_throw>  
 8020: 0000af48  .word   0x0000af48
```



Breaking Barrier #3

Overriding the allocator

Breaking Barrier #3: Override the allocator

```
extern "C"  
{  
    std::array<std::uint8_t, 256> storage{};  
    void* __cxa_allocate_exception(unsigned int p_size) noexcept  
    {  
        static constexpr size_t offset = 128;  
        return storage.data() + offset;  
    }  
    void __cxa_free_exception(void*) noexcept  
    {  
        // Do nothing here.  
    }  
} // extern "C"
```

Barrier #4

The Spectre of RTTI



Breaking Barrier #4

Disabling **RTTI**

Breaking Barrier #4: Disabling RTTI

Action: Replace `-f rtti` with `-fno-rtti`

What happens if you disable it?

Note: C++ Exceptions REQUIRE RTTI to work.

1.  Compiler tells you that you cannot use C++ exceptions with RTTI disabled
2.  Compiler strips all RTTI information not involved in exception handling?
3.  Compiler is silent and removes RTTI, using C++ exceptions becomes UB

Breaking Barrier #4: Disabling RTTI

Action: Replace `-f rtti` with `-fno-rtti`

What happens if you disable it?

1.  Compiler tells you that you cannot use C++ exceptions with RTTI disabled
2.  Compiler strips all RTTI information not involved in exception handling?
3.  Compiler is silent and removes RTTI, using C++ exceptions becomes UB

So why do firmware developers avoid C++ exceptions?

SPACE

Requires Dynamic Memory/Heap

Requires Malloc

Increases Binary Size

Requires whole
~~C++ STL~~

Run Time Type Info
(RTTI)

Unbounded Memory Usage

Exception Tables

Exception Code

TIME

Nondeterministic

Type Comparison (dynamic_cast)



Slow

Binary Search

Frame Unwinding

Frame Evaluation

But at this point,
I had seen enough!

Refactoring my project to use exceptions

Migrate SJSU-Dev2 completely to exception handling #1367

[Edit](#)[Code ▾](#)

kammce merged 1 commit into [master](#) from [exception-testing](#) on Sep 18, 2020

[Conversation 0](#)[Commits 1](#)[Checks 0](#)[Files changed 141](#)[+1,555 -2,339](#) 

kammce commented on Sep 14, 2020 • edited

[Member](#)[...](#)

- Remove Returns from codebase
- Remove SJ2_RETURN_ON_ERROR() macro and all of its usage
- Remove unit test sections that check for proper return code error propagation

Reviewers

No reviews



Assignees

No one—[assign yourself](#)



Result<T> PR closed

[WIP] Make all interfaces adopt Returns<> #1363

Closed kammce wants to merge 1 commit into `master` from `error-all-the-things`

Conversation 1 Commits 1 Checks 0 Files changed 20 +239 -132

kammce commented on Sep 8, 2020

It can not be determined that that any particular implementation of a hardware driver will or will not throw an error, thus every interface must be equip to return an error if need be.

[WIP] Make all interfaces adopt Returns<> ... 81ebc49

kammce commented on Oct 17, 2020

It was decided to transition to exception handling rather than return error types. For large enough projects exceptions are smaller in code size and faster when exceptions are not running.

Reviewers: No reviews

Assignees: No one—assign yourself

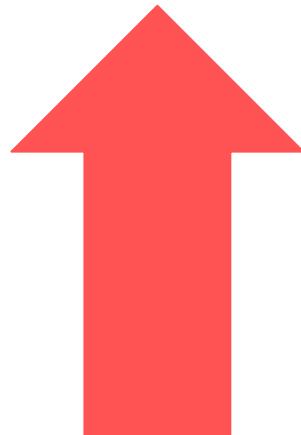
Labels: None yet

Projects: None yet

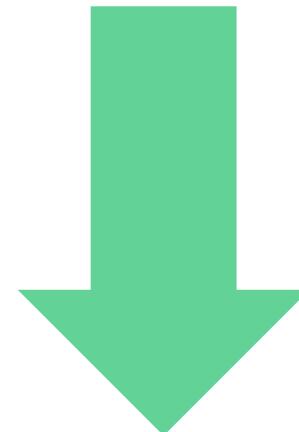
Milestone: No milestone

What I had observed in terms of Binary Size

Small Projects/Demos

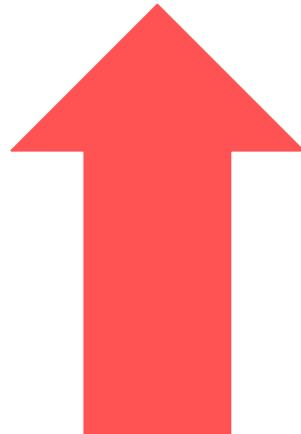


Large Projects

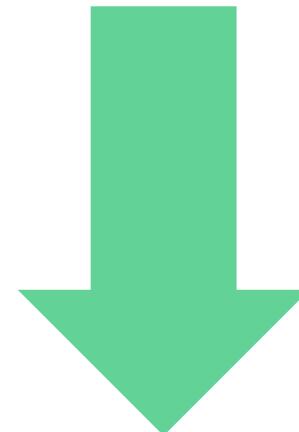


What I had observed in terms of Binary Size

Small Projects/Demos

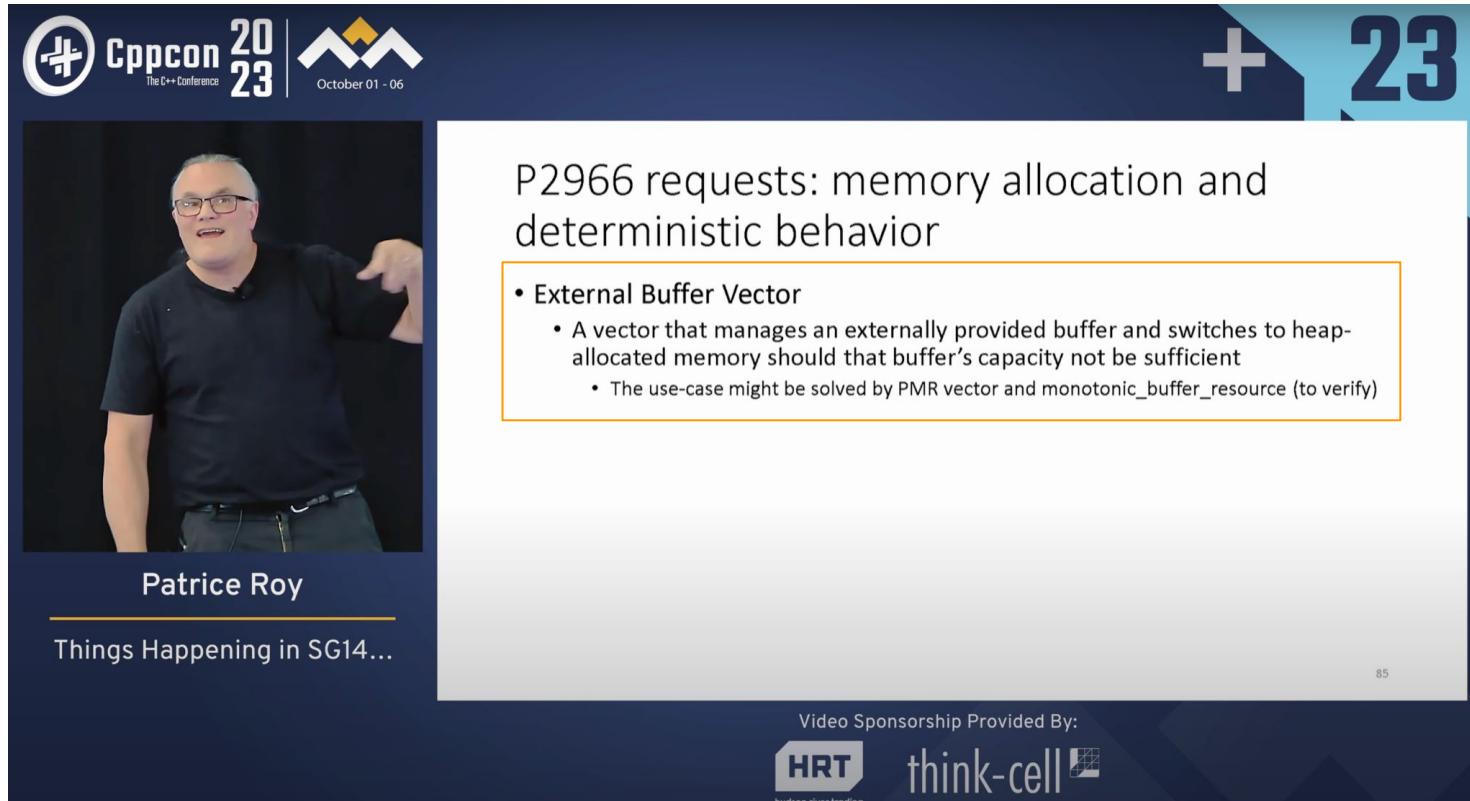


Large Projects



ISO C++ Meeting 2023 @ CppCon

Things Happening in SG14... - Patrice Roy - CppCon 2023



The screenshot shows a video player interface. At the top left is the CppCon 2023 logo with the text "The C++ Conference" and the date "October 01 - 06". To the right is a graphic with a large white plus sign and the number "23". The main video frame on the left shows a man with glasses and a black t-shirt, identified as Patrice Roy, gesturing while speaking. Below his name is the text "Things Happening in SG14...". The right side of the screen contains a white box with a yellow border containing text and a bulleted list.

P2966 requests: memory allocation and deterministic behavior

- External Buffer Vector
 - A vector that manages an externally provided buffer and switches to heap-allocated memory should that buffer's capacity not be sufficient
 - The use-case might be solved by PMR vector and monotonic_buffer_resource (to verify)

Video Sponsorship Provided By:

HRT hudson river trading

think-cell

Concerns with std::pmr::vector

`std::vector<T,Allocator>::push_back`

`void push_back(const T& value);` (1) `(constexpr since C++20)`

`void push_back(T&& value);` (2) `(since C++11)`
`(constexpr since C++20)`

How do you know it worked without exceptions?

"It's a shame embedded devs
don't use exceptions..."

It would make their binaries
smaller."

My Hypothesis

- Cost of using result types: $S(N) = aN$
 - a = average cost of a check per call
 - N = number of checked function calls

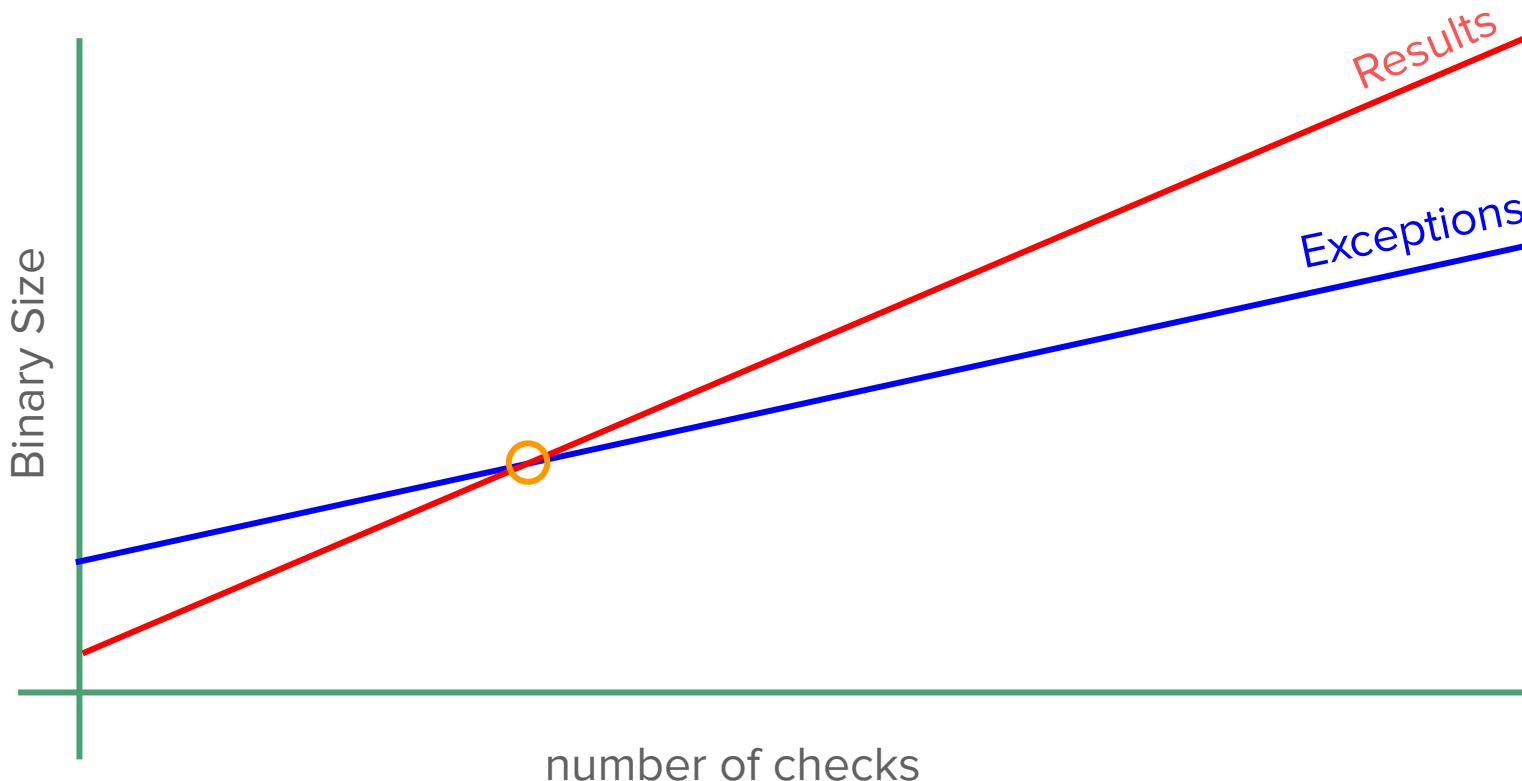
My Hypothesis

- Cost of using result types: $S(N) = aN$
 - a = average cost of a check per call
 - N = number of checked function calls
- Cost of using exception: $S(M) = bM + E$
 - b = average cost of exception metadata
 - M = number of non-leaf functions
 - E = cost of exception machinery (libunwind, etc)
 - **NOTE: this is grossly simplified**

My Hypothesis

- Cost of using result types: $S(N) = aN$
 - a = average cost of a check per call
 - N = number of checked function calls
- Cost of using exception: $S(M) = bM + E$
 - b = average cost of exception metadata
 - M = number of non-leaf functions
 - E = cost of exception machinery (libunwind, etc)
 - **NOTE: this is grossly simplified**
- I assert that: $aN > bM$
 - Let's assume we have an application of 100 error reporting functions
 - **N (# of functions called) $\geq M$ (total number of functions)**

My Hypothesis



The Experiment: Generating C++ code

Due to time constraints, I cannot go over this program in full.

- Defines classes with and without trivial destructors
- Create sequences of functions that call other functions
- Those functions would also create class objects and call their class functions

I wrote a program to compare **std::expected** to **exception handling**.

My requirements were:

1. Generate two C++ programs (**except.cpp** & **result.cpp**) that are isomorphic to each other.
2. Focus on propagation vs handling.

Sample Generated C++ code

```

int funct_group11_8()
{
    volatile static std::uint32_t inner_side_effect = 0;
    inner_side_effect = inner_side_effect + 1;
    class_0 instance_0(side_effect);
    instance_0.trigger();
    instance_0.trigger();
    instance_0.trigger();
    instance_0.trigger();

    class_1 instance_1(side_effect);
    instance_1.trigger();

    class_2 instance_2(side_effect);
    instance_2.trigger();
    instance_2.trigger();
    side_effect = side_effect - funct_group11_9();
    return side_effect;
}

```

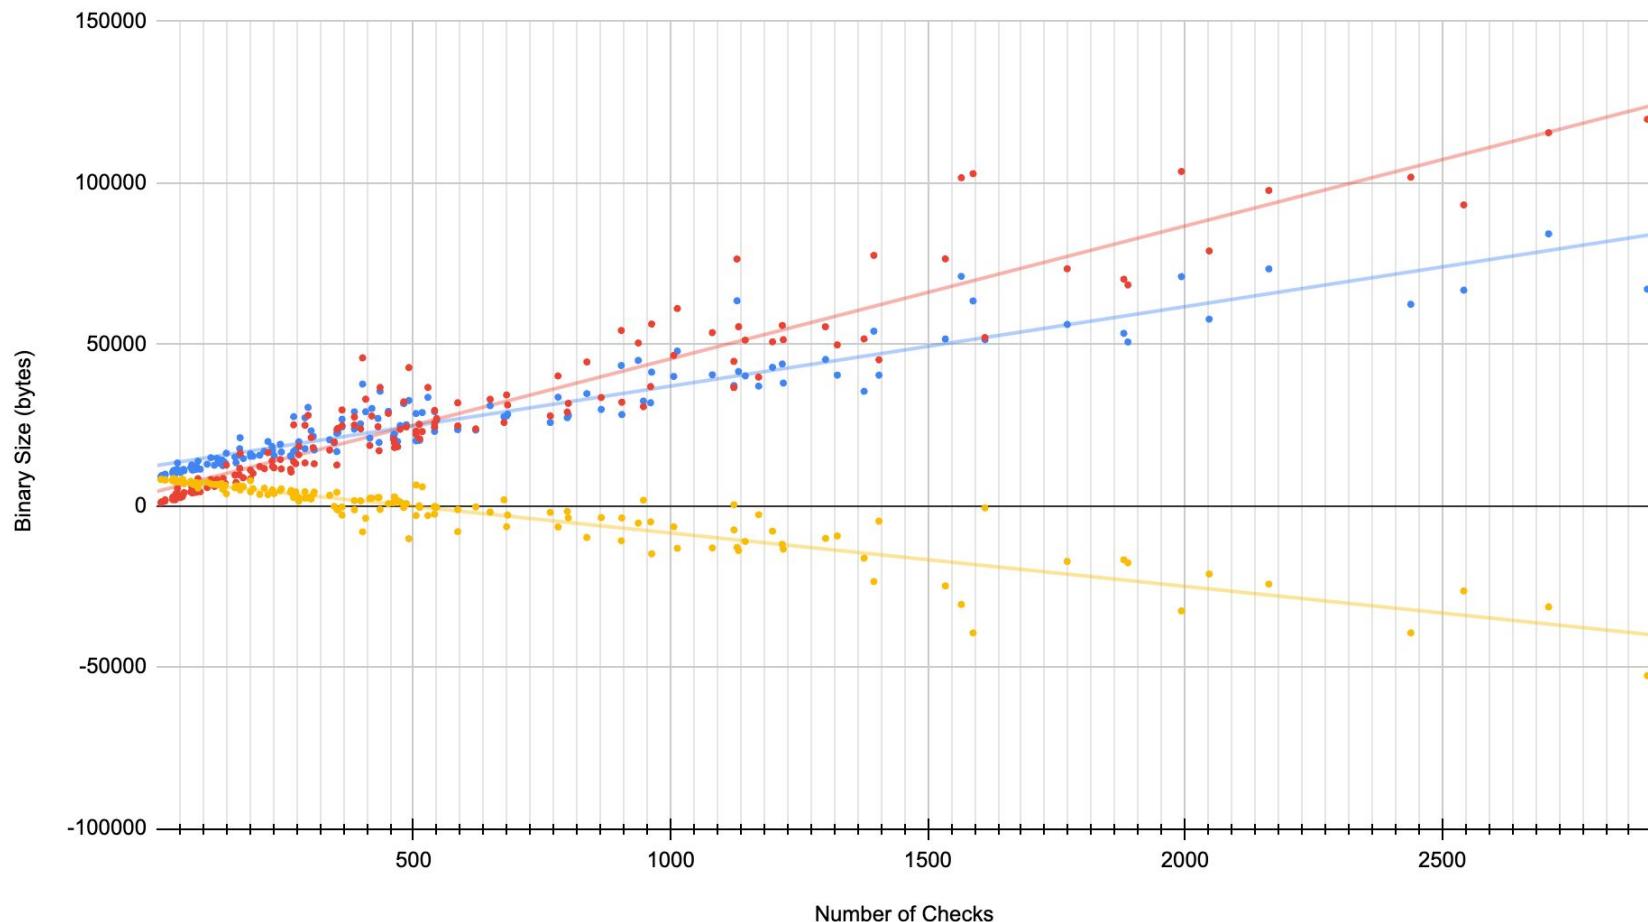
```

std::expected<int, my_error_t> funct_group11_8()
{
    volatile static std::uint32_t inner_side_effect = 0;
    inner_side_effect = inner_side_effect + 1;

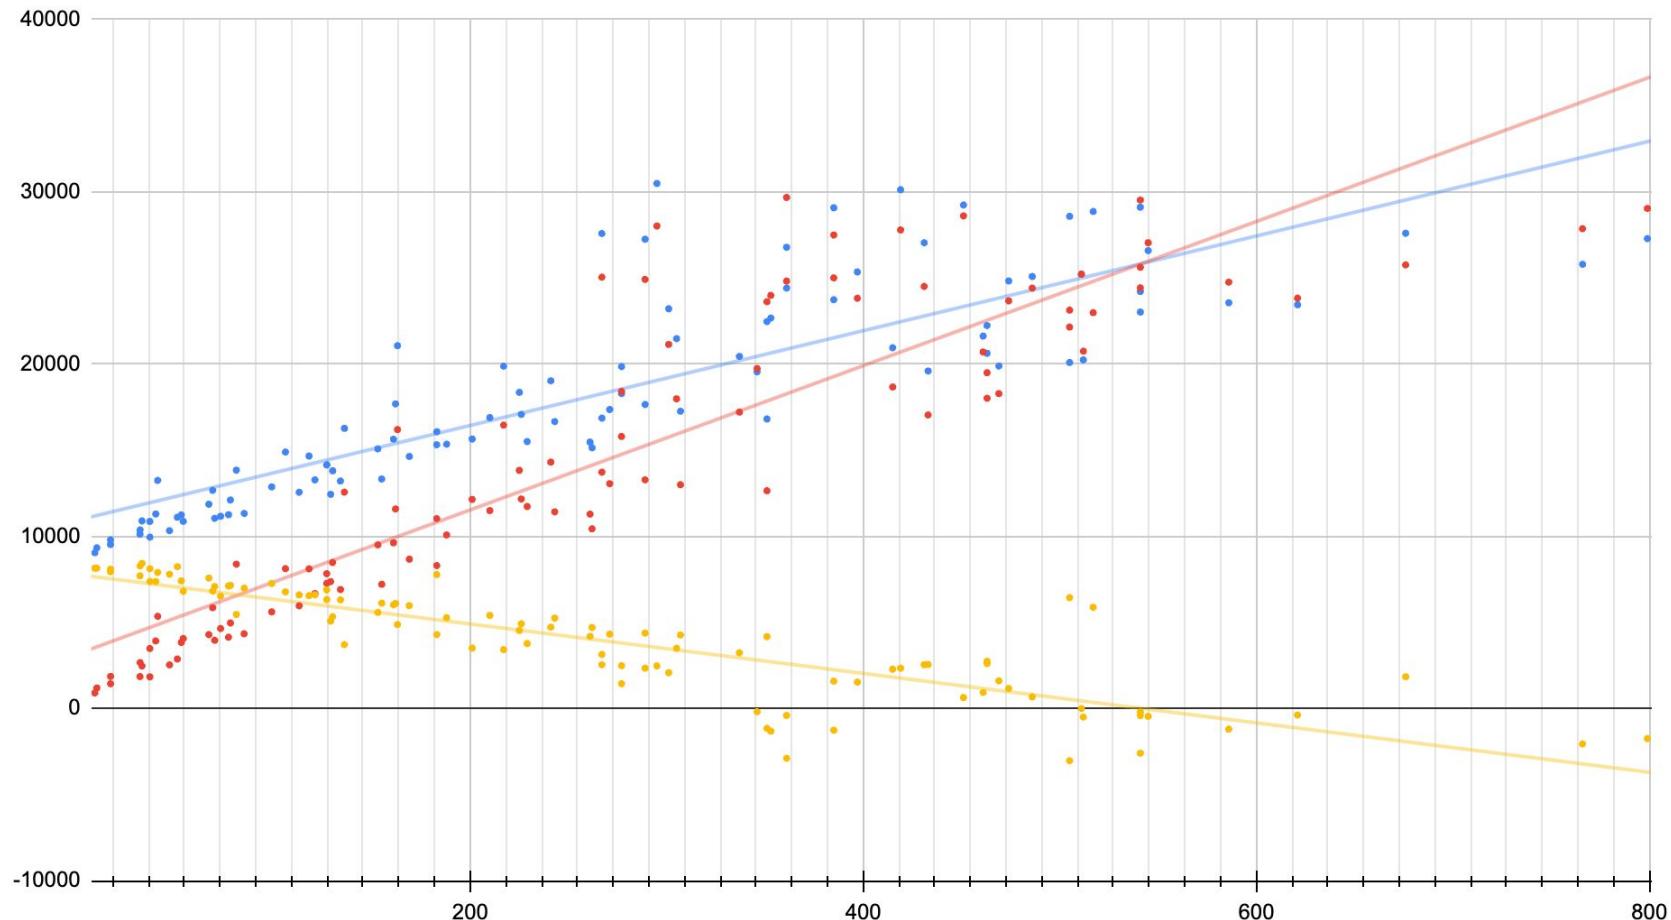
    auto instance_0 = class_0::make(side_effect);
    if (!instance_0) {
        return std::unexpected(instance_0.error());
    }
    {
        auto scoped_result = instance_0.value().trigger();
        if (!scoped_result) {
            return std::unexpected(scoped_result.error());
        }
    }
    {
        auto scoped_result = instance_0.value().trigger();
        if (!scoped_result) {
            return std::unexpected(scoped_result.error());
        }
    }
    {
        auto scoped_result = instance_0.value().trigger();
        if (!scoped_result) {
            return std::unexpected(scoped_result.error());
        }
    }
    // etc...
    if (auto result = funct_group11_9(); !result) {
        return std::unexpected(result.error());
    } else {
        side_effect = side_effect - result.value();
    }
    return side_effect;
}

```

● except.text ● 24.6*x + 12447 R² = 0.885 ● result.text ● 41.1*x + 4380 R² = 0.905 ● Size Delta ● -16.5*x + 8067 R² = 0.844



● except.text ● 27.5*x + 10929 R² = 0.731 ● result.text ● 41.9*x + 3154 R² = 0.807 ● Size Delta ● -14.4*x + 7775 R² = 0.733



But that wasn't to convince
you

That was to show you that
something is here.

C++ Exceptions from throw to catch on GCC ARM

Things that will NOT be covered here

- Nested exceptions
- Anything other than table based exceptions

Consider the following

```
struct error { };

void foo() {
    try {
        bar();
    } catch (error const& p_error) {
        // end up here...
    }
}

void bar() {
    destructible_t obj;
    baz();
} // ~destructible_t(&my_object)

void baz() { throw error{}; }
```

The diagram illustrates the execution flow in the code. A vertical orange arrow points upwards from the bottom of the `baz()` function back towards the `catch` block. From the top of the `catch` block, another vertical orange arrow points down to the `// ~destructible_t(&my_object)` comment. A curved orange arrow then branches off from the vertical line at the `// ~destructible_t(&my_object)` comment, pointing towards the `throw error{ };` statement in the `baz()` function.

The 3 types of functions

Leaf Functions

Transparent-Trivial Functions

Reentrant Functions

The 3 types of functions

Leaf Functions



- Calls other functions
- Contains destructors
- try/catch

Transparent-Trivial Functions

Reentrant Functions

The 3 types of functions

Leaf Functions



- Calls other functions
- Contains destructors
- try/catch

Transparent-Trivial Functions



- Calls other functions
- Contains destructors
- try/catch

Reentrant Functions



- Calls other functions
- Contains destructors
- try/catch

The 3 types of functions

Leaf Functions



- Calls other functions
- Contains destructors
- try/catch

Transparent-Trivial Functions



- Calls other functions
- Contains destructors
- try/catch

Reentrant Functions



- Calls other functions
- Contains destructors
- try/catch

Classifying the functions

Foo

Bar

Baz



```
struct error { };

void foo() {
    try {
        bar();
    } catch (error const& p_error) {
        // end up here...
    }
}

void bar() {
    destructible_t obj;
    baz();
} // ~destructible_t(&my_object)

void baz() { throw error{}; }
```

ARM Cortex-M Registers & Calling Convention

Registers	Info
R0	Scratch Registers: Used to pass parameters & return values Additional parameters will be passed on the stack if the inputs exceed what can be passed in the registers
R1	
R2	
R3	
R4	Preserved Registers: Must be restored before a subroutine returns.
R5	
R6	
R7	
R8	
R9	
R10	
R11	
R12	Scratch Register
R13 / SP	Stack Pointer
R14 / LR	Link Register (Updated with the return address of a subroutine call) Pushed on the stack if the function calls other functions
R15 / PC	Program Counter (current location in the program)

ARM Cortex-M Registers & Calling Convention

```
00003720 <foo ()>:  
    3720:   b510          push    {r4, lr} ; save R4 & LR  
    ;  
    ; Do stuff...  
    ; Call a function at some point...  
    ;  
    372c:   4620          mov r0, r4 ; Set return value  
    ; Restore R4 back to R4 & LR into PC return to the caller  
    372e:   bd10          pop {r4, pc}
```

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	XXXX	0x4000	0x??????????
R1	XXXX	0x3FFC	0x??????????
R2	XXXX	0x3FF8	0x??????????
R3	XXXX	0x3FF4	0x??????????
R4	?	0x3FF0	0x??????????
R5	?	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	?	0x3FCC	obj[3/3]
R14/LR	?	0x3FC8	bar() +14
R15/PC	?	0x3FC4	[bar] r4

```

struct error { };

void foo() {
    try {
        bar();
    } catch (error const& p_error) {
        // end up here...
    }
}

void bar() {
    destructible_t obj;
    baz();
} // ~destructible_t(&my_object)

void baz() { throw error{}; }

```

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	XXXX	0x4000	0x??????????
R1	XXXX	0x3FFC	0x??????????
R2	XXXX	0x3FF8	0x??????????
R3	XXXX	0x3FF4	0x??????????
R4	?	0x3FF0	0x??????????
R5	?	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	?	0x3FCC	obj[3/3]
R14/LR	?	0x3FC8	bar() +14
R15/PC	?	0x3FC4	[bar] r4

We are here

```

struct error { };

void foo() {
    try {
        bar();
    } catch (error const& p_error) {
        // end up here...
    }
}

void bar() {
    destructible_t obj;
    baz();
} // ~destructible_t(&my_object)

void baz() { throw error{}; }

```

Itanium C++ ABI: What are throws?

```
struct my_error {};  
  
void throws_my_error()  
{  
    throw my_error{};  
}
```



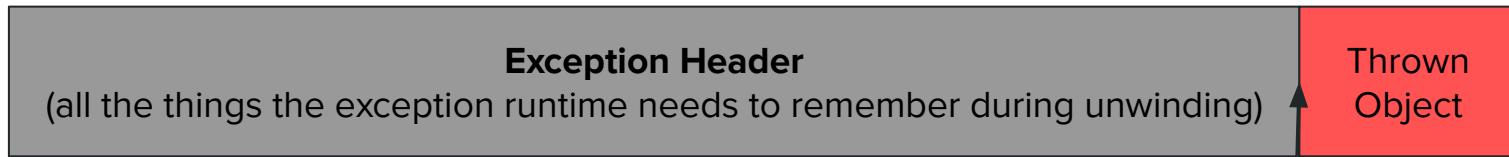
```
throws_my_error():  
    push    {r7, lr}  
    add     r7, sp, #0  
    movs   r0, #1  
    bl      __cxa_allocate_exception  
    movs   r3, r0  
    movs   r0, r3  
    ldr    r3, .L18  
    movs   r2, #0  
    movs   r1, r3  
    bl      __cxa_throw  
.L18:  
    .word   typeinfo_for_my_error
```

Itanium C++ ABI: Throw guidelines

```
void __cxa_throw(
    void* thrown_exception,
    std::type_info* type_info,
    void (*dtor)(void*));
```

Itanium C++ ABI: How __cxa_allocate_exception works?

```
std::array<std::uint8_t, 1024> exception_storage;  
void* __cxa_allocate_exception(size_t pThrownSize)  
{  
    if (exception_storage.size() < pThrownSize + sizeof(exception_header)) {  
        std::terminate();  
    }  
    return exception_storage.data() + sizeof(exception_header);  
}
```



Returned address

Itanium C++ ABI: What is the Exception Header?

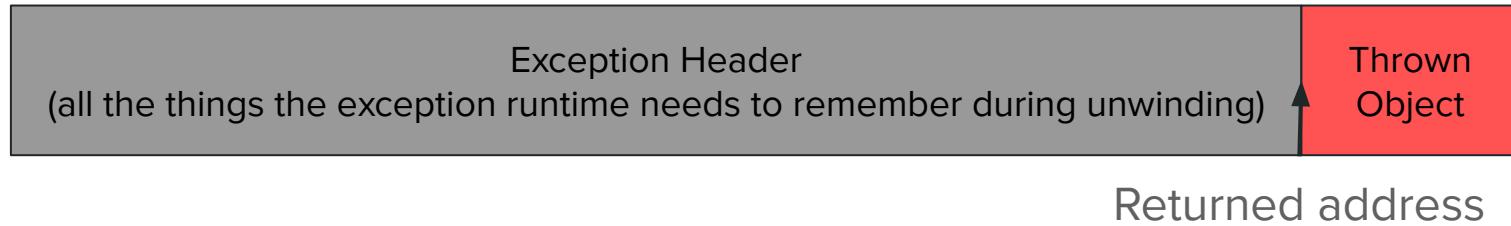
```
struct exception_header
{
    cortex_m_cpu cpu{};
    std::type_info* type_info;
    destructor_t* destructor;
    cache_t cache{};
};
```

```
enum class runtime_state : std::uint8_t
{
    get_next_frame = 0,
    enter_function = 1,
    unwind_frame = 2,
    handled = 3,
};

struct cache_t
{
    std::uint32_t state_and_rel_address;
    const index_entry_t* entry_ptr;
}
```

```
struct cortex_m_cpu
{
    register_t r0; // Remove?
    register_t r1; // Remove?
    register_t r2; // Remove?
    register_t r3; // Remove?
    register_t r4;
    register_t r5;
    register_t r6;
    register_t r7;
    register_t r8;
    register_t r9;
    register_t r10;
    register_t r11;
    register_t ip;
    register_t sp;
    register_t lr;
    register_t pc;
};
```

Why not call `malloc`?



Because of **ABI Compatibility & Flexibility**.

- Allows the size of the exception header to grow or shrink without an ABI break
- Reduces the allocations to 1

Itanium C++ ABI: How `__cxa_throw` works

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	XXXX	0x4000	0x?????????
R1	XXXX	0x3FFC	0x?????????
R2	XXXX	0x3FF8	0x?????????
R3	XXXX	0x3FF4	0x?????????
R4	?	0x3FF0	0x?????????
R5	?	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	?	0x3FCC	obj[3/3]
R14/LR	?	0x3FC8	bar() +14
R15/PC	?	0x3FC4	[bar] r4

```

void __cxa_throw(exception_ptr p_thrown_exception,
                  std::type_info* p_type_info,
                  destructor_t p_destructor)
{
    // TLS variable for std::current_exception
    active_exception = p_thrown_exception;

    auto& exception_header = extract_exception_header(
        p_thrown_exception);

    exception_header.type_info = p_type_info;
    exception_header.destructor = p_destructor;

    capture_cpu_core_and_unwind1(exception_header.cpu);
    raise_exception(exception_header);
    std::terminate();
}

```

Itanium C++ ABI: How `__cxa_throw` works

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'C000	0x4000	0x?????????
R1	type_info	0x3FFC	0x?????????
R2	0x0	0x3FF8	0x?????????
R3	0x?????????	0x3FF4	0x?????????
R4	0x?????????	0x3FF0	0x?????????
R5	0x?????????	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() + 10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FC4	0x3FCC	obj[3/3]
R14/LR	baz() + 4	0x3FC8	bar() + 14
R15/PC	__cxa_throw	0x3FC4	[bar] r4

```

void __cxa_throw(exception_ptr p_thrown_exception,
                  std::type_info* p_type_info,
                  destructor_t p_destructor)
{
    // TLS variable for std::current_exception
    active_exception = p_thrown_exception;
    auto& exception_header = extract_exception_header (
        p_thrown_exception);
    exception_header.type_info = p_type_info;
    exception_header.destructor = p_destructor;

    capture_cpu_core_and_unwind1(exception_header.cpu);
    raise_exception(exception_header);
    std::terminate();
}

```

Itanium C++ ABI: How `__cxa_throw` works

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'C000	0x4000	0x?????????
R1	type_info	0x3FFC	0x?????????
R2	0x0	0x3FF8	0x?????????
R3	0x?????????	0x3FF4	0x?????????
R4	0x?????????	0x3FF0	0x?????????
R5	0x?????????	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() + 10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FC4	0x3FCC	obj[3/3]
R14/LR	baz() + 4	0x3FC8	bar() + 14
R15/PC	baz() + 4	0x3FC4	[bar] r4

```

void __cxa_throw(exception_ptr p_thrown_exception,
                  std::type_info* p_type_info,
                  destructor_t p_destructor)
{
    // TLS variable for std::current_exception
    active_exception = p_thrown_exception;

    auto& exception_header = extract_exception_header(
        p_thrown_exception);

    exception_header.type_info = p_type_info;
    exception_header.destructor = p_destructor;

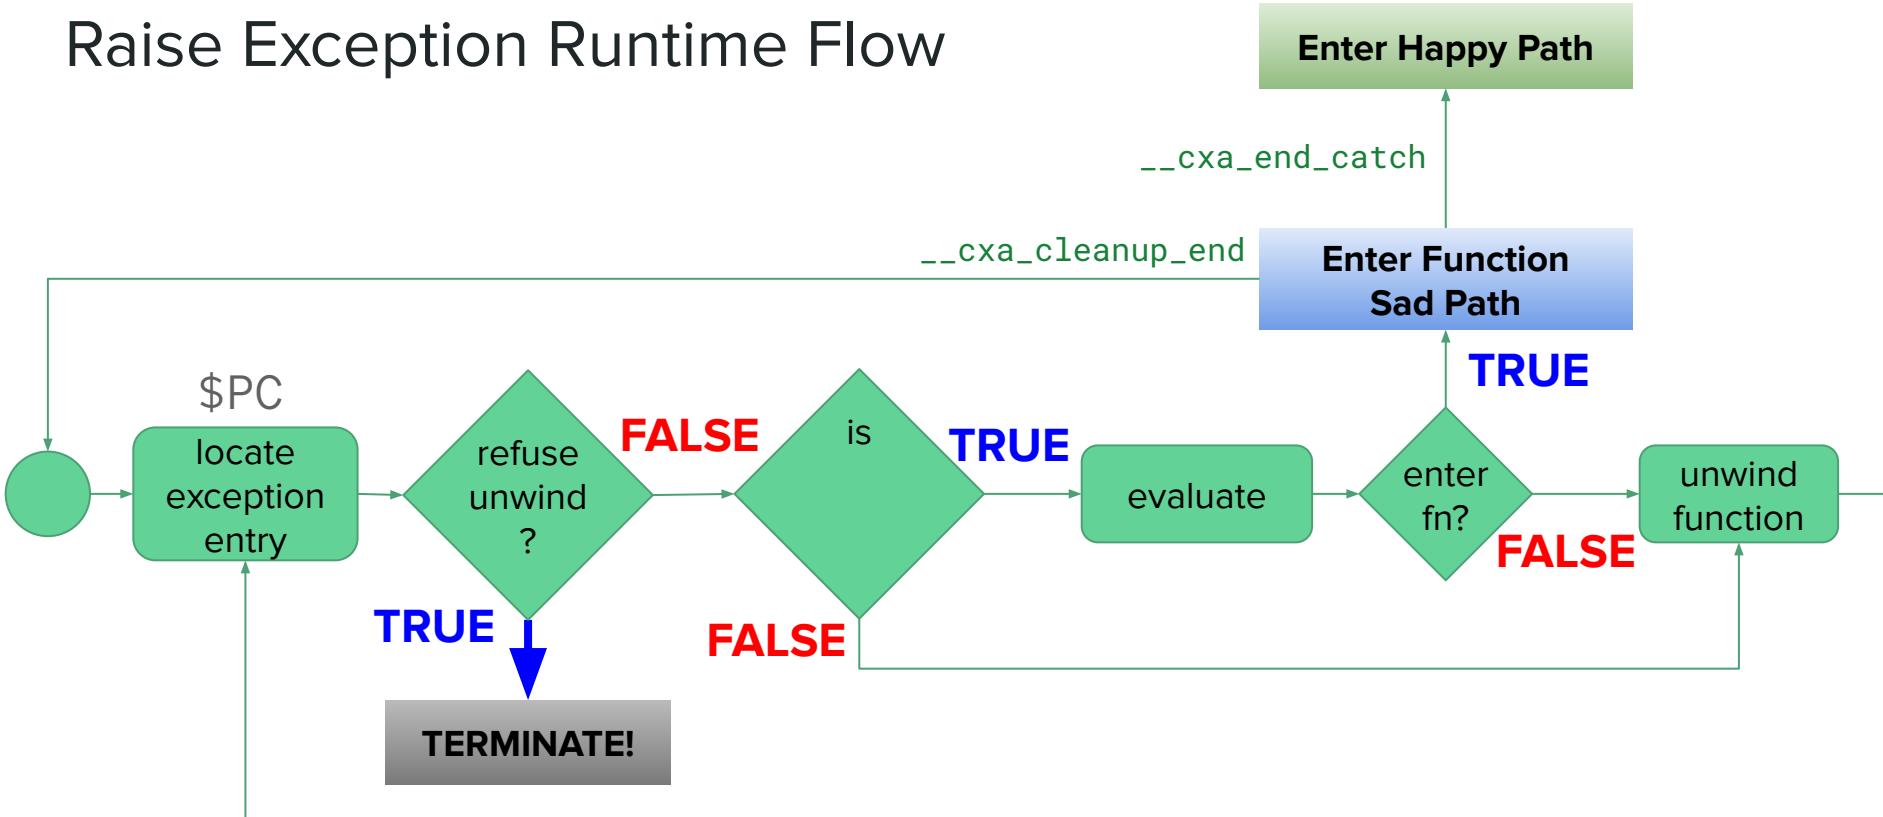
    capture_cpu_core_and_unwind1(exception_header.cpu);

    raise_exception(exception_header);

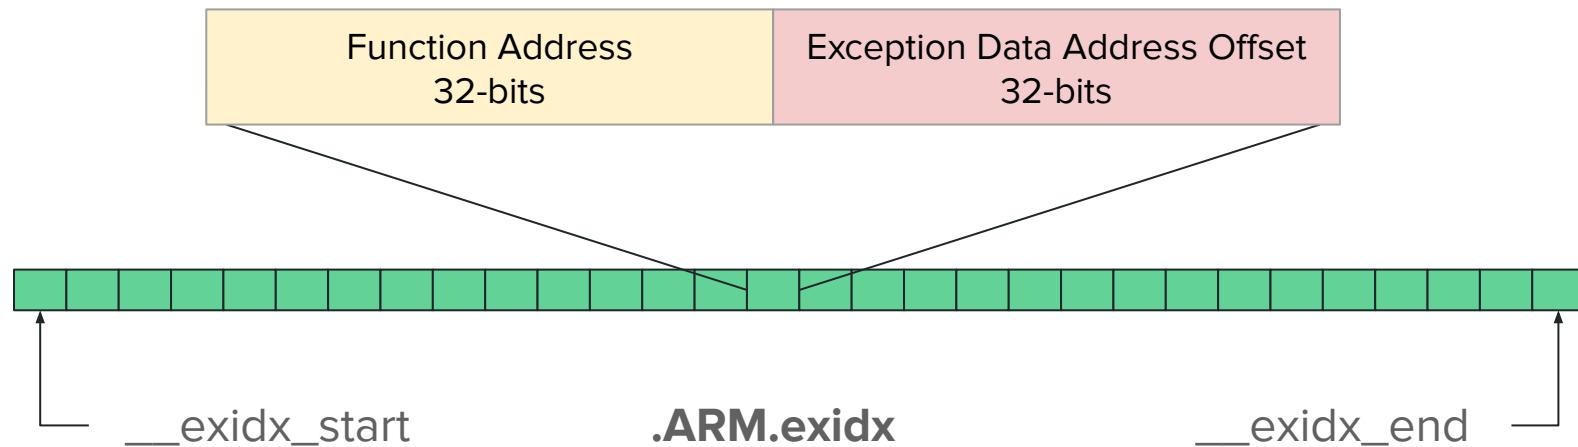
    std::terminate();
}

```

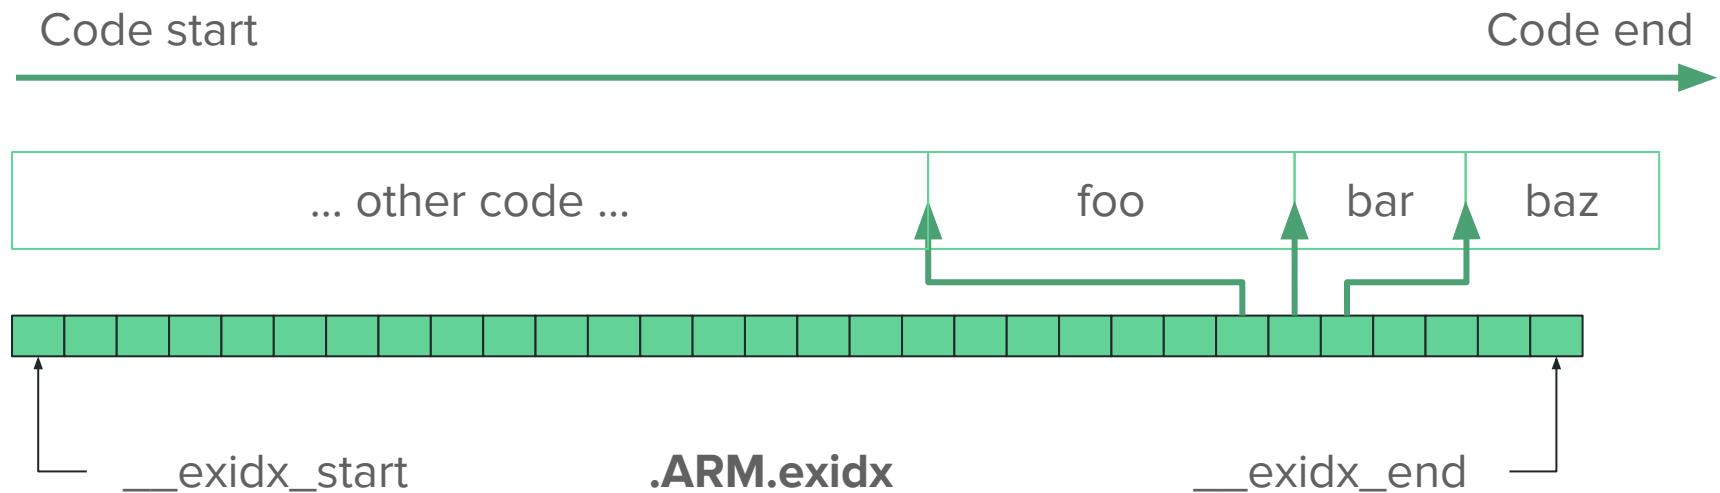
Raise Exception Runtime Flow



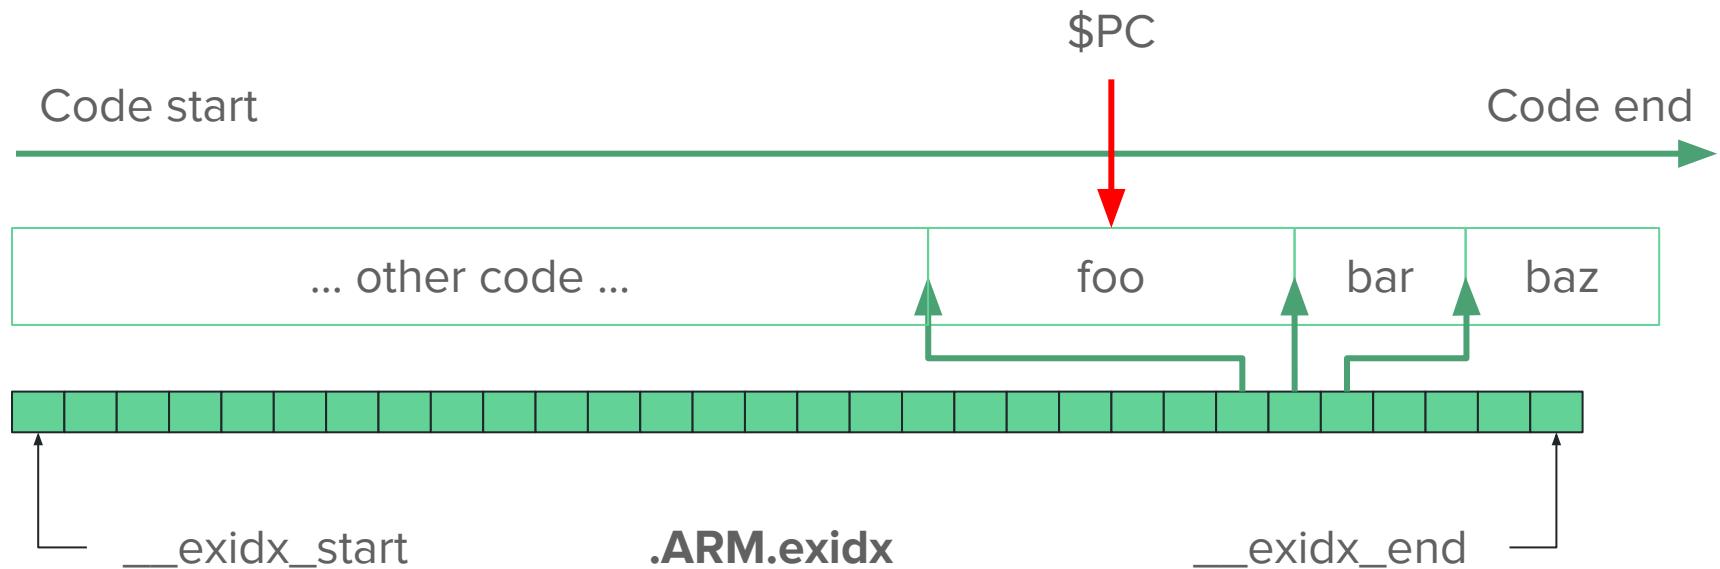
ARM Exception Index: Index Entry



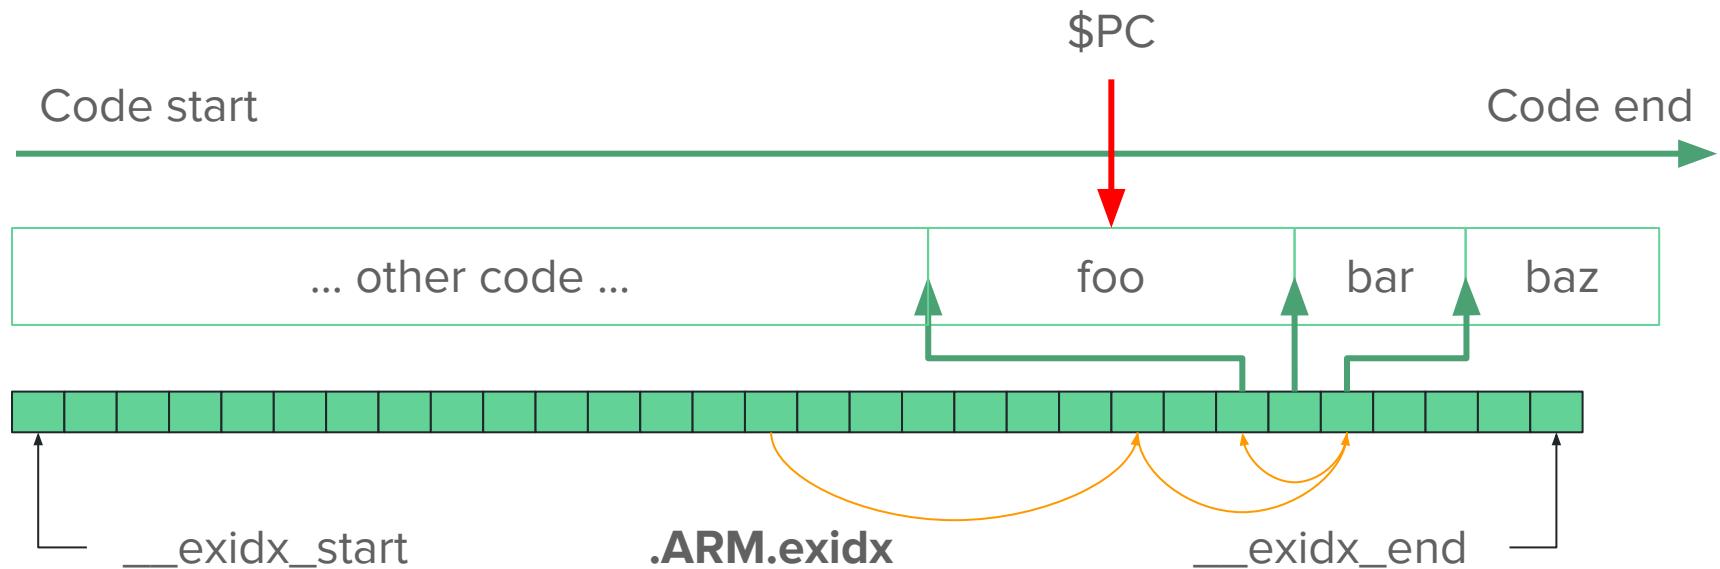
ARM Exception Index: Ordering



ARM Exception Index: Ordering



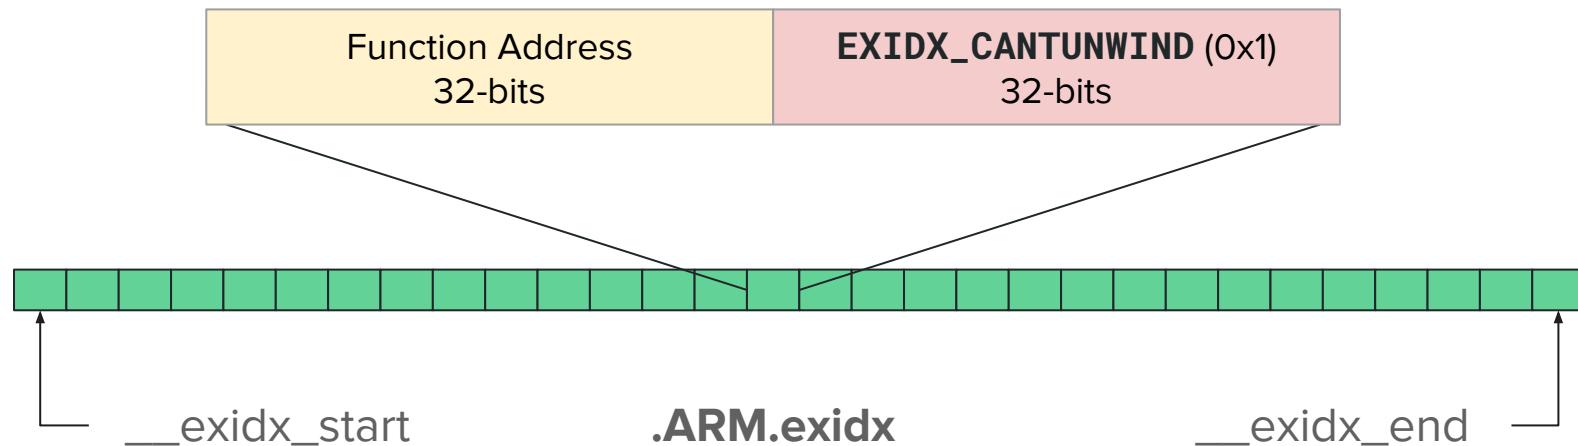
ARM Exception Index: Ordering



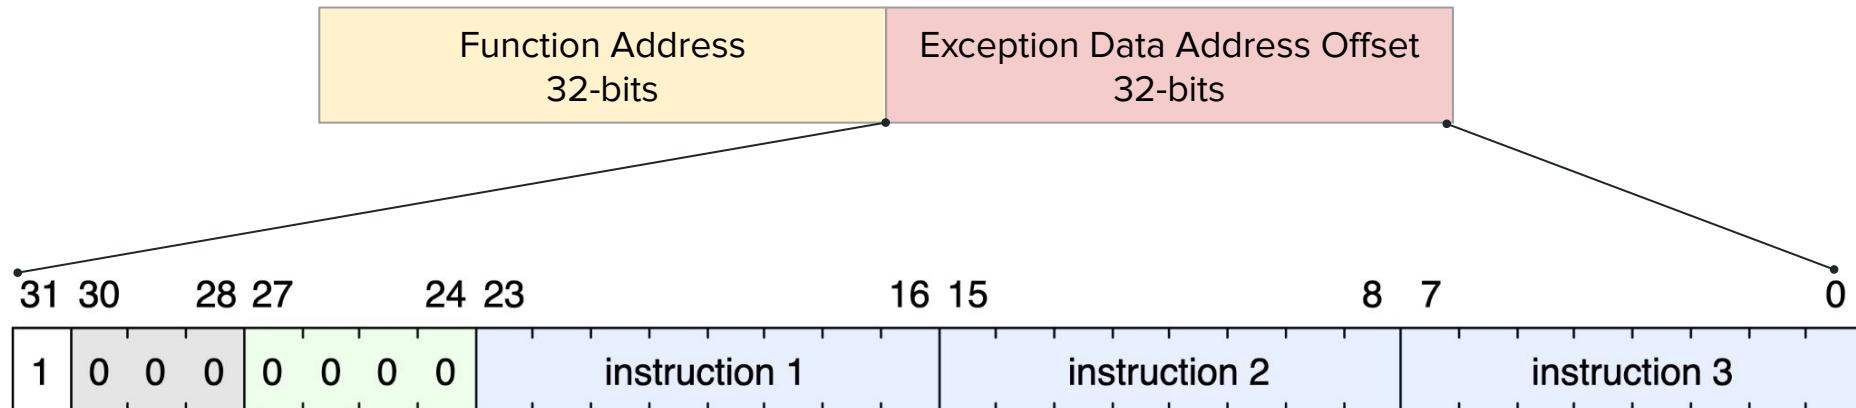
ARM Exception Index: Index Entry

```
struct arm_exception_index_t
{
    std::uint32_t prel31_address;
    union content
    {
        std::uint32_t prel31_metadata;
        std::uint32_t inlined_metadata;
    };
};
```

ARM Exception Index: 😐 "noexcept" Index Entry



ARM Exception Index: Inlined Personality Data



- For transparent 😊 functions needing only 3 unwind instructions.
- Saves memory

ARM unwind instructions

Instruction	Explanation
00xxxxxx	$vsp = vsp + (xxxxxx \ll 2) + 4$. Covers range 0x04-0x100 inclusive
01xxxxxx	$vsp = vsp - (xxxxxx \ll 2) - 4$. Covers range 0x04-0x100 inclusive
10100nnn	Pop r4-r[4+nnn]
10101nnn	Pop r4-r[4+nnn], r14 (lr)
10110010 uleb128	$vsp = vsp + 0x204 + (\text{uleb128} \ll 2)$ (for vsp increments of 0x104-0x200, use 00xxxxxx twice)
1001nnnn (nnnn != 13,15)	Set $vsp = r[nnnn]$
10110000	Finish
10000000 00000000	refuse to unwind / call std::terminate (for example, out of a cleanup)

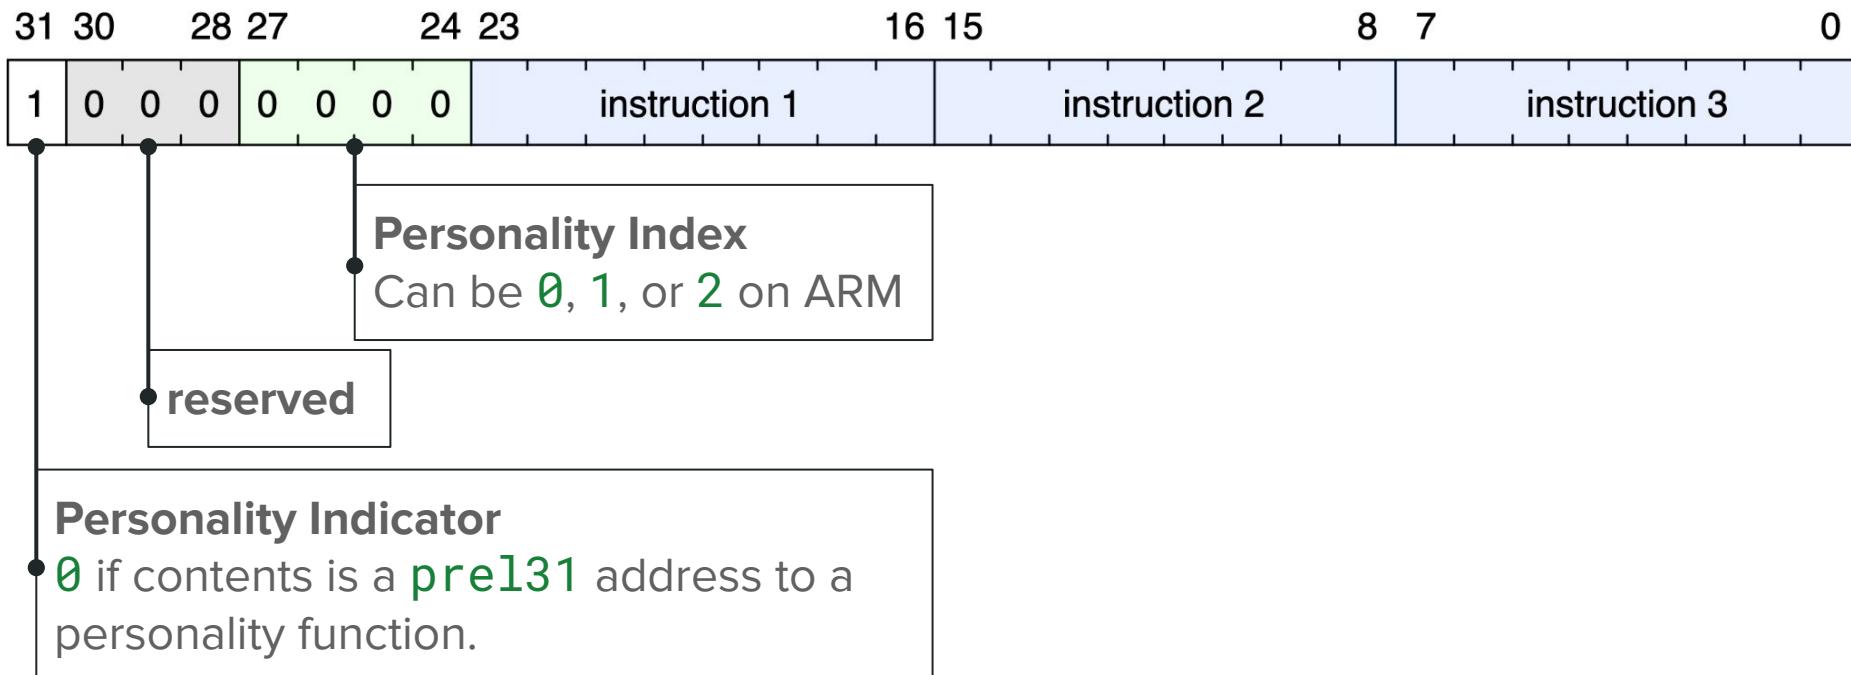
Note: Instructions are in 8-bit segments

3 bytes typically

Limit is 7 bytes

ARM Personality Data: Unwind Instructions

SU16 = Short Personality with 16-bit scopes (always 4-bytes)



ARM Personality Data: Unwind Instructions

LU16 = Long Personality with 16-bit scopes (always 8-bytes to 12-bytes)

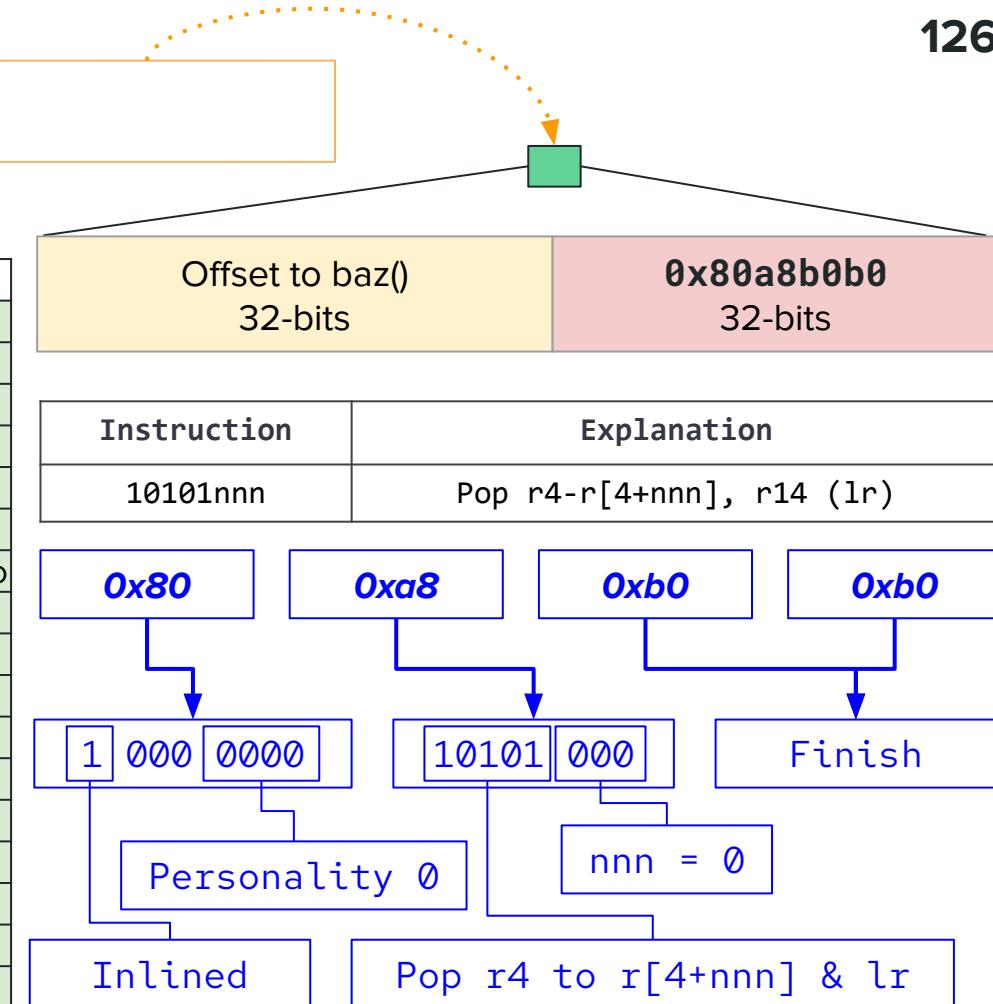
31	30	28	27	24	23	16	15	8	7	0		
1	0	0	0	0	0	0	1	length (1 or 2)		instruction 0		
instruction 2				instruction 3			instruction 4			instruction 5		
instruction 6				0xB0			0xB0			0xB0		

LU32 = Long Personality with 32-bit scopes (always 8-bytes to 12-bytes)

31	30	28	27	24	23	16	15	8	7	0		
1	0	0	0	0	0	1	0	length (1 or 2)		instruction 0		
instruction 2				instruction 3			instruction 4			instruction 5		
instruction 6				0xB0			0xB0			0xB0		

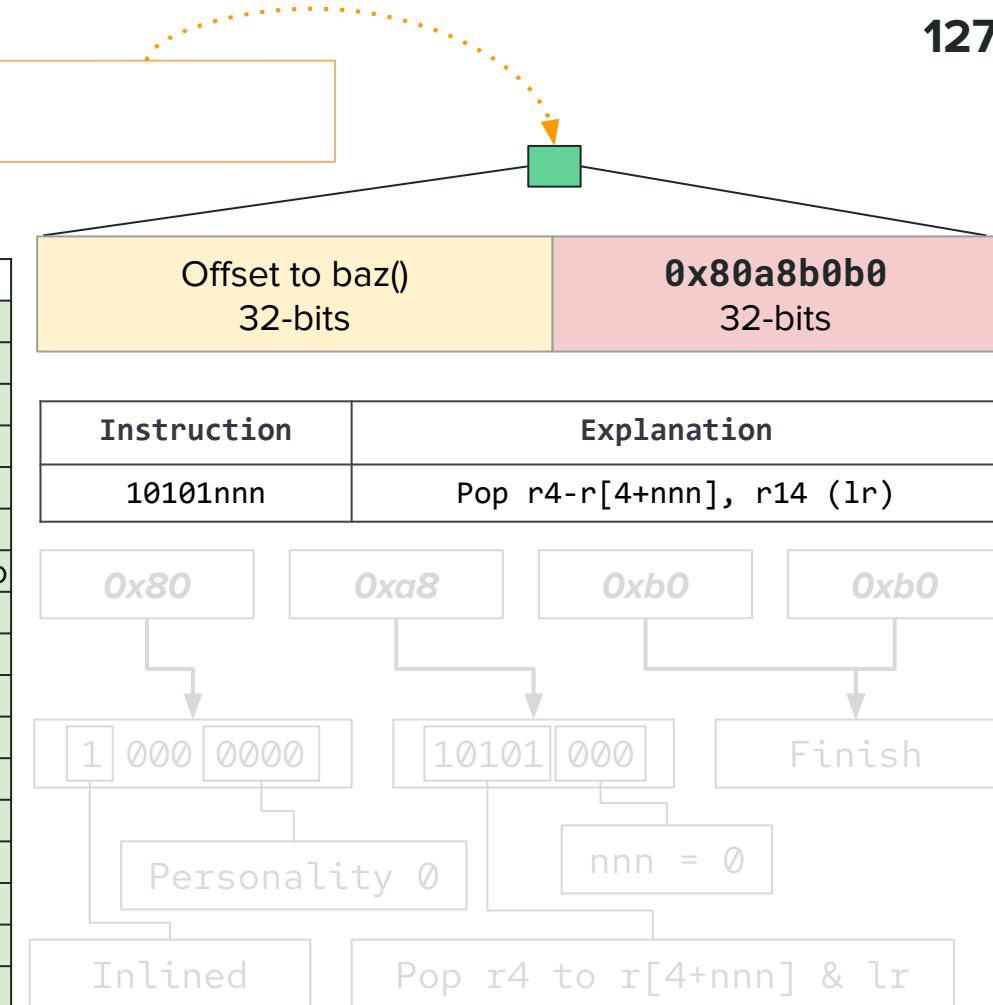
Unwinding baz()

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'C000	0x4000	0x?????????
R1	type_info	0x3FFC	0x?????????
R2	0x0	0x3FF8	0x?????????
R3	0x?????????	0x3FF4	0x?????????
R4	0x?????????	0x3FF0	0x?????????
R5	0x?????????	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FC4	0x3FCC	obj[3/3]
R14/LR	baz() +4	0x3FC8	bar() +14
R15/PC	baz() +4	0x3FC4	[bar] r4



Unwinding baz()

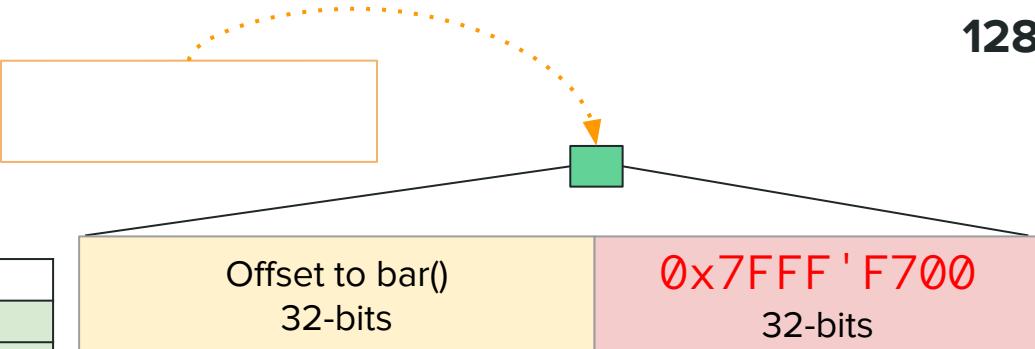
Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'C000	0x4000	0x?????????
R1	type_info	0x3FFC	0x?????????
R2	0x0	0x3FF8	0x?????????
R3	0x?????????	0x3FF4	0x?????????
R4	[bar] r4	0x3FF0	0x?????????
R5	0x?????????	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FCC	0x3FCC	obj[3/3]
R14/LR	bar() +14	0x3FC8	bar() +14
R15/PC	bar() +14	0x3FC4	[bar] r4



bar entry

Unwinding bar()

Virtual	ARM	CPU	SRAM	
Reg	Value		Address	Value
R0	0x2000'C000		0x4000	0x?????????
R1	type_info		0x3FFC	0x?????????
R2	0x0		0x3FF8	0x?????????
R3	0x?????????		0x3FF4	0x?????????
R4	[bar] r4		0x3FF0	0x?????????
R5	0x?????????		0x3FEC	caller of foo
R6	-		0x3FE8	caller r4
R7	-		0x3FE4	caller r5
R8	-		0x3FE0	foo() + 10
R9	-		0x3FDC	[foo] r5
R10	-		0x3FD8	[foo] r4
R11	-		0x3FD4	obj[1/3]
R12	-		0x3FD0	obj[2/3]
R13/SP	0x3FCC		0x3FCC	obj[3/3]
R14/LR	bar() + 14		0x3FC8	bar() + 14
R15/PC	bar() + 14		0x3FC4	[bar] r4



Note: Data isn't inlined because Bar is



Calling bar()'s Destructors...

We are here

```
struct error { };

void foo() {
    try {
        bar();
    } catch (error const& p_error) {
        // end up here...
    }
}

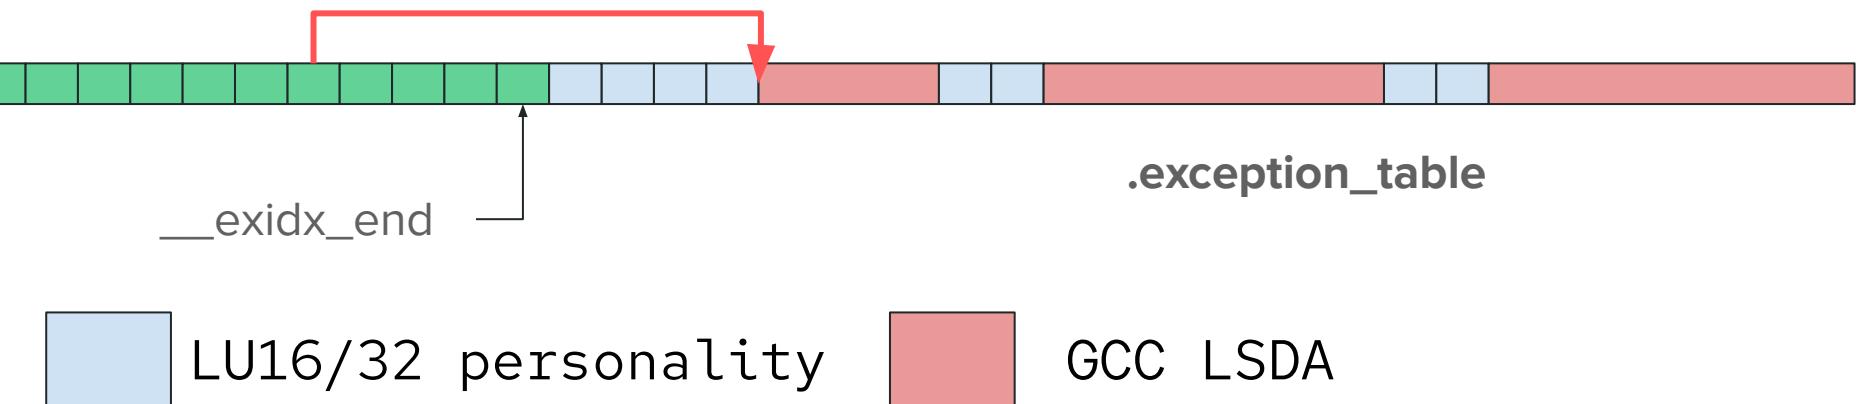
void bar() {
    destructible_t obj;
    baz();
} // ~destructible_t(&my_object)

void baz() { throw error{}; }
```

ARM Exception Table

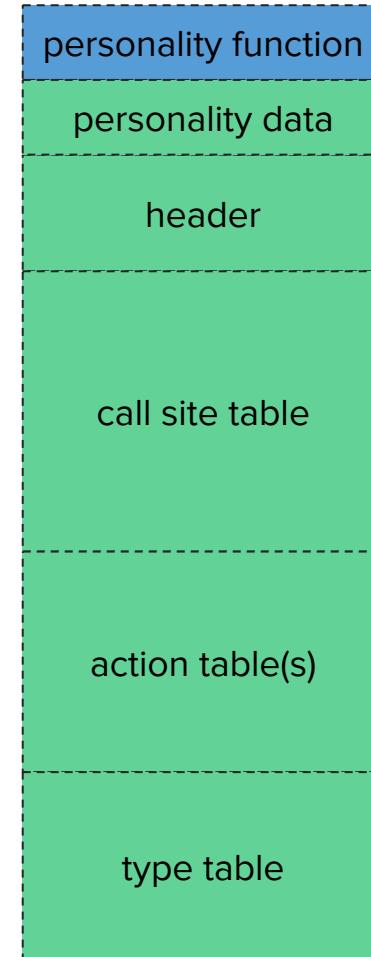
Note: Exception table is 32-bit aligned

Function Address 32-bits	Exception Data Address Offset 0x7FFF 'F700	= -2304
-----------------------------	--	---------



What is a LSDA?

- L.S.D.A. = **L**anguage **S**pecific **D**ata **A**rea
- LSDA is included in the Itanium exception ABIs in order to support other languages such as **Java**, **C#**, etc.
- Provides a means to allows exceptions to propagate from language to language.



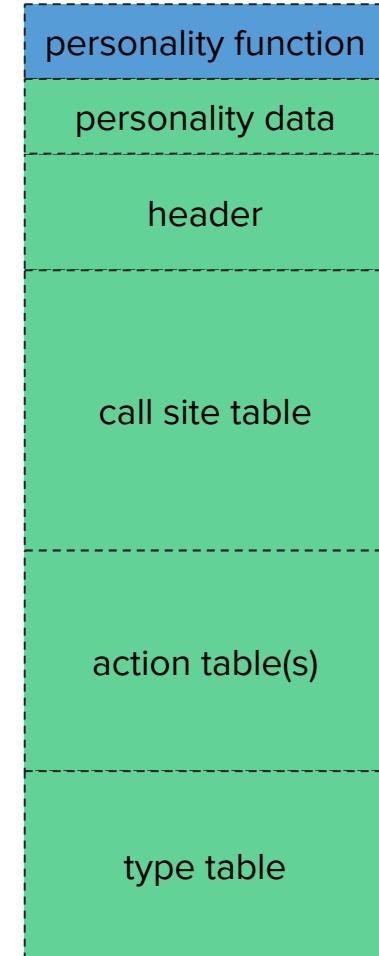
What is a LSDA?

- L.S.D.A. = **L**anguage **S**pecific **D**ata **A**rea
- LSDA is included in the Itanium exception ABIs in order to support other languages such as **Java**, **C#**, etc.
- Provides a means to allows exceptions to propagate from language to language.
- GCC leverages this to make platform agnostic exception data



GCC LSDA: Sections

- **Personality Function:** Function that can handle this data
- **Personality Data:** Unwind information
- **Header:** Describes the encoding of the tables below
- **Call Site Table:** Contains important frame scope areas
- **Action Table:** Provides catch order
- **Type table:** Holds the caught types in the function



GCC LSDA: Personality Function

31-bit offset to **gcc_personality_v0**
0x7FFF'F320



GCC LSDA: Personality Data

SU16 = Short Personality

31	30	28	27	24	23	16	15	8	7	0
1	0	0	0	0	0	instruction 1	instruction 2	instruction 3		

OR

LU16/32 = Long Personality

31	30	28	27	24	23	16	15	8	7	0
1	0	0	0	0	0	length (1 or 2)	instruction 0	instruction 1		
						instruction 2	instruction 3	instruction 4	instruction 5	
						instruction 6	0xB0	0xB0	0xB0	

personality function

personality data

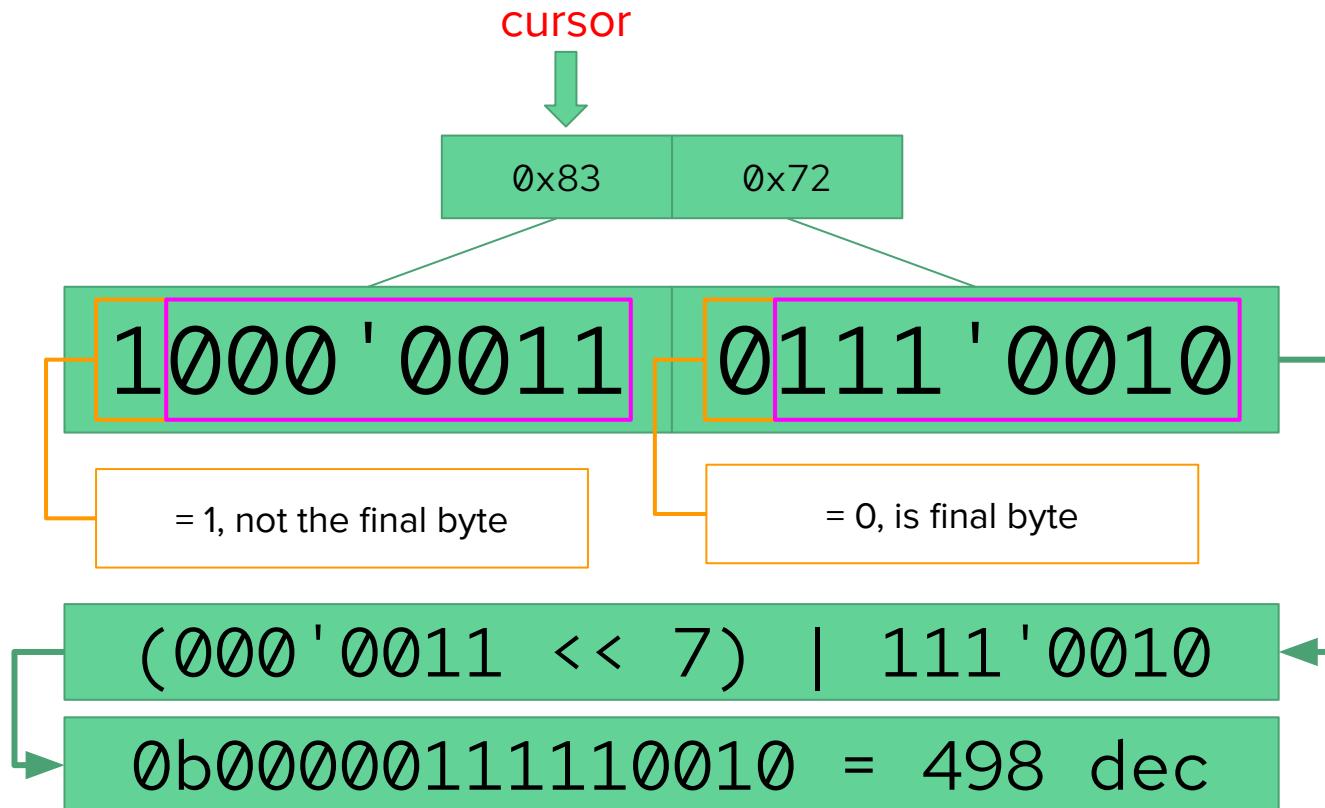
header

call site table

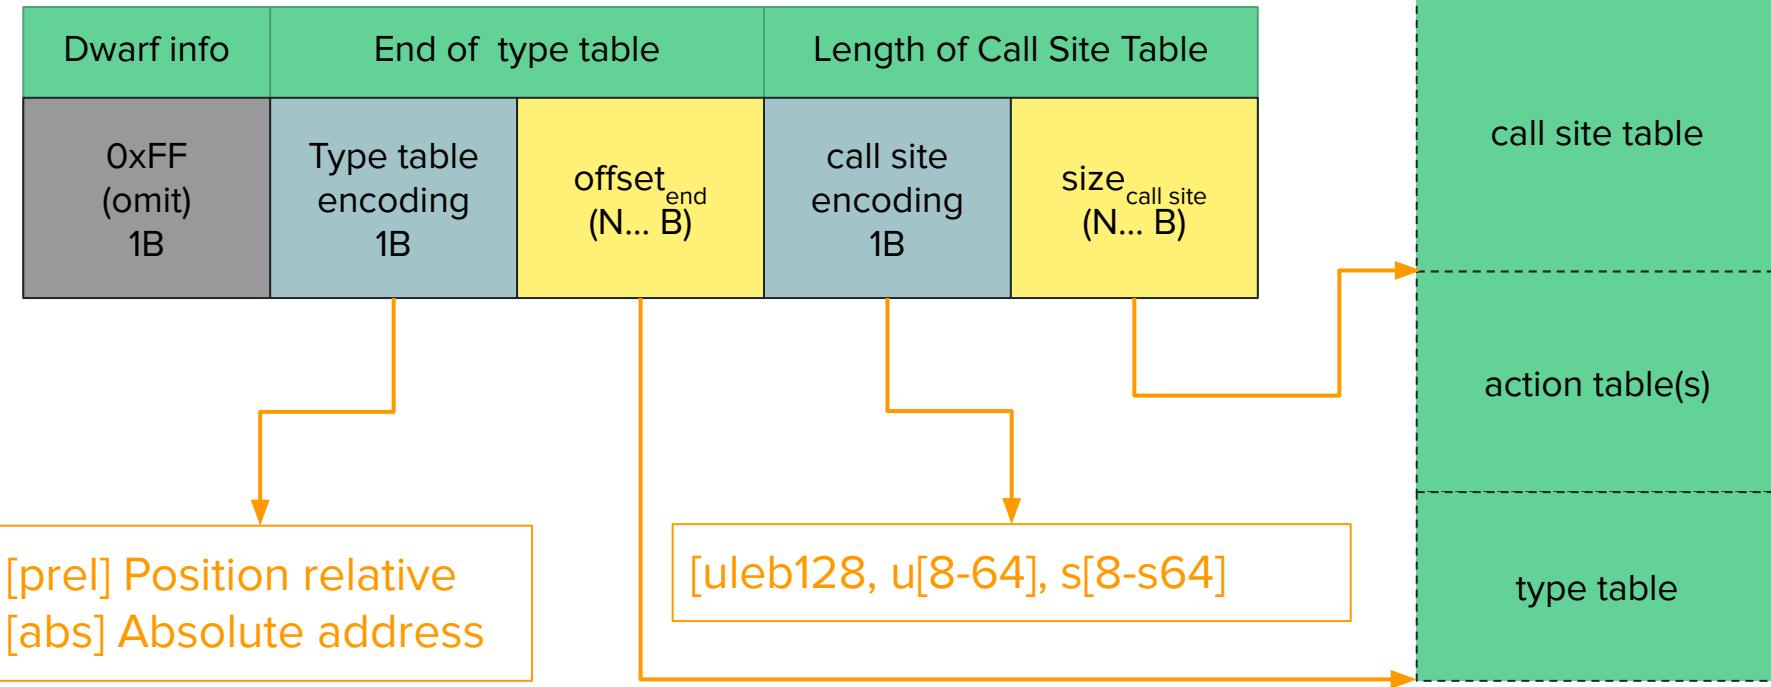
action table(s)

type table

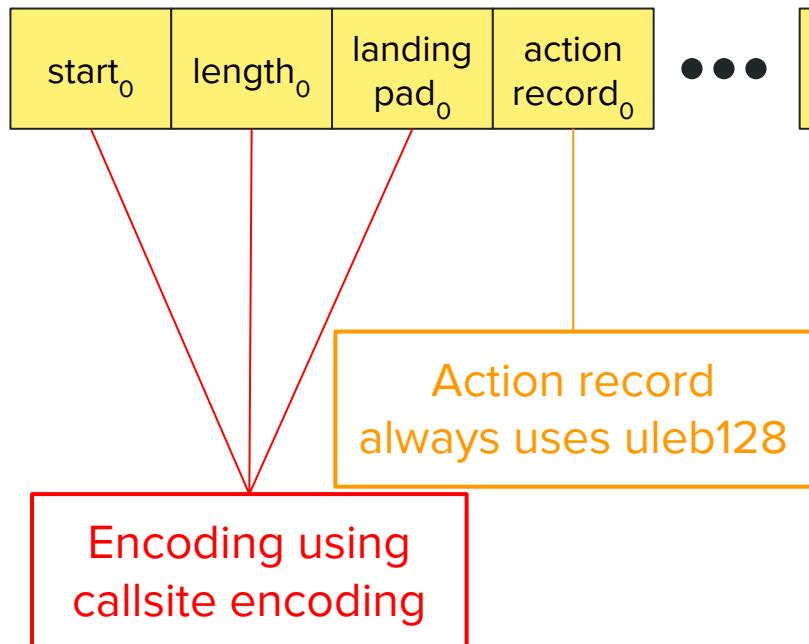
GCC LSDA: LEB128 variable length encoding



GCC LSDA: Typical Header



GCC LSDA: Call Site Table



GCC LSDA: Decoding the Call Site Table

```
do {  
    auto start = read_encoded_data(&lsda_cursor, call_site_format);  
    auto length = read_encoded_data(&lsda_cursor, call_site_format);  
    landing_pad = read_encoded_data(&lsda_cursor, call_site_format);  
    action = read_uleb128(&lsda_cursor);  
  
    if (start <= rel_pc && rel_pc < start + length) {  
        if (landing_pad == 0) {  
            p_exception_object.cache.state(runtime_state::unwind_frame);  
            return;  
        }  
        break;  
    }  
} while (lsda_cursor < call_site_end);
```

personality function

personality data

header

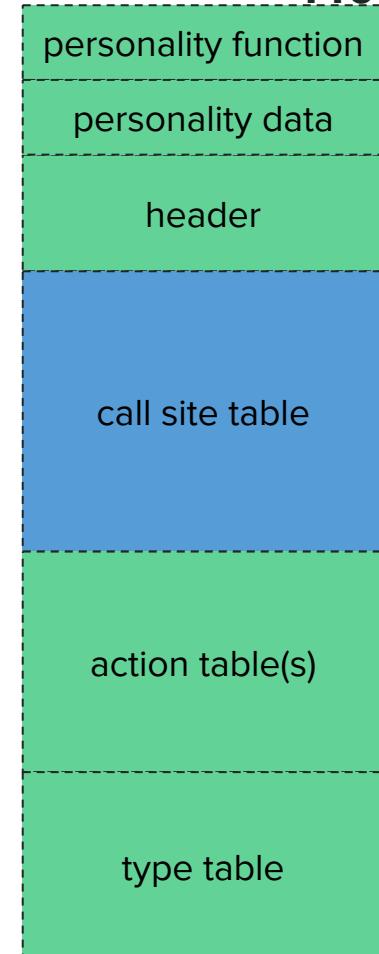
call site table

action table(s)

type table

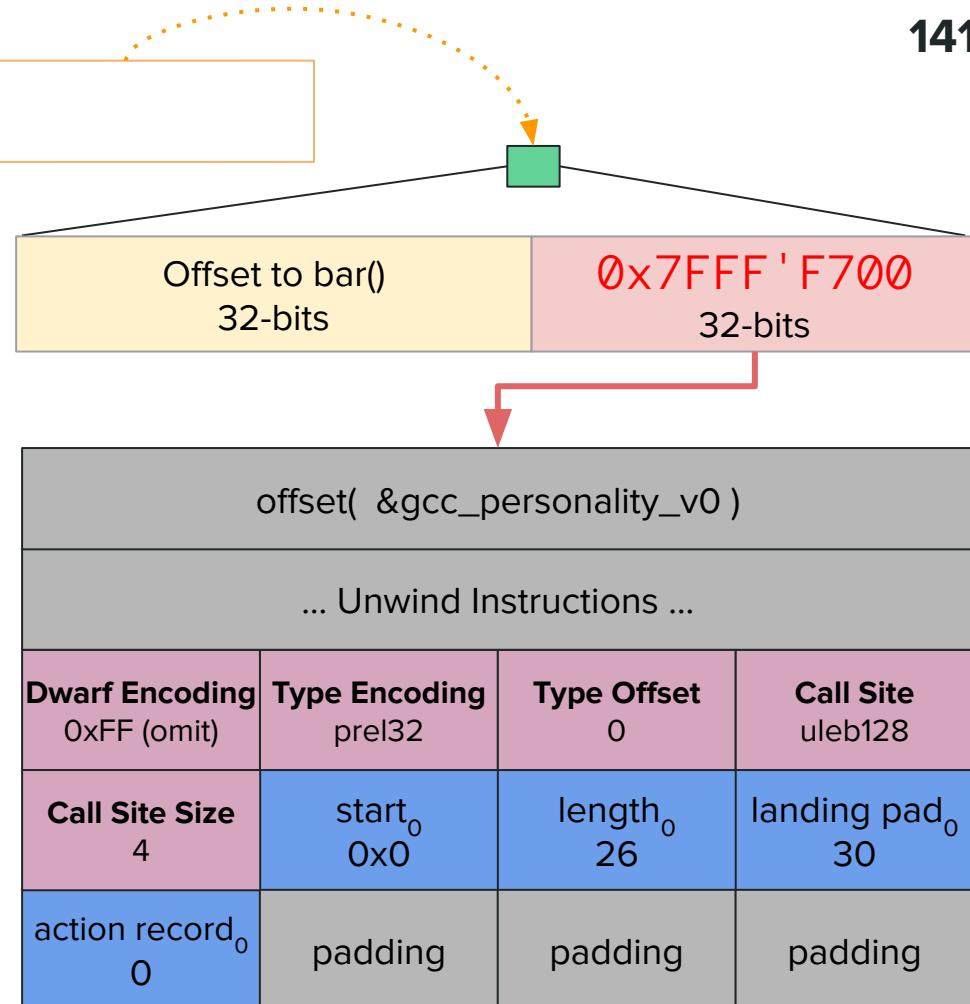
GCC LSDA: Executing Destructors

```
if (type_table_end < call_site_end ||  
    action == 0) {  
    // LSB must be set to 1 to jump to an address  
    auto const destination = (entry.function() + landing_pad) | 0b1;  
    // Set PC to the cleanup destination  
    virtual_cpu.pc = destination;  
    // Install CPU state  
    restore_cpu_core(virtual_cpu);  
}
```



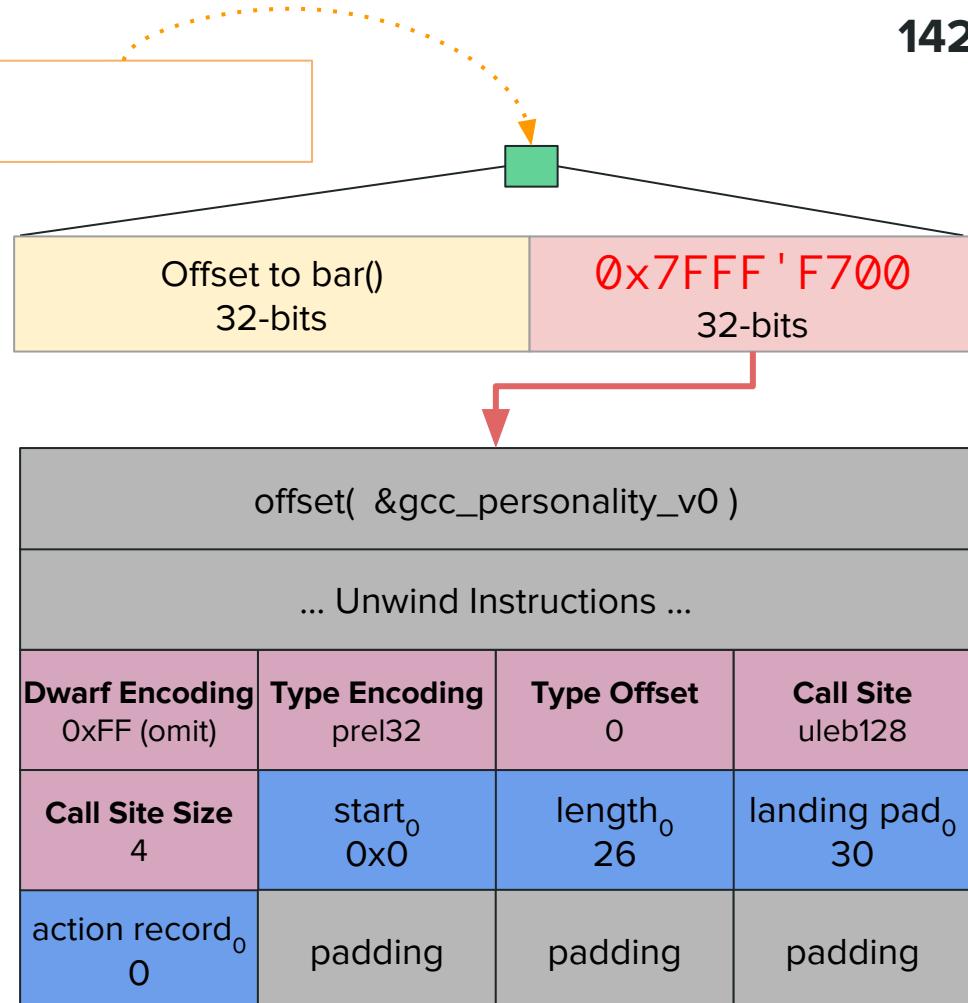
Entering bar()

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'C000	0x4000	0x?????????
R1	type_info	0x3FFC	0x?????????
R2	0x0	0x3FF8	0x?????????
R3	0x?????????	0x3FF4	0x?????????
R4	[bar] r4	0x3FF0	0x?????????
R5	0x?????????	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FCC	0x3FCC	obj[3/3]
R14/LR	bar() +14	0x3FC8	bar() +14
R15/PC	bar() +14	0x3FC4	[bar] r4



Entering bar()

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'C000	0x4000	0x?????????
R1	type_info	0x3FFC	0x?????????
R2	0x0	0x3FF8	0x?????????
R3	0x?????????	0x3FF4	0x?????????
R4	[bar] r4	0x3FF0	0x?????????
R5	0x?????????	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FCC	0x3FCC	obj[3/3]
R14/LR	bar() +30	0x3FC8	bar() +14
R15/PC	bar() +30	0x3FC4	[bar] r4



Executing bar()'s Destructor

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'C000	0x4000	0x?????????
R1	type_info	0x3FFC	0x?????????
R2	0x0	0x3FF8	0x?????????
R3	0x?????????	0x3FF4	0x?????????
R4	[bar] r4	0x3FF0	0x?????????
R5	0x?????????	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FCC	0x3FCC	obj[3/3]
R14/LR	bar() +30	0x3FC8	bar() +14
R15/PC	bar() +30	0x3FC4	[bar] r4

```

000088b4 <bar ()>:
    88b4: b500      push {lr}
    88b6: 4b09      ldr   r3, [pc, #36] ; (88dc <bar ()+0x28>)
    88b8: b083      sub   sp, #12
    88ba: 791b      ldrb  r3, [r3, #4]
    88bc: 220f      movs  r2, #15
    88be: 9201      str   r2, [sp, #4]
    88c0: b92b      cbnz  r3, 88ce <bar ()+0x1a>
    88c2: a801      add   r0, sp, #4
    88c4: f000 f832  bl    892c <destructible_t::~destructible_t ()>
    88c8: b003      add   sp, #12
    88ca: f85d fb04  ldr.w pc, [sp], #4
    88ce: f7f7 fc09  bl    e4 <baz () [clone .part.0]>
    88d2: a801      add   r0, sp, #4
    88d4: f000 f82a  bl    892c <destructible_t::~destructible_t ()>
    88d8: f000 fcc6  bl    9268 <__cxa_end_cleanup>
    88dc: 1000065c  .word   0x1000065c

```



GCC LSDA: Executing Destructors

personality function

personality data

header

call site table

action table(s)

type table

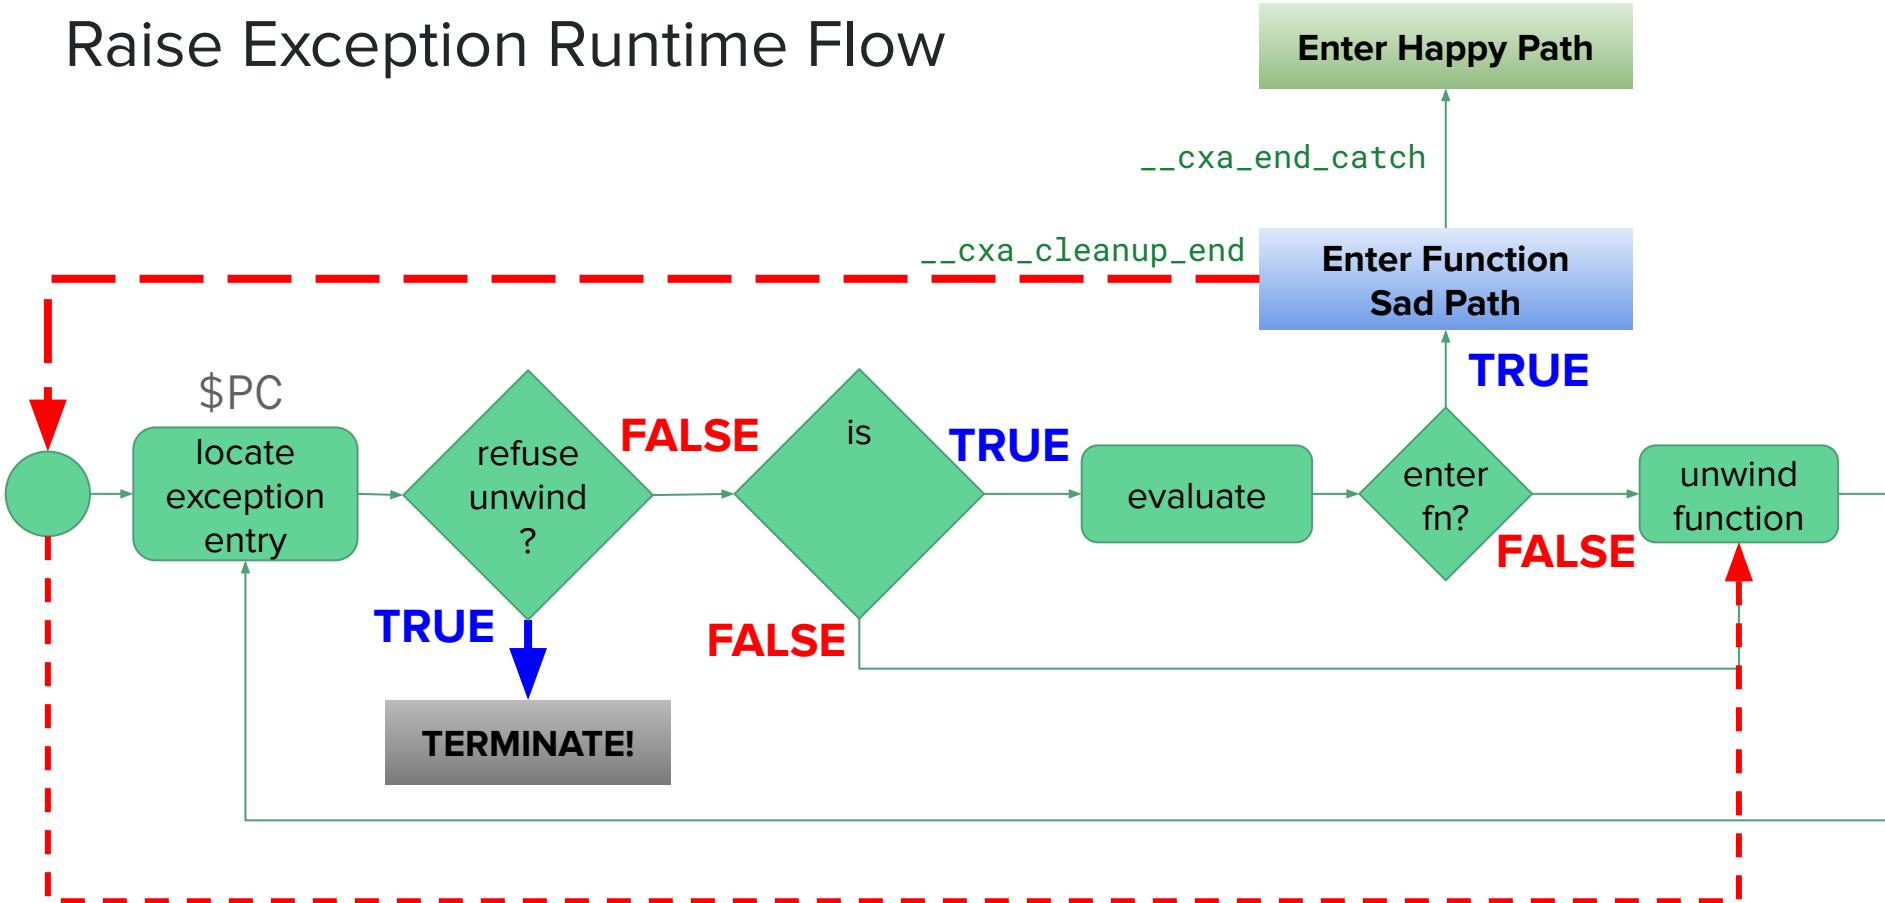
Load address of object on stack into R0

```
000088b4 <bar ()>:
    ; happy path instructions
    88d2: a801      add r0, sp, #4
    88d4: f000 f82a  bl  892c <destructible_t::~destructible_t()>
    88d8: f000 fcc6  bl  9268 <__cxa_end_cleanup>
```

Call Destructor

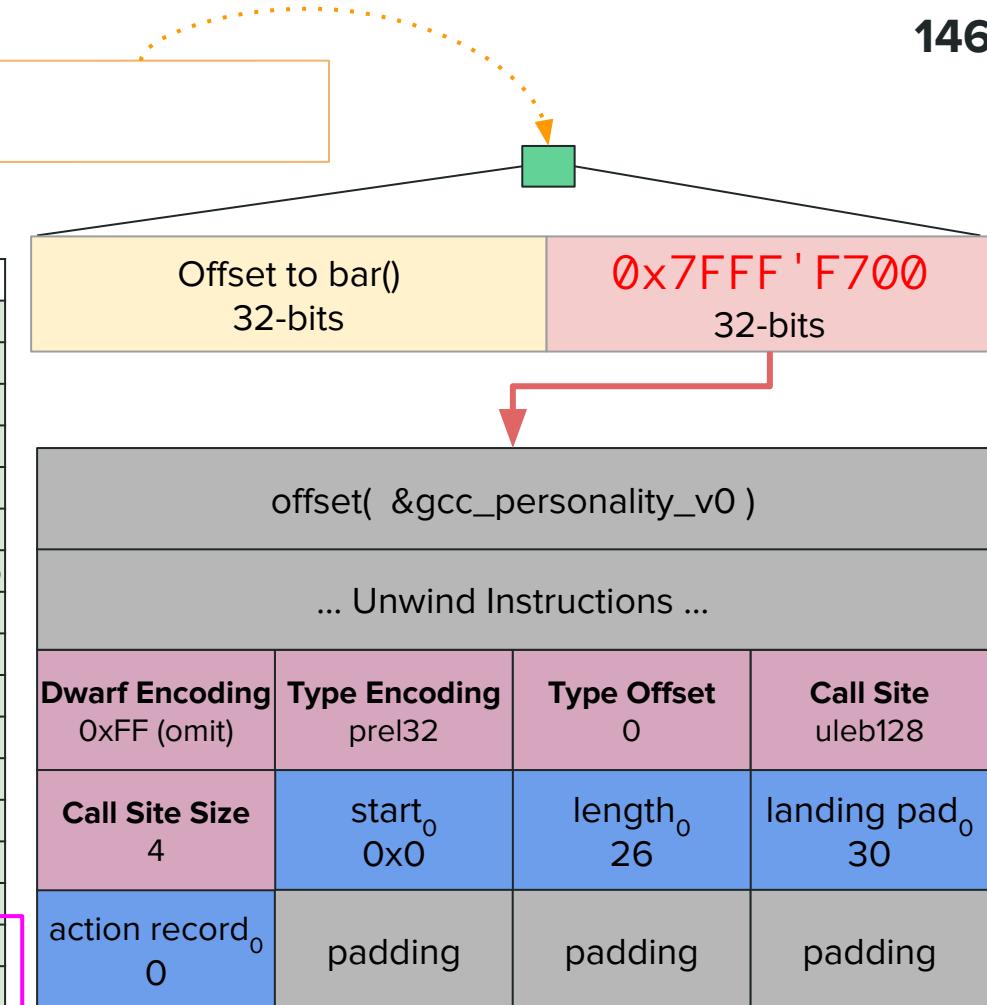
1. Recover exception object
2. Set state to "unwind frame"
3. raise_exception()

Raise Exception Runtime Flow



After Bar's Destructors

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'C000	0x4000	0x?????????
R1	type_info	0x3FFC	0x?????????
R2	0x0	0x3FF8	0x?????????
R3	0x?????????	0x3FF4	0x?????????
R4	[bar] r4	0x3FF0	0x?????????
R5	0x?????????	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FCC	0x3FCC	obj[3/3]
R14/LR	bar() +30	0x3FC8	0x?????????
R15/PC	bar() +30	0x3FC4	0x?????????



GCC LSDA: Executing Destructors

personality function

personality data

header

call site table

action table(s)

type table

Call Site 0 Landing Pad

```
000088b4 <bar():>
    ; happy path instructions
    88da: a801      add r0, sp, #4
    88dc: f000 f834  bl  8948 <destructible_t::~destructible_t()>
    88e0: 4668      mov r0, sp
    88e2: f000 f831  bl  8948 <destructible_t::~destructible_t()>
    88e6: f000 fcc4  bl  9284 <__cxa_end_cleanup>
```

Call Site 1 Landing Pad

Unwinding bar()

We are here

```
struct error { };

void foo() {
    try {
        bar();
    } catch (error const& p_error) {
        // end up here...
    }
}

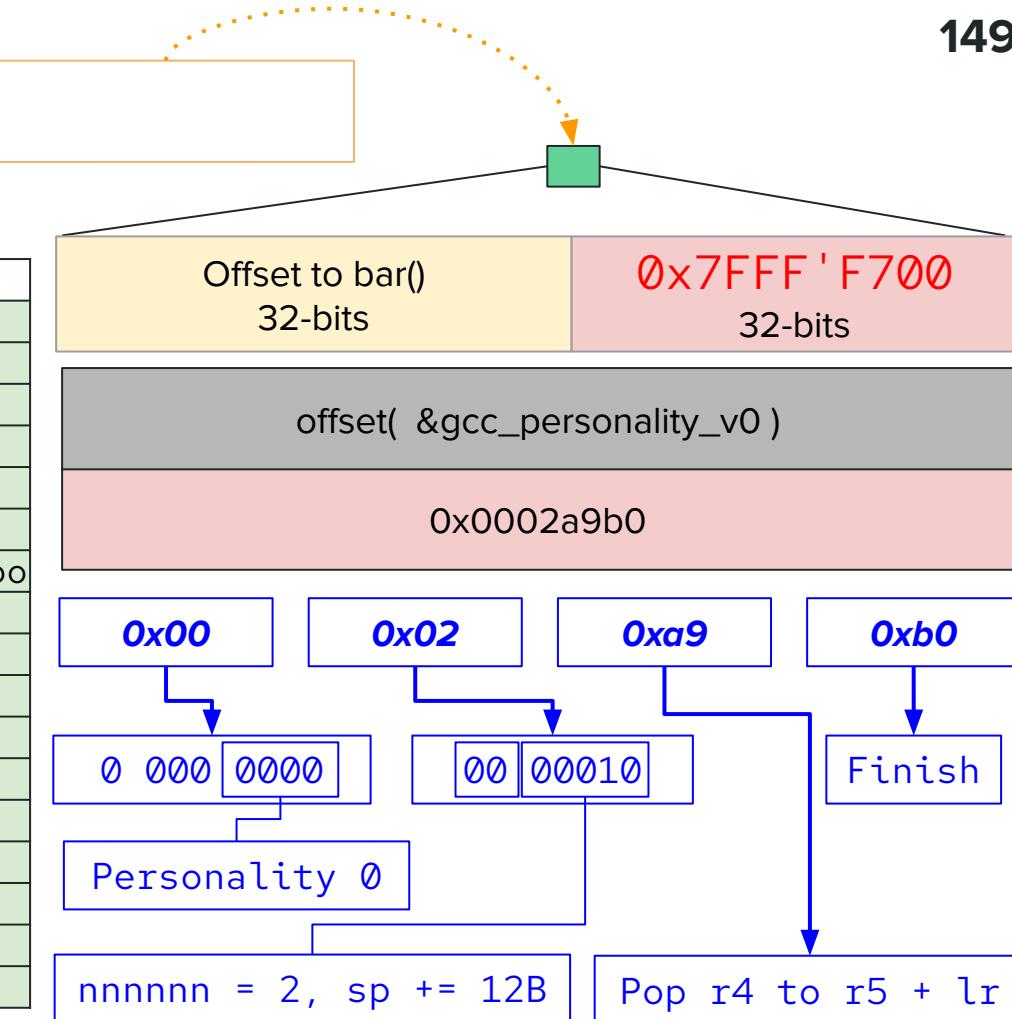
void bar() {
    destructible_t obj;
    baz();
} // ~destructible_t(&my_object)

void baz() { throw error{}; }
```

bar entry

Unwinding bar()

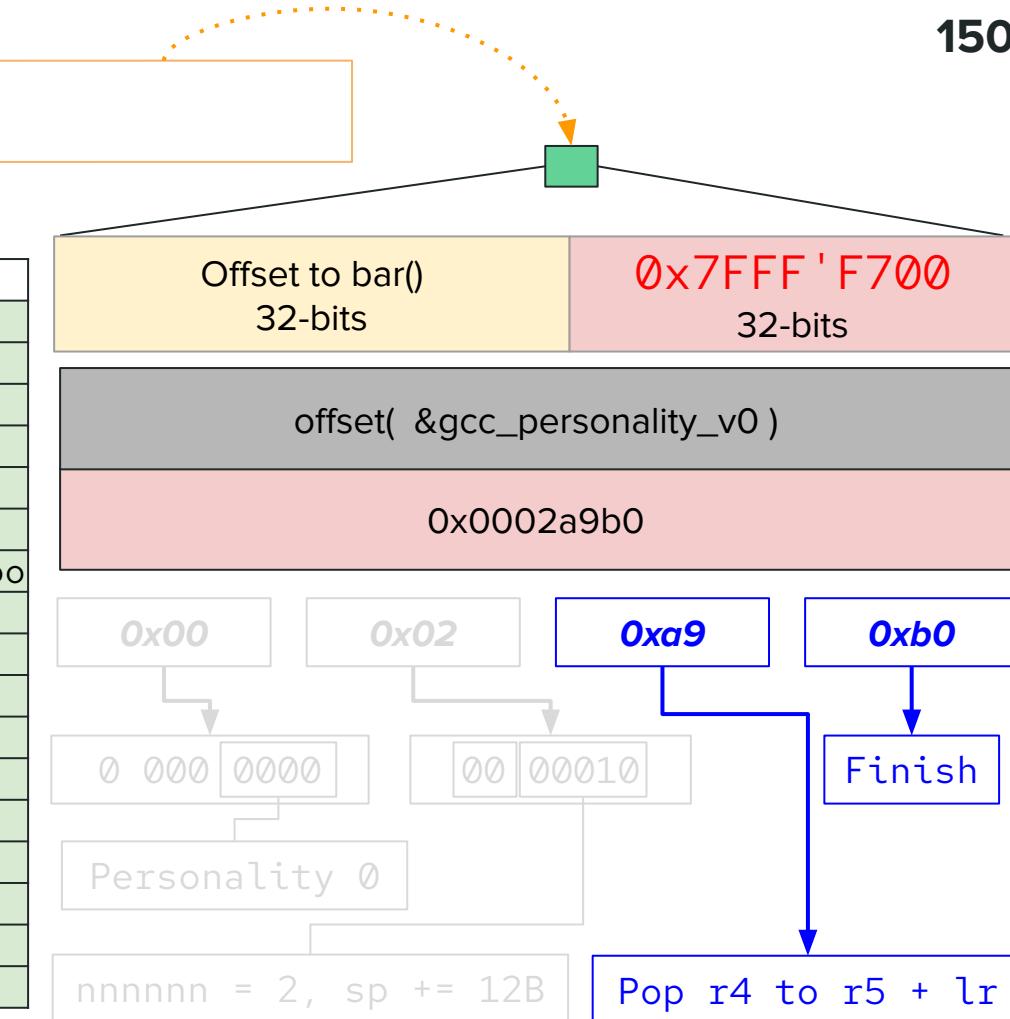
Virtual	ARM	CPU	SRAM	
Reg	Value		Address	Value
R0	0x2000'C000		0x4000	0x?????????
R1	type_info		0x3FFC	0x?????????
R2	0x0		0x3FF8	0x?????????
R3	0x?????????		0x3FF4	0x?????????
R4	[bar] r4		0x3FF0	0x?????????
R5	0x?????????		0x3FEC	caller of foo
R6	-		0x3FE8	caller r4
R7	-		0x3FE4	caller r5
R8	-		0x3FE0	foo() +10
R9	-		0x3FDC	[foo] r5
R10	-		0x3FD8	[foo] r4
R11	-		0x3FD4	obj[1/3]
R12	-		0x3FD0	obj[2/3]
R13/SP	0x3FCC		0x3FCC	obj[3/3]
R14/LR	bar() +14		0x3FC8	0x?????????
R15/PC	bar() +14		0x3FC4	0x?????????



bar entry

Unwinding bar()

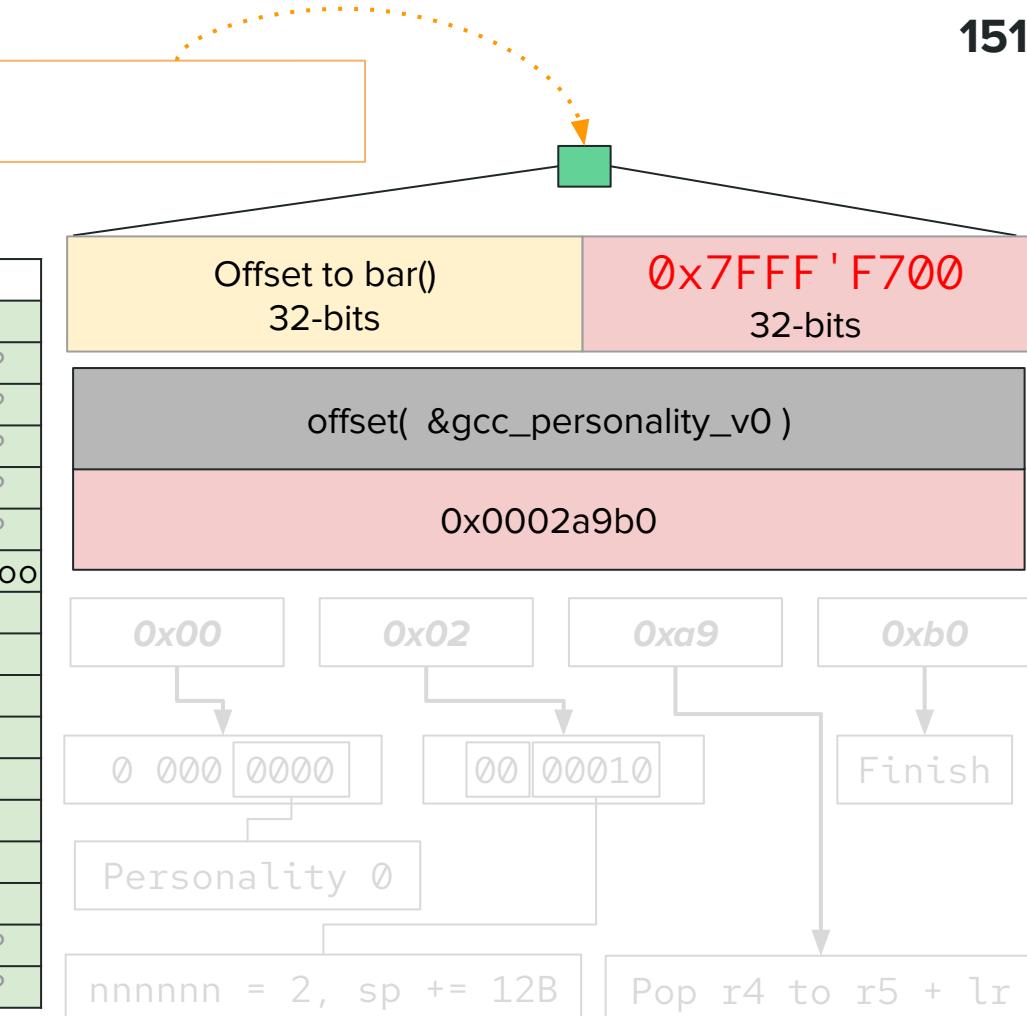
Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'C000	0x4000	0x????????
R1	type_info	0x3FFC	0x????????
R2	0x0	0x3FF8	0x????????
R3	0x????????	0x3FF4	0x????????
R4	[bar] r4	0x3FF0	0x????????
R5	0x????????	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FD8	0x3FCC	obj[3/3]
R14/LR	bar() +14	0x3FC8	0x????????
R15/PC	bar() +14	0x3FC4	0x????????



bar entry

Unwinding bar()

Virtual	ARM	CPU	SRAM	
Reg	Value		Address	Value
R0	0x2000'C000		0x4000	0x?????????
R1	type_info		0x3FFC	0x?????????
R2	0x0		0x3FF8	0x?????????
R3	0x?????????		0x3FF4	0x?????????
R4	[foo] r4		0x3FF0	0x?????????
R5	[foo] r5		0x3FEC	caller of foo
R6	-		0x3FE8	caller r4
R7	-		0x3FE4	caller r5
R8	-		0x3FE0	foo() +10
R9	-		0x3FDC	[foo] r5
R10	-		0x3FD8	[foo] r4
R11	-		0x3FD4	obj[1/3]
R12	-		0x3FD0	obj[2/3]
R13/SP	0x3FE0		0x3FCC	obj[3/3]
R14/LR	foo() +10		0x3FC8	0x?????????
R15/PC	foo() +10		0x3FC4	0x?????????



Entering foo()'s Catch block...

We are here

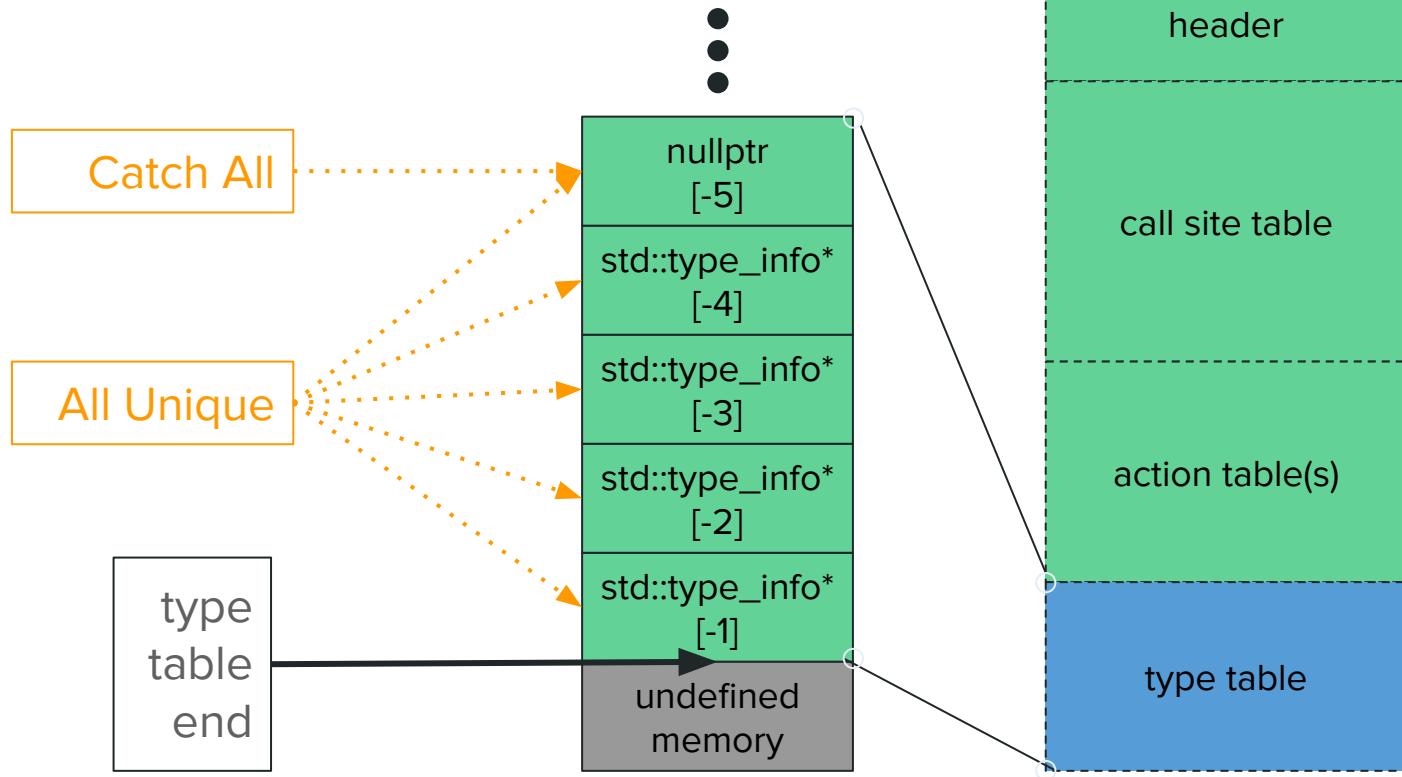
```
struct error { };

void foo() {
    try {
        bar();
    } catch (error const& p_error) {
        // end up here...
    }
}

void bar() {
    destructible_t obj;
    baz();
} // ~destructible_t(&my_object)

void baz() { throw error{}; }
```

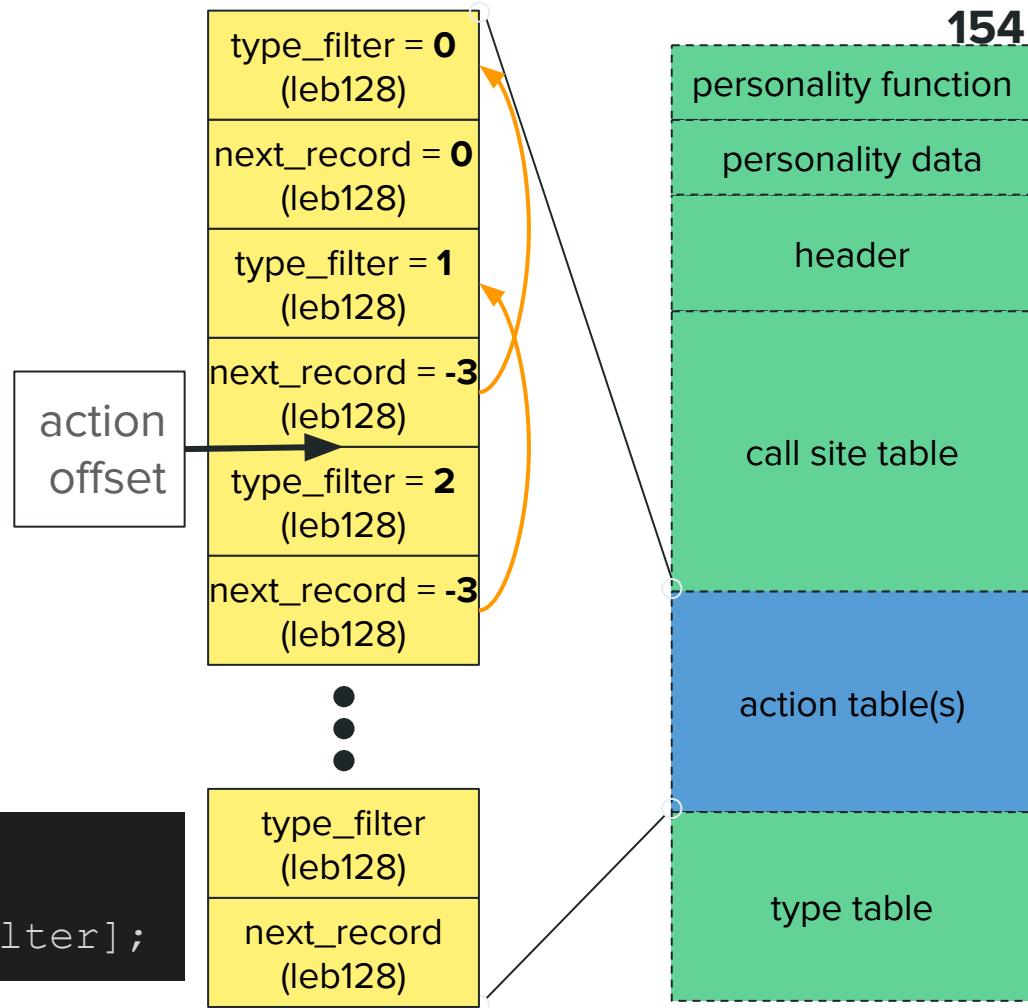
GCC LSDA: Type Table



GCC LSDA: Action Table

- The action table represents
 - The sequence of catch blocks attached to a try scope
 - provides the mapping between catch block, switch case value, and the `std::type_info` of the caught object.
- The type filter is the negative **index** into the **type table**.

```
std::type_info* current_type =
    type_table_end[-type_filter];
```



0 = catchbackrecords

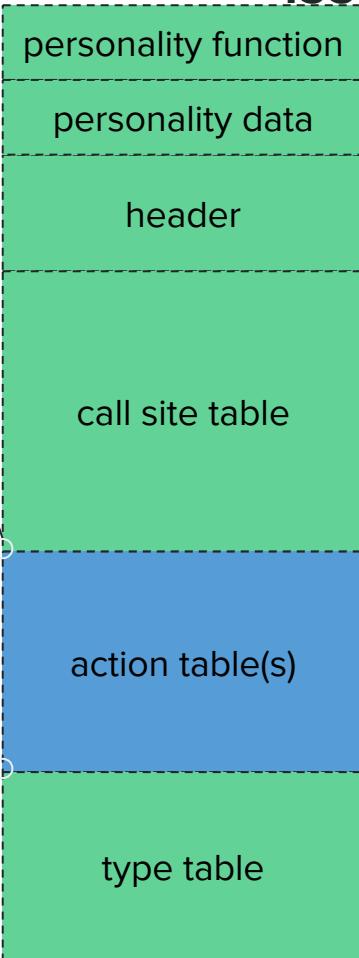
GCC LSDA: Action Table



If the runtime reaches a **next_record** value of 0 and with no types match, then the current frame should be unwound.

type_filter = 0 (leb128)
next_record = 0 (leb128)
type_filter = 1 (leb128)
next_record = -3 (leb128)
type_filter = 2 (leb128)
next_record = -3 (leb128)
⋮
type_filter (leb128)
next_record (leb128)

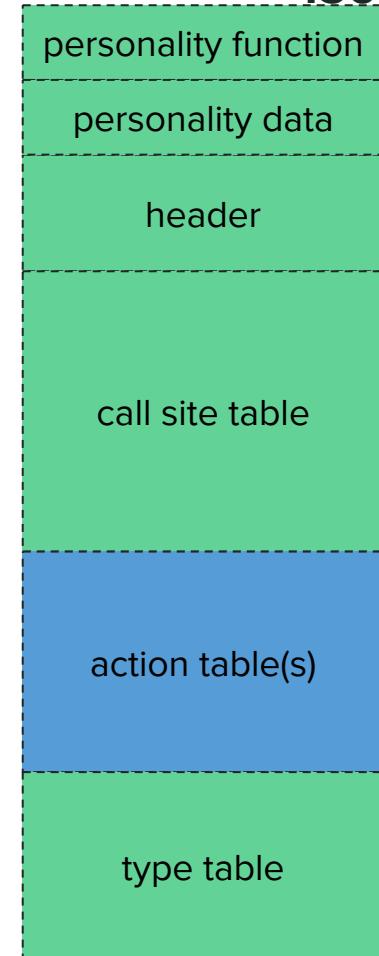
type_filter (leb128)
next_record (leb128)



GCC LSDA: Entering the Catch Block

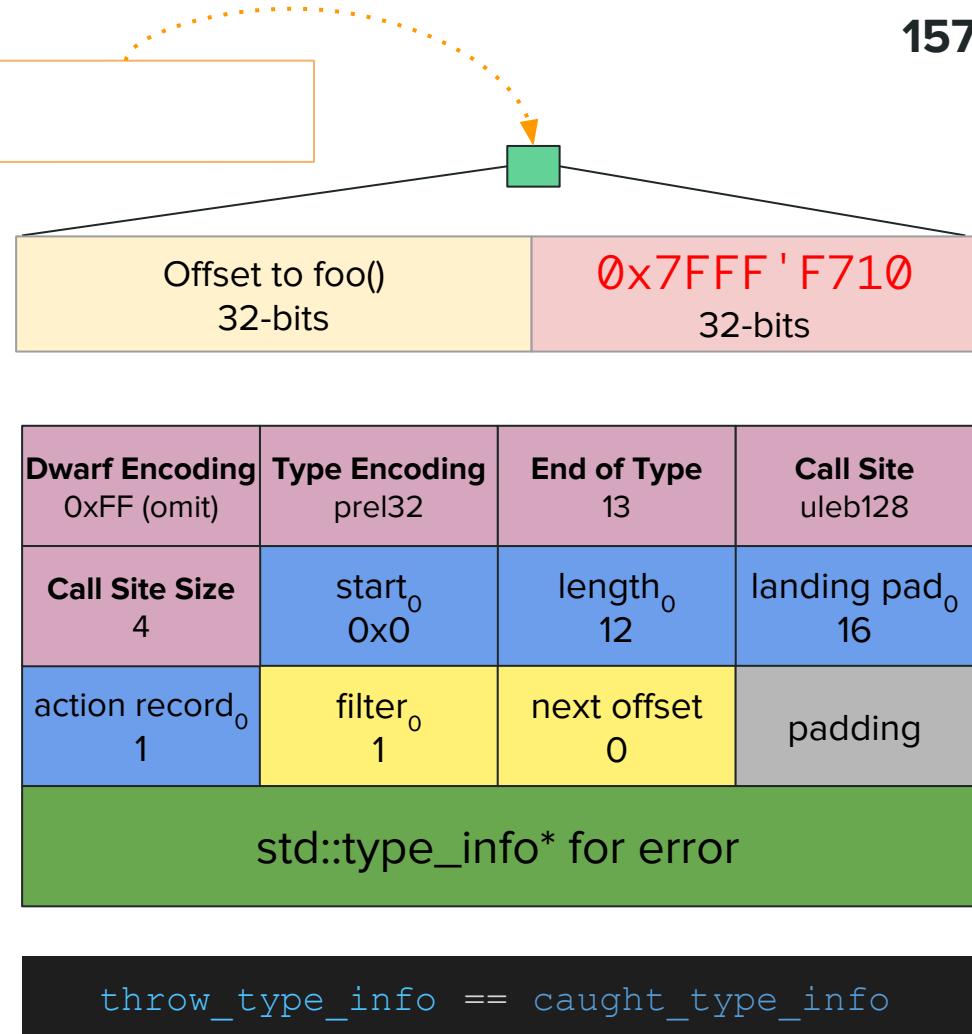
If a match between the type that was thrown and a catch type perform the following steps:

1. Set **R0** \leftarrow exception object's address
2. Set **R1** \leftarrow type filter number
3. Set **PC** \leftarrow `function_address + landing_pad | 1`
4. Set **actual CPU** to **virtual CPU**.



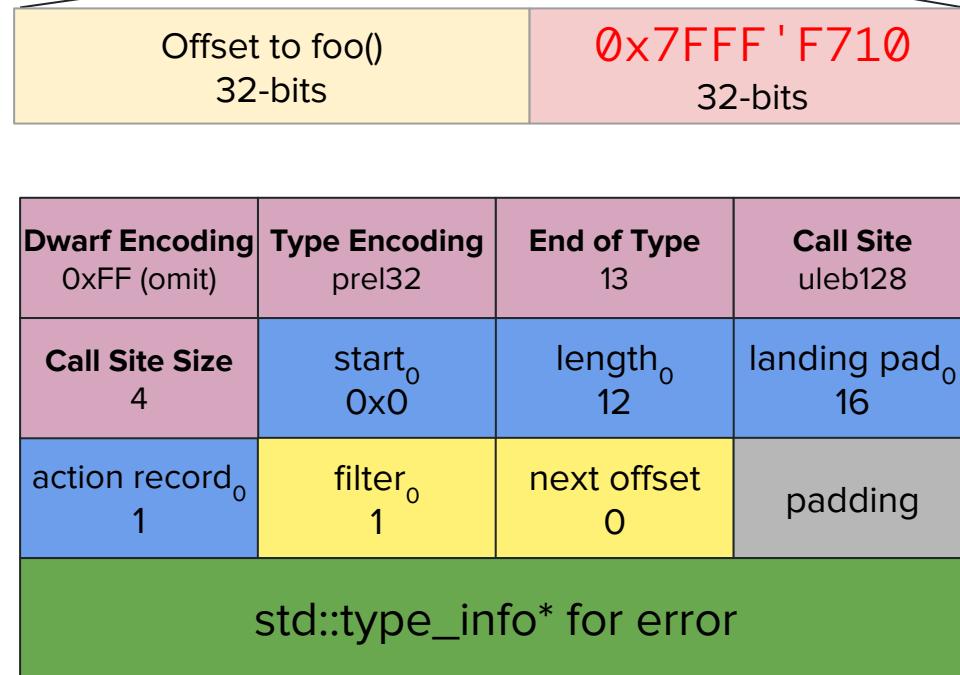
Entering foo()

Virtual	ARM	CPU	SRAM	
Reg	Value		Address	Value
R0	0x2000'C000		0x4000	0x?????????
R1	type_info		0x3FFC	0x?????????
R2	0x0		0x3FF8	0x?????????
R3	0x?????????		0x3FF4	0x?????????
R4	[foo] r4		0x3FF0	0x?????????
R5	[foo] r5		0x3FEC	caller of foo
R6	-		0x3FE8	caller r4
R7	-		0x3FE4	caller r5
R8	-		0x3FE0	foo() +10
R9	-		0x3FDC	[foo] r5
R10	-		0x3FD8	[foo] r4
R11	-		0x3FD4	obj[1/3]
R12	-		0x3FD0	obj[2/3]
R13/SP	0x3FE0		0x3FCC	obj[3/3]
R14/LR	foo() +10		0x3FC8	0x?????????
R15/PC	foo() +10		0x3FC4	0x?????????



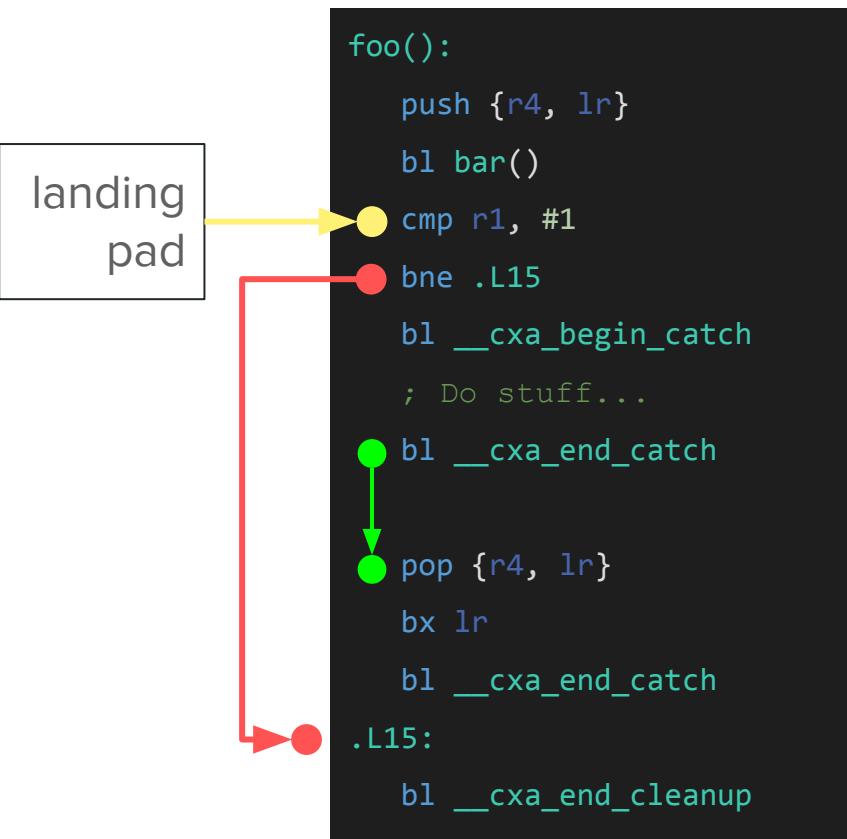
Entering foo()

Virtual ARM CPU		SRAM	
Reg	Value	Address	Value
R0	0x2000'A100	0x4000	0x?????????
R1	1	0x3FFC	0x?????????
R2	0x0	0x3FF8	0x?????????
R3	0x?????????	0x3FF4	0x?????????
R4	[foo] r4	0x3FF0	0x?????????
R5	[foo] r5	0x3FEC	caller of foo
R6	-	0x3FE8	caller r4
R7	-	0x3FE4	caller r5
R8	-	0x3FE0	foo() +10
R9	-	0x3FDC	[foo] r5
R10	-	0x3FD8	[foo] r4
R11	-	0x3FD4	obj[1/3]
R12	-	0x3FD0	obj[2/3]
R13/SP	0x3FE0	0x3FCC	obj[3/3]
R14/LR	foo() +16	0x3FC8	0x?????????
R15/PC	foo() +16	0x3FC4	0x?????????



```
throw_type_info == caught_type_info
```

Within the Catch Block



● `__cxa_begin_catch`

- R0 contains the exception object which will be the first parameter to this function.
- Set exception state to **handled**
- Takes the exception object and returns the adjusted pointer to the exception object.

● `__cxa_end_catch`

- Call destructor of thrown object
- Frees thrown object memory

And that's

ARM GCC Exceptions!



Space Cost of Exception Handling

We are going to size up the following

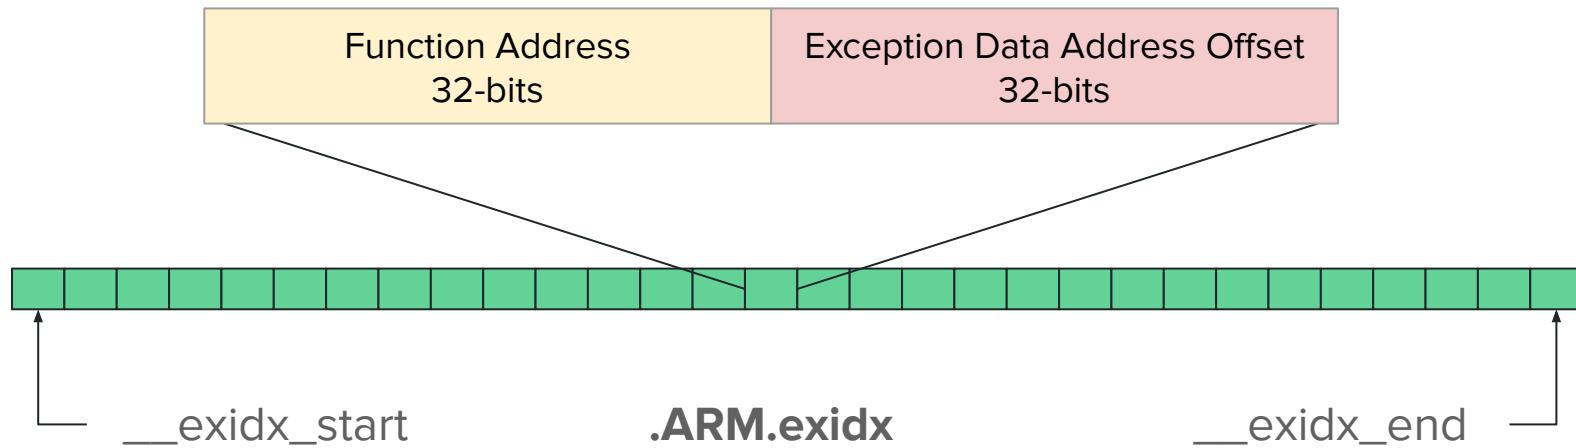
Structure	Size	Notes
Exception Index Entry	???	
SU16 Personality Data	???	
LU16/32 Personality Data	???	
LSDA header	???	
LSDA Call Site Table	???	
Cleanup Landing Area	???	
LSDA Action Table	???	
LSDA Type Table	???	
Catch Block	???	

Assumptions

- Assume the code is being compiled to minimize code size
- Assume that functions are generally at or below **16 kB** in size each
 - [Firefox Debug Symbols Spreadsheet](#)
- Assume that no function has more than a total of **63 catch blocks** and **destructor scopes** in it.
- Assume everything uses LEB128 & ULEB128 encodings.

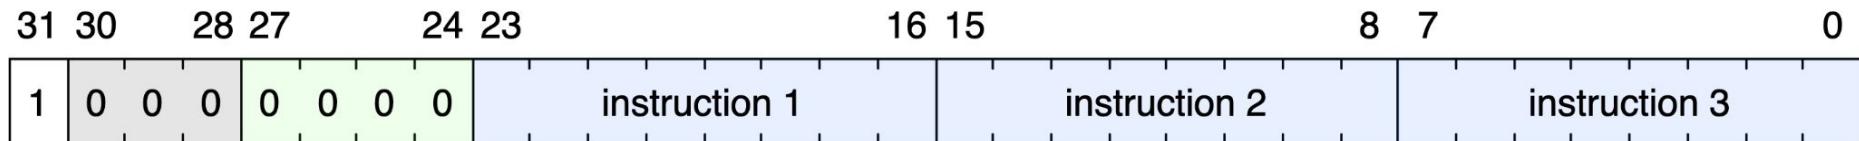
Firefox Function Distribution	
242	functions above (16kB)
325108	functions below (16kB)
157946	functions above (127B)
167404	functions below (127B)

`sizeof(Exception Index Entry) = 8 bytes`



As defined by the ABI.

`sizeof(SU16 Personality Data) = 4 bytes`



As defined by the ABI.

`sizeof(LU16/32 Personality Data) = 12 bytes`

31	30	28	27	24	23	16	15	8	7	0
1	0	0	0	0	0	length (1 or 2)		instruction 0	instruction 1	
								instruction 4	instruction 5	
								0xB0	0xB0	

Worst case size as defined in the ABI.

sizeof(Catch Block) = 12 bytes

- `sizeof(cmp)` = 2 bytes
- `sizeof(beq)` = 2 bytes
- `sizeof(bl __cxa_begin_catch)` = 4 bytes
- `sizeof(bl __cxa_end_catch)` = 4 bytes

```
00000000 <catch_my_errors():>
 0: b508      push    {r3, lr}
 2: f7ff fffe  bl <may_throw_exception()>
 6: e00f      b.n    1e <catch_my_errors()+0x28>
 8: 2901      cmp    r1, #1
 a: d001      beq.n   10 <catch_my_errors()+0x10>
 c: f7ff fffe  bl <__cxa_end_cleanup>
10: f7ff fffe  bl <__cxa_begin_catch>
14: 4805      ldr    r0, [pc, #12]
16: f7ff fffe  bl <printf>
1a: f7ff fffe  bl <__cxa_end_catch>
1e: bd08      pop    {r3, pc}
26: 00000000  .word   0x00000000
```

sizeof(Cleanup Landing Area) = **6D + 4**

```
000088b4 <bar ()>:  
    ; happy path instructions  
  
    88da: a801      add r0, sp, #4  
    88dc: f000 f834  bl  8948 <destructible_t::~destructible_t()>  
    88e0: 4668      mov r0, sp  
    88e2: f000 f831  bl  8948 <destructible_t::~destructible_t()>  
    88e6: f000 fcc4  bl  9284 <__cxa_end_cleanup>
```

D = Each objects within a frame requiring cleanup

- `sizeof(mov/add/sub)` = 2 bytes
- `sizeof(bl to destructor)` = 4 bytes
- `sizeof(bl __cxa_end_cleanup)` = 4 bytes

$$6(2) + 4 = 16B$$

sizeof(LSDA Action Table) = **2A**

A = catch block count

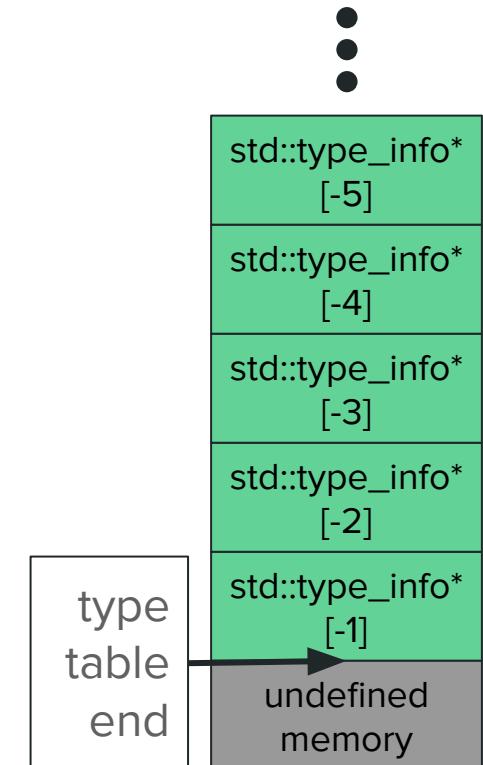
- There is an entry for each
 - catch block
 - cleanup scope
- The range of a signed byte of LEB128 integer is -64 to 63
- Due to our constraints:
 - "Type filter" will not exceed 63 selectable types
 - "Next record" should not exceed 63 because entries are packed next to each other

type_filter = 0 (leb128)
next_record = 0 (leb128)
type_filter = 1 (leb128)
next_record = -3 (leb128)
type_filter = 2 (leb128)
next_record = -3 (leb128)

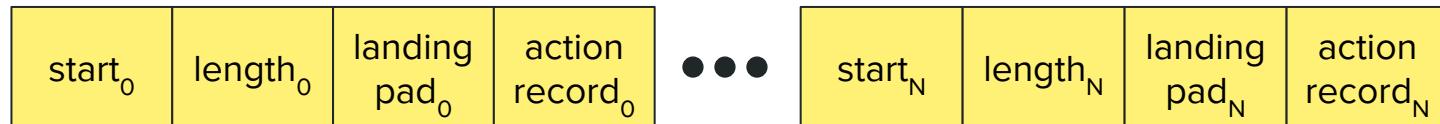


`sizeof(LSDA Type Table) = 4T`

$T = \text{number of unique types caught in the function}$



`sizeof(LSDA Call Site) = 4S ↔ 7S`



$S = \text{number of try scopes} + \text{scopes w/ cleanup}$

- 16kB function size constraint constraints **Start, length & landing pad** to 16kB or 2-bytes of ULEB128.
- The action table shouldn't exceed a size of 2^*A (or 2^*63) which is 126 bytes, thus the **Action record offset** should stay as a single ULEB128 byte

`sizeof(LSDA header) = 5 ↔ 7`

0xFF (omit)	Type table encoding	offset _{end} (N... B)	call site encoding	size _{call site} (N... B)
Dwarf info	End of type table		Length of Call Site Table	

- Minimum size of 5 bytes
- What could grow in it?
 - Type table size?
 - Call site length?

`sizeof(LSDA header) = 5 ↔ 7`

0xFF (omit)	Type table encoding	offset _{end} (N... B)	call site encoding	size _{call site} (N... B)
Dwarf info	End of type table		Length of Call Site Table	

- With our constraints

- **Type table** cannot exceed 63 types so the max size is $4*63 = \mathbf{252\ bytes}$
- **Action table** cannot exceed 63 catch blocks + cleanup regions so the max is $2*(63) = \mathbf{126\ Bytes}$
- Total size without the call site of **372 bytes**, leaving 16,004 bytes for the call site.
- Assuming the worst case of 7S, this would require a function to have **~2,286** call site scopes, which seems impractical for most code.

I assert that with our constraints the size range is 5 bytes to 7 bytes.

We are going to size up the following

Structure	Size (bytes)	Notes
Exception Index Entry	8	
SU16 Personality Data	4	
LU16/32 Personality Data	12	
LSDA header	5 ↔ 7	
LSDA Call Site Table	4S ↔ 7S	where S = count of try & cleanup scopes
Cleanup Landing Area	6D + 4	where D = count of objects that need cleanup
LSDA Action Table	2A	where A = the number of catch blocks & cleanup scopes
LSDA Type Table	4T	where T = number of unique types caught in the function
Catch Block	12	

Examples! Trivially Unwindable Function

```
void blink_n_times(
    output_pin& p_led,
    uint8_t p_count)
{
    for (uint8_t i = 0; i < p_count; i++) {
        led.level(false);
        delay(500ms);
        led.level(true);
        delay(500ms);
    }
}
```

- **Metadata Components:**
 - Index Entry: 8-bytes
- **Total: 8 bytes**
- The size will not increase unless:
 - Stack size exceeds 512 (+8 bytes)
 - Object with destructor added
 - try/catch added

Examples! Function with single destructor scope

```
void foo(i2c& p_i2c,
          std::span<std::byte> p_data)
{
    std::lock_guard<my_mutex_t> lock(my_mutex);
    write(i2c, 0x44, p_data);
}
```

- **Metadata Components**
 - Index Entry: 8 bytes
 - GCC LSDA:
 - Header: 5 bytes
 - Personality: 4 bytes (SU16)
 - Call site table: 4 bytes
 - Padding: 3 bytes
 - Cleanup Landing Area: 10 bytes
- **Total: 34 bytes**
- The size will not increase unless:
 - Stack size exceeds 512 (+8 bytes)
 - Object with destructor added (+6 each)
 - try/catch added (Uses LSDA)

Examples! Function with try/catch

```
[nodiscard] inline bool probe(i2c& p_i2c,
    hal::byte p_address) {
    std::array<std::byte, 1> data_in{};

    bool device_acknowledged = true;
    try {
        p_i2c.transaction(p_address, {}, data_in);
    } catch (const hal::no_such_device&) {
        device_acknowledged = false;
    }
    return device_acknowledged;
}
```

$$8 + (5 + 4 + 4S + 2A + 4T) + 12C \rightarrow 8 + (5 + 4 + 4(1) + 2(1) + 4(1)) + 12(1)$$

Total: **47 bytes**

- **Metadata Components**

- Index Entry: 8 bytes
- GCC LSDA:
 - Header: 5 bytes (min sized)
 - Personality: 4 bytes (SU16)
 - Try/Destructor Scopes: 1
 - Catch Count: 1
 - Unique Types Caught: 1
 - Catch Blocks: 1

Space Cost of Return Type Error Handling

Basis Function: A function calling other functions

Consider the following with no error handling/using exception

```
namespace {  
    uint32_t func1();  
    uint32_t func2();  
    uint32_t func3();  
} // namespace  
  
uint32_t none()  
{  
    return func1() + func2() +  
    func3();  
}
```

```
000004c0 <none ()>:  
    4c0:   b508      push   {r3, lr}  
    4c2:   f7ff fff7  bl    4b4 <func1 ()>  
    4c6:   4603      mov    r3, r0  
    4c8:   f7ff fff6  bl    4b8 <func2 ()>  
    4cc:   4403      add    r3, r0  
    4ce:   f7ff fff5  bl    4bc <func3 ()>  
    4d2:   4418      add    r0, r3  
    4d4:   bd08      pop    {r3, pc}  
    4d6:   bf00      nop
```

GCC ARM 12.3.0

-O3
-std=c++23
-fexceptions
-frtti
-mthumb
-mcpu=cortex-m4
-mfloating-abi=hard

Returning std::expected

```
enum class my_error_t : std::uint8_t
{
    error1,
    error2,
    error3,
    error4
};
```

```
namespace {
std::expected<uint32_t, my_error_t> func1();
// the rest . .
} // namespace

std::expected<uint32_t, my_error_t> returning_expected()
{
    uint32_t sum = 0;
    auto result1 = func1();
    [[unlikely]] if (!result1) {
        return std::unexpected(result1.error());
    }
    sum += result1.value();
    auto result2 = func2();
    [[unlikely]] if (!result2) {
        return std::unexpected(result2.error());
    }
    sum += result2.value();
    auto result3 = func3();
    [[unlikely]] if (!result3) {
        return std::unexpected(result3.error());
    }
    sum += result3.value();
    return sum;
}
```

Returning std::expected: Disassembled!

```

00000468 <returning_expected()>:
    468: b510      push   {r4, lr}
    46a: b084      sub sp, #16
    46c: 4603      mov r3, r0
    46e: 4668      mov r0, sp
    470: f7ff ffe8  bl 444 <func1()>
    474: f89d 2004  ldrb.w r2, [sp, #4]
    478: b1e2      cbz r2, 4b4 <using_expected()+0x4c>
    47a: a802      add r0, sp, #8
    47c: 9d00      ldr r4, [sp, #0]
    47e: f7ff ffe7  bl 450 <func2()>
    482: f89d 200c  ldrb.w r2, [sp, #12]
    486: b172      cbz r2, 4a6 <using_expected()+0x3e>
    488: 9a02      ldr r2, [sp, #8]
    48a: 4414      add r4, r2
    48c: f7ff ffe6  bl 45c <func3()>
    490: f89d 200c  ldrb.w r2, [sp, #12]
    494: b13a      cbz r2, 4a6 <using_expected()+0x3e>
    496: 9902      ldr r1, [sp, #8]
    498: 2201      movs r2, #1
    49a: 440c      add r4, r1
    49c: 4618      mov r0, r3
    49e: 601c      str r4, [r3, #0]
    4a0: 711a      strb r2, [r3, #4]

```

4a2:	b004	add sp, #16
4a4:	bd10	pop {r4, pc}
4a6:	f89d 1008	ldr.w r1, [sp, #8]
4aa:	7019	strb r1, [r3, #0]
4ac:	4618	mov r0, r3
4ae:	711a	strb r2, [r3, #4]
4b0:	b004	add sp, #16
4b2:	bd10	pop {r4, pc}
4b4:	f89d 1000	ldr.w r1, [sp]
4b8:	7019	strb r1, [r3, #0]
4ba:	4618	mov r0, r3
4bc:	711a	strb r2, [r3, #4]
4be:	b004	add sp, #16
4c0:	bd10	pop {r4, pc}
4c2:	bf00	nop

Cost of **6 bytes** per function call

Returning `bool`

```
namespace {
    bool func1(std::uint32_t* p_value);
    // The rest...
} // namespace
```

```
bool returns_bool (std::uint32_t* p_value)
{
    uint32_t sum = 0;
    uint32_t temporary = 0;

    auto success = func1(&temporary);
    [[unlikely]] if (not success) {
        return false;
    }
    sum += temporary;

    success = func2(&temporary);
    [[unlikely]] if (not success) {
        return false;
    }
    sum += temporary;

    success = func3(&temporary);
    [[unlikely]] if (not success) {
        return false;
    }
    sum += temporary;

    *p_value = sum;

    return true;
}
```

Return `bool`: Disassembled!

```

00000510 <returns_bool(unsigned long*)>:
    510: b510      push   {r4, lr}
    512: b082      sub sp, #8
    514: 4602      mov r2, r0
    516: a801      add r0, sp, #4
    518: f7ff ffee  bl 4f8 <func1(unsigned long*)>
51c: b910      cbnz r0, 524 <returns_bool+0x14>
    51e: 2000      movs r0, #0
    520: b002      add sp, #8
    522: bd10      pop {r4, pc}
    524: a801      add r0, sp, #4
    526: 9901      ldr r1, [sp, #4]
    528: f7ff ffea  bl 500 <func2(unsigned long*)>
52c: 2800      cmp r0, #0
52e: d0f6      beq.n 51e <returns_bool+0xe>
    530: a801      add r0, sp, #4
    532: 9c01      ldr r4, [sp, #4]
    534: f7ff ffe8  bl 508 <func3(unsigned long*)>
538: 2800      cmp r0, #0
53a: d0f0      beq.n 51e <returns_bool+0xe>
    53c: 9b01      ldr r3, [sp, #4]
    53e: 4421      add r1, r4
    540: 440b      add r3, r1
    542: 6013      str r3, [r2, #0]
    544: e7ec      b.n 520 <returns_bool+0x10>
    546: bf00      nop

```

Cost of **4 bytes** per function call

Using Rust

```
rustc  
--target thumbv7em-none-eabi  
-C opt-level="s"  
--emit=obj  
rust.rs
```

```
pub fn combined() -> Result<i32, u32> {  
    let mut num: i32 = 0;  
    num += func1()?;
    num += func2()?;
    num += func3()?;
    return Ok(num);
}  
  
#[inline(never)]  
pub fn func1() -> Result<i32, u32> { return Ok(5); }  
#[inline(never)]  
pub fn func2() -> Result<i32, u32> { return Err(88); }  
#[inline(never)]  
pub fn func3() -> Result<i32, u32> { return Ok(13); }
```

Return Rust: Disassembled!

```

00000000 <rust::combined::h7a7b0cc556a62e18>:
 0: b5f0      push {r4, r5, r6, r7, lr}
 2: af03      add r7, sp, #12
 4: f84d bd04 str r11, [sp, #-4]!
 8: b082      sub sp, #8
 a: 4604      mov r4, r0
 c: 4668      mov r0, sp
 e: f7ff fffe bl 0xe
12: e9dd 0500 ldrd r0, r5, [sp]
16: b108      cbz r0, 0x1c
18: 2001      movs r0, #1
1a: e014      b 0x46
1c: 4668      mov r0, sp
1e: f7ff fffe bl 0x1e
22: e9dd 0600 ldrd r0, r6, [sp]
26: b110      cbz r0, 0x2e
28: 2001      movs r0, #1
2a: 4635      mov r5, r6
2c: e00b      b 0x46
2e: 4668      mov r0, sp

```

```

30: f7ff fffe bl 0x30
34: e9dd 0100 ldrd r0, r1, [sp]
38: b110      cbz r0, 0x40
3a: 2001      movs r0, #1
3c: 460d      mov r5, r1
3e: e002      b 0x46
40: 1970      adds r0, r6, r5
42: 1845      adds r5, r0, r1
44: 2000      movs r0, #0
46: e9c4 0500 strd r0, r5, [r4]
4a: b002      add sp, #8
4c: f85d bb04 ldr r11, [sp], #4
50: bdf0      pop {r4, r5, r6, r7, pc}

```

`objdump --disassemble --demangle
rust.o > rust.S`

Cost of **6 bytes** per function call

Using **Rust** (u64)

```
rustc  
--target thumbv7em-none-eabi  
-C opt-level="s"  
--emit=obj  
rust.rs
```

```
#[inline(never)]  
pub fn combined() -> Result<i32, u64> {  
  
    let mut num: i32 = 0;  
  
    num += func1()?;
    num += func2()?;
    num += func3()?;
    Ok(num)  
}  
  
#[inline(never)]  
pub fn func1() -> Result<i32, u64> { Ok(5) }  
#[inline(never)]  
pub fn func2() -> Result<i32, u64> { Err(8800) }  
#[inline(never)]  
pub fn func3() -> Result<i32, u64> { Ok(13) }
```

Return Rust (u64): Disassembled

```

00000000 <rust::combined::hece36e1d2889c9fd>:
    0: b5f0      push   {r4, r5, r6, r7, lr}
    2: af03      add    r7, sp, #12
    4: f84d bd04 str    r11, [sp, #-4]!
    8: b084      sub    sp, #16
    a: 4604      mov    r4, r0
    c: 4668      mov    r0, sp
    e: f7ff fffe bl    0xe
   12: 9800      ldr    r0, [sp]
   14: b958      cbnz  r0, 0x2e
   16: 4668      mov    r0, sp
   18: 9d01      ldr    r5, [sp, #4]
   1a: f7ff fffe bl    0x1a
   1e: 9800      ldr    r0, [sp]
   20: b928      cbnz  r0, 0x2e
   22: 4668      mov    r0, sp
   24: 9e01      ldr    r6, [sp, #4]
   26: f7ff fffe bl    0x26
   2a: 9800      ldr    r0, [sp]
   2c: b148      cbz   r0, 0x42
   2e: 9903      ldr    r1, [sp, #12]
   30: 9802      ldr    r0, [sp, #8]
   32: e9c4 0102 strd   r0, r1, [r4, #8]

```

```

   36: 2001      movs   r0, #1
   38: 6020      str    r0, [r4]
   3a: b004      add    sp, #16
   3c: f85d bb04 ldr    r11, [sp], #4
   40: bdf0      pop    {r4, r5, r6, r7, pc}
   42: 9901      ldr    r1, [sp, #4]
   44: 1970      adds   r0, r6, r5
   46: 4408      add    r0, r1
   48: 6060      str    r0, [r4, #4]
   4a: 2000      movs   r0, #0
   4c: e7f4      b     0x38

```

`objdump --disassemble --demangle
rust.o > rust.S`

Cost of **4 bytes** per function call

Distributed vs Centralized Error Handling

Distributed Error Handling

- A form of error handling where the work & code for propagating error information up the stack to a handler is distributed across all functions that participate in error handling.
 - If a function wants to call a function that can emit an error, it is the calling functions job to either:
 - Handle the error
- OR
- Propagate (return) the error
 - Using a return type or an out parameter as your error handling is an example of this.

Centralized Error Handling

- A form of error handling where a singular mechanism along with some metadata about the system is capable of propagating error information from the detection site to the appropriate handler.
- Calling code does not have to perform any work during normal operation in order to participate in error handling.
- Normal code & error handling code are in separate domains.
- Exception handling is a form of centralized error handling

My Claim:

Centralized Error handling schemes have an advantage in terms of code size compared to distributed error handling.

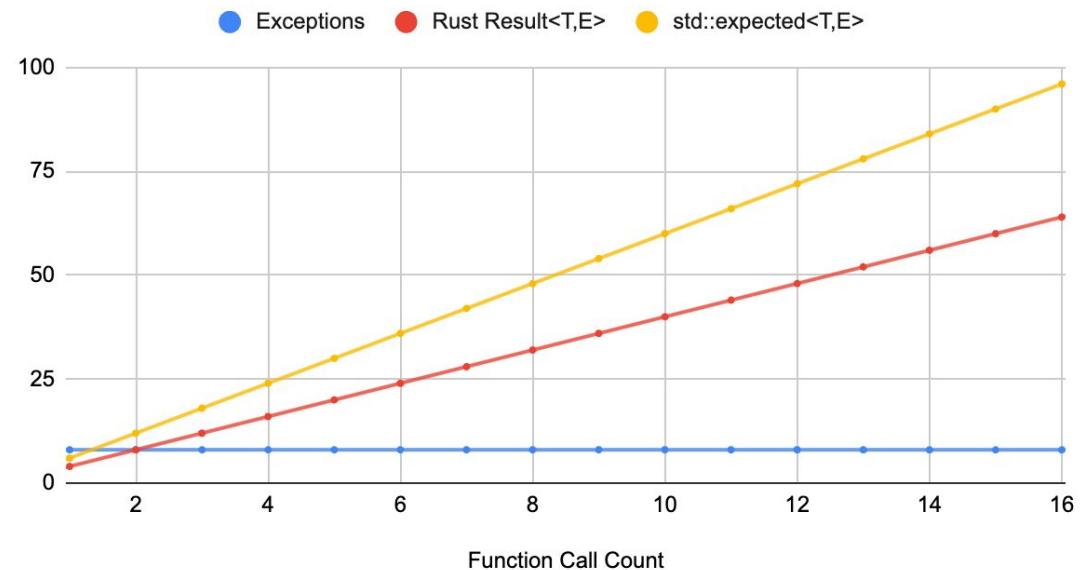
Point 1:
Branching isn't FREE!

Point 1: Branching isn't free!

Consider the following function using any distributed type error mechanism...

```
uint32_t func1();  
uint32_t func2();  
uint32_t func3();  
uint32_t func4();  
uint32_t func5();  
  
uint32_t combined_function() {  
    return func1() + func2() +  
           func3() + func4() +  
           func5();  
}
```

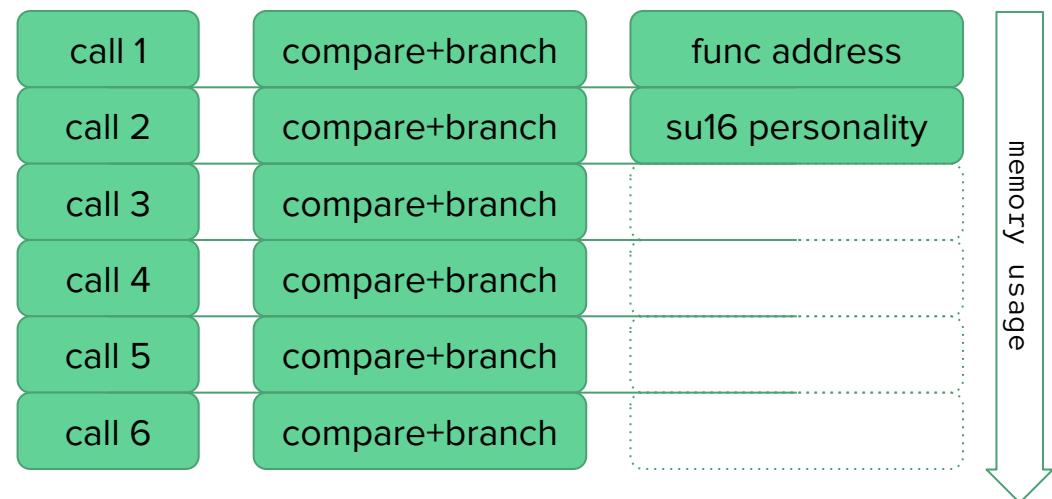
Exceptions, Rust Result<T,E> and std::expected<T,E>



Point 1: Branching isn't free!

Consider the following function using any return type error mechanism...

```
uint32_t func1();  
uint32_t func2();  
uint32_t func3();  
uint32_t func4();  
uint32_t func5();  
uint32_t func6();  
  
uint32_t combined_function() {  
    return func1() + func2() + func3() +  
        func4() + func5() + func6();  
}
```



Point 1: Branching isn't free!

```
[nodiscard] inline bool probe(i2c& p_i2c,  
    hal::byte p_address)  
{  
    std::array<std::byte, 1> data_in{};  
    bool device_acknowledged = true;  
  
    try {  
        p_i2c.transaction(p_address, {}, data_in);  
    } catch (const hal::no_such_device&) {  
        device_acknowledged = false;  
    }  
    return device_acknowledged;  
}
```

Consider this function from before.
Adds size is **47 bytes**

How many checked functions would
cost the same?

About **11 to 12 checked** function calls,
ignoring the additional return paths.

Point 2:

Most functions are trivially
unwindable

Point 2: Most functions are trivially unwindable

1. What types of functions do you write?
2. What kinds of functions do you normally see/read?
3. How many functions in your code base **handle errors** vs **propagating it up?**
4. How much of your functions create objects with **non-trivial destructors**?

Point 2: Most functions are trivially unwindable

Consider **Audacity**:

- **1054** C++ files **only 39 files** (or 3.7% of files) have a catch block.
- Total of **60** try scopes and **69** catch blocks were found in the source code.
- Out of **177433** functions **50** functions (0.028%) had a try/catch blocks
- Exception sections add **+8%** memory relative to the **.text** section (15,415 kB)

Point 2: Most functions are trivially unwindable

Let's consider code with this distribution of functions:

1. **~5%** of functions catch errors
2. **~25%** allocate objects with non-trivial destructors
3. **~55%** of functions are trivially unwindable
4. **~10%** of functions is leaf functions

NOTE:**C code is always leaf/trivial**

Point 2: Most functions are trivially unwindable

So if we had some code with 1,000 functions:

1. 50 try/catch functions require +47B → lets use **64 bytes**
2. 250 cleanup functions require +34B → lets use **46 bytes**
3. 550 trivial functions require 8 to 24 → lets use **8 bytes** (LU should be rare)
4. 100 leaf functions use **0 bytes** of exception info

NOTE:**C code is always leaf/trivial**

Point 2: Most functions are trivially unwindable

So if we had some code with 1,000 functions:

1. 50 try/catch functions require +47B → lets use **64 bytes**
2. 250 cleanup functions require +34B → lets use **46 bytes**
3. 550 trivial functions require 8 to 24 → lets use **8 bytes** (LU should be rare)
4. 100 leaf functions use **0 bytes** of exception info

$$8000 + 50*64 + 250*46 + 550*8 + 100*0 = \mathbf{27,100\ bytes}$$

Point 2: Most functions are trivially unwindable

So if we had some code with 1,000 functions:

1. 50 try/catch functions require +47B → lets use **64 bytes**
2. 250 cleanup functions require +34B → lets use **46 bytes**
3. 550 trivial functions require 8 to 24 → lets use **8 bytes** (LU should be rare)
4. 100 leaf functions use **0 bytes** of exception info

$$8000 + 50*64 + 250*46 + 550*8 + 100*0 = \mathbf{27,100 \text{ bytes}}$$

$$27,100 \text{ bytes} * (1 \text{ check} / 4 \text{ bytes}) = \mathbf{6,775 \text{ checked function calls}}$$

Point 2: Most functions are trivially unwindable

So if we had some code with 1,000 functions:

1. 50 try/catch functions require +47B → lets use **64 bytes**
2. 250 cleanup functions require +34B → lets use **46 bytes**
3. 550 trivial functions require 8 to 24 → lets use **8 bytes** (LU should be rare)
4. 100 leaf functions use **0 bytes** of exception info

$$8000 + 50*64 + 250*46 + 550*8 + 100*0 = \mathbf{27,100 \text{ bytes}}$$

$$27,100 \text{ bytes} * (1 \text{ check} / 4 \text{ bytes}) = \mathbf{6,775 \text{ checked function calls}}$$

$$6,775 \text{ checked} / 900 \text{ functions} = \mathbf{\textcolor{red}{\sim 7.5 \text{ checked function calls per function}}}$$

NOTE: Should be less, because we aren't considering additional epilogs

Point 2: Most functions are trivially unwindable

So if we had some code with 1,000 functions:

1. 50 try/catch functions require +47B → lets use **64 bytes**
2. 250 cleanup functions require +34B → lets use **46 bytes**
3. 550 trivial functions require 8 to 24 → lets use **8 bytes** (LU should be rare)
4. 100 leaf functions use **0 bytes** of exception info

$$8000 + 50*64 + 250*46 + 550*8 + 100*0 = \mathbf{27,100 \text{ bytes}}$$

$$27,100 \text{ bytes} * (1 \text{ check} / 6 \text{ bytes}) = \mathbf{4,516 \text{ checked function calls}}$$

$$4,516 \text{ checked} / 900 \text{ functions} = \mathbf{\sim 5.0 \text{ checked function calls per function}}$$

NOTE: Should be less, because we aren't considering additional epilogs

Point 2: Most functions are trivially unwindable

So if we had some code with 1,000 functions:

1. 50 try/catch functions require +47B → lets use **64 bytes**
2. 250 cleanup functions require +34B → lets use **46 bytes**
3. 550 trivial functions require 8 to 24 → lets use **8 bytes** (LU should be rare)
4. 100 leaf functions use **0 bytes** of exception info

$$8000 + 50*64 + 250*46 + 550*8 + 100*0 = \mathbf{27,100 \text{ bytes}}$$

$$27,100 \text{ bytes} * (1 \text{ check} / 8 \text{ bytes}) = \mathbf{3387.5 \text{ checked function calls}}$$

$$3387.5 \text{ checked} / 900 \text{ functions} = \mathbf{\textcolor{red}{3.76 \text{ checked function calls per function}}}$$

NOTE: Assuming the average if 8-bytes per check w/ return path

How does C++ exceptions enable smaller binaries?

- Reduced codegen when calling fallible functions.
- The unwind information **scales** with the **number of functions you have** NOT the **number of functions you call**.
- Eliminates additional return paths in functions.
- The unwind information is a compressed representation of normal instructions.
- You only pay for the exception mechanism code once.

THIS ISN'T EVEN ITS
FINAL FORM!

Exceptions can be smaller!!

Why not do this?

`__gcc_personality_v0`

Dwarf info	End of type table	Length of Call Site Table
0xFF (omit) 1B	Type table encoding 1B	offset _{end} (N... B)



`__new_personality`

End of type table	Length of Call Site Table
offset _{end} (N... B)	size _{call site} (N... B)

Type table encoding = prel
Call site encoding = uleb128

THIS IS BACKWARDS
COMPATIBLE!

Got old code? Why not use a tool to convert?

`__gcc_personality_v0`

Dwarf info	End of type table	Length of Call Site Table
0xFF (omit) 1B	Type table encoding 1B	offset _{end} (N... B)



`__new_personality?`

End of type table	Length of Call Site Table
offset _{end} (N... B)	size _{call site} (N... B)

Saves 3 bytes for each reentrant function

Why not do this?

```
catch_chain():

    push    {lr}
    bl      foo()
    pop    {pc}
    cmp    r1, #1
    bne   .L3
    bl     __cxa_begin_catch
    ldr   r0, .L13
    bl     printf
    bl     __cxa_end_catch

.L3:
    bl     __cxa_begin_catch
    ldr   r0, .L13+4
    bl     printf

.L4:
    bl     __cxa_end_catch
    pop    {lr}
```



Saves 2B for
each
attached
catch.

```
catch_chain():

    push    {lr}
    bl      foo()
    pop    {pc}
    cmp    r1, #1
    bne   .L3
    bl     __cxa_begin_catch
    ldr   r0, .L13
    bl     printf
    b     .L4

.L3:
    bl     __cxa_begin_catch
    ldr   r0, .L13+4
    bl     printf

.L4:
    bl     __cxa_end_catch
    pop    {lr}
```

Exceptions can be so much more than they are now

What are C++ exceptions really?

What are C++ exceptions really?
They are a **code compression**

```

00000000 <rust::combined::hece36e1d2889c9fd>:
 0: b5f0      push {r4, r5, r6, r7, lr}
 2: af03      add r7, sp, #12
 4: f84d bd04 str r11, [sp, #-4]!
 8: b084      sub sp, #16
 a: 4604      mov r4, r0
 c: 4668      mov r0, sp
 e: f7ff fffe bl 0xe
12: 9800      ldr r0, [sp]
14: b958      cbnz r0, 0x2e
16: 4668      mov r0, sp
18: 9d01      ldr r5, [sp, #4]
1a: f7ff fffe bl 0x1a
1e: 9800      ldr r0, [sp]
20: b928      cbnz r0, 0x2e
22: 4668      mov r0, sp
24: 9e01      ldr r6, [sp, #4]
26: f7ff fffe bl 0x26
2a: 9800      ldr r0, [sp]
2c: b148      cbz r0, 0x42
2e: 9903      ldr r1, [sp, #12]
30: 9802      ldr r0, [sp, #8]
32: e9c4 0102 strd r0, r1, [r4, #8]

```



```

36: 2001      movs r0, #1
38: 6020      str r0, [r4]
3a: b004      add sp, #16
3c: f85d bb04 ldr r11, [sp], #4
40: bdf0      pop {r4, r5, r6, r7, pc}
42: 9901      ldr r1, [sp, #4]
44: 1970      adds r0, r6, r5
46: 4408      add r0, r1
48: 6060      str r0, [r4, #4]
4a: 2000      movs r0, #0
4c: e7f4      b 0x38

```

+8 bytes

```

000004c0 <none ()>:
 4c0:   b508      push {r3, lr}
 4c2:   f7ff ffff7 bl 4b4 <func1 ()>
 4c6:   4603      mov r3, r0
 4c8:   f7ff ffff6 bl 4b8 <func2 ()>
 4cc:   4403      add r3, r0
 4ce:   f7ff ffff5 bl 4bc <func3 ()>
 4d2:   4418      add r0, r3
 4d4:   bd08      pop {r3, pc}
 4d6:   bf00      nop

```

So why do firmware developers avoid C++ exceptions?

SPACE

Requires Dynamic Memory/Heap

Requires Malloc

Increases Binary Size

Requires whole
~~C++ STL~~

Run Time Type Info
(RTTI)

Unbounded Memory Usage

Exception Tables

Exception Code

TIME

Nondeterministic

Type Comparison (dynamic_cast)



Slow

Binary Search

Frame Unwinding

Frame Evaluation

wHaT aBoUt tHe
SaD pATH pErFoRmAnCe!?

Cutting C++ Exception time by ~~+76%~~ +88%

By: Khalil Estell

Some performance numbers

Function Depth	std::expected <u32, u32>	GCC & libunwind	My implementation
6	556	15344	2652
Multiplier for depth 6	1x	27x slower	4.77x slower
96	8566	184454	22102
Multiplier for depth 96	1x	21.53x slower	2.58x slower

But let's be honest with
ourselves...

EXCEPTIONS

F*****ing

SUCK!

What does **foo** throw???

```
void foo();
```

What does `foo` throw???

```
void foo();
```

Docs say, "throws `some_error`"

But does it actually ONLY throw that?

What does `foo` throw???

```
void foo();
```

Docs say, "throws `some_error`"

But does it actually ONLY throw that?

For version 5.1.0 this is true.

What does `foo` throw???

```
void foo();
```

Docs say, "throws `some_error`"

But does it actually ONLY throw that?

For version 5.1.0 this is true.

```
self.requires("foo-lib/[^5.1.0]")
self.requires("foo-lib/[^5.1.1]")
```

YOU DON'T KNOW!

THE **ULTIMATE** PROBLEM WITH EXCEPTIONS!

Luckily it's not impossible
to know

Exception Insights Tool

By: Khalil Estell

Special Thanks

C++'s evolution priorities

Historical

add things
fix things
simplify

A future worth considering?

CppCon
The C++ Conference
2019
cppcon.org

De-fragmenting C++: Making Exceptions and RTTI More Affordable and Usable -
Herb Sutter CppCon 2019

87K views • 4 years ago

CppCon

A fundamental reason why C++ is successful and loved is its adherence to Stroustrup's zero-overhead principle: You don't pay for ...

9 moments C++'s evolution priorities | Part 1: Exception Handling (EH) | Code review... | Pathology 101 | Core... ▾

Video Sponsorship Provided By:
ansa 1:33:12

Doc. No. P1947

Date: 2019-11-18

Audience: EWG and LEWG

Reply to: Bjarne Stroustrup (bs@ms.com)

C++ exceptions and alternatives

Bjarne Stroustrup

There has been little serious research on the topics of performance of exceptions and reliability of the resulting code in C++.

✨ Supporting
this work



Github sponsorship page
For inquiry: estell.exceptions@gmail.com

Questions?

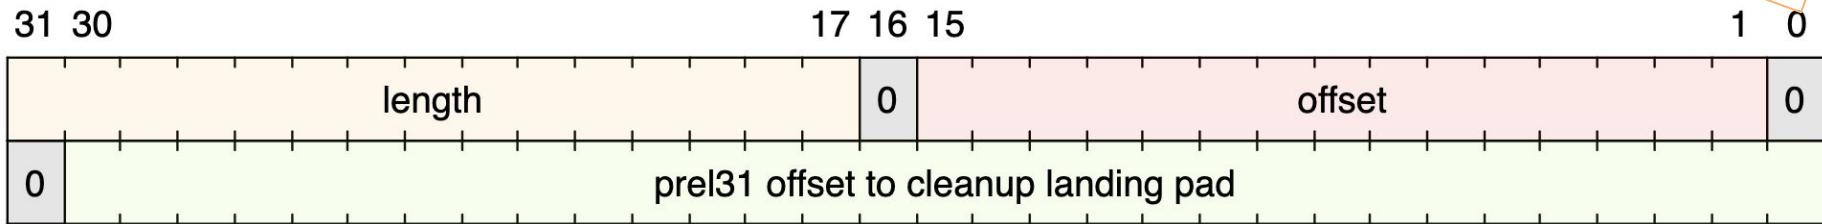


Github sponsorship page
For inquiry: estell.exceptions@gmail.com

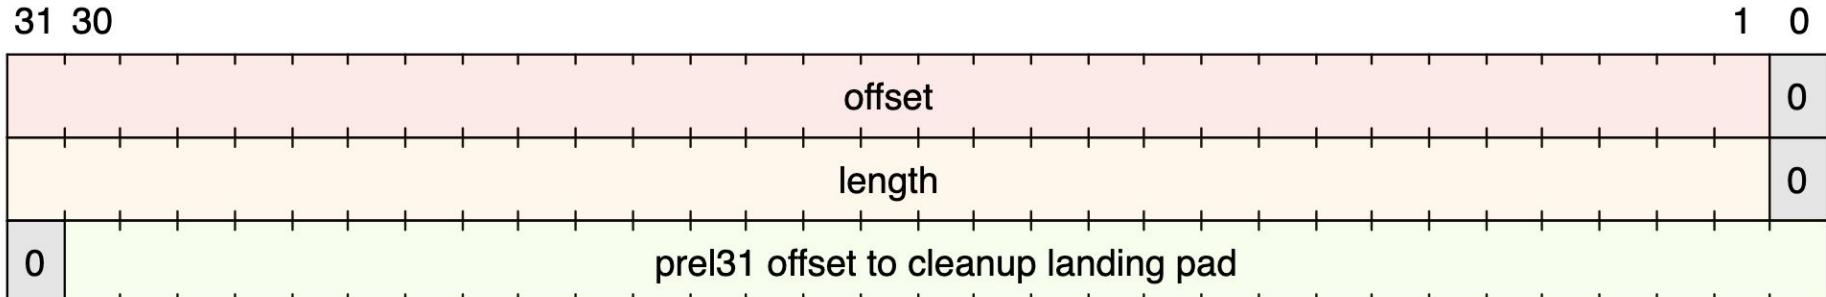
ARM Personality Data: cleanup descriptors

**NOT
IMPORTANT!**

16-bit cleanup descriptor



32-bit cleanup descriptor



ARM Personality Data: catch descriptors

**NOT
IMPORTANT!**

Catch by reference descriptor

31 30

17 16 15

1 0

length

1

offset

0

1

prel31 offset to catch landing pad

std::type_info*

Catch by "anything else"

31 30

17 16 15

1 0

length

1

offset

0

0

prel31 offset to catch landing pad

std::type_info*

ARM Personality Data: catch descriptors

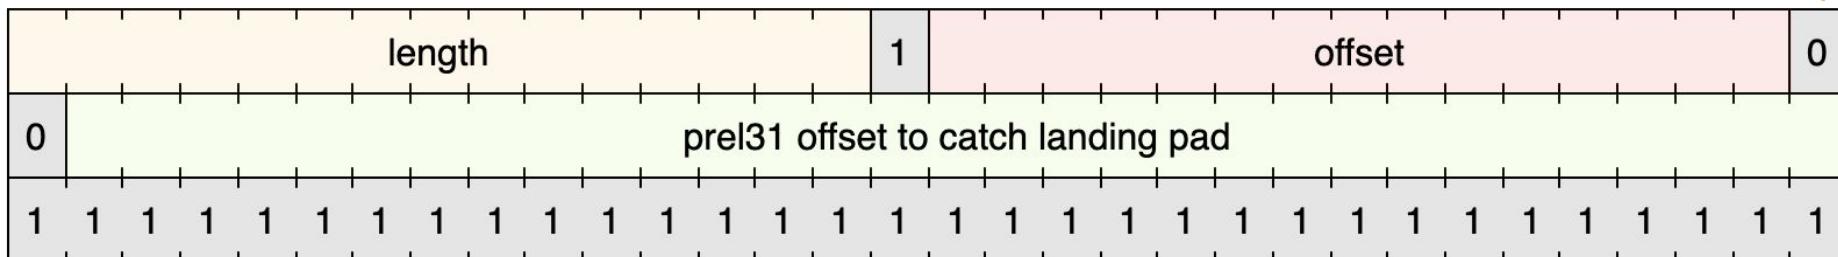
NOT IMPORTANT!

Catch All

31 30

17 16 15

10

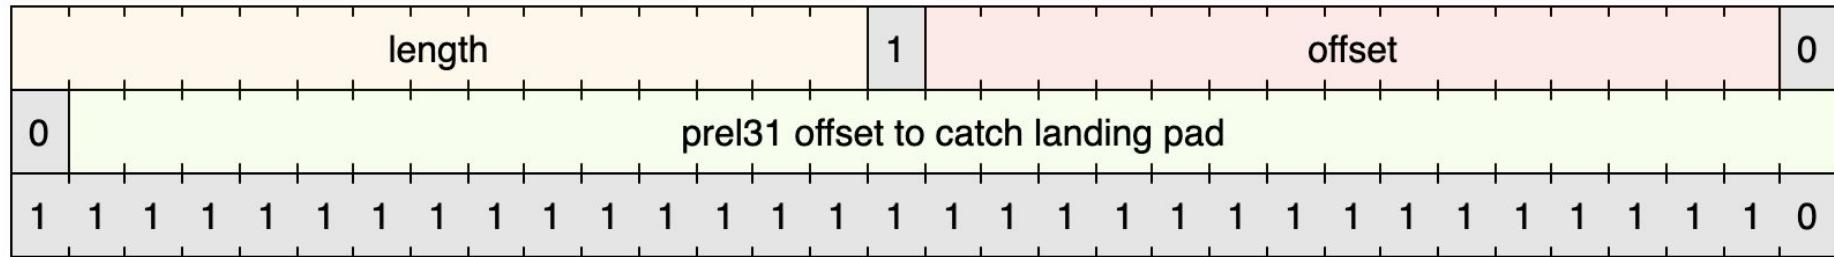


Catch All & terminate

31 30

17 16 15

1 0



ARM Personality Data: Example Complete data

data
**NOT
IMPORTANT**

EOD (End of Descriptors)