

+ 24

Interesting Upcoming Features from Low Latency, Parallelism and Concurrency

From Kona 2023, Tokyo 2024, and St. Louis 2024

PAUL E. MCKENNEY,
MAGED MICHAEL
& MICHAEL WONG



20
24



Agenda

1. Improving C++20 Atomic Min/Max(P0493; Michael)
2. Hazard pointer extensions (P3135; Maged)
3. Pointer tagging (P3125; Maged)
4. Parallel Range algorithms (P3179; Michael), may be Parallel Algorithms (P2500)

C++26: Atomic Min/Max

C++26: Improving C++20 Concurrency primitives

atomic min/max ([P0493](#))

Atomic min/max motivation (P0493)

Long history - almost as old as atomic addition

Multithreaded applications often involve scenarios where multiple threads need to concurrently update a shared variable to track the minimum or maximum value.

Without atomic operations, race conditions can occur, leading to data corruption and unpredictable behavior.

Atomic min/max operations provide a solution by ensuring that updates to the shared variable are performed atomically, guaranteeing data integrity.

Efficient and safe concurrent updates to shared variables.

Enable the implementation of lock-free data structures, leading to improved performance and scalability.

Useful in various applications, including reductions in data-parallel algorithms, statistics collection, and optimization processes.

Useful for:

- Lock-free data structures
- Parallel reductions (OpenMP)
- Optimization algorithms
- Statistics collection

Proposed interface

```
namespace std {
```

```
template<class T>  
T atomic_fetch_max(atomic<T>*,  
                  typename atomic<T>::value_type) noexcept;
```

```
template<class T>  
T atomic_fetch_max_explicit(atomic<T>*,  
                           typename atomic<T>::value_type,  
                           memory_order) noexcept;
```

```
// Similar for atomic_fetch_min  
// Member functions on atomic<T>
```


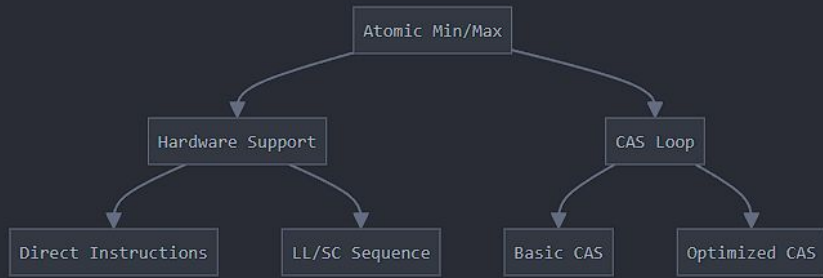
```
}
```

Performance Considerations & Benchmark Results

Hardware support can be significantly faster

CAS loops have higher contention

But optimized CAS can be competitive



Example Usage

```
#include <atomic>
#include <thread>
#include <vector>
```

```
std::atomic<int> max_value{0};
```


```
void find_max_in_range(int
start, int end) {
    for (int i = start; i <
end; ++i) {
```

```
        max_value.fetch_max(i,
std::memory_order_relaxed)
    ;
}
```

```
int main() {
    std::vector<std::thread> threads;
    for (int i = 0; i < 10; ++i) {
```

```
        threads.emplace_back(find_max_in_range
, i * 100, (i + 1) * 100);
    }
    for (auto& thread : threads) {
        thread.join();
    }
    std::cout << "Maximum value: " <<
max_value << std::endl;
    return 0;
}
```


Memory Ordering



`std::memory_order_relaxed`: No synchronization or ordering constraints.

`std::memory_order_acquire`: Ensures that subsequent loads see the effects of this operation and any prior release operations.

`std::memory_order_release`: Makes the effects of this operation and any prior operations visible to subsequent acquire operations.

`std::memory_order_acq_rel`: Combines acquire and release semantics.

`std::memory_order_seq_cst`: Provides the strongest ordering guarantees, ensuring a globally consistent view of memory across all threads.

Current Status, Next Steps, & Conclusion

Proposal in flight since 2016, aiming for C++26

Current revision: P0493R5 (February 2024)

Implementation experience in Clang

Open questions:

- Floating point support?
- New-value-returning variants?

- Atomic min/max operations are valuable additions to the C++ standard library.
- They enable efficient and safe concurrent updates to shared variables, facilitating the development of high-performance and scalable multithreaded applications.
- By understanding the semantics and proper usage of atomic min/max, you can leverage their power to build robust and reliable concurrent code.

Atomic Floating-Point Min/Max in C++ (P3008)

Atomic min/max operations are crucial for concurrent algorithms

Integer versions added in C++20

Floating-point versions were removed due to concerns

The Challenge of Floating-Point Corner Cases

- NaN (Not a Number): Represents undefined or unrepresentable values.
- Signed Zero: Both +0 and -0 exist, but their ordering can be ambiguous

Evolution of Floating-Point Min/Max in C++

- `std::min`, `std::max`: Undefined behavior with NaNs.
 - a. `std::min(-0.0f, +0.0f);` // Returns -0.0f
 - b. `std::min(NaN, 2.0f);` // Undefined behavior
- `fmin`, `fmax`: Treat NaNs as missing data, signed zero behavior can vary.
- `fminimum`, `fmaximum`: Treat NaNs as errors, -0 < +0.
- `fminimum_num`, `fmaximum_num`: Treat NaNs as missing data, -0 < +0.

Hardware Support

- Many GPUs have native atomic float min/max
- CPU support growing (e.g. ARM v8.1)
- Hardware treats -0 < +0 and NaN as "missing data"

Atomic Floating-Point Min/Max

`atomic<T>::fetch_min`,
`atomic<T>::fetch_max`: Proposed atomic operations.

```
float fetch_min(float, memory_order =  
memory_order::seq_cst) noexcept;
```

```
float fetch_max(float, memory_order =  
memory_order::seq_cst) noexcept;
```

Semantics:

- Unspecified behavior with NaNs.
- Treats $-0 < +0$ (recommended, not required)
- Based on C23's `fminimum_num`/`fmaximum_num`

```
float fetch_fminimum(float, memory_order =  
memory_order::seq_cst) noexcept;
```

```
float fetch_fmaximum(float, memory_order  
= memory_order::seq_cst) noexcept;
```

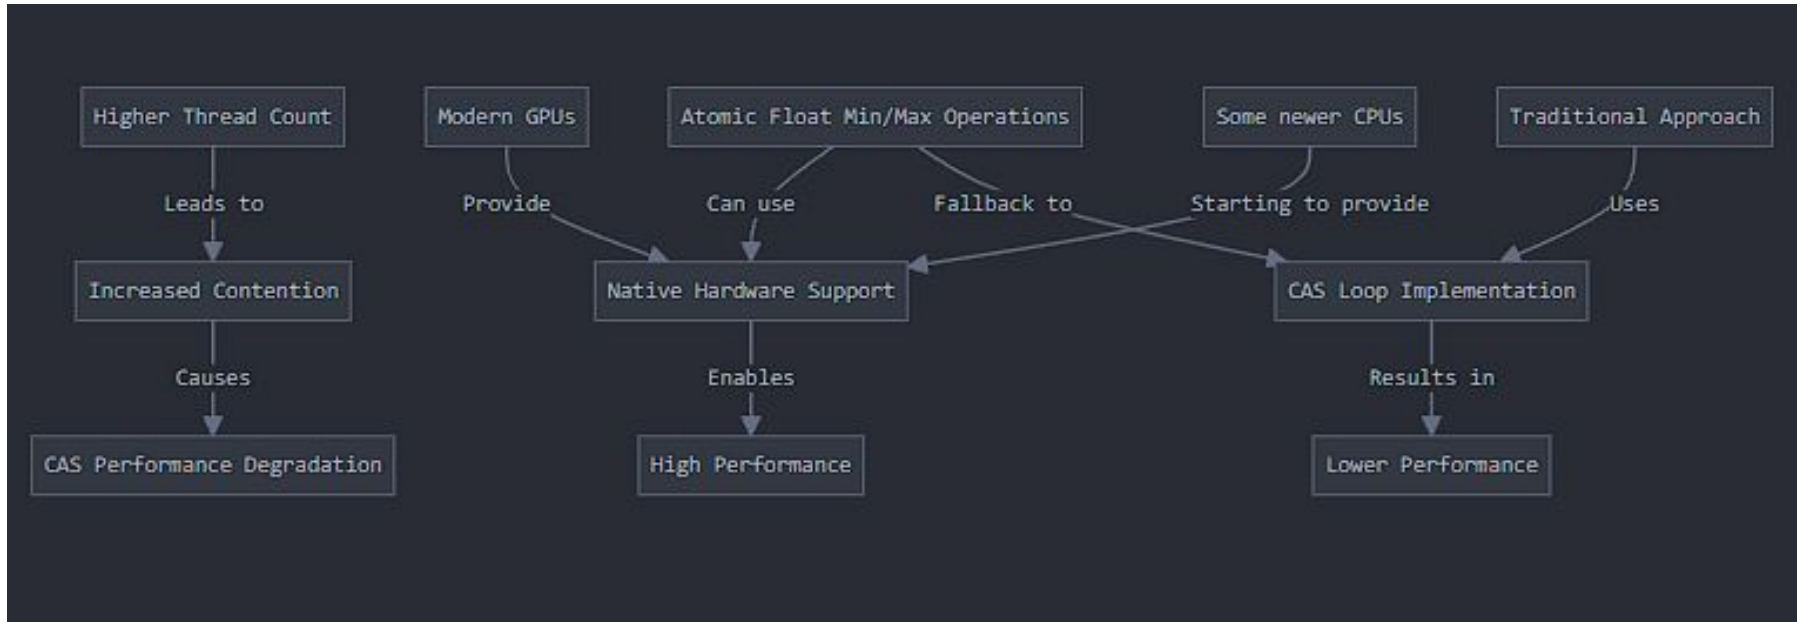
```
float fetch_fminimum_num(float,  
memory_order = memory_order::seq_cst)  
noexcept;
```

```
float fetch_fmaximum_num(float,  
memory_order = memory_order::seq_cst)  
noexcept;
```

Performance Considerations

Native atomic operations vs. CAS loops.

Significant performance difference in highly concurrent systems.



Code Examples

```
#include <atomic>
#include <cmath>

#include <iostream>

int main() {

    std::atomic<float> atomic_value(10.0f);

    float new_max = 15.0f;

    float result_max = atomic_value.fetch_max(new_max);

    std::cout << "Old max: " << result_max << ", New max: "
    << atomic_value << std::endl;

    float new_min = 5.0f;

    float result_min = atomic_value.fetch_min(new_min);

    std::cout << "Old min: " << result_min << ", New min: "
    << atomic_value << std::endl;

    return 0;

}
```

Expected output:

```
Old max: 10, New max: 15
```

```
Old min: 15, New min: 15
```

Conclusion

Atomic floating-point min/max: Essential for modern concurrent systems

Careful consideration of corner cases: NaNs and signed zeros

Performance implications: Native atomic operations vs. CAS loops

New atomic float min/max operations

Semantics close to hardware capabilities

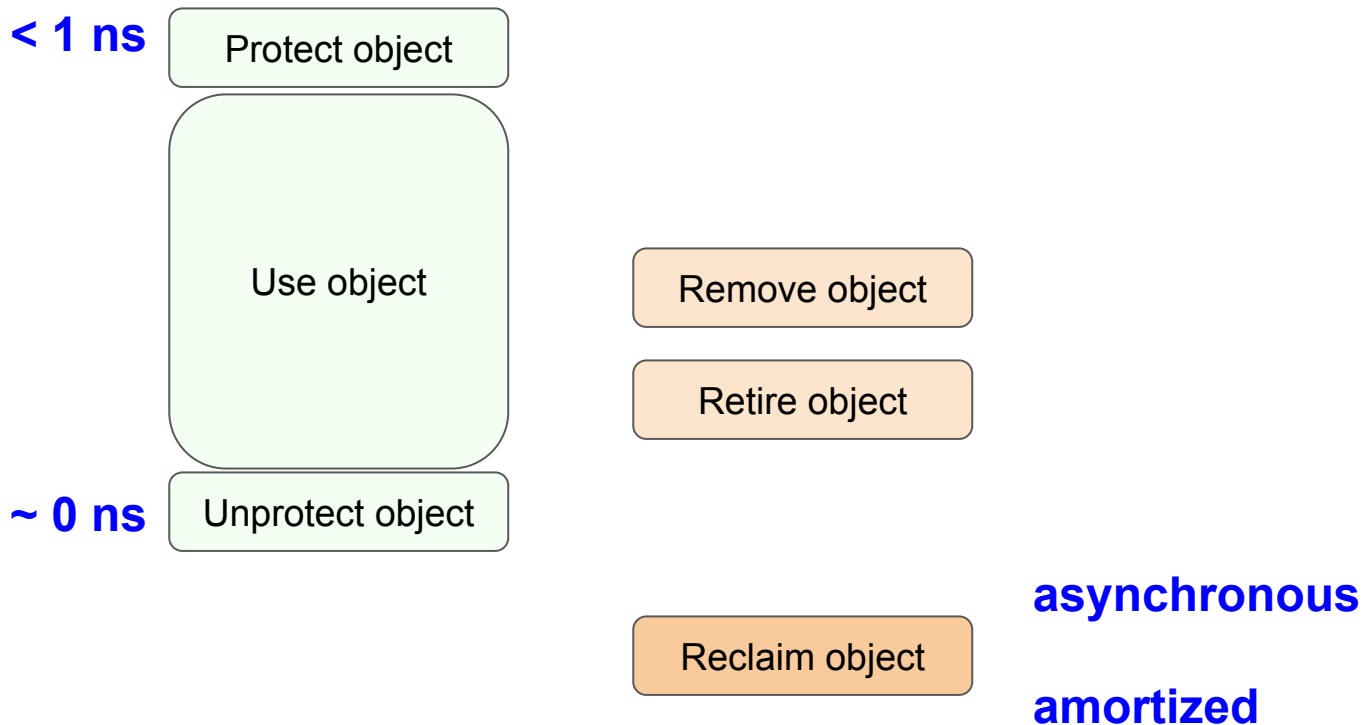
Significant performance benefits

Additional functions for more specific needs

Hazard Pointer Extensions beyond C++26

Hazard Pointers in C++26 -- Background

Hazard pointers **protect** dynamic **objects** from being **reclaimed**, allowing **safe access** to protected objects **without** additional **synchronization**.



Hazard Pointers C++26

```
template <typename T, typename D = default_delete<T>>
class hazard_pointer_obj_base {
    void retire(D d = D()) noexcept;
};
```

```
class hazard_pointer {
    hazard_pointer() noexcept; // construct empty hazard_pointer
    hazard_pointer(hazard_pointer&&) noexcept;
    hazard_pointer& operator=(hazard_pointer&&) noexcept;
    ~hazard_pointer();
    [[nodiscard]] bool empty() const noexcept;
    template <typename T> T* protect(const atomic<T*>& src) noexcept;
    template <typename T> bool try_protect(T*& ptr, const atomic<T*>& src) noexcept;
    template <typename T> void reset_protection(const T* ptr) noexcept;
    void reset_protection(nullptr_t = nullptr) noexcept;
    void swap(hazard_pointer&) noexcept;
};
```

```
hazard_pointer make_hazard_pointer(); // construct nonempty hazard_pointer
void swap(hazard_pointer&, hazard_pointer&) noexcept;
```

Example Using C++26 Hazard Pointers

```
class T : public hazard_pointer_obj_base<T> { /* T members */ };
```

```
std::atomic<T*> src_;
```

```
U readAndAccess(Func userFn) { // Called frequently
    hazard_pointer hp = make_hazard_pointer(); // Construct hazard pointer.
    T* ptr = hp.protect(src_);                // Get pointer to a protected object.
    return userFn(ptr);
}
```

```
Void update(T* newptr) { // Called infrequently
    T* oldptr = src_.exchange(newptr);
    oldptr->retire(); // Pass to hazard pointer library for safe reclamation.
}
```

P3135R1: Hazard Pointer Extensions (beyond C++26)

P3135R1: Hazard Pointer Extensions (wg21.link/p3135r1)

No need for extending C++26:

- **Protection Counting** (can be a topic for a future talk)
- **Execution of Asynchronous Reclamation**

Proposed standard extensions:

- **Synchronous reclamation**
- **Batch creation and destruction**

Hazard Pointer Execution of Asynchronous Reclamation

Hazard Pointer Execution of Asynchronous Reclamation

Using C++26

Possible inline asynchronous reclamation

```
void worker() {  
    /* ... */  
    obj->retire();  
    // Possible inline reclamation.  
}
```

Using a separate asynchronous reclamation executor

```
void worker() {  
    /* ... */  
    ex_.submit([obj] { obj->retire(); });  
    // No inline reclamation.  
}
```

Hazard Pointer Synchronous Reclamation

Hazard Pointer Synchronous Reclamation

C++26 (Asynchronous Reclamation Only)

```
template <class T> class Container {
    class Obj : hazard_pointer_obj_base<Obj>
    { T data; /* etc */ };
    void insert(T data) {
        Obj* obj = new Obj(data); /* etc */ }
    void erase(Args args) {
        Obj* obj = find(args);
        /* Remove obj from container */
        obj->retire();
    }
};

class A {
    // Deleter does not depend on resources
    // with independent lifetime.
    ~A();
};

{ Container<A> container;
  container.insert(a);
  container.erase(a); }
// Obj containing 'a' may be not deleted yet.
```

OK

Need Synchronous Reclamation

```
template <class T> class Container {
    class Obj : hazard_pointer_obj_base<Obj>
    { T data; /* etc */ };
    void insert(T data) {
        Obj* obj = new Obj(data); /* etc */ }
    void erase(Args args) {
        Obj* obj = find(args);
        /* Remove obj from container */
        obj->retire();
    }
};

class B {
    // Deleter may depend on resources
    // with independent lifetime.
    ~B() { use_resource_XYZ(); }
};

make_resource_XYZ();
{ Container<B> container;
  container.insert(b);
  container.erase(b); }
// Obj containing 'b' may be not deleted yet.
destroy_resource_XYZ();
```

ERROR

Hazard Pointer Synchronous Reclamation

Cohorts:

- **Folly:** folly/synchronization/Hazptr.h
- **CPPCON 2021:** *Hazard Pointer Synchronous Reclamation*

Possible API

```
class hazard_pointer_cohort {
    hazard_pointer_cohort() noexcept;
    hazard_pointer_cohort(const hazard_pointer_cohort&) = delete;
    hazard_pointer_cohort(hazard_pointer_cohort&&) = delete;
    hazard_pointer_cohort& operator=(const hazard_pointer_cohort&) = delete;
    hazard_pointer_cohort& operator=(hazard_pointer_cohort&&) = delete;
    ~hazard_pointer_cohort();
};

template <class T, class D = default_delete<T>>
class hazard_pointer_obj_base {
    void retire_to_cohort(hazard_pointer_cohort&, D d = D()) noexcept;
};

void asynchronous_reclamation() noexcept;
```

synchronous

asynchronous

Hazard Pointer Synchronous Reclamation

C++26 (Asynchronous Reclamation Only)

```
template <class T> class Container {
    class Obj : hazard_pointer_obj_base<Obj>
    { T data; /* etc */ };

    void insert(T data) {
        Obj* obj = new Obj(data); /* etc */ }
    void erase(Args args) {
        Obj* obj = find(args);
        /* Remove obj from container */
        obj->retire();
    }
};

class A {
    // Deleter does not depend on resources
    // with independent lifetime.
    ~A();
};

{ Container<A> container;
  container.insert(a);
  container.erase(a); }
// Obj containing 'a' may be not deleted yet.
```

OK

Synchronous Reclamation

```
template <class T> class Container {
    class Obj : hazard_pointer_obj_base<Obj>
    { T data; /* etc */ };
    hazard_pointer_cohort cohort_;
    void insert(T data) {
        Obj* obj = new Obj(data); /* etc */ }
    void erase(Args args) {
        Obj* obj = find(args);
        /* Remove obj from container */
        obj->retire_to_cohort(cohort_);
        ex_.submit([] {asynchronous_reclamation(); });
    }
};

class B {
    // Deleter may depend on resources
    // with independent lifetime.
    ~B() { use_resource_XYZ(); }
};

make_resource_XYZ();
{ Container<B> container;
  container.insert(b);
  container.erase(b); }
// Obj containing 'b' must be already deleted.
destroy_resource_XYZ();
```

OK

Batch Creation and Destruction

Hazard Pointer Batch Creation and Destruction

C++26 (one at a time)

```
{ hazard_pointer hp[3];  
  /* Three hazard pointers are made nonempty  
     separately. */  
  hp[0] = make_hazard_pointer();  
  hp[1] = make_hazard_pointer();  
  hp[2] = make_hazard_pointer();  
  assert(!hp[0].empty());  
  assert(!hp[1].empty());  
  assert(!hp[2].empty());  
  /* src is atomic<T*> */  
  T* ptr = hp[0].protect(src);  
  /* etc */  
} /* Three nonempty hazard pointers are  
   destroyed separately. */
```

e.g., ~6 ns

Batch creation and destruction

```
{ hazard_pointer hp[3];  
  /* Three hazard pointers are made nonempty  
     together. */  
  make_hazard_pointer_batch(std::span{hp});  
  SCOPE_EXIT {  
    destroy_hazard_pointer_batch(std::span{hp});  
  };  
  assert(!hp[0].empty());  
  assert(!hp[1].empty());  
  assert(!hp[2].empty());  
  /* src is atomic<T*> */  
  T* ptr = hp[0].protect(src);  
  /* etc */  
} /* Three nonempty hazard pointers are emptied  
   together, and then destroyed separately. */
```

e.g., ~2 ns

Possible API


```
void make_hazard_pointer_batch(std::span<hazard_pointer>);  
void destroy_hazard_pointer_batch(std::span<hazard_pointer>) noexcept;
```

Pointer Tagging

Pointer Tagging

P3125R0: Pointer Tagging (author: Hana Dusíková) (wg21.link/p3125r0)

Pointer to aligned object T with `alignof(T) = 4`:



Motivation (P3125R0):

"Pointer tagging is not allowed in standard C++ as manipulating pointer bits is UB. Because of this limitation, some advanced data structures are not implementable or sub-optimally implementable."

This technique is an existing practice, is widely used in the field, and standardising it would lower the bar for its safe usage among C++'s users."

P3125R0: Pointer Tagging API

```
template <typename T, size_t Alignment = alignof(T)> class tagged_pointer;
```

```
template <typename T, size_t Alignment = alignof(T)>  
constexpr auto tag_bit_mask() noexcept -> uintptr_t;
```

available bits for tagging

```
template <typename T, size_t Alignment = alignof(T)>  
constexpr auto tag_pointer(T* original, uintptr_t tag) noexcept  
-> tagged_pointer<T, Alignment>;
```

pointer --> tagged_pointer

Precondition: tag == (tag & tag_bit_mask<T, Alignment>)

```
template <typename T, size_t Alignment = alignof(T)>  
constexpr auto untag_pointer(tagged_pointer<T, Alignment> ptr) noexcept -> T*;
```

tagged_pointer --> pointer

```
template <typename T, size_t Alignment = alignof(T)>  
constexpr auto tag_value(tagged_pointer<T, Alignment> ptr) noexcept -> uintptr_t;
```

tagged_pointer --> tag

Pointer Tagging Example

```
using TaggedPointer = tagged_pointer<T,2>;
```

```
bool try_tag_untagged_pointer(atomic<TaggedPointer>& src) {  
    TaggedPointer current = src.load();  
    assert(tag_value(current) == 0);           // Assert src was untagged  
    assert(1 & tag_bit_mask<T,2>() == 1);      // Assert that 1 is a valid tag  
    T* ptr = untag_pointer(current);           // Extract pointer  
    TaggedPointer newval = tag_pointer(ptr, 1); // Tag the pointer  
    return src.compare_exchange_weak(current, newval); // CAS src  
}
```


Extending Hazard Pointers to Tagged Pointers

C++26

```
template <typename T> T* protect(const atomic<T*>& src) noexcept;  
template <typename T> bool try_protect(T*& ptr, const atomic<T*>& src) noexcept;
```

Possible Extensions

```
template <typename T, size_t Alignment = alignof(T)>  
tagged_pointer<T, Alignment> protect(  
    const atomic<tagged_pointer<T, Alignment>>& src) noexcept;
```

```
template <typename T, size_t Alignment = alignof(T)>  
bool try_protect(  
    tagged_pointer<T, Alignment>& ptr,  
    const atomic<tagged_pointer<T, Alignment>>& src) noexcept;
```

Example

```
atomic<tagged_pointer<T>> src_;
```

```
hazard_pointer hp = make_hazard_pointer();  
tagged_pointer<T> tagged = hp.protect(src_);  
/* Safe to use ptr, where ptr == untag_pointer(tagged). */
```

Bringing parallelism to `std::ranges` algorithms

ISO C++ Parallelism/Concurrency Programming Language (based on Gonzalo's ISC C++ BoF)

Parallel Algorithms [many]


Concurrency++

Memory Model++ [MichaelW,
Maged, Paul M]

Forward Progress

Ranges++ many

Multi-dimensional Spans [Bryce, Christian]
operator[i, j, k]



Bringing parallelism to std::ranges algorithms (P3179)

// C++03

```
template<class RandomAccessIterator, class Compare>  
void sort(RandomAccessIterator first, RandomAccessIterator last, Compare comp);
```

// C++17

```
template<class ExecutionPolicy, class RandomAccessIterator, class Compare>  
void sort(ExecutionPolicy&& exec, RandomAccessIterator first, RandomAccessIterator last, Compare comp);
```

// C++20

```
template<random_access_range R, class Comp = ranges::less, class Proj = identity>  
requires sortable<iterator_t<R>, Comp, Proj>  
constexpr borrowed_iterator_t<R> ranges::sort(R&& r, Comp comp = {}, Proj proj = {});
```

// C++26? (<https://wg21.link/P3179>)

```
template<class ExecutionPolicy, random_access_range R, class Comp = ranges::less, class Proj = identity>  
requires sortable<iterator_t<R>, Comp, Proj>  
constexpr borrowed_iterator_t<R> ranges::sort(ExecutionPolicy&& exec, R&& r, Comp comp = {}, Proj proj = {});
```

C++ Parallel Range Algorithms: Unifying Parallelism and Ranges

Why Combine Parallelism with Ranges?

- Ranges offer a productive API with opportunities for optimization.
- Users are already using ranges with non-range parallel algorithms; integrating execution policies simplifies and streamlines code.

The Power of Ranges and Parallelism

- The C++ Ranges library provides a powerful way to express and compose computations lazily.
- C++17 introduced parallel algorithms, but they don't integrate seamlessly with ranges.
- This paper proposes adding parallel algorithms that work directly with ranges, combining the benefits of both worlds.

The Need for Parallel Range Algorithms

- Users often combine ranges and parallel algorithms, but the current approach is verbose and error-prone.
- The proposed parallel range algorithms offer a more natural and expressive way to parallelize range-based computations

Design Overview: Key Modifications

- Execution policy parameter added to range algorithms.
- Introduction of bounded ranges for better parallel performance.
-
- **Execution policies:** Parallel range algorithms accept execution policies to control parallelism.
- **Random access ranges:** Algorithms require random-access ranges for efficient parallelization.
- **Bounded ranges:** At least one input and the output range must be bounded for safety and performance.
- **Algorithm return types:** Consistent with serial range algorithms for easy migration.
-
- Enable single-call fusion of multiple operations
- Preserve the expressiveness of ranges

Key Design Decisions

1. Return types match serial range algorithms
2. Require `random_access_range` (for now)
3. Take range as output
4. Require bounded ranges
5. Preserve callable requirements from C++17 parallel algorithms

```
template <class ExecutionPolicy,  
          random_access_range R,  
          class Proj = identity,
```

```
indirectly_unary_invocable<projected<i  
iterator_t<R>, Proj>> Fun>
```

```
requires
```

```
sized_sentinel_for<ranges::sentinel_t<  
R>, ranges::iterator_t<R>>
```

```
ranges::borrowed_iterator_t<R>
```

```
ranges::for_each(ExecutionPolicy&&  
policy, R&& r, Fun f, Proj proj = {});
```

Differences to C++17 Parallel Algorithms

Key Differences:

- Parallel range algorithms require random access ranges.
- Output can now be a range instead of just an iterator.

Benefits

Parallel range algorithms offer a natural and efficient way to parallelize range-based computations.

The proposed design integrates seamlessly with the Ranges library and existing parallel algorithms.

This feature will enhance the expressiveness and performance of parallel code in C++.

More expressive code

Potential for better performance

Safer APIs (bounded ranges, range outputs)

Simplified migration from serial to parallel code

C++ Parallel Algorithms

P2500R2 C++ parallel algorithms and P2300

Overview:

- The evolution of parallelism in C++.
- Why P2300 is important for C++26.

Goal of the Talk:

- Discuss the integration of C++ parallel algorithms with the facilities introduced in P2300.
- <https://www.open-std.org/jtc1/sc22/wg21/docs/papers/2023/p2500r2.htm>

ISO C++ Parallelism/Concurrency Programming Language (based on Gonzalo's ISC C++ BoF)

Parallel Algorithms [many]


Concurrency++

Memory Model++ [MichaelW,
Maged, Paul M]

Forward Progress

Ranges++ many

Multi-dimensional Spans [Bryce, Christian]
operator[i, j, k]



Further evolution of parallel algorithms (P2500)

//C++ 11/14

```
namespace stde = std::execution;  
std::vector<int> vec{N};
```

// <https://wg21.link/P3179>: parallel range algorithms

```
std::ranges::generate( vec, std::minstd_rand{} );  
std::ranges::sort( stde::par, vec );
```

// <https://wg21.link/P2500>: support schedulers to specify where to execute, targeted to cover synchronous parallel algorithms integration with P2300.

```
std::ranges::generate(stde::execute_on(some_sched, stde::seq), vec, std::minstd_rand{} );  
std::ranges::sort( stde::execute_on(some_sched, stde::par), vec );
```

Motivation

Why P2300?

- The need for a flexible abstraction that answers "where" the code should be executed.
- Limitations of current execution policies in specifying hardware execution contexts.

C++17 parallel algorithms: A good start

Current limitation: No control over execution hardware

P2300 introduces flexible schedulers

Need: Integrate schedulers with parallel algorithms

The Need for Integration

- C++ parallel algorithms offer parallelism, but lack control over execution hardware ("where").
- P2300 introduces the "scheduler" concept, representing execution contexts, addressing the "where."
- The integration of these two is crucial for leveraging hardware capabilities effectively.

Design Overview

The Proposed Solution in P2500

- The paper proposes extending C++ algorithms to accept a "policy-aware scheduler."
- This scheduler combines an execution policy ("how") and a scheduler ("where").
- The `execute_on` function facilitates the creation of such policy-aware schedulers.

Design Goals:

- Extending C++ parallel algorithms with policy-aware schedulers.
- Allowing customization for different execution contexts.
-
- Preserve core semantics of algorithms and policies
- Cover both "classic" and range-based algorithms
-
- Minimal API changes: The design aims to preserve the existing usage patterns of C++ algorithms.
- Flexibility: It allows execution semantics to be adjusted based on the capabilities of the execution context.
- Customization: Implementers of execution contexts can customize the implementation of standard algorithms for optimal performance.
-

Key Features:

- Combining scheduler and policy.
- Minimal, incremental API changes.

Combining Scheduler with Policy

Why Use Schedulers?

- Schedulers represent execution contexts and provide flexibility.
- API overview

Key Concepts

1. `policy_aware_scheduler`
2. `execute_on`
3. Customizable functions
 - - The `execution_policy` concept defines the requirements for execution policies.
 - The `policy_aware_scheduler` concept represents an entity combining a scheduler and an execution policy.
 - The `execute_on` customization point binds a scheduler and an execution policy.
 - Parallel algorithms are defined as customizable functions, allowing customization for specific policy-aware schedulers.

policy_aware_scheduler Concept

API

```
template <typename S>
concept policy_aware_scheduler =
scheduler<S> && requires (S s) {
    typename
    S::base_scheduler_type;
    typename S::policy_type;
    { s.get_policy() } ->
    execution_policy;
};
```

Usage Example

```
struct MyScheduler {
    using base_scheduler_type = /* some scheduler
    type */;

    using policy_type = /* some execution policy
    type */;

    policy_type get_policy() const {
        return /* return the associated policy
        */;
    }
};

static_assert(policy_aware_scheduler<MySchedul
er>);
```

execute_on Function API and Usage example

```
inline namespace
__execute_on_fn_namespace
{

    inline constexpr
    __detail::__execute_on_fn
    execute_on;

}
```

```
auto
policy_aware_sched =
std::execute_on(my_scheduler,
std::execution::par);
```

Proposed API (Example with for_each)

// Existing API

```
template<class ExecutionPolicy, class It, class Fun>  
constexpr void for_each(ExecutionPolicy&& policy, It first, It last, Fun f);
```

// New Policy-based API

```
template<execution_policy Policy, input_iterator I, sentinel_for<I> S, class Proj = identity,
```

```
        indirectly_unary_invocable<projected<I, Proj>> Fun>
```

```
constexpr ranges::for_each_result<I, Fun>
```

```
    ranges::for_each(Policy&& policy, I first, S last, Fun f, Proj proj = {});
```

// New Scheduler-based API

```
template<policy_aware_scheduler Scheduler, input_iterator I, sentinel_for<I> S,
```

```
        class Proj = identity, indirectly_unary_invocable<projected<I, Proj>> Fun>
```

```
constexpr ranges::for_each_result<I, Fun>
```

```
    ranges::for_each(Scheduler sched, I first, S last, Fun f, Proj proj = {}) /*customizable*/;
```

Allowing schedulers with C++ algorithms.

Blocking behavior similar to C++17 parallel algorithms.

```
template<policy_aware_scheduler Scheduler,  
typename ForwardIterator, typename Function>  
void for_each(Scheduler&& sched, ForwardIterator  
first, ForwardIterator last, Function f) {  
    // Implementation using scheduler and policy  
    sched.execute([&]() {  
        for (; first != last; ++first) {  
            f(*first);  
        }  
    });  
}
```

Using the API

```
std::for_each(
```

```
std::execute_on(my_gpu_scheduler,  
std::execution::par),  
    begin(data),  
    end(data),  
    [](auto& item) { item.process();  
}  
);
```

Customization and Extensibility

Parallel Algorithms as Customizable Functions:

- Customization through policy-aware schedulers.
- Flexibility to support platform-specific optimizations.

Example Customization:

- Using CUDA-specific scheduler for `std::for_each`.

```
namespace cuda {
    struct scheduler {
        friend constexpr auto
        tag_invoke(std::tag_t<ranges::for_each>, scheduler, /*...*/) {
            // CUDA-optimized implementation
            cuda_kernel<<<blocks, threads>>>(/*...*/);
            return std::ranges::for_each_result{/*...*/};
        }
    };
}
```

Summary

P2300 offers significant flexibility and control over execution contexts.

P2500 Integration with parallel algorithms is crucial for modern C++ development.

The future of C++ parallelism lies in customizable and extensible algorithms.

Conclusion and Questions

We covered a few interesting proposals that are coming in C++ parallelism and concurrency based on the last 3 C++ Std meetings since the last CPPCON:

1. Improving C++20 Atomic Min/Max(P0493, Michael)
2. Hazard pointer extensions (P3135; Maged)
3. Pointer tagging (P3125; Maged)
4. Parallel Range algorithms (P3179; Michael), maybe Parallel Algorithms (P2500),

1 and 4 are likely to make C++26. 2 and 3 will be beyond C++ 26.

These are just some of the likely features that have progressed sufficiently and fit in a talk format, though they still may enter the standard in slightly different forms.

If you like this talk, we will continue to cover more parallelism and concurrency proposals in future CPPCON talks to prepare you for the future, giving our take on their significance, and usefulness in parallel concurrent programming.