



Performance Engineering

Being Friendly to Your Hardware

IGNAS BAGDONAS



20
24



Being Friendly to Your Hardware

Performance Engineering

A gentle introduction to hardware for software engineers

Where does C++ run?


On an abstract C++ machine

On an abstract C++ machine?




C++ runs on a computer platform

Computer platform




Computer platform

Processor + Interconnect + Memory





Computer platform

Processor + Interconnect + Memory





Computer platform

Processor + Interconnect + Memory






Computer platform

Processor + Interconnect + Memory




Computer platform

Processor + Interconnect + Memory




Computer platform

Processor + Interconnect + Memory





Computer system

Processor cores + Interconnect + Memory




Computer system

Processor cores + Interconnect + Memory




Computer system

Processor cores + Interconnect + Memory




Complex and scary

And mostly out of
scope of our
discussion anyway




Memory

- Addressable capacitor array
- Short term nonvolatile
- Speeds – capacitor array vs bus interface




Memory

- Open a row
- Send part of address - one command



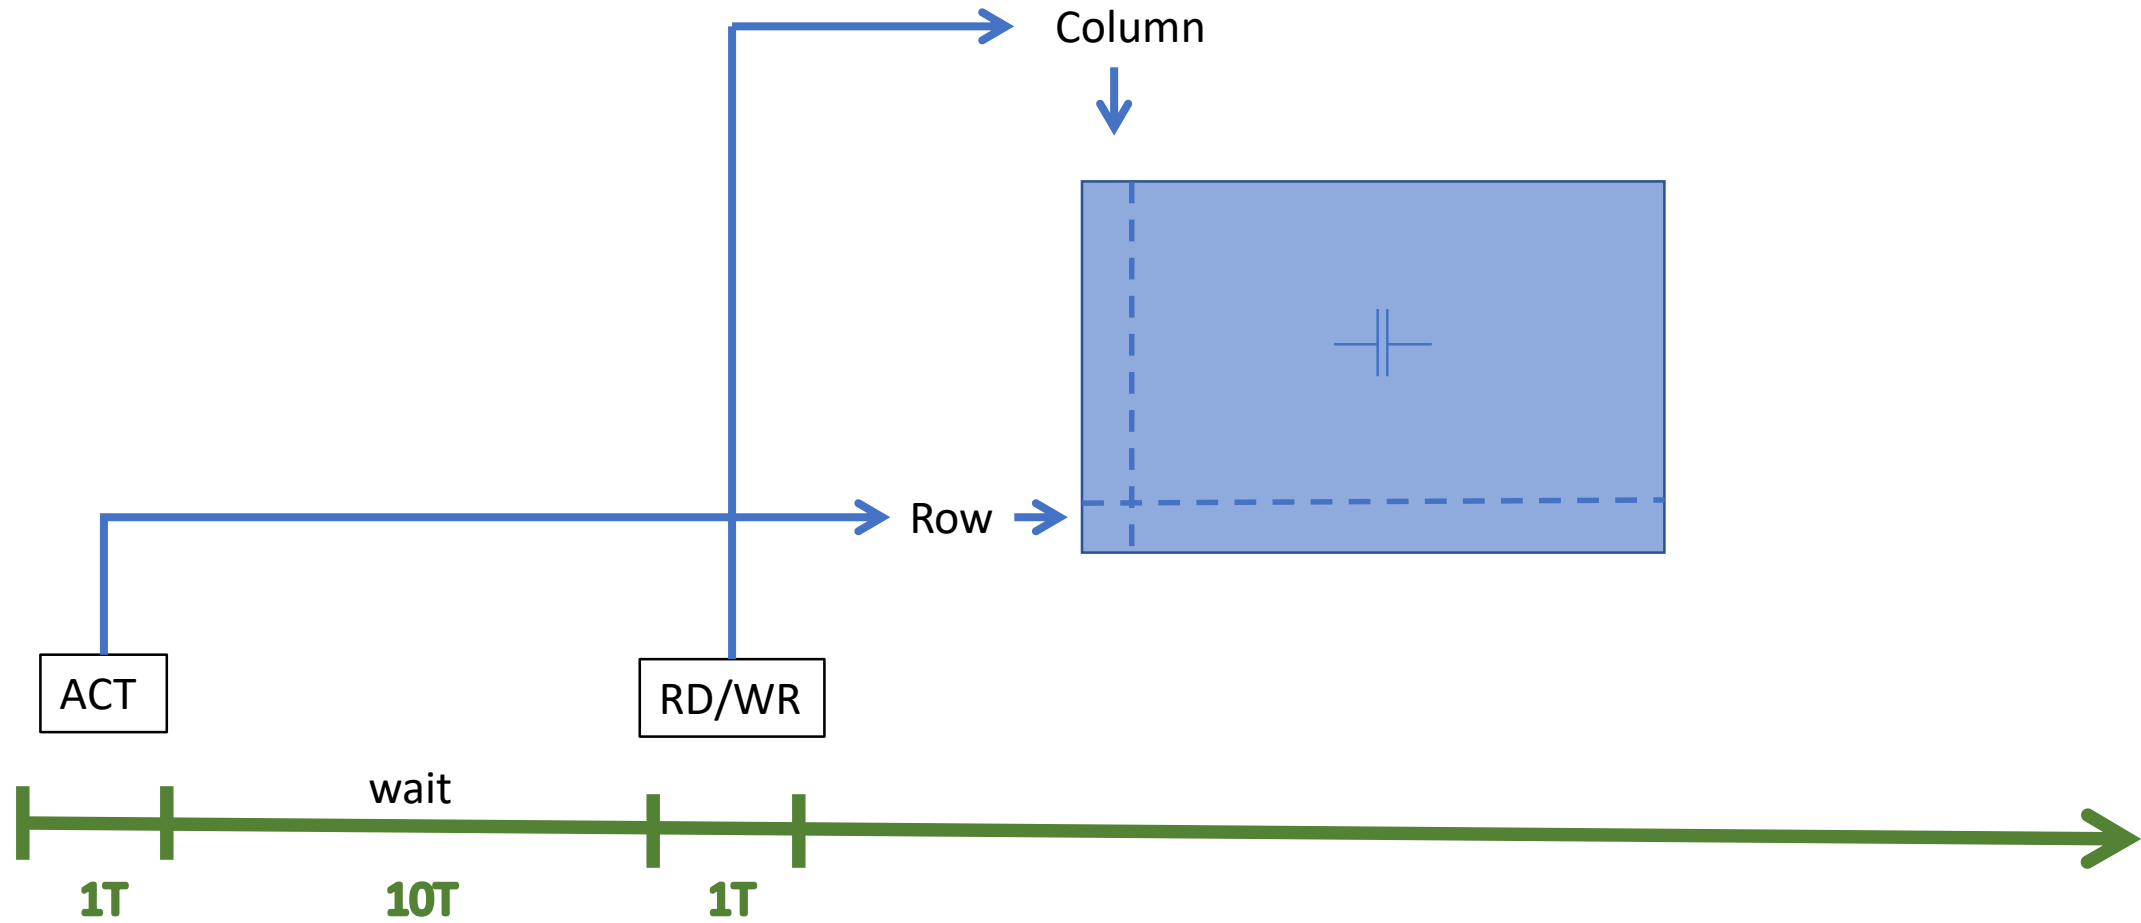
Memory

- Wait for a portion of array to get selected.



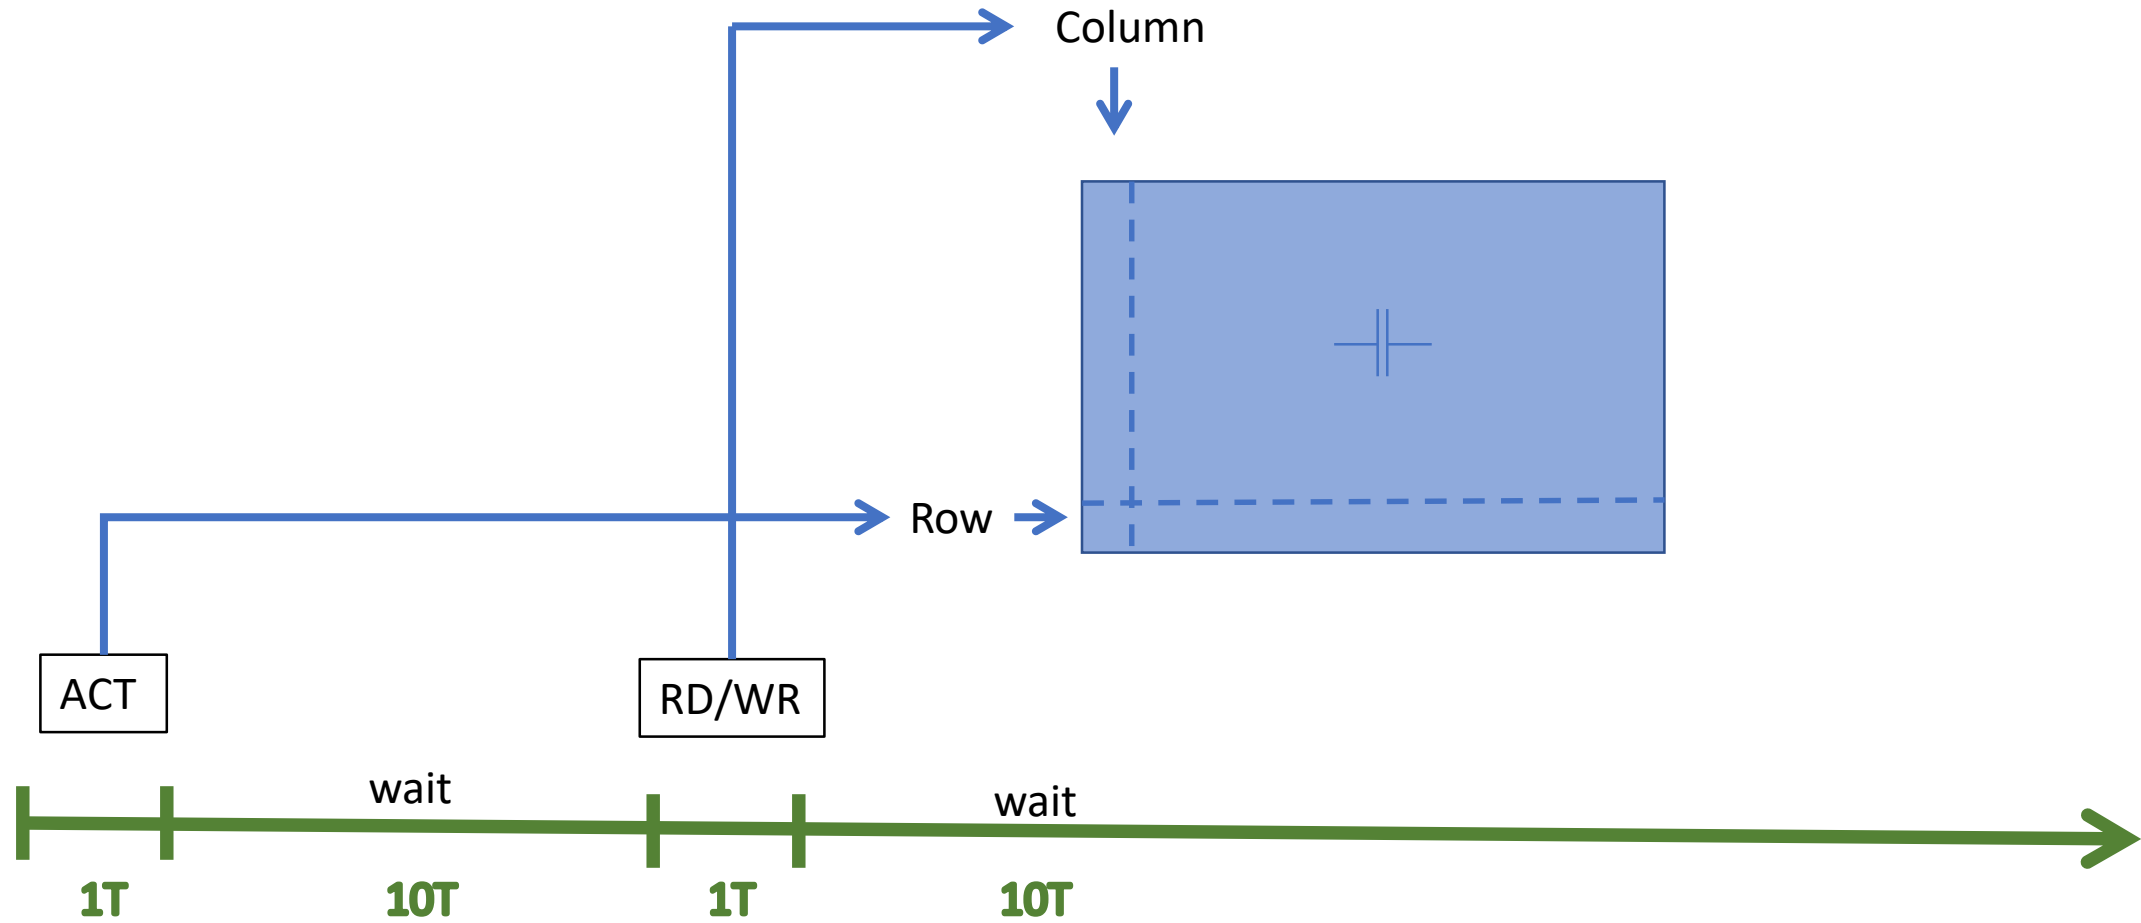
Memory

- Send column address and actual memory operation – another command



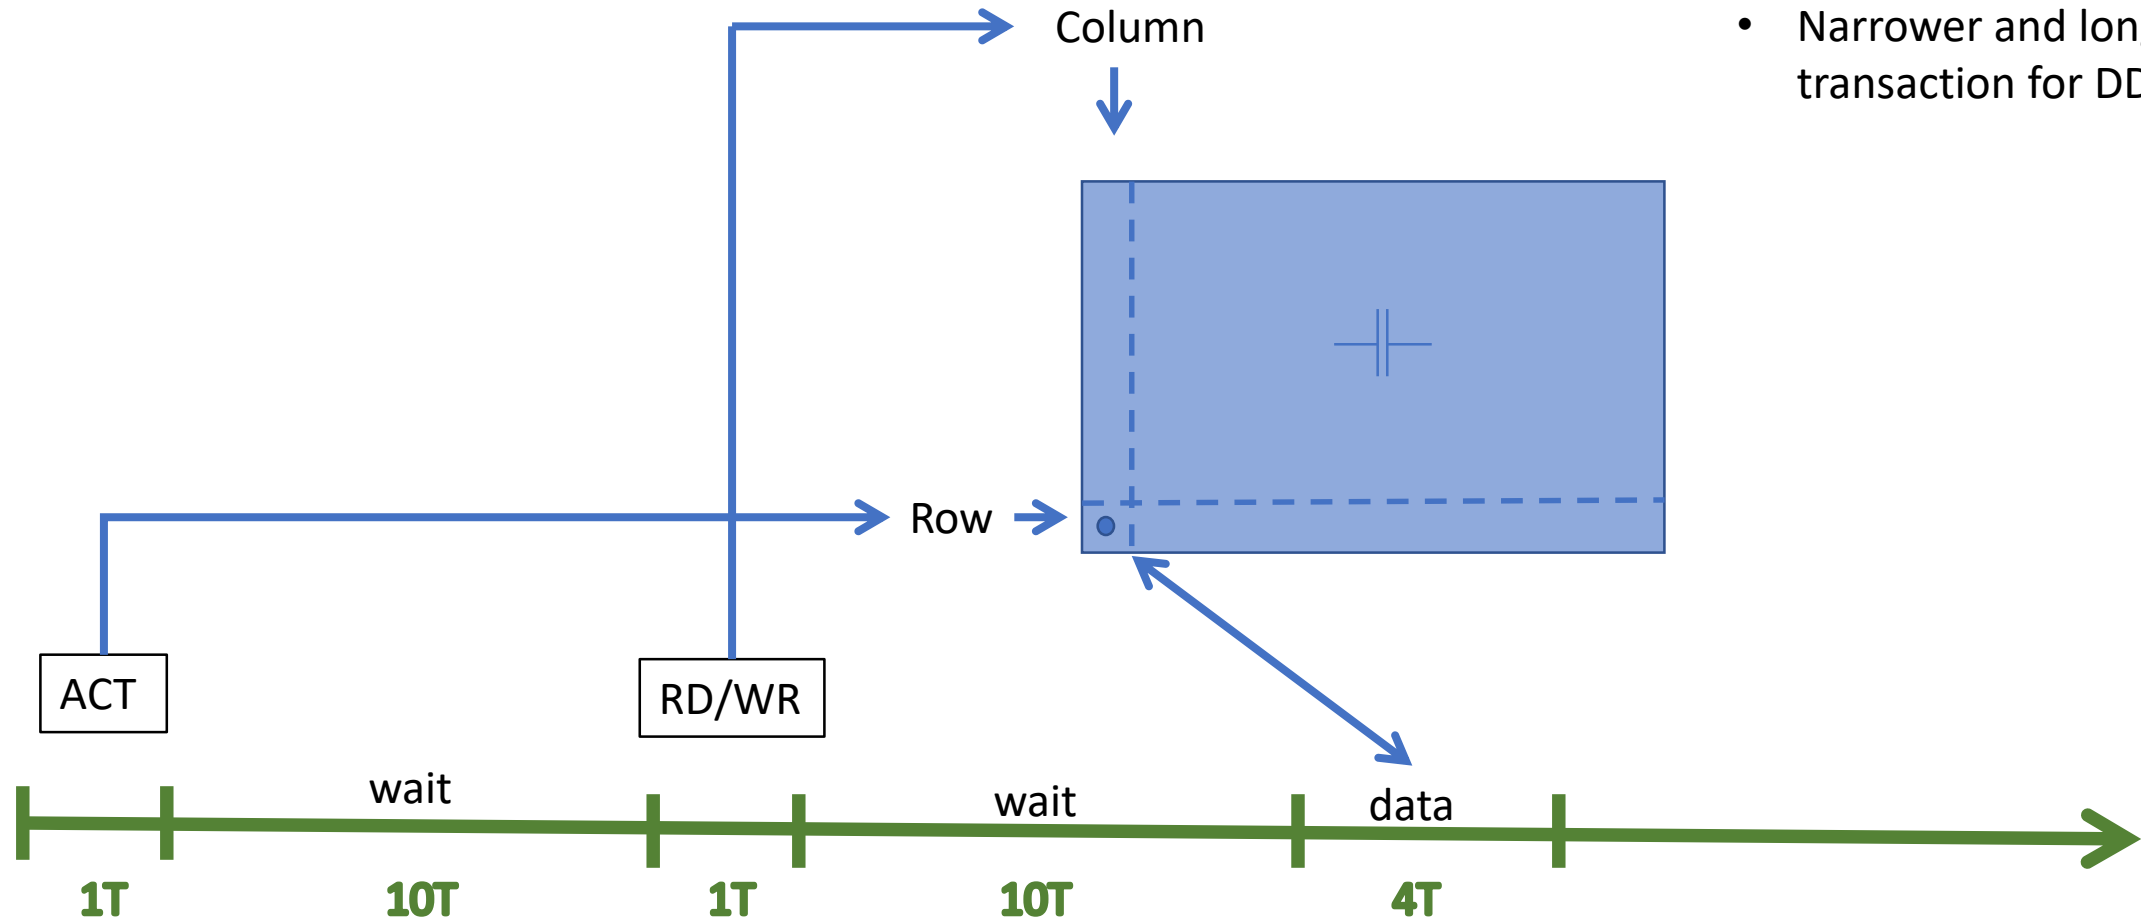
Memory

- Wait for a portion of row to get selected and capacitors sensed.



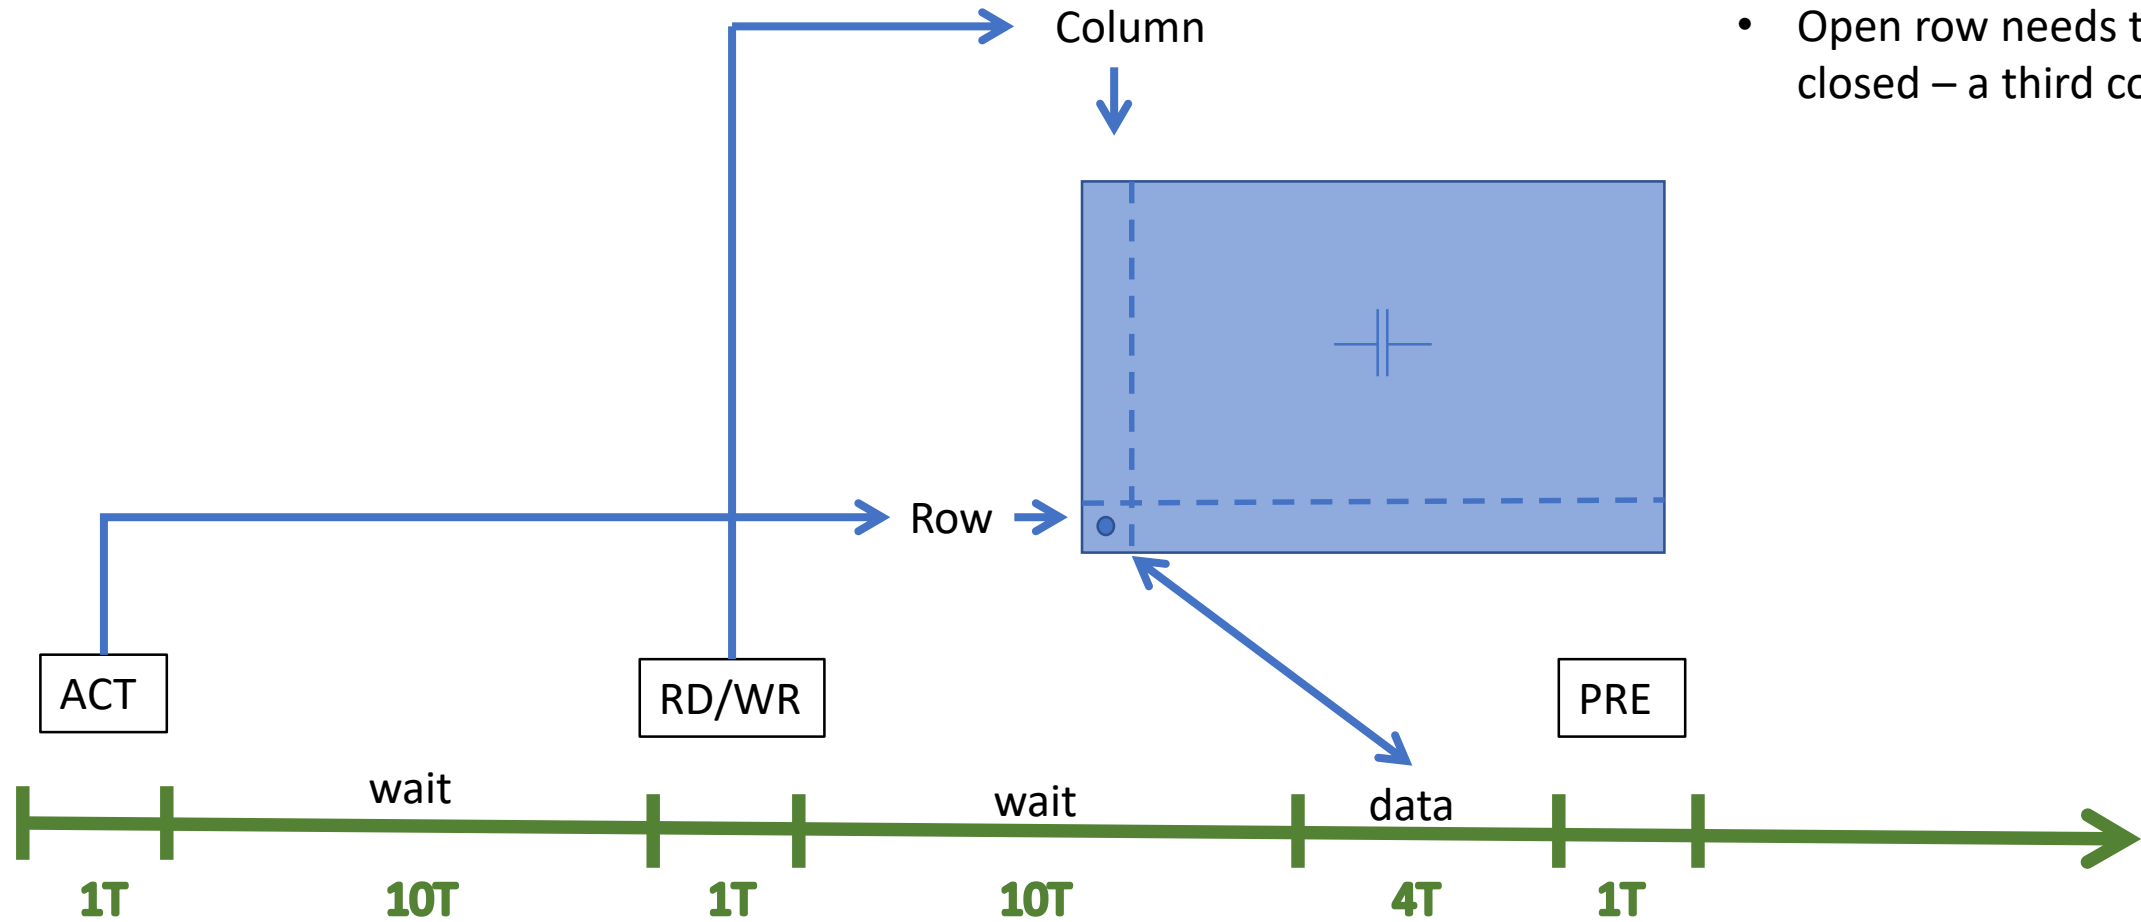
Memory

- Actual data transfer.
- Wider and shorter transaction for DDR4
- Narrower and longer transaction for DDR5



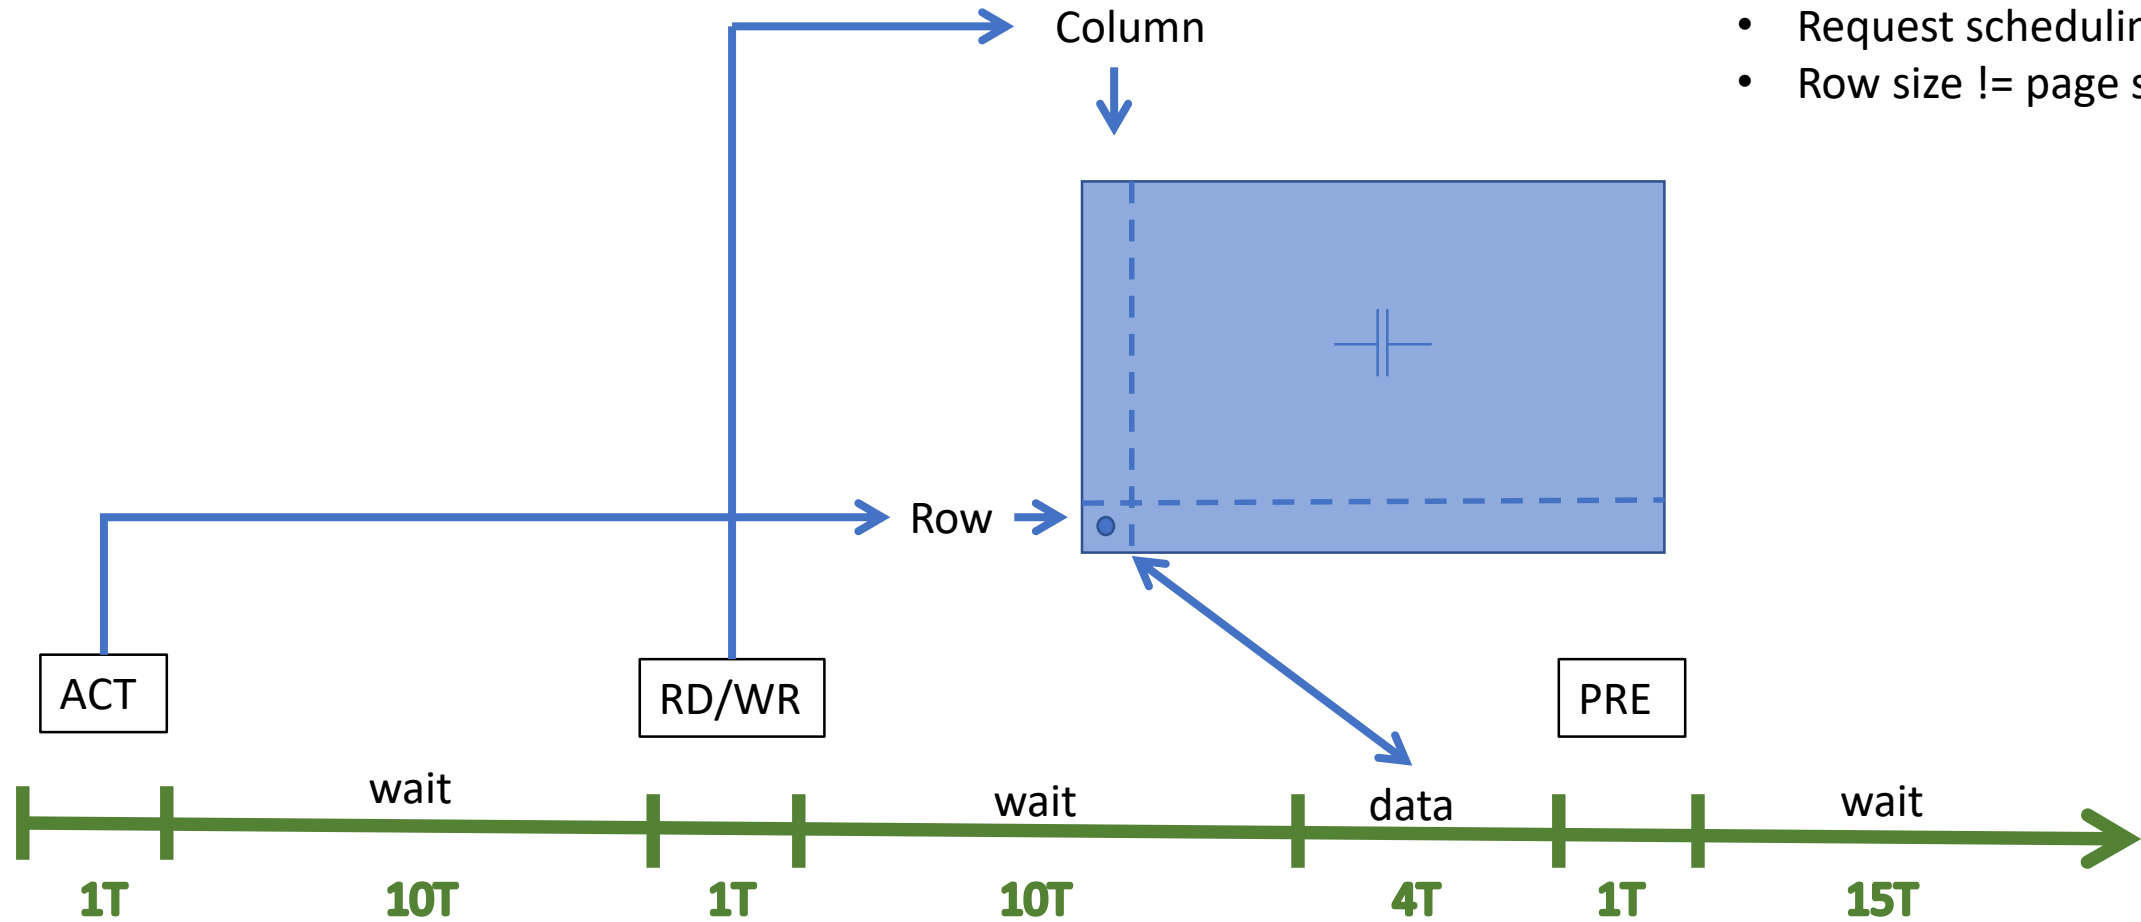
Memory

- Capacitor array access is destructive
- Even for a read
- Open row needs to be closed – a third command




Memory

- One row active at a time
- All rows need to be refreshed periodically
- Request scheduling
- Row size \neq page size



Memory

- Usable utilization is quite far away from theoretical limit
- Bandwidth vs latency



Memory


- DDR architecture – it is only a fraction of bus transaction that is in fact DDR
- Row addressing
- Column addressing
- Generation to generation capacitor array operates at mostly the same frequency, bus transfer speed in fact increases
- Multiple memory modules (sockets)
- Page size – CPU vs memory
- Address mapping

Memory




- One row open per bank
- Logical partitioning

Memory




- One row open per bank
- Logical partitioning

Memory




- One row open per bank
- Logical partitioning

Memory




- One row open per bank
- Logical partitioning

Memory




- Bank group – a grouping of logically addressable banks
- Multiple open rows per bank group
- Multiple outstanding commands in progress

Memory




- Bank group – a grouping of logically addressable banks
- Multiple open rows per bank group
- Multiple outstanding commands in progress

Memory




- Bank group – a grouping of logically addressable banks
- Multiple open rows per bank group
- Multiple outstanding commands in progress

Memory




- Bank group – a grouping of logically addressable banks
- Multiple open rows per bank group
- Multiple outstanding commands in progress

Memory



- Multiple bank groups for a rank
- A rank is just a CS line from the electrical domain perspective
- It is part of addressing scheme too

Memory



- Multiple bank groups for a rank
- A rank is just a CS line from the electrical domain perspective
- It is part of addressing scheme too
- Multiple ranks physically equal to multiple separate groups of memory components assembled on the same PCB
- MDP and vendor marketing names for multilayer components

Memory

- What about HBM?
 - Not 1024 bits, 8 x 128 instead.
 - Bandwidth will only be high when there is sufficient stream of commands inflight
- What about remote memory?
 - CXL and vendor proprietary interconnects
 - Coherency is expensive
 - Single physical address space is expensive
 - Single logical cache coherent address space is very expensive
- What about nonvolatile memory?
 - NVDIMM-* variants, vendor specifics.
 - Useful for debugging
- System address mapping design and control

Addressable memory units

- SW visible page size is (typically) 4KB
- DDR memory page (row) size depends on the actual components used
- Mapping is controllable by low level firmware
- In most practical cases at boot time only

Configuration		1 Gb x4	512 Mb x8	256 Mb x16
Bank Address	# of Bank Groups	4	4	2
	BG Address	BG0~BG1	BG0~BG1	BG0
	Bank Address in a BG	BA0~BA1	BA0~BA1	BA0~BA1
Row Address		A0~A15	A0~A14	A0~A14
Column Address		A0~A9	A0~A9	A0~A9
Page size		512B	1KB	2KB


Same capacity, different composition => different performance profile

From JESD 79-4 DDR4 specification


Memory

- Memory system is in the uncore
- Cores act as clients
- Remote socket cores act as clients
- Distributed cache coherence is hard and expensive
- Writing implies reading first in most of cases
- Reading has a load on cache hierarchy
- Nontemporal writes may help in some cases
- Cacheability parameters can be tuned, this is highly platform dependent though.

Processor (core)




Instruction fetch




- Gets next block of (partial) instructions
- Linear fetch
- Incoming branch
- Instruction alignment
- Instruction fusing

Branch prediction



- Governs fetching of next instruction blocks
- A set of tables
- Branch history
- Branch site
- Branch target
- Repetitive patterns
- Really complex
- High cost of error

Instruction decoding

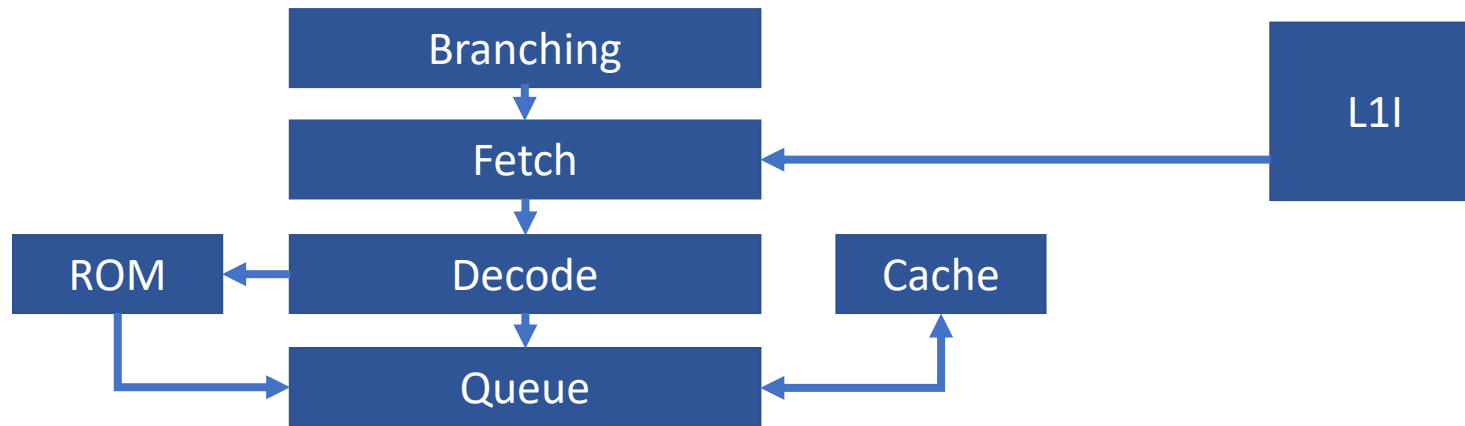


- Multiple instructions get decoded in parallel
- Even for variable length encoding
- Not that complex in HW domain
- Fusion

Instruction input

- Has to be logically sequential
- Can be physically parallel
- HW is inherently parallel
- Myth: decoding variable length instructions is:
 - Complex
 - Serial
 - Slow
- Fetch block size
- Linear fetch vs incoming branch

Instruction decoding



- Decoded operations may get cached
- There is a library of operations for complex instructions and events
- Decoded operations may get fused

Code density

```
uint64_t v = 0x123456789abcdef0;
```

```
x86  
movabs r10, 0x123456789abcdef0  
49 ba f0 de bc 9a 78 56 34 12
```

Code density

```
uint64_t v = 0x123456789abcdef0;
```

```
x86
movabs r10, 0x123456789abcdef0
49 ba f0 de bc 9a 78 56 34 12
```

```
MIPS
li      $2, 38141952
ori     $2, $2, 0x8acf
sll     $2, $2, 17
daddiu  $2, $2, 9903
sll     $2, $2, 18
ori     $2, $2, 0xdef0
46 02 02 3c cf 8a 42 34 78 14 02 00 af 26 42 64 b8 14 02 00 f0 de 42 34
```

Code density

```
uint64_t v = 0x123456789abcdef0;
```

```
x86
movabs r10, 0x123456789abcdef0
49 ba f0 de bc 9a 78 56 34 12
```

```
MIPS
li      $2, 38141952
ori     $2, $2, 0x8acf
dsll    $2, $2, 17
daddiu  $2, $2, 9903
dsll    $2, $2, 18
ori     $2, $2, 0xdef0
46 02 02 3c cf 8a 42 34 78 14 02 00 af 26 42 64 b8 14 02 00 f0 de 42 34
```

```
RISC-V
li      a5, 305418240
addi    a5, a5, 1657
li      a0, -1698897920
slli    a5, a5, 32
addi    a0, a0, -272
add     a0, a5, a0
12 34 57 b7 67 97 87 93 9a bc e5 37 02 07 97 93 ef 05 05 13 00 a7 85 33
```


Code density

```
uint64_t v = 0x123456789abcdef0;
```

```
x86
movabs r10, 0x123456789abcdef0
49 ba f0 de bc 9a 78 56 34 12
```

```
MIPS
li      $2, 38141952
ori     $2, $2, 0x8acf
dsll    $2, $2, 17
daddiu  $2, $2, 9903
dsll    $2, $2, 18
ori     $2, $2, 0xdef0
46 02 02 3c cf 8a 42 34 78 14 02 00 af 26 42 64 b8 14 02 00 f0 de 42 34
```

```
RISC-V
li      a5, 305418240
addi    a5, a5, 1657
li      a0, -1698897920
slli    a5, a5, 32
addi    a0, a0, -272
add     a0, a5, a0
12 34 57 b7 67 97 87 93 9a bc e5 37 02 07 97 93 ef 05 05 13 00 a7 85 33
```

```
SPARC
sethi    %hi(0x12345400), %g1
sethi    %hi(0x9ABCD00), %o0
or       %g1, 0x278, %g1
or       %o0, 0x2F0, %o0
sllx     %g1, 32, %g1
add      %g1, %o0, %o0
15 8d 04 03 37 af 26 11 78 62 10 82 f0 22 12 90 20 70 28 83 08 40 00 90
```

Code density

```
uint64_t v = 0x123456789abcdef0;
```

x86

```
movabs r10, 0x123456789abcdef0
49 ba f0 de bc 9a 78 56 34 12
```

MIPS

```
li      $2, 38141952
ori     $2, $2, 0x8acf
dsll    $2, $2, 17
daddiu  $2, $2, 9903
dsll    $2, $2, 18
ori     $2, $2, 0xdef0
46 02 02 3c cf 8a 42 34 78 14 02 00 af 26 42 64 b8 14 02 00 f0 de 42 34
```

ARM

```
mov     x0, 0xdef0
movk    x0, 0x9abc, lsl 16
movk    x0, 0x5678, lsl 32
movk    x0, 0x1234, lsl 48
00 de 9b d2 80 57 b3 f2 00 cf ca f2 80 46 e2 f2
```

RISC-V

```
li      a5, 305418240
addi    a5, a5, 1657
li      a0, -1698897920
slli    a5, a5, 32
addi    a0, a0, -272
add     a0, a5, a0
12 34 57 b7 67 97 87 93 9a bc e5 37 02 07 97 93 ef 05 05 13 00 a7 85 33
```

SPARC

```
sethi   %hi(0x12345400), %g1
sethi   %hi(0x9ABCD00), %o0
or      %g1, 0x278, %g1
or      %o0, 0x2F0, %o0
sllx    %g1, 32, %g1
add     %g1, %o0, %o0
15 8d 04 03 37 af 26 11 78 62 10 82 f0 22 12 90 20 70 28 83 08 40 00 90
```

Code density - fusion

```
uint64_t v = 0x123456789abcdef0;
```

x86
`movabs r10, 0x123456789abcdef0`
49 ba f0 de bc 9a 78 56 34 12


ARM
`mov x0, 0xdef0`
`movk x0, 0x9abc, lsl 16`
`movk x0, 0x5678, lsl 32`
`movk x0, 0x1234, lsl 48`
00 de 9b d2 80 57 b3 f2 00 cf ca f2 80 46 e2 f2



Imaginary ARM
`mov r20, 0x123456789abcdef0`

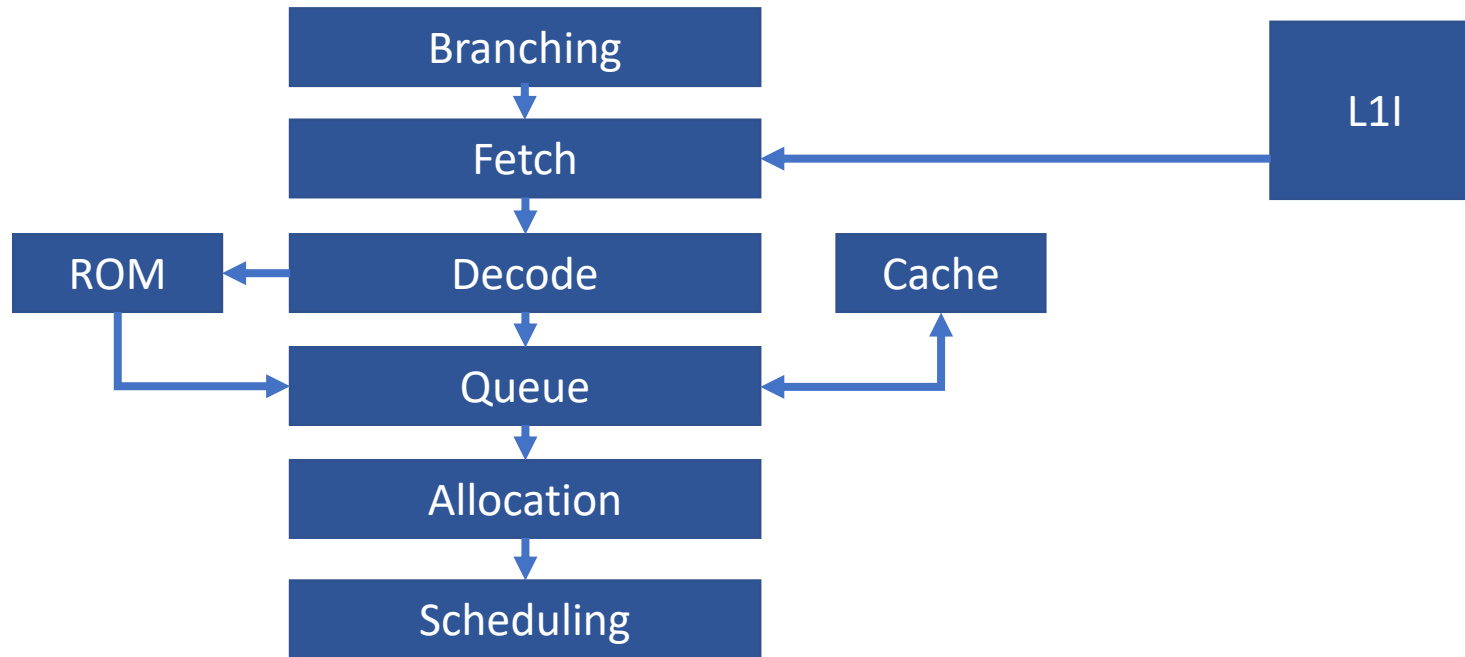
- Implementation may detect a sequence of instructions as a pattern
- And logically combine them
- Multiple instructions resulting in fewer operations
- ISA restrictions may have impact to performance

Register renaming



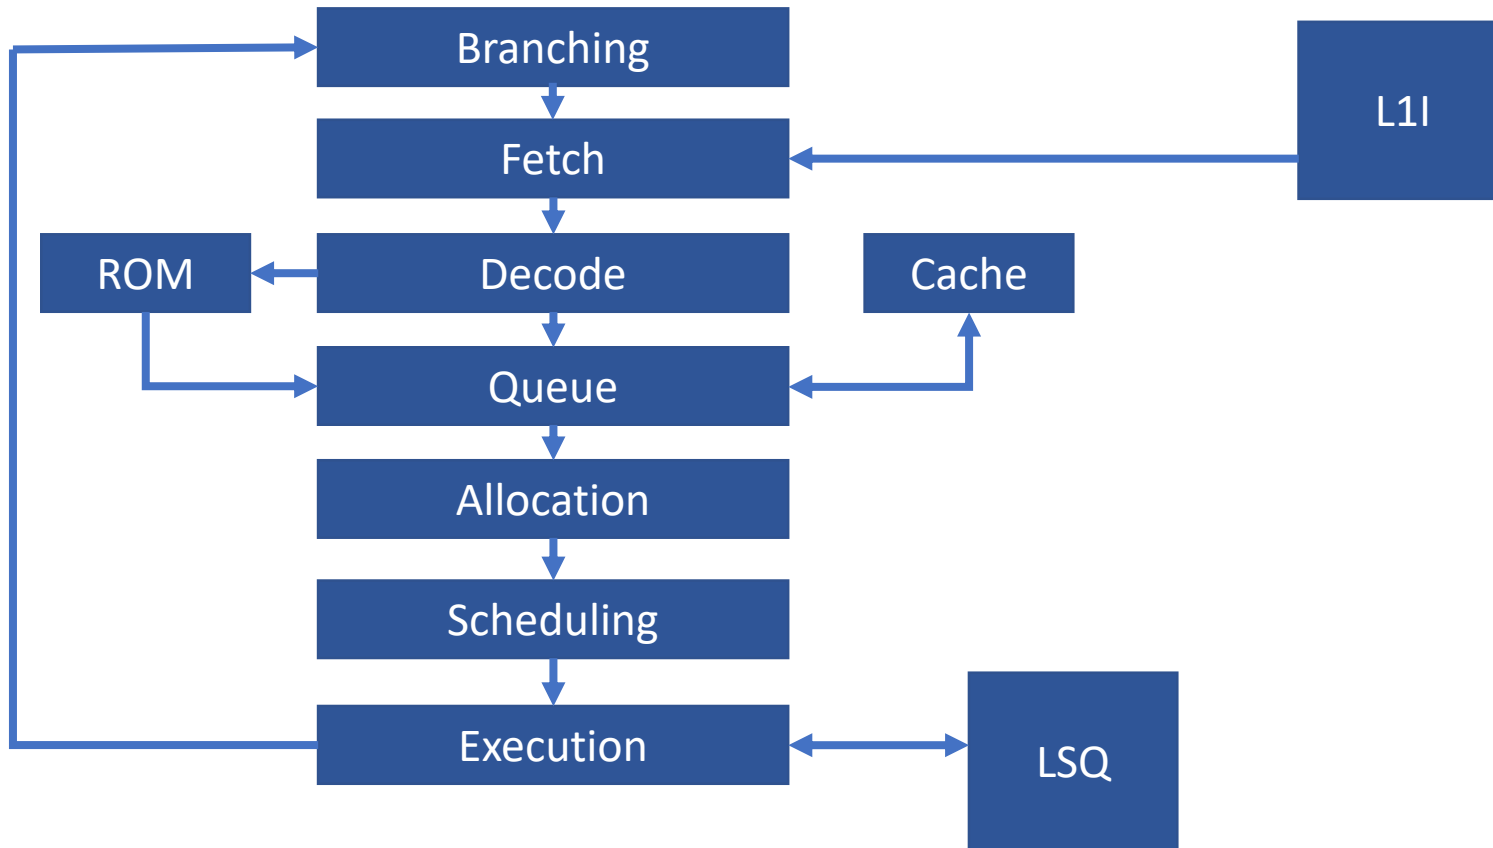
- ABI registers implement a SW contract
- They do not correspond to the actual execution
- Conversion of control flow to a variant of data flow
- Really complex
- Some operations end here
- Can a smart compiler help here?
- WLIV history

Scheduling



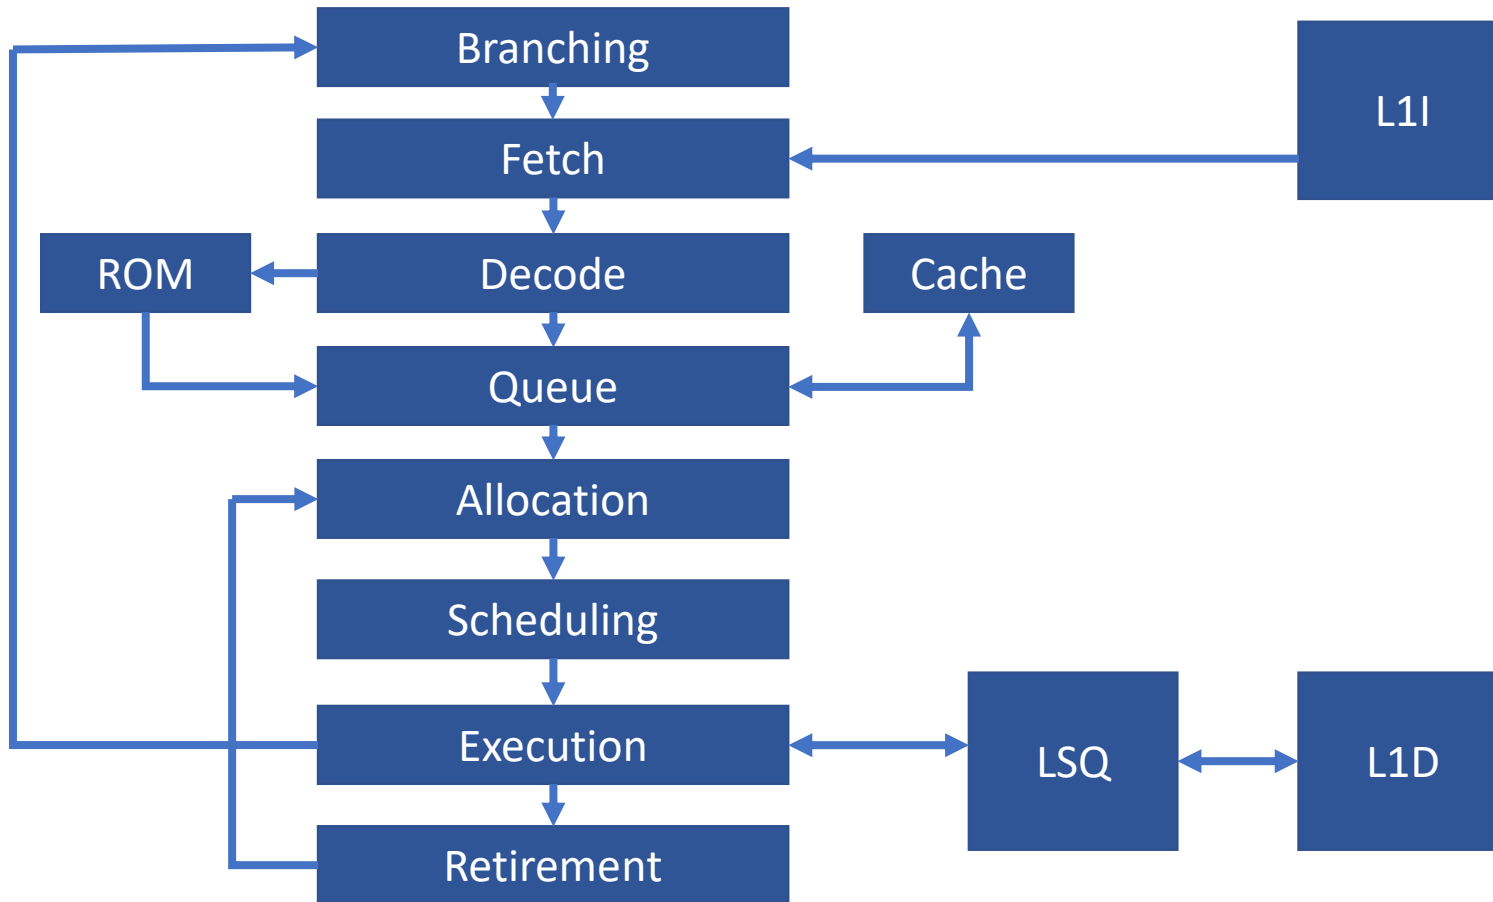
- Not all operations are equal
- Not all combinations of operations are equal

Execution



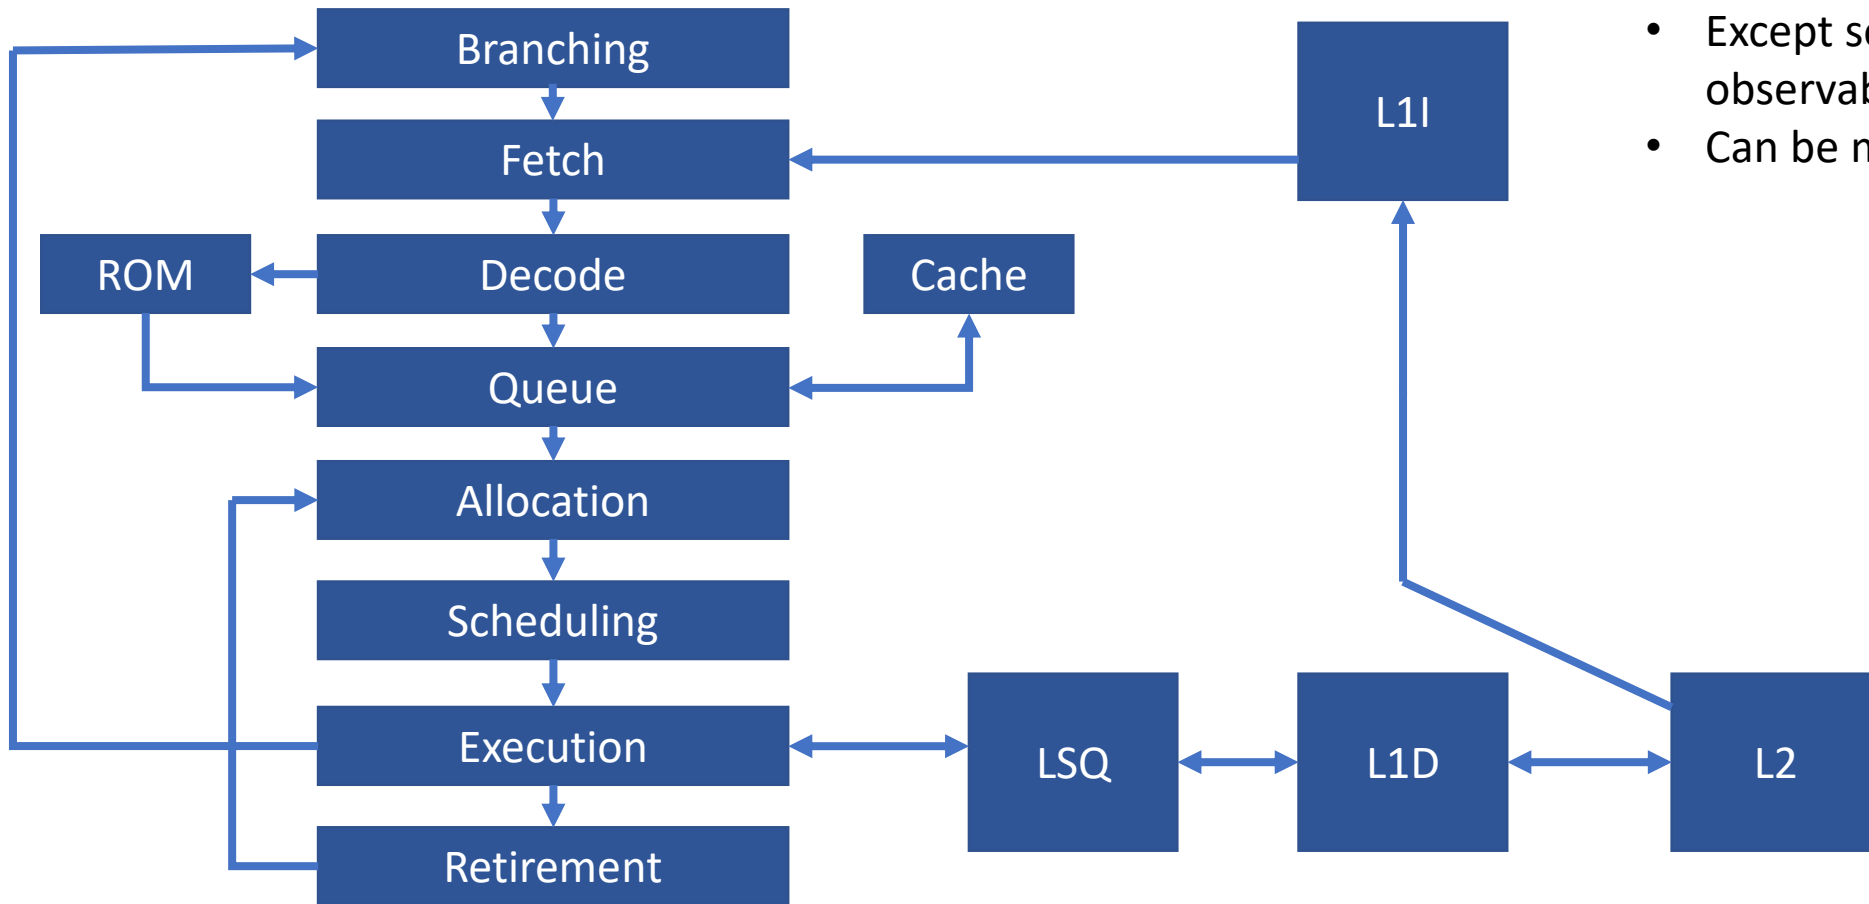
- Multiple specialized functional units
- Performs the actual (eventually) externally visible work
- Cycles
- Latency vs throughput

Retirement



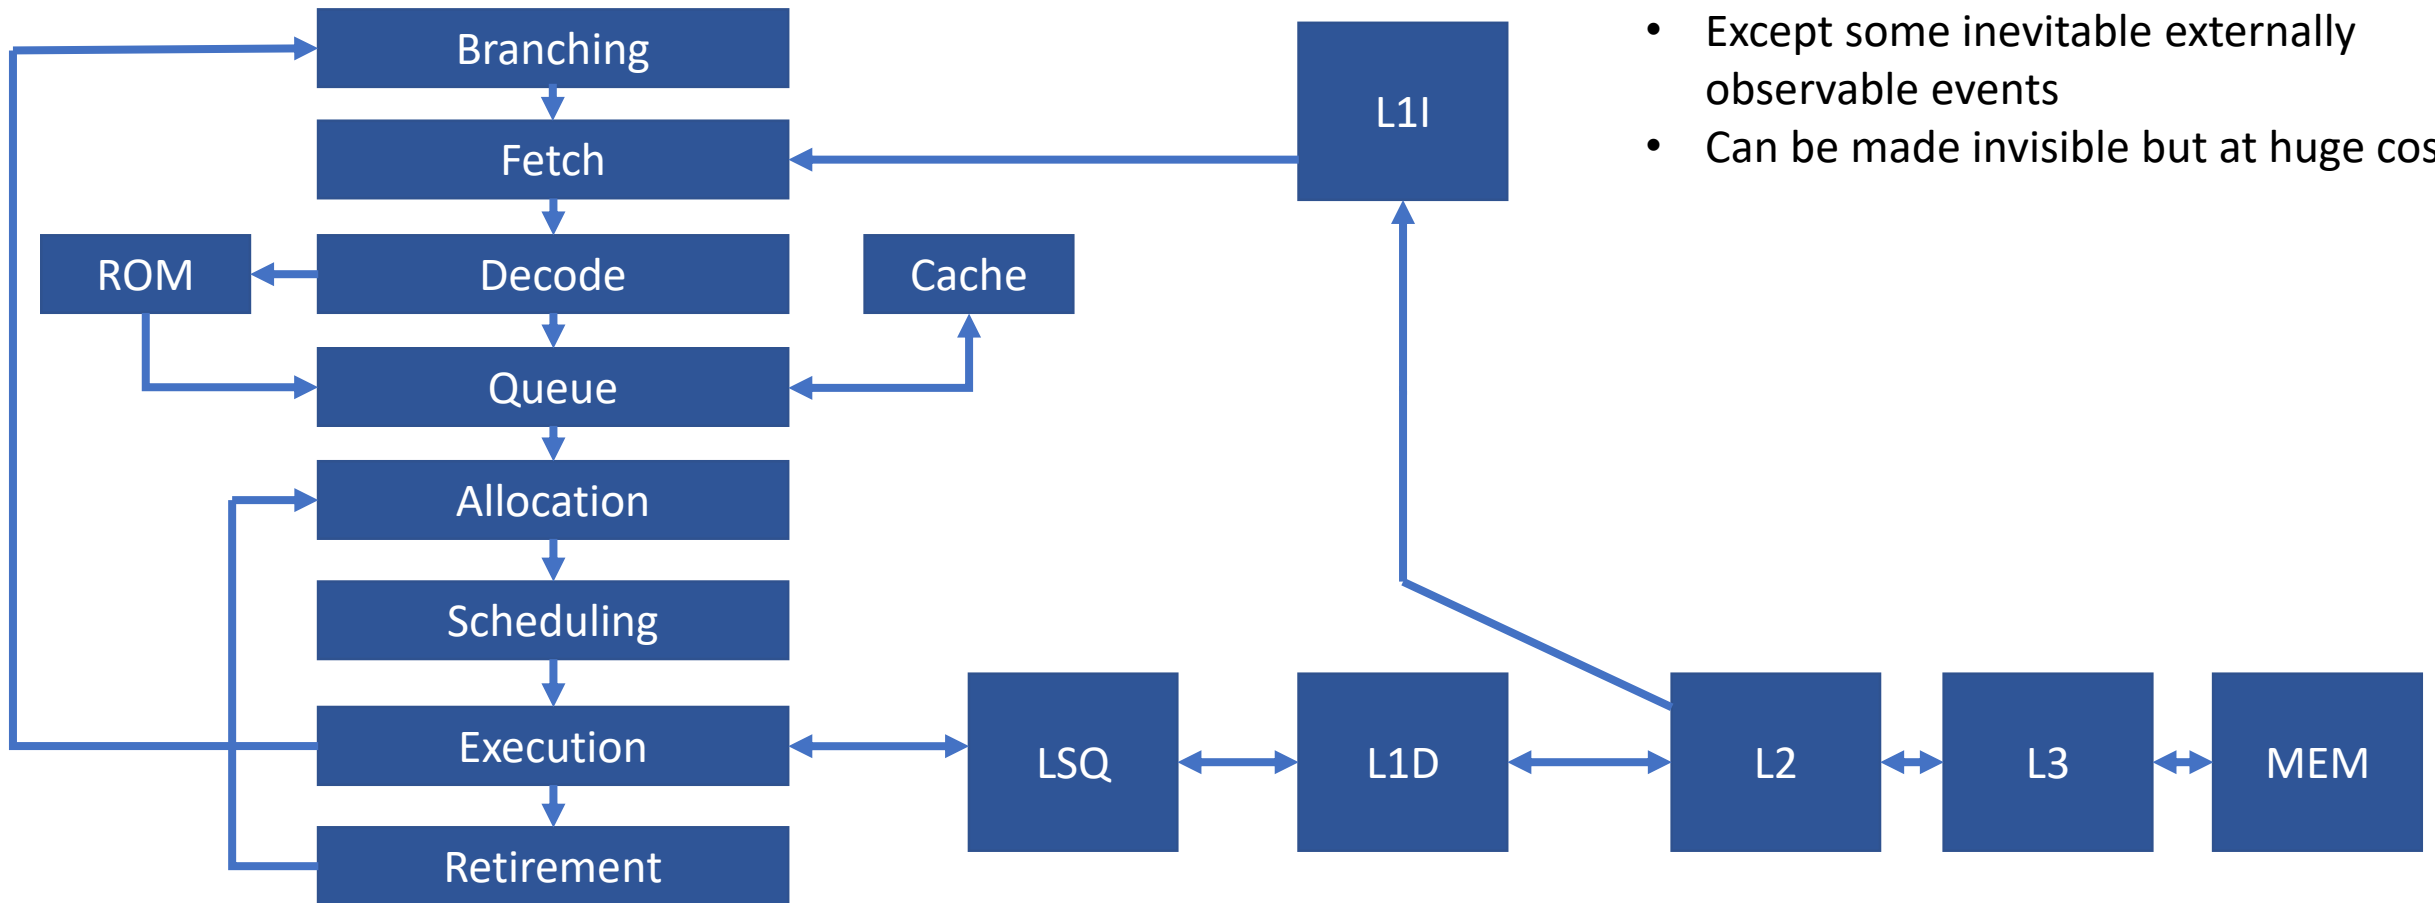
- All operations until now are speculative
- Except some inevitable externally observable events
- Can be made invisible but at huge cost

Retirement



- All operations until now are speculative
- Except some inevitable externally observable events
- Can be made invisible but at huge cost

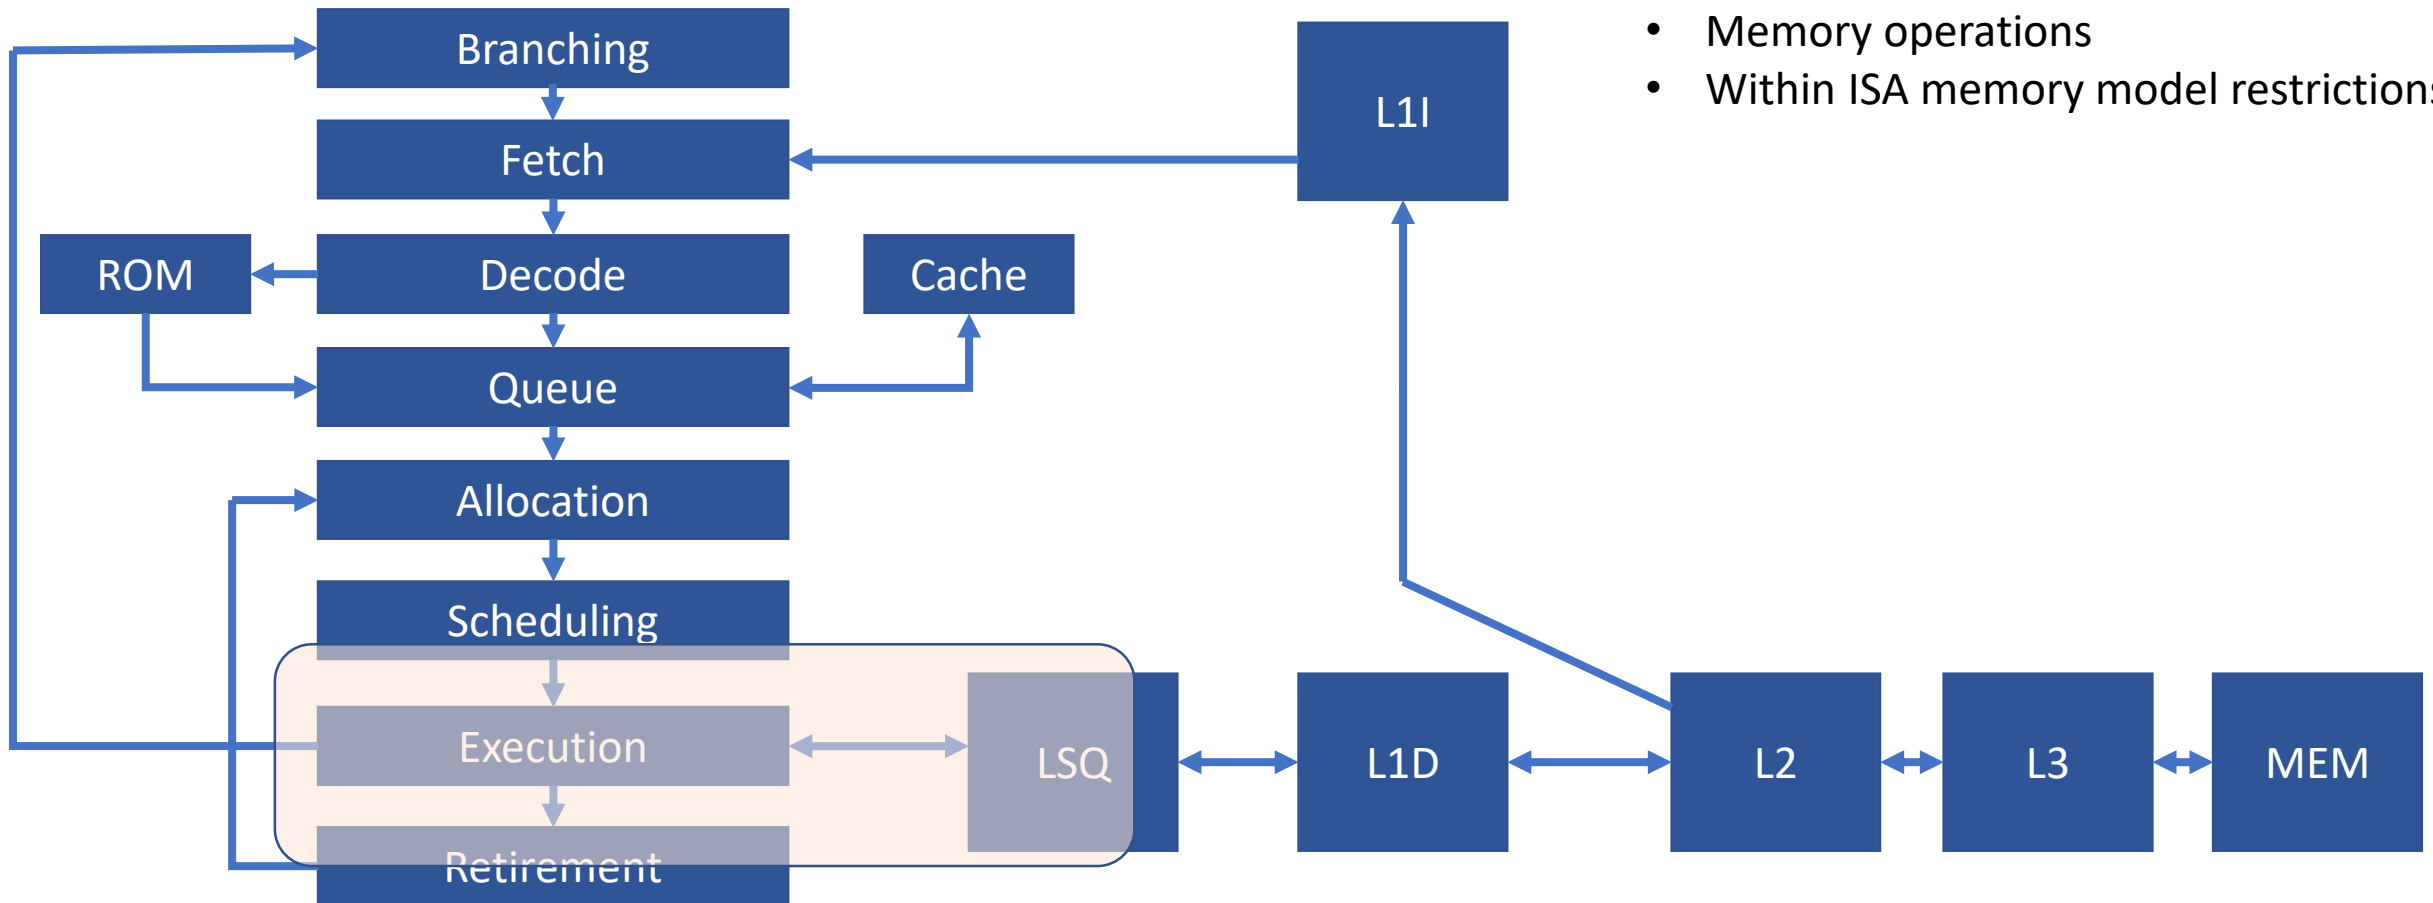
Retirement



- All operations until now are speculative
- Except some inevitable externally observable events
- Can be made invisible but at huge cost

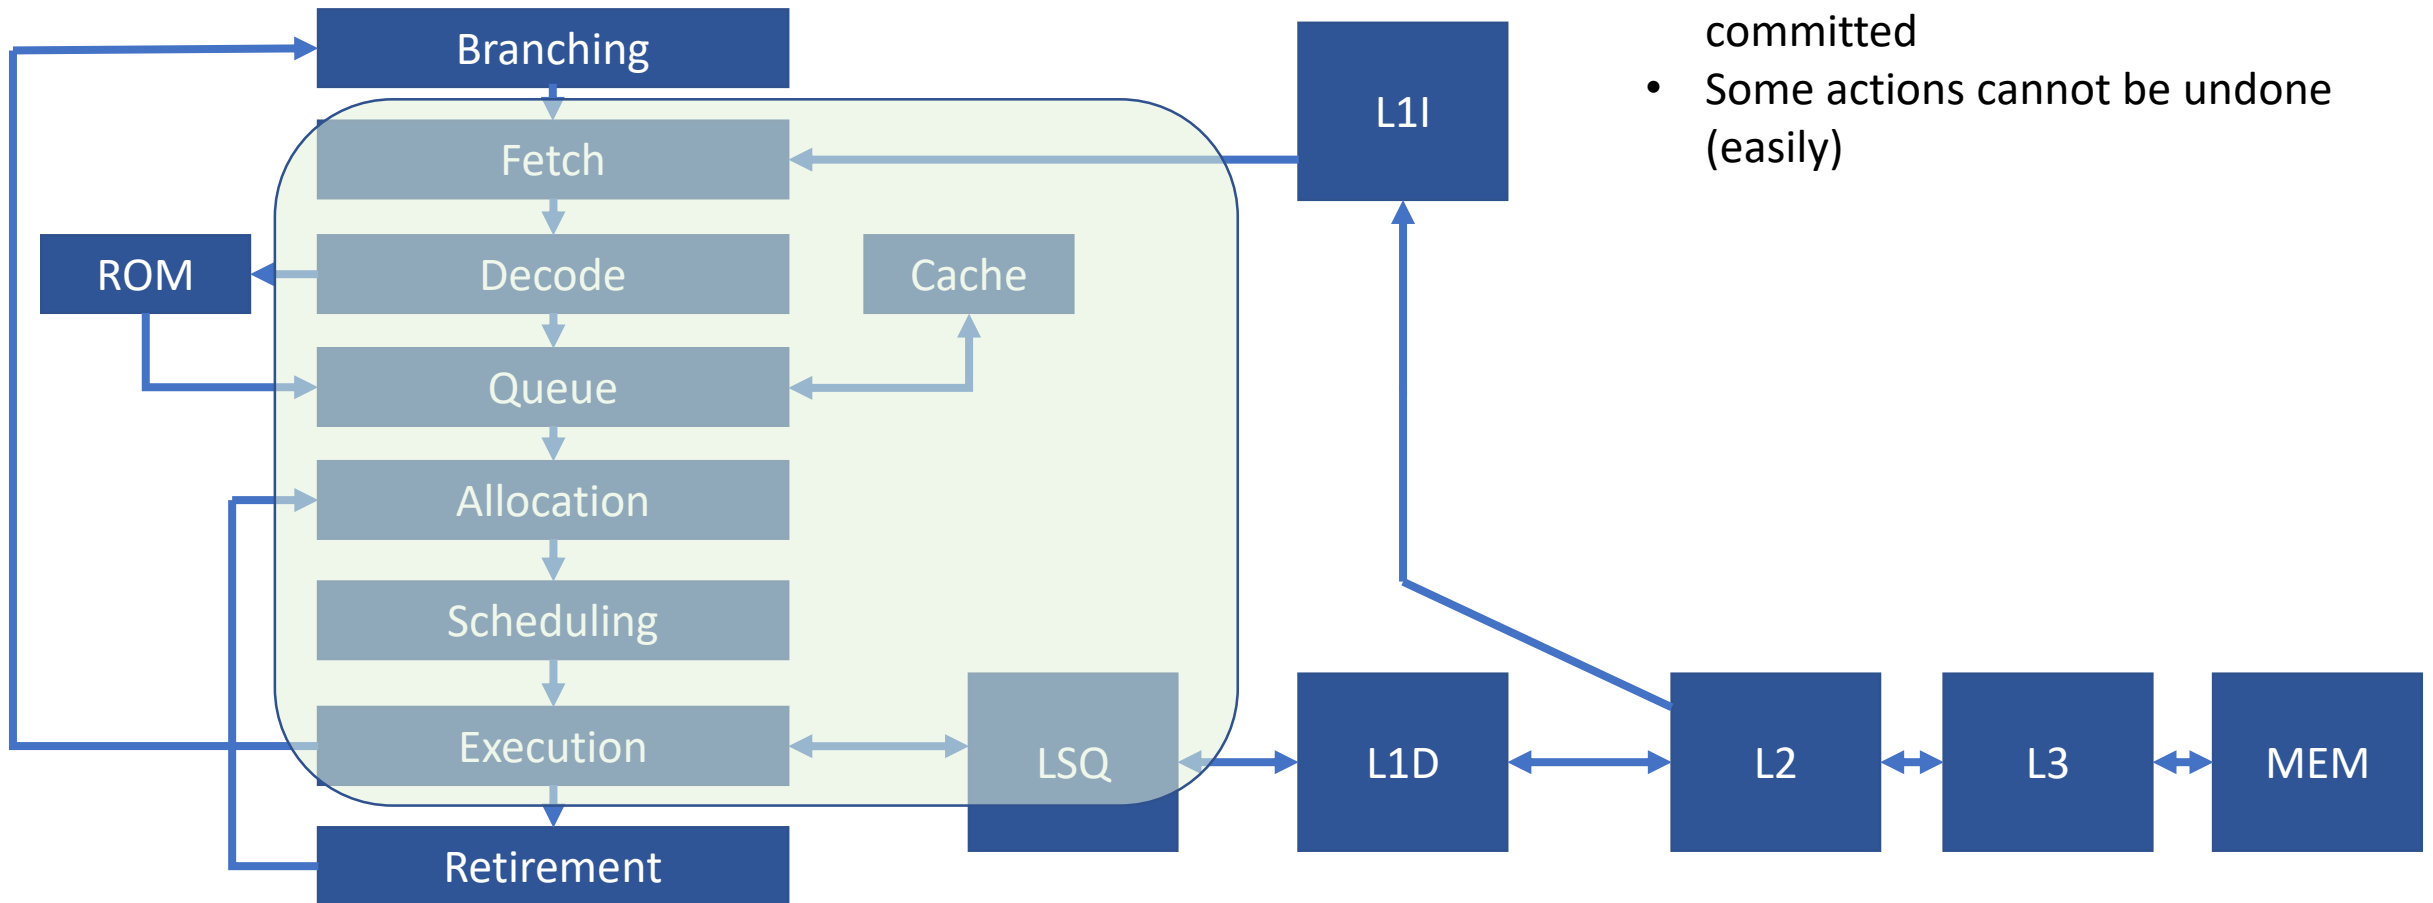
What is out of order?

- Execution of operations
- With some exceptions
- Memory operations
- Within ISA memory model restrictions

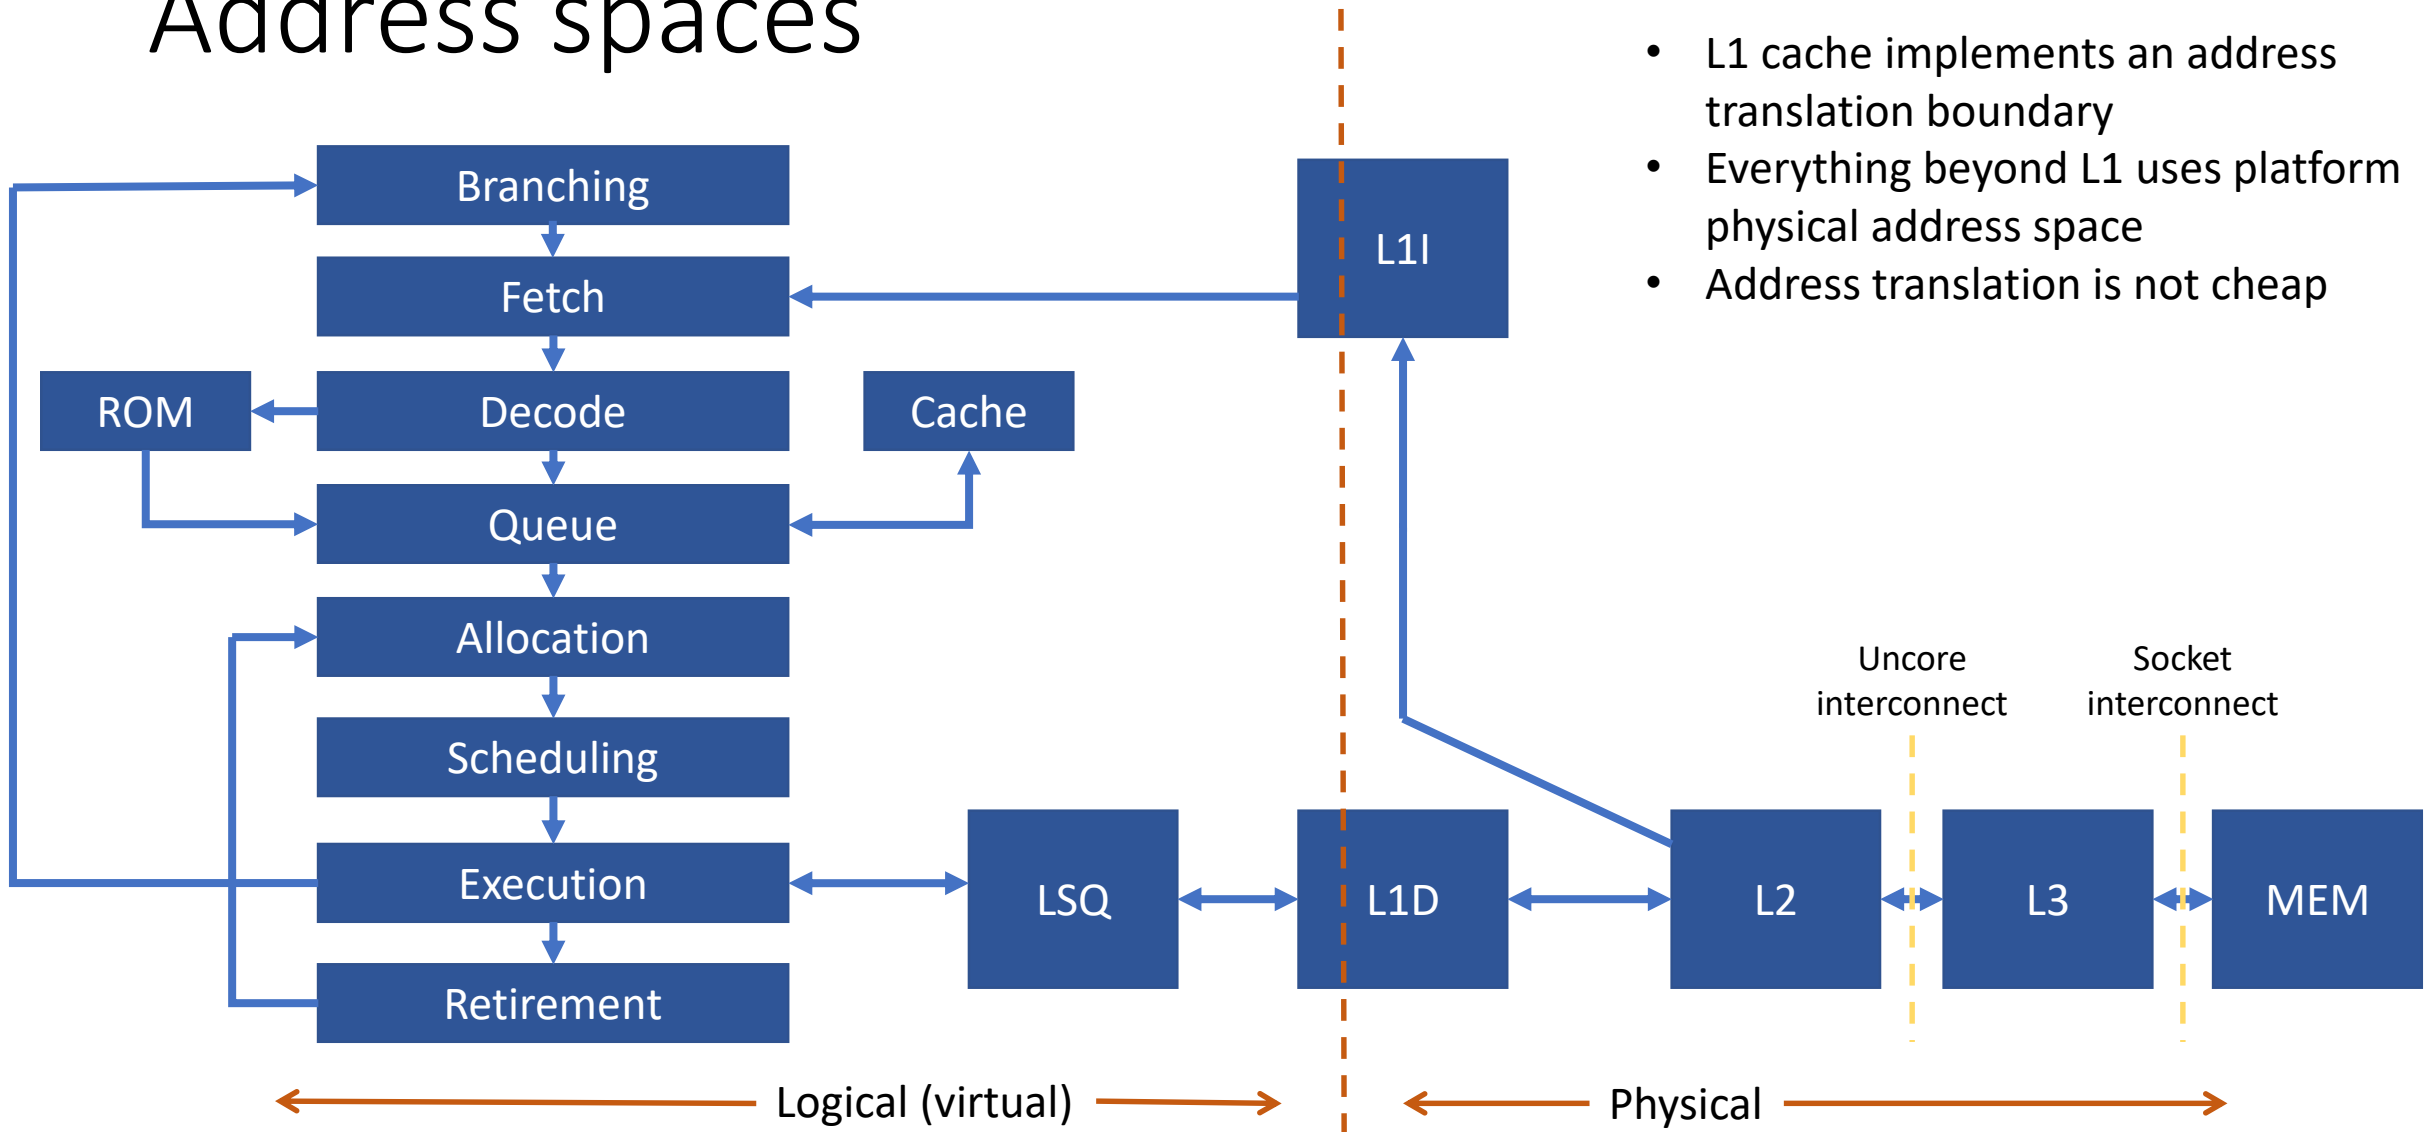


What is speculative?

- Mostly everything
- Until the results can be safely committed
- Some actions cannot be undone (easily)



Address spaces



- L1 cache implements an address translation boundary
- Everything beyond L1 uses platform physical address space
- Address translation is not cheap

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

```
48 83 FA 1B 75 09 89 F8 C1 E0 04 01 F0 EB 06 89  
F8 D1 E8 29 F0 C3 66 66 0F 1F 84 00 00 00 00 00
```

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
    cmp rdx, 27  
    jne L_0F  
  
    mov eax, edi  
    shl eax, 4  
    add eax, esi  
    jmp L_15  
  
L_0F:  
    mov eax, edi  
    shr eax, 1  
    sub eax, esi  
  
L_15:  
    ret  
    nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
    cmp rdx, 27  
    jne L_0F  
  
    mov eax, edi  
    shl eax, 4  
    add eax, esi  
    jmp L_15  
  
L_0F:  
    mov eax, edi  
    shr eax, 1  
    sub eax, esi  
  
L_15:  
    ret  
    nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
    cmp rdx, 27  
    jne L_0F  
  
    mov eax, edi  
    shl eax, 4  
    add eax, esi  
    jmp L_15  
  
L_0F:  
    mov eax, edi  
    shr eax, 1  
    sub eax, esi  
  
L_15:  
    ret  
    nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
    cmp rdx, 27  
    jne L_0F  
  
    mov eax, edi  
    shl eax, 4  
    add eax, esi  
    jmp L_15  
  
L_0F:  
    mov eax, edi  
    shr eax, 1  
    sub eax, esi  
  
L_15:  
    ret  
    nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

```
    cmp rdx, 27  
    jne +0x9  
    mov eax, edi  
    shl eax, 4
```

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

```
cmp rdx, 27  
jne +0x9  
mov eax, edi  
shl eax, 4
```

No history entry,
forward direction =>
not taken

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

```
cmp rdx, 27  
jne +0x9  
mov eax, edi  
shl eax, 4
```

```
add eax, esi  
jmp +0x6  
mov eax, edi  
shr eax, 1
```

Unaligned fetch

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
    cmp rdx, 27  
    jne L_0F  
  
    mov eax, edi  
    shl eax, 4  
    add eax, esi  
    jmp L_15  
  
L_0F:  
    mov eax, edi  
    shr eax, 1  
    sub eax, esi  
  
L_15:  
    ret  
    nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

```
    cmp rdx, 27  
    jne +0x9  
    mov eax, edi  
    shl eax, 4  
  
    add eax, esi  
    jmp +0x6  
    mov eax, edi  
    shr eax, 1
```

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

48 83 FA 1B 75 09 89 F8 C1 E0 04 01 F0 EB 06 89
F8 D1 E8 29 F0 C3 66 66 0F 1F 84 00 00 00 00 00

```
cmp rdx, 27  
jne 0x0f  
mov eax, edi  
shl eax, 4  
  
add eax, esi  
jmp 0x15  
mov eax, edi  
shr eax, 1
```

Unconditional branch
=> flush subsequent
decoded instructions

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 42);
```

48 83 FA 1B 75 09 89 F8 C1 E0 04 01 F0 EB 06 89
F8 D1 E8 29 F0 C3 66 66 0F 1F 84 00 00 00 00

```
cmp rdx, 27  
jne 0x0f  
mov eax, edi  
shl eax, 4
```

```
add eax, esi  
jmp 0x15  
mov eax, edi  
shr eax, 1
```

```
ret  
nop 10  
...
```

Unaligned fetch

Example - branching

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
    cmp rdx, 27  
    jne L_0F  
  
    mov eax, edi  
    shl eax, 4  
    add eax, esi  
    jmp L_15  
  
L_0F:  
    mov eax, edi  
    shr eax, 1  
    sub eax, esi  
  
L_15:  
    ret  
    nop 10
```


```
uint32_t res = fn1(x, y, 42);
```

48 83 FA 1B 75 09 89 F8 C1 E0 04 01 F0 EB 06 89
F8 D1 E8 29 F0 C3 66 66 0F 1F 84 00 00 00 00

```
cmp rdx, 27  
jne 0x0f  
mov eax, edi  
shl eax, 4
```

```
add eax, esi  
jmp 0x15  
mov eax, edi  
shr eax, 1
```

```
ret  
nop 10  
...
```



```
cmp rdx, 27  
jne 0x0f  
mov eax, edi  
shl eax, 4  
add eax, esi  
ret  
...
```

Flush, update branch
history, refetch

Example – branching, 2nd attempt

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
uint32_t res = fn1(x, y, 27);
```

```
48 83 FA 1B 75 09 89 F8 C1 E0 04 01 F0 EB 06 89  
F8 D1 E8 29 F0 C3 66 66 0F 1F 84 00 00 00 00 00
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```


Example – branching, 2nd attempt

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
    cmp rdx, 27  
    jne L_0F  
  
    mov eax, edi  
    shl eax, 4  
    add eax, esi  
    jmp L_15  
  
L_0F:  
    mov eax, edi  
    shr eax, 1  
    sub eax, esi  
  
L_15:  
    ret  
    nop 10
```

```
uint32_t res = fn1(x, y, 27);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

Example – branching, 2nd attempt

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 27);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

```
cmp rdx, 27  
jne L_0F  
mov eax, edi  
shl eax, 4
```

Example – branching, 2nd attempt

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 27);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

```
cmp rdx, 27  
jne L_0F  
mov eax, edi  
shl eax, 4
```

Branch history is
present from last run
=> taken

Example – branching, 2nd attempt

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 27);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

```
cmp rdx, 27  
jne L_0F  
mov eax, edi  
shl eax, 4
```

Start fetching from
branch target

Example – branching, 2nd attempt

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 27);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00

```
cmp rdx, 27  
jne L_0F  
mov eax, edi  
shl eax, 4  
  
mov eax, edi  
shr eax, 1  
sub eax, esi  
ret
```

Example – branching, 2nd attempt

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
    cmp rdx, 27  
    jne L_0F  
  
    mov eax, edi  
    shl eax, 4  
    add eax, esi  
    jmp L_15  
  
L_0F:  
    mov eax, edi  
    shr eax, 1  
    sub eax, esi  
  
L_15:  
    ret  
    nop 10
```

```
uint32_t res = fn1(x, y, 27);
```

48 83 FA 1B 75 09 89 F8 C1 E0 04 01 F0 EB 06 89
F8 D1 E8 29 F0 C3 66 66 0F 1F 84 00 00 00 00 00

```
cmp rdx, 27  
jne L_0F  
mov eax, edi  
shl eax, 4
```

```
mov eax, edi  
shr eax, 1  
sub eax, esi  
ret
```

Example – branching, 2nd attempt

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 27);
```

48 83 FA 1B 75 09 89 F8 C1 E0 04 01 F0 EB 06 89
F8 D1 E8 29 F0 C3 66 66 0F 1F 84 00 00 00 00 00

```
cmp rdx, 27  
jne L_0F  
mov eax, edi  
shl eax, 4
```

```
mov eax, edi  
shr eax, 1  
sub eax, esi  
ret
```

Flush, update
branch history,
and restart

Example – branching, two-way

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
    cmp rdx, 27  
    jne L_0F  
  
    mov eax, edi  
    shl eax, 4  
    add eax, esi  
    jmp L_15  
  
L_0F:  
    mov eax, edi  
    shr eax, 1  
    sub eax, esi  
  
L_15:  
    ret  
    nop 10
```

```
uint32_t res = fn1(x, y, 27);
```

48	83	FA	1B	75	09	89	F8	C1	E0	04	01	F0	EB	06	89
F8	D1	E8	29	F0	C3	66	66	0F	1F	84	00	00	00	00	00


```
cmp rdx, 27  
jne L_0F  
mov eax, edi  
shl eax, 4
```


Example – branching, two-way

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 27);
```




```
cmp rdx, 27  
jne L_0F  
mov eax, edi  
shl eax, 4
```

Example – branching, two-way

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
cmp rdx, 27  
jne L_0F  
  
mov eax, edi  
shl eax, 4  
add eax, esi  
jmp L_15  
  
L_0F:  
mov eax, edi  
shr eax, 1  
sub eax, esi  
  
L_15:  
ret  
nop 10
```

```
uint32_t res = fn1(x, y, 27);
```



```
cmp rdx, 27  
jne L_0F  
mov eax, edi  
shl eax, 4
```

```
mov eax, edi  
shr eax, 1  
sub eax, esi  
ret
```


Example – branching, factual

```
uint32_t fn(uint32_t x, uint32_t y, size_t cond) {  
    if (cond == 27)  
        return (x << 4) + y;  
    else  
        return (x >> 1) - y;  
}
```

```
uint32_t res = fn1(x, y, 0);
```


```
89 f8 c1 e0 04 01 f0 d1 ef 29 f7 48 83 fa 1b 0f  
45 c7 c3 66 66 66 66 2e 0f 1f 84 00 00 00 00 00
```

```
mov eax, edi  
shl eax, 4  
add eax, esi  
shr edi, 1  
sub edi, esi  
cmp rdx, 27  
cmovne eax, edi  
ret  
nop 13
```



Example - memcpy

- Problem space
- Performance requirements
- Scalar, various vectors, specialty instructions, on-core and off-core accelerators
- Data layout: both software and hardware characteristics



- Alignment: source and destination
- Size
- Direction
- Linearity

Example – memcpy: scalar naive

```
void *memcpy_scalar(char *dst, const char *src, size_t n) {  
    if(n) {  
        while (n--) {  
            *dst++ = *src++;  
        }  
    }  
    return dst;  
}
```

Scalar base ISA only, no vectorization

```
0000000000001270 <_Z13memcpy_scalarPcPKcm>:  
1270:    48 89 f8      mov     rax,rdi  
1273:    48 85 d2      test    rdx,rdx  
1276:    74 1b         je      1293 <_Z13memcpy_scalarPcPKcm+0x23>  
1278:    31 c9         xor     ecx,ecx  
127a:    66 0f 1f 44 00 00 nop     WORD PTR [rax+rax*1+0x0]  
1280:    0f b6 3c 0e   movzx   edi,BYTE PTR [rsi+rcx*1]  
1284:    40 88 3c 08   mov     BYTE PTR [rax+rcx*1],dil  
1288:    48 ff c1      inc     rcx  
128b:    48 39 ca      cmp     rdx,rcx  
128e:    75 f0         jne     1280 <_Z13memcpy_scalarPcPKcm+0x10>  
1290:    48 01 c8      add     rax,rcx  
1293:    c3           ret  
1294:    66 66 66 2e 0f 1f 84 data16  data16 nop WORD PTR cs:[rax+rax*1+0x0]  
129b:    00 00 00 00 00
```

Memory operations noticeably suboptimal
Caching system will step in, with some impact

Example – memcpy: scalar++

```
void *memcpy_scalar(char *dst, const char *src, size_t n) {
    if(n) {
        while (n--) {
            *dst++ = *src++;
        }
    }
    return dst;
}
```

Autovectorized, loops not unrolled

```
0000000000001280 <_Z13memcpy_scalarPcPKcm>:
1280: 48 85 d2          test    rdx,rdx
1283: 74 26             je      12ab <_Z13memcpy_scalarPcPKcm+0x2b>
1285: 48 83 fa 20       cmp     rdx,0x20
1289: 0f 92 c0          setb    al
128c: 48 89 f9          mov     rcx,rdi
128f: 48 29 f1          sub     rcx,rsi
1292: 48 81 f9 80 00 00 00 cmp     rcx,0x80
1299: 0f 92 c1          setb    cl
129c: 08 c1            or      cl,al
129e: 74 0f             je      12af <_Z13memcpy_scalarPcPKcm+0x2f>
12a0: 48 89 d1          mov     rcx,rdx
12a3: 49 89 f0          mov     r8,rsi
12a6: 48 89 f8          mov     rax,rdi
12a9: eb 3f            jmp     12ea <_Z13memcpy_scalarPcPKcm+0x6a>
12ab: 48 89 f8          mov     rax,rdi
12ae: c3              ret
12af: 49 89 d1          mov     r9,rdx
12b2: 49 83 e1 e0       and     r9,0xfffffffffffffffe0
12b6: 89 d1            mov     ecx,edx
12b8: 83 e1 1f         and     ecx,0x1f
12bb: 4e 8d 04 0e       lea     r8,[rsi+r9*1]
12bf: 4a 8d 04 0f       lea     rax,[rdi+r9*1]
12c3: 45 31 d2          xor     r10d,r10d
12c6: 66 2e 0f 1f 84 00 00 nop     WORD PTR cs:[rax+rax*1+0x0]
12cd: 00 00 00
12d0: c4 a1 7c 10 04 16 vmovups ymm0,YMMWORD PTR [rsi+r10*1]
12d6: c4 a1 7c 11 04 17 vmovups YMMWORD PTR [rdi+r10*1],ymm0
12dc: 49 83 c2 20       add     r10,0x20
12e0: 4d 39 d1          cmp     r9,r10
12e3: 75 eb            jne     12d0 <_Z13memcpy_scalarPcPKcm+0x50>
12e5: 49 39 d1          cmp     r9,rdx
12e8: 74 1a            je      1304 <_Z13memcpy_scalarPcPKcm+0x84>
12ea: 31 d2            xor     edx,edx
12ec: 0f 1f 40 00       nop     DWORD PTR [rax+0x0]
12f0: 41 0f b6 34 10    movzx   esi,BYTE PTR [r8+rdx*1]
12f5: 40 88 34 10       mov     BYTE PTR [rax+rdx*1],sil
12f9: 48 ff c2          inc     rdx
12fc: 48 39 d1          cmp     rcx,rdx
12ff: 75 ef            jne     12f0 <_Z13memcpy_scalarPcPKcm+0x70>
1301: 48 01 d0          add     rax,rdx
1304: c5 f8 77          vzeroupper
1307: c3              ret
1308: 0f 1f 84 00 00 00 00 nop     DWORD PTR [rax+rax*1+0x0]
130f: 00
```

Example – memcpy: handcrafted

```
void *memcpy_ch(void * __restrict dst_, const void * __restrict src_, size_t size) {  
  
    char * __restrict dst = reinterpret_cast<char * __restrict>(dst_);  
    const char * __restrict src = reinterpret_cast<const char * __restrict>(src_);  
  
    void * ret = dst;  
  
tail:  
  
    if (size <= 16)  
    {  
        if (size >= 8)  
        {  
            __builtin_memcpy(dst + size - 8, src + size - 8, 8);  
            __builtin_memcpy(dst, src, 8);  
        }  
        else if (size >= 4)  
        {  
            __builtin_memcpy(dst + size - 4, src + size - 4, 4);  
            __builtin_memcpy(dst, src, 4);  
        }  
        else if (size >= 2)  
        {  
            __builtin_memcpy(dst + size - 2, src + size - 2, 2);  
            __builtin_memcpy(dst, src, 2);  
        }  
        else if (size >= 1)  
        {  
            *dst = *src;  
        }  
    }  
    else  
    {  
        if (size <= 128)  
        {  
            __mm_storeu_si128(reinterpret_cast<__m128i *>(dst + size - 16), __mm_loadu_si128(reinterpret_cast<const __m128i *>(src + size - 16)));  
        }  
    }  
}
```

Factual Clickhouse
implementation as an
example

Example – memcpy: handcrafted

```
void *memcpy_ch(void * __restrict dst_, const void * __restrict src_, size_t size) {  
  
    char * __restrict dst = reinterpret_cast<char * __restrict>(dst_);  
    const char * __restrict src = reinterpret_cast<const char * __restrict>(src_);  
  
    while (size > 16)  
    {  
        _mm_storeu_si128(reinterpret_cast<__m128i *>(dst), _mm_loadu_si128(reinterpret_cast<const __m128i *>(src)));  
        dst += 16;  
        src += 16;  
        size -= 16;  
    }  
    else  
    {  
        size_t padding = (16 - (reinterpret_cast<size_t>(dst) & 15)) & 15;  
  
        if (padding > 0)  
        {  
            __m128i head = _mm_loadu_si128(reinterpret_cast<const __m128i*>(src));  
            _mm_storeu_si128(reinterpret_cast<__m128i*>(dst), head);  
            dst += padding;  
            src += padding;  
            size -= padding;  
        }  
  
        __m128i c0, c1, c2, c3, c4, c5, c6, c7;  
  
        while (size >= 128)  
        {  
            c0 = _mm_loadu_si128(reinterpret_cast<const __m128i*>(src) + 0);  
            c1 = _mm_loadu_si128(reinterpret_cast<const __m128i*>(src) + 1);  
            c2 = _mm_loadu_si128(reinterpret_cast<const __m128i*>(src) + 2);  
            c3 = _mm_loadu_si128(reinterpret_cast<const __m128i*>(src) + 3);  
            c4 = _mm_loadu_si128(reinterpret_cast<const __m128i*>(src) + 4);  
            c5 = _mm_loadu_si128(reinterpret_cast<const __m128i*>(src) + 5);  
            c6 = _mm_loadu_si128(reinterpret_cast<const __m128i*>(src) + 6);  
            c7 = _mm_loadu_si128(reinterpret_cast<const __m128i*>(src) + 7);  
            src += 128;  
        }  
    }  
}
```

Factual Clickhouse
implementation as an
example

Example – memcpy: handcrafted

```
void *memcpy_ch(void * __restrict dst_, const void * __restrict src_, size_t size) {  
    char * __restrict dst = reinterpret_cast<char * __restrict>(dst_);  
    const char * __restrict src = reinterpret_cast<const char * __restrict>(src_);  
    while (size > 16)  
    {  
        _mm_storeu_si128(reinterpret_cast<__m128i *>(dst), _mm_loadu_si128(reinterpret_cast<const __m128i *>(src)));  
        _mm_store_si128(reinterpret_cast<__m128i *>(dst) + 0, c0);  
        _mm_store_si128(reinterpret_cast<__m128i *>(dst) + 1, c1);  
        _mm_store_si128(reinterpret_cast<__m128i *>(dst) + 2, c2);  
        _mm_store_si128(reinterpret_cast<__m128i *>(dst) + 3, c3);  
        _mm_store_si128(reinterpret_cast<__m128i *>(dst) + 4, c4);  
        _mm_store_si128(reinterpret_cast<__m128i *>(dst) + 5, c5);  
        _mm_store_si128(reinterpret_cast<__m128i *>(dst) + 6, c6);  
        _mm_store_si128(reinterpret_cast<__m128i *>(dst) + 7, c7);  
        dst += 128;  
        size -= 128;  
    }  
    goto tail;  
}  
tail:  
    return ret;  
}  
}
```

```
c1 = _mm_loadu_si128(reinterpret_cast<const __m128i *>(src) + 1);  
c2 = _mm_loadu_si128(reinterpret_cast<const __m128i *>(src) + 2);  
c3 = _mm_loadu_si128(reinterpret_cast<const __m128i *>(src) + 3);  
c4 = _mm_loadu_si128(reinterpret_cast<const __m128i *>(src) + 4);  
c5 = _mm_loadu_si128(reinterpret_cast<const __m128i *>(src) + 5);  
c6 = _mm_loadu_si128(reinterpret_cast<const __m128i *>(src) + 6);  
c7 = _mm_loadu_si128(reinterpret_cast<const __m128i *>(src) + 7);  
src += 128;
```

Factual Clickhouse
implementation as an
example

Example – memcpy: handcrafted

```

0000000000001310 <_Z9memcpy_chPvPKvm>:
1310: 53          push    rbp
1311: 48 89 f8    mov     rax,rdi
1314: 48 83 fa 11 cmp     rdx,0x11
1318: 73 2c      jae     1346 <_Z9memcpy_chPvPKvm+0x36>
131a: 48 83 fa 08 cmp     rdx,0x8
131e: 0f 82 96 00 00 00 jbe     13ba <_Z9memcpy_chPvPKvm+0xaa>
1324: 48 8b 4c 16 f8 mov     rcx,QWORD PTR [rsi+rdx*1-0x8]
1329: 48 89 4c 17 f8 mov     QWORD PTR [rdi+rdx*1-0x8],rcx
132e: 48 8b 0e    mov     rcx,QWORD PTR [rsi]
1331: 48 89 0f    mov     QWORD PTR [rdi],rcx
1334: 5b        pop     rbp
1335: c5 f8 77   vzeroupper
1338: c3        ret
1339: 0f 1f 80 00 00 00 00 nop     DWORD PTR [rax+0x0]
1340: 48 83 fa 10 cmp     rdx,0x10
1344: 76 d4     jbe     131a <_Z9memcpy_chPvPKvm+0xa>
1346: 48 81 fa 80 00 00 00 cmp     rdx,0x80
134d: 0f 86 7e 00 00 00 jbe     13d1 <_Z9memcpy_chPvPKvm+0xc1>
1353: 89 f9     mov     ecx,edi
1355: f7 d9     neg     ecx
1357: 83 e1 0f   and     ecx,0xf
135a: 74 24     je      1380 <_Z9memcpy_chPvPKvm+0x70>
135c: c5 f8 10 06 vmovups xmm0,XMMWORD PTR [rsi]
1360: c5 f8 11 07 vmovups XMMWORD PTR [rdi],xmm0
1364: 48 01 cf   add     rdi,rcx
1367: 48 01 ce   add     rsi,rcx
136a: 48 29 ca   sub     rdx,rcx
136d: 48 81 fa 80 00 00 00 cmp     rdx,0x80
1374: 72 ca     jb      1340 <_Z9memcpy_chPvPKvm+0x30>
1376: 66 2e 0f 1f 84 00 00 nop     WORD PTR cs:[rax+rax*1+0x0]
137d: 00 00 00
1380: c5 fc 10 06 vmovups ymm0,YMMWORD PTR [rsi]
1384: c5 fc 10 4e 20 vmovups ymm1,YMMWORD PTR [rsi+0x20]
1389: c5 fc 10 56 40 vmovups ymm2,YMMWORD PTR [rsi+0x40]
138e: c5 fc 10 5e 60 vmovups ymm3,YMMWORD PTR [rsi+0x60]
1393: c5 fc 11 07 vmovups YMMWORD PTR [rdi],ymm0
1397: c5 fc 11 4f 20 vmovups YMMWORD PTR [rdi+0x20],ymm1
139c: c5 fc 11 57 40 vmovups YMMWORD PTR [rdi+0x40],ymm2
13a1: 48 83 ee 80 sub     rsi,0xfffffffffffffff80
13a5: c5 fc 11 5f 60 vmovups YMMWORD PTR [rdi+0x60],ymm3
13aa: 48 83 ef 80 sub     rdi,0xfffffffffffffff80
13ae: 48 83 c2 80 add     rdx,0xfffffffffffffff80
13b2: 48 83 fa 7f cmp     rdx,0x7f
13b6: 77 c8     ja      1380 <_Z9memcpy_chPvPKvm+0x70>
13b8: eb 86     jmp     1340 <_Z9memcpy_chPvPKvm+0x30>
13ba: 48 83 fa 04 cmp     rdx,0x4
13be: 72 39     jb      13f9 <_Z9memcpy_chPvPKvm+0xe9>
13c0: 8b 4c 16 f8 mov     ecx,DWORD PTR [rsi+rdx*1-0x4]
13c4: 89 4c 17 f8 mov     DWORD PTR [rdi+rdx*1-0x4],ecx
13c8: 8b 0e     mov     ecx,DWORD PTR [rsi]
13ca: 89 0f     mov     DWORD PTR [rdi],ecx
13cc: 5b        pop     rbp
13cd: c5 f8 77   vzeroupper
13d0: c3        ret
13d1: c5 f8 10 44 16 f0 vmovups xmm0,XMMWORD PTR [rsi+rdx*1-0x10]
13d7: c5 f8 11 44 17 f0 vmovups XMMWORD PTR [rdi+rdx*1-0x10],xmm0
13dd: 48 83 c2 ef add     rdx,0xfffffffffffffffef
13e1: 48 83 e2 f0 and     rdx,0xfffffffffffffff0
13e5: 48 83 c2 10 add     rdx,0x10
13e9: 48 89 c3    mov     rbx,rax
13ec: c5 f8 77   vzeroupper

```

```

13cd: c5 f8 77   vzeroupper
13d0: c3        ret
13d1: c5 f8 10 44 16 f0 vmovups xmm0,XMMWORD PTR [rsi+rdx*1-0x10]
13d7: c5 f8 11 44 17 f0 vmovups XMMWORD PTR [rdi+rdx*1-0x10],xmm0
13dd: 48 83 c2 ef add     rdx,0xfffffffffffffffef
13e1: 48 83 e2 f0 and     rdx,0xfffffffffffffff0
13e5: 48 83 c2 10 add     rdx,0x10
13e9: 48 89 c3    mov     rbx,rax
13ec: c5 f8 77   vzeroupper
13ef: e8 3c fc ff ff call    1030 <memcpy@plt>
13f4: 48 89 d8    mov     rax,rbx
13f7: 5b        pop     rbp
13f8: c3        ret
13f9: 48 83 fa 02 cmp     rdx,0x2
13fd: 72 15     jb      1414 <_Z9memcpy_chPvPKvm+0x104>
13ff: 0f b7 4c 16 fe movzx   ecx,WORD PTR [rsi+rdx*1-0x2]
1404: 66 89 4c 17 fe mov     WORD PTR [rdi+rdx*1-0x2],cx
1409: 0f b7 0e   movzx   ecx,WORD PTR [rsi]
140c: 66 89 0f   mov     WORD PTR [rdi],cx
140f: 5b        pop     rbp
1410: c5 f8 77   vzeroupper
1413: c3        ret
1414: 48 85 d2   test    rdx,rdx
1417: 74 05     je      141e <_Z9memcpy_chPvPKvm+0x10e>
1419: 0f b6 0e   movzx   ecx,BYTE PTR [rsi]
141c: 88 0f     mov     BYTE PTR [rdi],cl
141e: 5b        pop     rbp
141f: c5 f8 77   vzeroupper
1422: c3        ret
1423: 66 66 66 66 2e 0f 1f data16 data16 data16 nop WORD PTR cs:[rax+rax*1+0x0]
142a: 84 00 00 00 00 00

```

Re-autovectorized, loops unrolled
Still predominantly unaligned

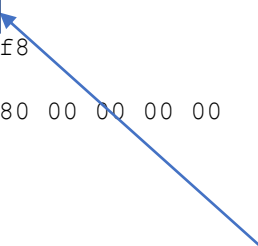
Factual Clickhouse
implementation as an
example

Example – memcpy: ERMS

```
void *memcpy_erms(void *dst, const void *src, size_t n) {  
    asm volatile ("rep movsb"  
        : "=D" (dst), "=S" (src), "=c" (n)  
        : "0" (dst), "1" (src), "2" (n)  
        : "memory");  
  
    return dst;  
}
```

Let the HW do the right thing

```
00000000000000001260 <_Z11memcpy_ermsPvPKvm>:  
1260: 48 89 d1                mov     rcx,rdx  
1263: f3 a4                  rep movs BYTE PTR es:[rdi],BYTE PTR ds:[rsi]  
1265: 48 89 f8                mov     rax,rdi  
1268: c3                     ret  
1269: 0f 1f 80 00 00 00 00    nop     DWORD PTR [rax+0x0]
```



- Microcoded subroutine
- Aware of platform specifics
- Has insight into internal machine state

Example – memcpy: ERMS

```
void *memcpy_erms(void *dst, const void *src, size_t n) {  
    asm volatile ("rep movsb"  
        : "=D" (dst), "=S" (src), "=c" (n)  
        : "0" (dst), "1" (src), "2" (n)  
        : "memory");  
  
    return dst;  
}
```

Let the HW do the right thing

```
00000000000001260 <_Z11memcpy_ermsPvPKvm>:  
1260: 48 89 d1                mov     rcx,rdx  
1263: f3 a4                  rep movs BYTE PTR es:[rdi],BYTE PTR ds:[rsi]  
1265: 48 89 f8                mov     rax,rdi  
1268: c3                     ret  
1269: 0f 1f 80 00 00 00 00    nop     DWORD PTR [rax+0x0]
```

- Microcoded subroutine
- Aware of platform specifics
- Has insight into internal machine state

Will this universally be the best option?
Only representative performance testing will show

Example - hashing

```
uint32_t hash_fnv1a(const char* key, const size_t len) {  
  
    uint32_t hash = 0x811c9dc5;  
  
    for(size_t i = 0; i < len; ++i) {  
        hash = hash ^ key[i];  
        hash *= 0x1000193;;  
    }  
  
    return hash;  
}
```

```
b8 c5 9d 1c 81      mov eax, 0x811c9dc5  
48 85 f6            test rsi, rsi  
74 1a              je L11c4  
31 c9              xor ecx, ecx  
0f 1f 40 00        nop 4  
                  L11b0:  
0f b6 14 0f        movzx edx, BYTE PTR [rdi+rcx*1]  
31 c2              xor edx, eax  
69 c2 93 01 00 01  imul eax, edx, 0x1000193  
48 ff c1           inc rcx  
48 39 ce           cmp rsi, rcx  
75 ec             jne L11b0  
                  L11c4:  
c3                ret  
66 66 2e 0f 1f 84 00  nop 11  
00 00 00 00
```

Example - hashing

```
uint32_t hash_fnv1a(const char* key, const size_t len) {  
  
    uint32_t hash = 0x811c9dc5;  
  
    for(size_t i = 0; i < len; ++i) {  
        hash = hash ^ key[i];  
        hash *= 0x1000193;;  
    }  
  
    return hash;  
}
```

```
b8 c5 9d 1c 81      mov eax, 0x811c9dc5  
48 85 f6            test rsi, rsi  
74 1a               je L11c4  
31 c9               xor ecx, ecx  
0f 1f 40 00          nop 4  
L11b0:  
0f b6 14 0f          movzx edx, BYTE PTR [rdi+rcx*1]  
31 c2               xor edx, eax  
69 c2 93 01 00 01    imul eax, edx, 0x1000193  
48 ff c1            inc rcx  
48 39 ce             cmp rsi, rcx  
75 ec               jne L11b0  
L11c4:  
c3                  ret  
66 66 2e 0f 1f 84 00  nop 11  
00 00 00 00
```

Candidate for fusion



Example - hashing

```
uint32_t hash_fnv1a(const char* key, const size_t len) {  
  
    uint32_t hash = 0x811c9dc5;  
  
    for(size_t i = 0; i < len; ++i) {  
        hash = hash ^ key[i];  
        hash *= 0x1000193;  
    }  
  
    return hash;  
}
```

```
b8 c5 9d 1c 81      mov eax, 0x811c9dc5  
48 85 f6            test rsi, rsi  
74 1a              je L11c4  
31 c9              xor ecx, ecx  
0f 1f 40 00        nop 4  
  
L11b0:  
0f b6 14 0f        movzx edx, BYTE PTR [rdi+rcx*1]  
31 c2              xor edx, eax  
69 c2 93 01 00 01  imul eax, edx, 0x1000193  
48 ff c1           inc rcx  
48 39 ce           cmp rsi, rcx  
75 ec             jne L11b0  
  
L11c4:  
c3                ret  
66 66 2e 0f 1f 84 00  nop 11  
00 00 00 00
```

```
mov eax, 0x811c9dc5  
test rsi, rsi  
je L11c4
```

```
movzx edx, BYTE PTR [rdi+rcx*1]  
xor edx, eax  
imul eax, edx, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
movzx edx, BYTE PTR [rdi+rcx*1]  
xor edx, eax  
imul eax, edx, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
movzx edx, BYTE PTR [rdi+rcx*1]  
xor edx, eax  
imul eax, edx, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
movzx edx, BYTE PTR [rdi+rcx*1]  
xor edx, eax  
imul eax, edx, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
movzx edx, BYTE PTR [rdi+rcx*1]  
xor edx, eax  
imul eax, edx, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
ret
```

Several iterations run
ahead speculatively

Example - hashing

```
uint32_t hash_fnv1a(const char* key, const size_t len) {  
  
    uint32_t hash = 0x811c9dc5;  
  
    for(size_t i = 0; i < len; ++i) {  
        hash = hash ^ key[i];  
        hash *= 0x1000193;;  
    }  
  
    return hash;  
}
```

```
b8 c5 9d 1c 81      mov eax, 0x811c9dc5  
48 85 f6            test rsi, rsi  
74 1a              je L11c4  
31 c9              xor ecx, ecx  
0f 1f 40 00        nop 4  
L11b0:  
0f b6 14 0f        movzx edx, BYTE PTR [rdi+rcx*1]  
31 c2              xor edx, eax  
69 c2 93 01 00 01  imul eax, edx, 0x1000193  
48 ff c1           inc rcx  
48 39 ce           cmp rsi, rcx  
75 ec             jne L11b0  
L11c4:  
c3                ret  
66 66 2e 0f 1f 84 00  nop 11  
00 00 00 00
```

```
mov eax, 0x811c9dc5  
test rsi, rsi  
je L11c4
```

```
movzx r100b, BYTE PTR [rdi+rcx*1]  
xor r100d, eax  
imul eax, r100d, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
movzx r101b, BYTE PTR [rdi+rcx*1]  
xor r101d, eax  
imul eax, r101d, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
movzx r102b, BYTE PTR [rdi+rcx*1]  
xor r102d, eax  
imul eax, r102d, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
movzx edx, BYTE PTR [rdi+rcx*1]  
xor edx, eax  
imul eax, edx, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
movzx edx, BYTE PTR [rdi+rcx*1]  
xor edx, eax  
imul eax, edx, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
ret
```

Several iterations run
ahead speculatively

And out of program
order

Example - hashing

```
uint32_t hash_fnv1a(const char* key, const size_t len) {  
  
    uint32_t hash = 0x811c9dc5;  
  
    for(size_t i = 0; i < len; ++i) {  
        hash = hash ^ key[i];  
        hash *= 0x1000193;;  
    }  
  
    return hash;  
}
```

```
b8 c5 9d 1c 81      mov eax, 0x811c9dc5  
48 85 f6            test rsi, rsi  
74 1a              je L11c4  
31 c9              xor ecx, ecx  
0f 1f 40 00        nop 4  
  
L11b0:  
0f b6 14 0f        movzx edx, BYTE PTR [rdi+rcx*1]  
31 c2              xor edx, eax  
69 c2 93 01 00 01  imul eax, edx, 0x1000193  
48 ff c1           inc rcx  
48 39 ce           cmp rsi, rcx  
75 ec             jne L11b0  
  
L11c4:  
c3                ret  
66 66 2e 0f 1f 84 00  nop 11  
00 00 00 00
```

```
mov eax, 0x811c9dc5  
test rsi, rsi  
je L11c4
```

```
movzx r100b, BYTE PTR [rdi+rcx_100*1]  
inc rcx_100  
movzx r101b, BYTE PTR [rdi+rcx_101*1] ← Possible OOB!  
inc rcx_101  
movzx r102b, BYTE PTR [rdi+rcx_102*1] ← Possible OOB!  
inc rcx_102
```

```
xor r100d, eax  
imul eax, r100d, 0x1000193  
cmpjne rsi, rcx_100, L11b0
```

```
xor r101d, eax  
imul eax, r101d, 0x1000193  
cmpjne rsi, rcx_101, L11b0
```

```
xor r102d, eax  
imul eax, r102d, 0x1000193  
cmpjne rsi, rcx_102, L11b0
```

```
movzx edx, BYTE PTR [rdi+rcx*1] ← Possible OOB!  
xor edx, eax  
imul eax, edx, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
movzx edx, BYTE PTR [rdi+rcx*1] ← Possible OOB!  
xor edx, eax  
imul eax, edx, 0x1000193  
inc rcx  
cmpjne rsi, rcx, L11b0
```

```
ret
```

Example - hashing

```
struct kv {
    uint32_t key1;
    uint16_t key2;
    uint64_t key3;
    void *value;
} kv;

uint32_t hash_ext(struct kv *kv) {

    uint32_t fnv1a_hash {};

    fnv1a_hash += hash_fnv1a_ext((const char *)&kv->key1, sizeof(kv->key1));
    fnv1a_hash += hash_fnv1a_ext((const char *)&kv->key2, sizeof(kv->key2));
    fnv1a_hash += hash_fnv1a_ext((const char *)&kv->key3, sizeof(kv->key3));

    return fnv1a_hash;
}
```

Hash function implementation in a separate translation unit.

```
000000000000012a0 <_Z4hash_extP2kv>:
    12a0:    55                push    rbp
    12a1:    41 56            push    r14
    12a3:    53                push    rbx
    12a4:    48 89 fb         mov     rbx,rdi
    12a7:    be 04 00 00 00   mov     esi,0x4
    12ac:    e8 7f fe ff ff   call    1130 <_Z14hash_fnv1a_extPKcm>
    12b1:    89 c5            mov     ebp,eax
    12b3:    48 8d 7b 04      lea     rdi,[rbx+0x4]
    12b7:    be 02 00 00 00   mov     esi,0x2
    12bc:    e8 6f fe ff ff   call    1130 <_Z14hash_fnv1a_extPKcm>
    12c1:    41 89 c6         mov     r14d,eax
    12c4:    41 01 ee         add     r14d,ebp
    12c7:    48 83 c3 08      add     rbx,0x8
    12cb:    be 08 00 00 00   mov     esi,0x8
    12d0:    48 89 df         mov     rdi,rbx
    12d3:    e8 58 fe ff ff   call    1130 <_Z14hash_fnv1a_extPKcm>
    12d8:    44 01 f0         add     eax,r14d
    12db:    5b                pop     rbx
    12dc:    41 5e            pop     r14
    12de:    5d                pop     rbp
    12df:    c3                ret
```

```
0000000000000000 <_Z14hash_fnv1a_extPKcm>:
    0:    b8 c5 9d 1c 81   mov     eax,0x811c9dc5
    5:    48 85 f6         test    rsi,rsi
    8:    74 1a            je      24 <_Z14hash_fnv1a_extPKcm+0x24>
    a:    31 c9            xor     ecx,ecx
    c:    0f 1f 40 00      nop     DWORD PTR [rax+0x0]
   10:    0f be 14 0f      movsx   edx,BYTE PTR [rdi+rcx*1]
   14:    31 c2            xor     edx,eax
   16:    69 c2 93 01 00 01 imul    eax,edx,0x1000193
   1c:    48 ff c1         inc     rcx
   1f:    48 39 ce         cmp     rsi,rcx
   22:    75 ec            jne     10 <_Z14hash_fnv1a_extPKcm+0x10>
   24:    c3                ret
```

Example - hashing

```
struct kv {
    uint32_t key1;
    uint16_t key2;
    uint64_t key3;
    void *value;
} kv;

uint32_t hash_int(struct kv *kv) {

    uint32_t fnv1a_hash {};

    fnv1a_hash += hash_fnv1a_int((const char *)&kv->key1, sizeof(kv->key1));
    fnv1a_hash += hash_fnv1a_int((const char *)&kv->key2, sizeof(kv->key2));
    fnv1a_hash += hash_fnv1a_int((const char *)&kv->key3, sizeof(kv->key3));

    return fnv1a_hash;
}
```

Hash function implementation in the same translation unit.

```
00000000000012a0 <_Z4hash_intP2kv>:
    12a0:    b9 c5 9d 1c 81      mov     ecx,0x811c9dc5
    12a5:    31 c0               xor     eax,eax
    12a7:    66 0f 1f 84 00 00 00 nop     WORD PTR [rax+rax*1+0x0]
    12ae:    00 00
    12b0:    0f b6 14 07      movzx   edx,BYTE PTR [rdi+rax*1]
    12b4:    31 ca             xor     edx,ecx
    12b6:    69 ca 93 01 00 01  imul    ecx,edx,0x1000193
    12bc:    48 ff c0         inc     rax
    12bf:    48 83 f8 04      cmp     rax,0x4
    12c3:    75 eb             jne     12b0 <_Z4hashP2kv+0x10>
    12c5:    b8 c5 9d 1c 81      mov     eax,0x811c9dc5
    12ca:    31 d2             xor     edx,edx
    12cc:    0f 1f 40 00      nop     DWORD PTR [rax+0x0]
    12d0:    0f b6 74 17 04      movzx   esi,BYTE PTR [rdi+rdx*1+0x4]
    12d5:    31 c6             xor     esi,eax
    12d7:    69 c6 93 01 00 01  imul    eax,esi,0x1000193
    12dd:    48 ff c2         inc     rdx
    12e0:    48 83 fa 02      cmp     rdx,0x2
    12e4:    75 ea             jne     12d0 <_Z4hashP2kv+0x30>
    12e6:    ba c5 9d 1c 81      mov     edx,0x811c9dc5
    12eb:    31 f6             xor     esi,esi
    12ed:    0f 1f 00      nop     DWORD PTR [rax]
    12f0:    44 0f b6 44 37 08  movzx   r8d,BYTE PTR [rdi+rsi*1+0x8]
    12f6:    41 31 d0         xor     r8d,edx
    12f9:    41 69 d0 93 01 00 01 imul    edx,r8d,0x1000193
    1300:    48 ff c6         inc     rsi
    1303:    48 83 fe 08      cmp     rsi,0x8
    1307:    75 e7             jne     12f0 <_Z4hashP2kv+0x50>
    1309:    01 c8             add     eax,ecx
    130b:    01 d0             add     eax,edx
    130d:    c3               ret
    130e:    66 90             xchg    ax,ax
```

Example - hashing

```
#define __packed __attribute__((packed))

struct kv {
    uint32_t key1;
    uint16_t key2;
    uint64_t key3;
    void *value;
} __packed kv;

uint32_t hash_packed(struct kv *kv) {

    uint32_t fnv1a_hash {};

    fnv1a_hash += hash_fnv1a_ref((const char *)&kv,
                                sizeof(kv->key1) +
                                sizeof(kv->key2) +
                                sizeof(kv->key3));

    return fnv1a_hash;
}
```

```
00000000000001310 <_Z11hash_packedP2kv>:
1310: 48 89 7c 24 f8      mov     QWORD PTR [rsp-0x8],rdi
1315: b8 c5 9d 1c 81      mov     eax,0x811c9dc5
131a: 31 c9              xor     ecx,ecx
131c: 0f 1f 40 00        nop     DWORD PTR [rax+0x0]
1320: 0f b6 54 0c f8      movzx   edx,BYTE PTR [rsp+rcx*1-0x8]
1325: 31 c2              xor     edx,eax
1327: 69 c2 93 01 00 01   imul    eax,edx,0x1000193
132d: 48 ff c1          inc     rcx
1330: 48 83 f9 0e        cmp     rcx,0xe
1334: 75 ea             jne     1320 <_Z11hash_packedP2kv+0x10>
1336: c3               ret
1337: 66 0f 1f 84 00 00 00 nop     WORD PTR [rax+rax*1+0x0]
133e: 00 00
```

Rearranged data structure

Measuring performance

- All of this is very platform specific
- Core and uncore performance counters
- Granularity, resolution, value limits
- Not everything at once
- Microarchitectural microbenchmarks are fun
- Although nowhere near being easy
- And might not matter that much in a global scope
- Fleetwide benchmarks do matter though

Universal ISA

- Comparing x86, ARM, R-V, some historical others
- RISC vs CISC debate, and VLIW as well
- 32 → 31 → even less registers
- L/S vs R/M
- Flags
- Operand count
- Destructive destination
- Spilling strategy
- ISA as a contract

Vectorization?

- The logic of a scalar algorithm applied to a multitude of separate sets of data.
 - Performance characterization in terms of latency and throughput.
 - SIMD as a specific instantiation of vectorization approach.
-
- Practical and Everyday => relies on the compiler. But we are not there yet. Not certain whether we will be at all.

Vectorization in 20 minutes


In order to vectorize scalar code successfully, horizontal operations are mandatory.
Control flow dependencies need to be translated to data flow dependencies.
Data structure layout should not conflict with native vector width and length.

- Controllable/maskable access to memory
- Assembling a vector register from scalar values (insert/extract)
- Mixing multiple vectors into one (blend)
- Swapping blocks and elements within a vector (permute, shuffle)
- Nonlinear and indexed memory access (scatter, gather)
- Conditional execution (predicates, masking)

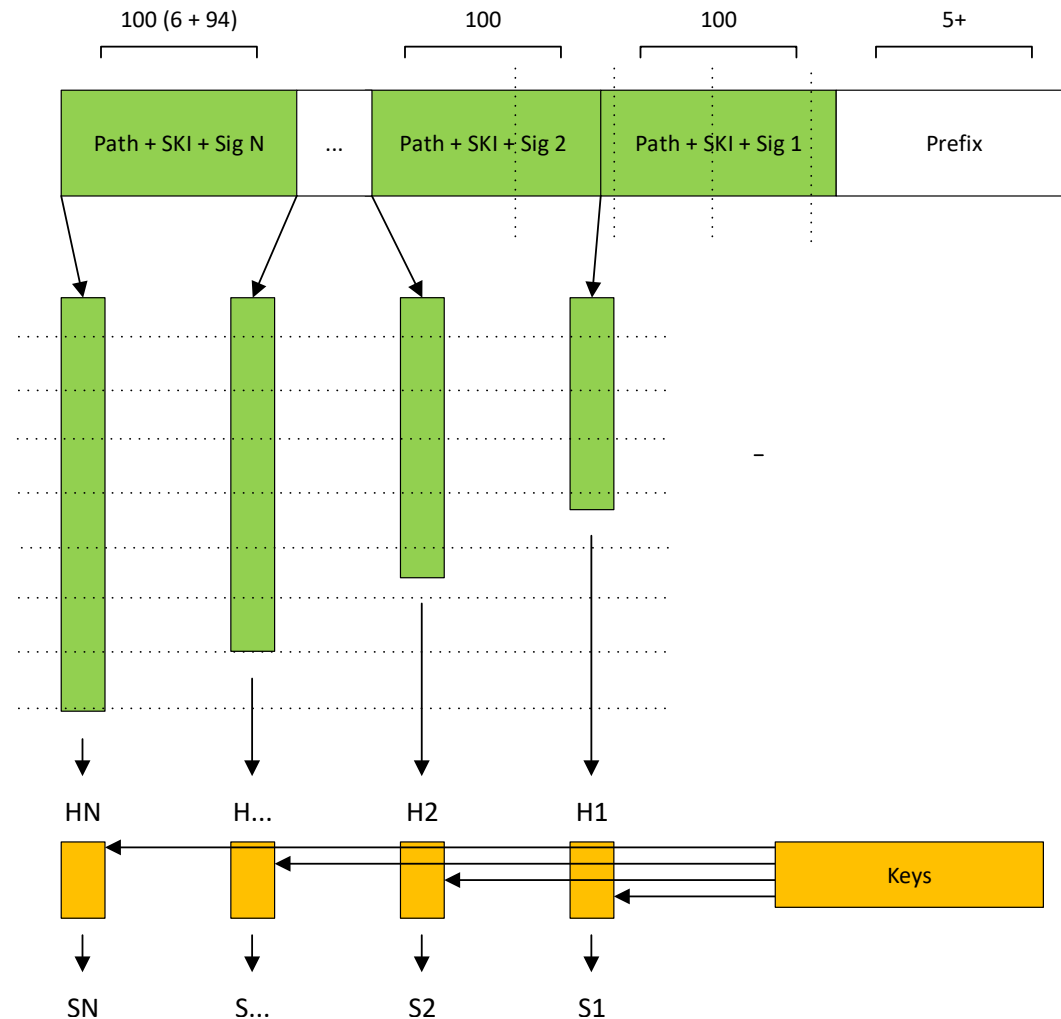
Vectorization Example – SHA2

- Block-based cryptographic hash function.
- Simple bitwise operations – many of them.
- Instruction set equivalence may not be assumed.
- Vertical vs horizontal data layout.

	Latency	Throughput
Scalar	1	1
Accelerated	4	2.2
Vectorized	0.85	14



Vectorized SHA2 and P-256 ECDSA



Linear code block operating on different data sets in parallel

Hash multiple blocks in parallel
Sign/verify multiple hashes/signatures in parallel

Vector lanes of fixed width

Gather operations place significant restrictions on data format

+20% latency results in +1500% throughput

Only if data structures allow!

Why bother?

Compiler is smart, one just needs to specify a correct command-line option, no?

```
for (int i = 0; i < m; i++) {  
    for (int j = 0; j < n; j++) {  
        for (int p = 0; p < k; p++) {  
            C(i, j) += A(i, p) * B(p, j);  
        }  
    }  
}
```

For vertical operations that is indeed not a complex task to do, and compilers perform just fine.

Why bother?

What if the level of triviality is reduced a little?

- Control flow dependencies.
- Different lane widths and parameters
- Branching
- Iterations of different length

```
missed: not vectorized: complicated access pattern.  
missed: not vectorized: no vectype for stmt: sum[0] = { 0, 0, 0, 0 };  
missed: couldn't vectorize loop
```

```
uint32_t CityHash32(const char *s, size_t len) {  
  
    if (len <= 24) {  
        return len <= 12 ? (len <= 4 ?  
            Hash32Len0to4(s, len) :  
            Hash32Len5to12(s, len)) :  
            Hash32Len13to24(s, len);  
    }  
  
    h ^= a0;  
    h = h * 5 + 0xe6546b64;  
    h ^= a2;  
  
    do {  
        h ^= a0;  
        h = Rotate32(h, 18);  
        h = h * 5 + 0xe6546b64;  
        f += a1;  
        s += 20;  
    } while (--iters != 0);  
  
    h = h * 5 + 0xe6546b64;  
    h = Rotate32(h, 17) * c1;  
    h = Rotate32(h + f, 19);  
    h = h * 5 + 0xe6546b64;  
    h = Rotate32(h, 17) * c1;  
  
    return h;  
}
```

Historical Perspective

MMX (1996)	8 64-bit integer registers 8/16/32/64, saturation, two-operand.
SSE (1999):	8 128-bit int and fp registers, three-operand, some domain-specific accelerations.
AVX (2011):	16 256-bit fp-only registers, 128-bit lanes, 32/64-bit elements.
AVX2 (2013):	Integer AVX version, horizontal operations, gather, 8/16/32/64-bit elements.
LRBNI (2009):	32 x 512-bit data and 8 x 16-bit predicate registers, i32, f32/f64, gather and scatter, flexible bit manipulation.
IMCI (2010):	32 + 8 registers, cache management, fp focus.
AVX-512 (2015)	IMCI backport to AVX2, 32 + 8 registers, int and fp focus.

Vectorize what?

- Historically the domain of HPC
- Differential equation systems and CFD simulations are important
- TLS termination is important too
- An open question remains – what is the ratio of cores doing the former and the latter?

Discussion