




Reflection Is Not Contemplation

Andrei Alexandrescu | CppCon 2024



“

Ba-dum tsss...

”

I'll Start With the Punchline

- Static reflection without code generation is incomplete
- The “reading” part of reflection generally agreed upon
- The “generation” part of reflection suffered of neglect
 - P2996 very gingerly sneaks in a foot in the door (`define_class`)
 - P3294 finally blows the door off its hinges
- The two facets of reflection are equally important
- Where do AI tools fit within this craze?

The Reflection Circularity Problem

Without generation, we're chasing our tails

- Recall $^{\wedge\wedge}x$ reflects on x , $[:y:]$ unreflects (splices) y ; $[:^{\wedge\wedge}x:]$ is x
- Large consensus on introspection query: contemplation is great
- Fear and loathing about code generation
 - Expansion of existing introspection objects deemed acceptable
- Consequence:
 - Severely limited: can't do 3D with 2D abilities
 - No consensus on “how much” generation is just enough
 - No clarity on necessary primitives

Where Do We Want To Be?

Reflection is ultimate code reuse

- Procedural: *I can write functions then call them*
- OOP: *I can call functions that haven't been written yet*
- Templates: *I can design with types that haven't been written yet*
- Reflection: *I can customize types that will not have been meant for customization*

How to Generate a Property

```
struct Book {  
    consteval {  
        property("author", ^^std::string);  
        property("title", ^^std::string);  
    }  
};
```

- A property introduces a data member, a getter, and a setter
- We'd like to be able to define such metafunctions

How to Generate a Property

// Equivalent hand-written code

```
struct Book {  
    std::string m_author;  
    const std::string& get_author() const { return m_author; }  
    void set_author(std::string s) { m_author = std::move(s); }  
  
    std::string m_title;  
    const std::string& get_title() const { return m_title; }  
    void set_title(std::string s) { m_title = std::move(s); }  
};
```

Early Observation

- We create identifiers from strings
 - Plus ability to concatenate “m_”, “get_”, “set_” with parameter
- Identifiers may be:
 - introduced (get_title); or
 - part of generated code (return m_title;)
- Ability to fluidly translate strings to identifiers *crucial*

Comparison of Splicing Models

- **The Spec API:** function calls create a “spec” of a type that is spliced
 - Complex; problematic; death of a thousand cuts
- **The CodeReckons API:** OOP interface for AST building
 - Verbose, loquacious, garrulous, prolix, long-winded, circumlocutory, and unceasingly inclined toward linguistic superfluity.
- **String Injection:** offer a primitive that splices CT strings into code
 - Horribly unstructured. Also... can't use macros?!?
- **Fragments:** C++ fragments of stylized code
 - Thorough early checking makes complexity explode

“

“There’s a representation of C++ code that is well-defined, complete, and easy to understand: C++ code.”

– Daveed Vandevoorde

”

The Game Changer: Token Sequences

P3294

- At this point complexity is a huge liability
- Adding yet another sublanguage to C++ deemed undesirable
- String-based generation best; *let's eliminate its disadvantages*
- Strings opaque/unwieldy? Use *token sequences* instead of strings
 - Cost: one added literal kind
- Injection risks and dangers? Restrict string expansion
 - Carefully controlled escapes, interpolation-style

Token Sequences

Status: prototype implementation

```
constexpr auto t1 = ^^{ a + /* hi! */ b }; // three tokens
static_assert(std::is_same_v<decltype(t1), >);
constexpr auto t2 = ^^{ a += ( );
constexpr auto t3 = ^^{ abc { def }; // Error, unpaired brace
```

- Concatenation, comparison, comments etc are token-level, not textual
 - Eliminates string and preprocessor dirt
- Macros work

Quoting & Splicing a Token Sequence

Status: prototype implementation

- Escapes inside a token sequence:
 - `\(expr)` evaluates `expr` eagerly during creation of the token sequence
 - `\id(e1, e2, ...)` concatenates strings and integrals, creates an identifier
 - `\tokens(expr)` expands another token sequence
- Inside any `constexpr` function:
 - `queue_injection(tokens_expr)` injects a token sequence into the *current declaration context*

Quoting in a Token Sequence

Status: prototype implementation

```
constexpr void property(std::meta::info type, std::string_view name) {  
    queue_injection(^^{  
        [:\(type):] \id("m_", name);  
        [:\(type):] const& \id("get_", name)() const {  
            return \id("m_", name);  
        }  
        void \id("set_", name)([:\(type):] x) {  
            \id("m_", name) = std::move(x);  
        }  
    });  
}
```


For Comparison: property Using The Spec API

```
template <class T> using getter_type = auto() const -> T const&;
template <class T> using setter_type = auto(T const&) -> void;

constexpr auto property(string_view name, meta::info type) -> void {

    auto member = inject(data_member_spec{
        .name=std::format("m_{}", name), .type=type
    });

    inject(function_member_spec{
        .name=std::format("get_{}", name),
        .signature=substitute(^getter_type, {^type}),
        .body=defer(member, ^[]<std::meta::info M>(auto const& self) -> auto const& {
            return self.[:M:];
        })
    });

    inject(function_member_spec{
        .name=std::format("set_{}", name),
        .signature=substitute(^setter_type, {^type}),
        .body=defer(member, ^[]<std::meta::info M>(auto& self, typename [:type_of(M):] const& x) -> void {
            self.[:M:] = x;
        })
    });
}
```

For Comparison: property Using The CodeReckons API

```
constexpr auto property(class_builder& b, type type, std::string name) -> void {
    auto member_name = identifier{"m_" + name}.c_str();
    append_field(b, member_name, type);

    method_prototype mp;
    object_type(mp, make_const(decl_of(b)));
    return_type(mp, make_lvalue_reference(make_const(type)));
    append_method(b, identifier{"get_" + name}.c_str(), mp,
    [member_name](method_builder& b){
        append_return(b,
            make_field_expr(
                make_deref_expr(make_this_expr(b)),
                member_name));
    });

    method_prototype mp1;
    append_parameter(mp1, "x", make_lvalue_reference(make_const(type)));
    object_type(mp1, decl_of(b));
    return_type(mp1, ^void);
    append_method(b, identifier{"set_" + name}.c_str(), mp1,
    [member_name](method_builder& b){
        append_expr(b,
            make_operator_expr(
                operator_kind::assign,
                make_field_expr(make_deref_expr(make_this_expr(b)), member_name),
                make_decl_ref_expr(parameters(decl_of(b))[1])
            ));
    });
}
```


“

“You can observe a lot by just
implementing *identity*”
– Yogi Berra

”

Notion of Identity/Copying/Cloning is Crucial

Allow me a little soapbox

- Identity fundamental concept in Computer Science
 - Different languages define it differently
- Philosophical underpinnings
 - Are objects fungible?
 - What is “same” object vs. “a copy”?
 - See Ship of Theseus Paradox/Hobbes variation
- Identifying, aliasing, copying, and moving objects central topic in C++
 - Alias analysis, self-assignment, self-move, double deletion, ...
- “Identity Function:” can insert in any (sub)expression without changing meaning

Identity Function Example

Allow me a little soapbox

FROM
ACCU
2008

- We had to add `&&` because C++98 couldn't implement `identity`
- The year was 2008 and this was C++0x:

```
template <typename T>  
T&& identity(T&& x) { return x; }
```

- Done! Or are we?

Identity Function Example

```
const int& a = identity(42); // dangling
int&& b = identity(42);      // dangling
auto&& c = identity(42);     // dangling
```

- When input is rvalue or rvalue reference, output should be a “true” rvalue
- When input is a reference or lvalue, output should be (the same) reference
- Early feedback:
 - *“I don’t understand his obsession with identity”*
 - *“Identity has no use in the real world”*
 - *“What’s the applicability here?”*

A Two-Body Problem

```
template <class T>
T& identity(T& x) { return x; }

template <class T>
T identity(T&& x) { return std::move(x); }
```

- All rvalue (reference)s go through the second overload, others through the first
- Key problem: scalability
 - Two definitions
 - Scales poorly to multiple arguments (T mentioned in the returned type)

One-Body identity()

```
template <class T>
T identity(T&& x) {
    return T(std::forward<T>(x));
}
```

- **T** morphs into either a value or reference (as before)
- **std::forward<T>(x)** makes sure a move happens if needed
- The additional **T()** is either a move ctor call OR a no-op, as needed
- Scalability problem remains
 - **T** still mentioned in the returned type
 - We can't return arbitrary expressions with fidelity

C++14 Fixes “Last” Problem: decltype(auto)

```
constexpr auto = ^int typename[:3] = 42 // Same as: int x = 42; typename[^char] = "" // Same as: char c = 'x';

template <class T>
decltype(auto) identity(T&& x) {
    return T(std::forward<T>(x));
}
```

- `decltype(auto)` teleports either a value type or reference type out
- Turns out it was important – see Howard Hinnant’s N2199
 - Rids `std::m(in|ax)` of dangling references with C++11 tech
 - Valiant effort: 210 lines, 21 struct definitions (7 distinct names)
 - Plagued by... `identity()` not implementable scalably
 - Rejected as too complex for what it does
- `min()` with `decltype(auto)` to N2199 spec in 8 lines, no helpers
 - godbolt.org/z/bGPnjM76e

std::min



C++ Safety

Just For Kicks

```
template <class A, class B>
decltype(auto)
min(A&& a, B&& b) {
    static_assert(... no dangerous comparisons ...);
    return b < a ? B(std::forward<B>(b)) : A(std::forward<A>(a));
}
```

- See also: [P1179](#), different take (also 12 years later)

identity Appears in the Darndest Places

Not Only for Functions!

- “The functionality [the identity metafunction] provides is both fundamental and surprisingly useful.” — Timur Doumler, P0887
- `std::identity` (C++20)
 - Default projection in constrained algorithms
 - (Intentionally does not create copies so it’s different from the one above)
- `thrust::identity`
 - Templated, works on CPU and GPU, does not accept rvalues
- `std::type_identity`, `std::type_identity_t` (C++20)
 - Introduce non-deduced contexts in template argument deduction

So What Is the Identity of Reflection?

Assembling a Type From Parts of Another is the Core Business of Reflection



“

Identity of reflection is being able
to deconstruct an entity and
construct a new, identical one
piecewise.

”

“

Please note: that entails (a) complete deconstruction of the entity; (b) construction from scratch of the entity.

”

Once You Have Identity of Reflection...

Everything becomes easy!

- Herb keynote's interface example?
 - “Deconstruct type, clone inserting `virtual` and `=0` for each memfun”
 - “Complain if you find data members or other suspect items”
- `polymorphic_base`:
 - “Deconstruct type, reassemble making sure no copying allowed”
 - “Ensure the dtor is public/virtual or protected/nonvirtual”
- `ordered`:
 - “Deconstruct type, reassemble and add operator<=>”
- *The point is (re)assembly from small, replaceable pieces*

The Real Challenge

- Cloning a class template is *much more difficult* than cloning a class
- `clone<MyStruct>` easy, `clone<std::vector>` absolutely crazy
 - Reflection must preserve order of declaration
 - Introspect function templates (e.g. `emplace_back`)
 - Signatures have constraints, noexcept clauses, attributes...
 - Clone inner classes: `iterator`, `const_iterator`
 - Manipulate (template) function signatures, e.g. `erase()`
 - Probably a million things we haven't thought of
- We can't avoid it! `clone<std::vector<int>>` equally hard
- Important part: *We know where the destination is*

Angle of Attack

Tokens, tokens everywhere...

- P3294 (Injection with Token Sequences) a game changer for generation
 - Lennon/McCartney of Daveed Vandevoorde, Barry Revzin, and yours truly
 - Flexible, loosely-structured “atoms” for code generation
 - Perfect currency for querying and writing code
- Repurposing P3157 (Generative Extensions for Reflection)
 - Retrieve/set template signature via multiple token strings:
 - `info get_template parameters(info);`
 - `info get_template_constraints(info);`
 - `info get_parameters(info);`
 - `info get_constraints(info);`
 - `info get_attributes(info);`
 - `info get_body(info); // (!)`



“

Suddenly, an AI

”

We Can't Not Talk About This

- How is generative AI changing language (features) design?
- [jippity.pro](#): 350 lines of working code written by a 9 yo in 3h
- Entire project passed to an LLM alongside a change request
- Will we really need in the future:
 - Boilerplate generation? AI could do that
 - Read all that boilerplate? AI editors can fold that away
 - Refactor/change large swaths of code? AI could read/modify/write code

“

Query (o1): *Please write a functionally equivalent clone of `std::vector` that uses a private member of type `std::vector` and forwards all methods to it. It should count calls to each method separately.*

”

“

(After 35 seconds) Generated a
fragment, gave me homework.

”

“

Me: *I noticed you wrote this:
// Additional methods forwarded similarly...
Can you please generate the FULL code?
This is very important for my career.*

”



“

(9 seconds later) Generated 447 lines of
working code.

”

“

Me: *Can you please do the same for
`std::unordered_map`? Otherwise, the
Taliban will have a word with me.*

”



“

(6 seconds later) Generated 470 lines of working code. No uncomfortable questions asked.

”



“

Don't forget: today's AI coding tools are absolutely the dumbest, most laughable, least capable they'll ever be. Everything now is at best alpha quality.

”

Where Is this Going?

- We need to assess how AI will influence programming language design
 - We now assume human production and consumption; that is being challenged
- We need to figure:
 - What is the spec? Is it the code? Is it some initial code and a set of queries?
 - Will we look at source code, stylized summaries, or something else?
 - “This is a proxy, this is a memoized fetcher, this is a façade...”
- How do we define technical debt?
 - Today: entropy created in order to deliver something quickly
 - Tomorrow: entropy *that is difficult for humans+AI to reduce*

However

- *Current LLM technology akin to an extra engineer, not a language feature*
 - 100 lines/s way faster than human, but 100x slower than reflection
 - Not as reliable as vetted code that generates code
- Reflection > 10x faster than TMP
 - 100x slower than compiled code
 - Close to parity with JIT or DLL tech
- Caching of reflection output will be a major focus
 - Dramatic improvement of compilation times

Don't Forget

- Procedural: *I can write functions then call them*
- OOP: *I can call functions that haven't been written yet*
- Templates: *I can design with types that haven't been written yet*
- Reflection: *I can customize types that will not have been meant for customization!*



Thank You!