

+ 24

Shared Libraries and Where To Find Them

LUIS CARO CAMPOS



September 15 - 20

Shared Libraries and Where to Find Them



CONAN

C/C++ Package Manager



Luis Caro Campos
R&D Team Lead, JFrog



Introduction

From the point of view of a C++ developer, **compiled code** typically ends up in either:

- The executables themselves
- Static libraries
- **Shared libraries**

Libraries* are a vehicle for “reusable” code that can be invoked by other libraries or applications (even from other languages)

* libraries: not defined by the standard! A word of many meanings...

How developers deal with shared libraries

LNK1104 wl

Asked 16 days ago

Id: warning: Could not find or use auto-linked framework 'CoreAudioTypes': framework 'CoreAudioTypes' not found Undefined symbols for architecture arm64: "_main", referenced from: Id: symbol(s) not found for architecture arm64 clang: error: linker command failed with exit code 1 (use -v to see invocation)

I have a c++ DLL proj
to use this library in n

Here is what I did so f

1. Added MyDll project as a reference in the console app.
2. Added MyDll.lib to the console's **Linker > Input**
3. Added MyDll's header path to the **Additional Include Directories**

Both projects hav
\$(SolutionDir):
are being genera

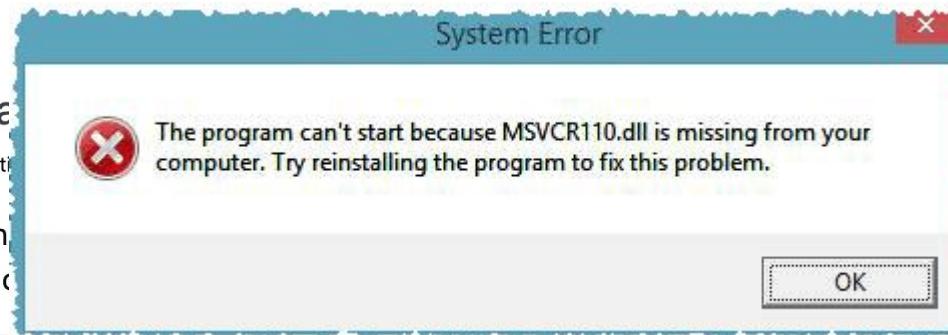
Why do recompile attempts fa

Asked 9 days ago Modified 9 days ago Viewed 36 t

I was able to h
directory. This
for library? Ho

When attempting to recompile an app th
running AlmaLinux 8 the following error o

```
/usr/bin/ld: cannot find -lgfortran  
collect2: error: ld returned 1 exit status  
make: [pest.mak:307: pest] Error 1
```



About this talk

- What this talk covers:
 - WHEN shared libraries are needed
 - WHICH programs needs to locate shared libraries
 - TOOLS to query, troubleshoot
 - WHERE to find shared libraries
- What this talk is NOT about:
 - Shared vs static linking
 - Internal structures of shared libraries
 - Linker specifics, relocations, advanced techniques

WHEN shared libraries are needed

When creating other shared libraries and executables

When running executables:

- Launching them
- Loading loadable plugins

When configuring the build

When packaging application for distribution

WHEN shared libraries are needed

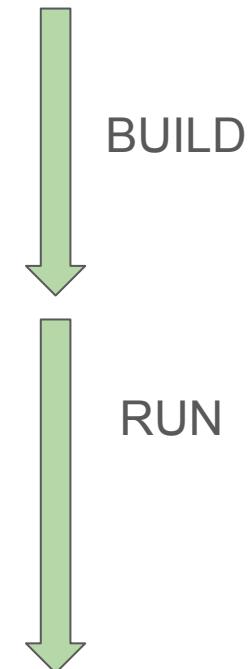
When configuring the build

When creating other shared libraries and executables

When running executables:

- Launching them
- Loading loadable plugins

When packaging application for distribution



WHICH programs need to locate shared libraries

When configuring the build



Build system generator

When creating other shared libraries and executables



(Static) **LINKER**

When running executables:

- Launching them
- Loading loadable plugins



(Dynamic) **LINKER/LOADER**

When packaging application for distribution



“Packaging tools”

GNU Linux

Linux example

We want to create an executable that links to libfoo

```
#include <foo.h>

int main() {
    i_am_foo();
    return 0;
}
```

my_program.cpp

Library exists and is installed:

- /usr/include/foo.h
- /usr/lib/libfoo.so



```
$ g++ my_program.cpp -o my_program
/usr/bin/ld: /tmp/cc8r2SFr.o: in function `main':
my_program.cpp:(.text+0x8): undefined reference to `i_am_foo()'
```

Linux example

We want to create an executable

```
#include <foo.h>

int main() {
    i_am_foo();
    return 0;
}
```

my_program.cpp

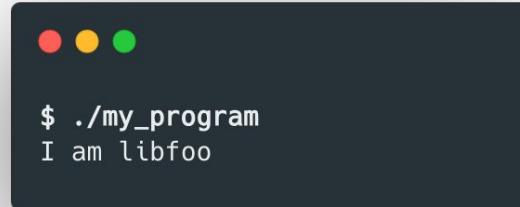
Library exists and is installed:

- /usr/include/foo.h
- /usr/lib/libfoo.so



```
$ g++ my_program.cpp -o my_program -lfoo
```

Linux example



We've shown both stages:

- The **linker (ld)** finds `libfoo.so` when creating the executable
 - The linker is indirectly invoked by `gcc`
 - Part of the compiler toolchain
- The **runtime linker-loader (ld.so)** found the library when launching it
 - Launched the application
 - Part of the system

The linker [ld]

- Program is typically called ld
 - other implementations: ld.gold, lld, mold
- On Linux, typically invoked by gcc itself via the collect2 tool

```
g++ my_program.cpp -o my_program -lfoo -v
```

```
4-linux-gnu/13:/usr/lib/gcc/aarch64-linux-gnu/
LIBRARY_PATH=/usr/lib/gcc/aarch64-linux-gnu/13:/usr/lib/gcc/aarch64-linux-gnu/13/../../aarch64-linux-gnu:/usr/lib/gcc/aarch64-linux-gnu/13/../../../../../lib/:/lib/aarch64-linux-gnu:/lib/..lib:/usr/lib/aarch64-linux-gnu:/usr/lib/..lib:/usr/lib/gcc/aarch64-linux-gnu/13/../../../../../lib/:/usr/lib/
COLLECT_GCC_OPTIONS='-o' 'my_program' '-v' '-shared-libgcc' '-mlittle-endian' '-mabi=lp64' '-dumpdir' 'my_program.'
/usr/libexec/gcc/aarch64-linux-gnu/13/collect2 -plugin /usr/libexec/gcc/aarch64-linux-gnu/13/liblto_plugin.so -plugin-opt=/usr/libexec/gcc/aarch64-linux-gnu/13/lto-wrapper -plugin-opt=-tresolution=/tmp/ccKKJULW.res -plugin-opt=-pass-through=-lgcc_s -plugin-opt=-pass-through=-lgcc -plugin-opt=-pass-through=-lc -plugin-opt=-pass-through=-lgcc_s -plugin-opt=-pass-through=-lgcc --build-id --eh-frame-hdr --hash-style=gnu --as-needed -dynamic-linker /lib/ld-linux-aarch64.so.1 -X -EL -maarch64linux --fix-cortex-a53-843419 -pie -z now -z relro -o my_program /usr/lib/gcc/aarch64-linux-gnu/13/../../aarch64-linux-gnu/Scrt1.o /usr/lib/gcc/aarch64-linux-gnu/13/../../aarch64-linux-gnu/crti.o /usr/lib/gcc/aarch64-linux-gnu/13/crtbeginS.o -L/usr/lib/gcc/aarch64-linux-gnu/13 -L/usr/lib/gcc/aarch64-linux-gnu/13/../../aarch64-linux-gnu -L/usr/lib/gcc/aarch64-linux-gnu/13/../../../../../lib -L/lib/aarch64-linux-gnu -L/lib/..lib -L/usr/lib/aarch64-linux-gnu -L/usr/lib/..lib -L/usr/lib/gcc/aarch64-linux-gnu/13/../../../../../ /tmp/cced6axp.o -lfoo -lstdc++ -lm -lgcc_s -lgcc -lc -lgcc_s -lgcc /usr/lib/gcc/aarch64-linux-gnu/13/crtendS.o /usr/lib/gcc/aarch64-linux-gnu/13/../../aarch64-linux-gnu/crtn.o
COLLECT_GCC_OPTIONS='-o' 'my_program' '-v' '-shared-libgcc' '-mlittle-endian' '-mabi=lp64' '-dumpdir' 'my_program.'
```

The linker [ld] - search procedure

--library-path=searchdir

Add path searchdir to the list of paths that **ld** will search for archive libraries and **ld** control scripts. You may use this option any number of times. The directories are searched in the order in which they are specified on the command line. Directories specified on the command line are searched before the default directories. All **-L** options apply to all **-l** options, regardless of the order in which the options appear. **-L** options do not affect how **ld** searches for a linker script unless **-T** option is specified.

If searchdir begins with "=" or **\$SYSROOT**, then this prefix will be replaced by the **sysroot** prefix, controlled by the **--sysroot** option, or specified when the linker is configured.

The default set of paths searched (without being specified with **-L**) depends on which emulation mode **ld** is using, and in some cases also on how it was configured.

The paths can also be specified in a link script with the "SEARCH_DIR" command. Directories specified this way are searched at the point in which the linker script appears in the command line.

-m emulation

The linker [ld] - search procedure

Ubuntu Manpage: ld - The GNU Linker

`--library-path=searchdir`

Add path `searchdir` to the list of paths that `ld` will search for archive libraries and `ld` control scripts. You may use this option any number of times. The directories are searched in the order in which they are specified on the command line. Directories specified on the command line are searched before the default directories. All `-L` options apply to all `-l` options, regardless of the order in which the options appear.

⚠

Some caution:

- Each distro/toolchain may be configured differently
- `ld` defaults may be overridden by the invocation by `gcc`
 - Check the `collect2` invocation with `-v`
- GNU Linker is typically the default, there's others:
 - `gold`, `lld`, `mold`
 - They may have slightly different behaviors

The linker [ld] - search procedure (cont'd)

```
g++ my_program.cpp -o my_program -lfoo -Xlinker --verbose

[...]
using internal linker script:
=====
/* Script for -pie -z combreloc -z relro -z now */
/* Copyright (C) 2014-2024 Free Software Foundation, Inc.
   Copying and distribution of this script, with or without modification,
   are permitted in any medium without royalty provided the copyright
   notice and this notice are preserved.  */
OUTPUT_FORMAT("elf64-littleaarch64", "elf64-bigaarch64",
              "elf64-littleaarch64")
OUTPUT_ARCH(aarch64)
ENTRY(_start)
SEARCH_DIR( "/usr/local/lib/aarch64-linux-gnu" );
SEARCH_DIR( "/lib/aarch64-linux-gnu" );
SEARCH_DIR( "/usr/lib/aarch64-linux-gnu" );
SEARCH_DIR( "/usr/local/lib" );
SEARCH_DIR( "/lib" );
SEARCH_DIR( "/usr/lib" );
SEARCH_DIR( "/usr/aarch64-linux-gnu/lib" );
SECTIONS
[...]
```

The dynamic linker/loader [ld-linux.so]

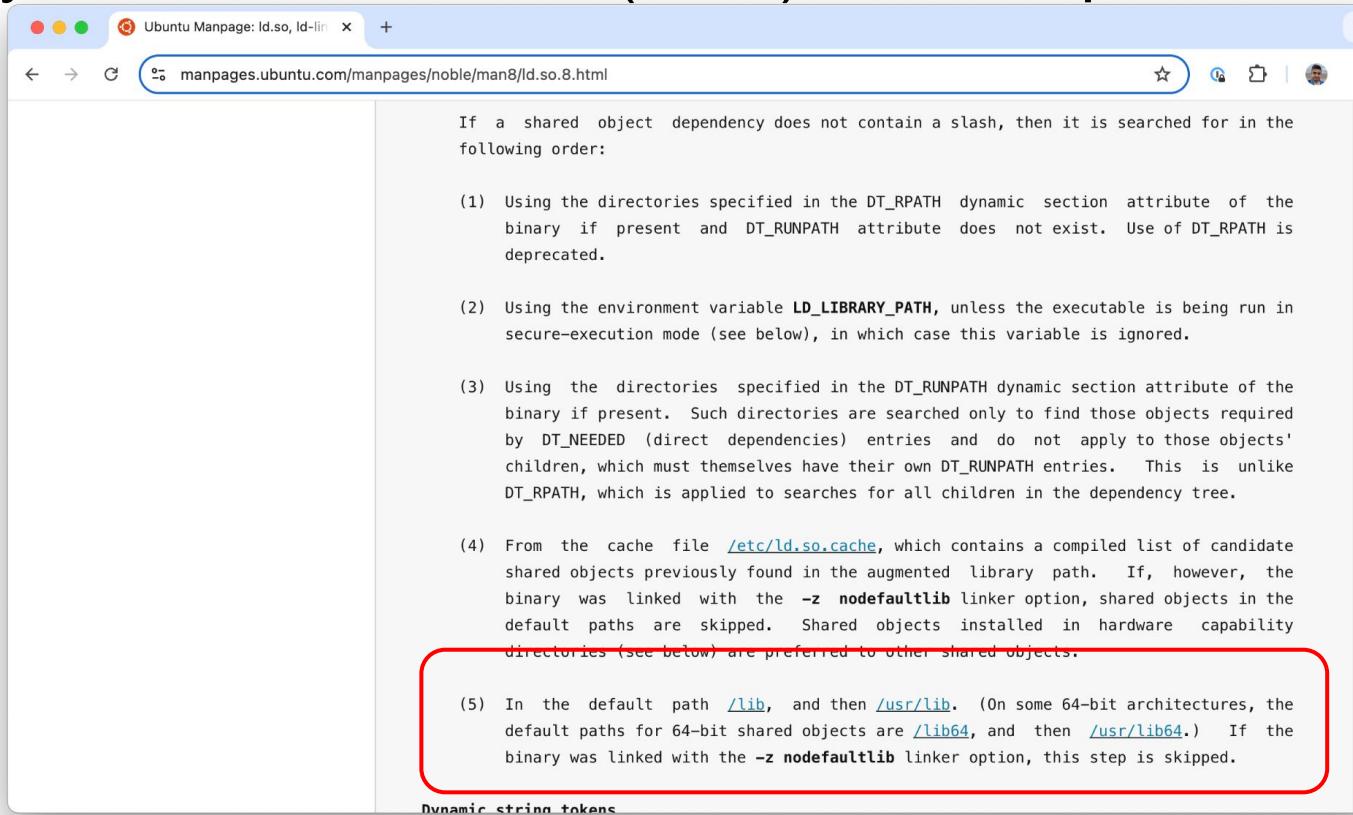
ELF executables are “interpreted” by an application: ld-linux.so

The full path to it is embedded in the executable ELF headers:

```
$ file my_program  
  
my_program: ELF 64-bit LSB pie executable, ARM aarch64, version 1 (SYSV),  
dynamically linked, interpreter /lib/ld-linux-aarch64.so.1,  
BuildID[sha1]=e9c0fb22d483b33ba9fac2b1574dbdf09e3421ce, for GNU/Linux 3.7.0, not  
stripped
```

```
$ readelf -l my_program | grep interpreter  
[Requesting program interpreter: /lib/ld-linux-aarch64.so.1]
```

The dynamic linker loader (ld.so) - search procedure



If a shared object dependency does not contain a slash, then it is searched for in the following order:

- (1) Using the directories specified in the DT_RPATH dynamic section attribute of the binary if present and DT_RUNPATH attribute does not exist. Use of DT_RPATH is deprecated.
- (2) Using the environment variable `LD_LIBRARY_PATH`, unless the executable is being run in secure-execution mode (see below), in which case this variable is ignored.
- (3) Using the directories specified in the DT_RUNPATH dynamic section attribute of the binary if present. Such directories are searched only to find those objects required by DT_NEEDED (direct dependencies) entries and do not apply to those objects' children, which must themselves have their own DT_RUNPATH entries. This is unlike DT_RPATH, which is applied to searches for all children in the dependency tree.
- (4) From the cache file `/etc/ld.so.cache`, which contains a compiled list of candidate shared objects previously found in the augmented library path. If, however, the binary was linked with the `-z nodefaultlib` linker option, shared objects in the default paths are skipped. Shared objects installed in hardware capability directories (see below) are preferred to other shared objects.
- (5) In the default path `/lib`, and then `/usr/lib`. (On some 64-bit architectures, the default paths for 64-bit shared objects are `/lib64`, and then `/usr/lib64`.) If the binary was linked with the `-z nodefaultlib` linker option, this step is skipped.

Dynamic string tokens

The dynamic linker loader (ld.so) - search procedure

```
$ objdump -p my_program | grep NEEDED
NEEDED          libfoo.so
NEEDED          libc.so.6
```

```
$ readelf -d my_program

Dynamic section at offset 0xfd90 contains 28 entries:
  Tag      Type           Name/Value
  0x0000000000000001 (NEEDED)    Shared library: [libfoo.so]
  0x0000000000000001 (NEEDED)    Shared library: [libc.so.6]
  [...]
```

The dynamic linker loader (ld.so) - search procedure

```
$ LD_DEBUG=libs ./my_program
    186: find library=libfoo.so [0]; searching
    186:   search cache=/etc/ld.so.cache
    186:   search path=/lib/aarch64-linux-gnu:/usr/lib/aarch64-linux-gnu:/lib:/usr/lib
(system search path)
    186:     trying file=/lib/aarch64-linux-gnu/libfoo.so
    186:     trying file=/usr/lib/aarch64-linux-gnu/libfoo.so
    186:     trying file=/lib/libfoo.so
```

The ldd tool

```
$ ldd my_program
    linux-vdso.so.1 (0x0000ffff98184000)
    libbar.so => /opt/foobar/lib/libbar.so (0x0000ffff980f0000)
    libstdc++.so.6 => /lib/aarch64-linux-gnu/libstdc++.so.6 (0x0000ffff97e70000)
    libc.so.6 => /lib/aarch64-linux-gnu/libc.so.6 (0x0000ffff97cb0000)
    libfoo.so => /opt/foobar/lib/libfoo.so (0x0000ffff97c80000)
    libm.so.6 => /lib/aarch64-linux-gnu/libm.so.6 (0x0000ffff97bd0000)
    /lib/ld-linux-aarch64.so.1 (0x0000ffff98147000)
    libgcc_s.so.1 => /lib/aarch64-linux-gnu/libgcc_s.so.1 (0x0000ffff97b90000)
```

Summary so far

Static linker [ld]

- Default search path
 - Default locations in SEARCH_DIR in linker script (pass -Xlinker –verbose to g++)
 - Compiler may be configured to pass additional locations (check -v output of g++)

Dynamic linker/loader [ld-linux .so]

- Default search path
 - Man ld.so
 - /etc/ld.so.conf.d/
 - (a cache)
- LD_DEBUG=libs ./executable (for search procedure)
- ldd ./executable (to list locations)

Linux example: library in non-default location

```
/opt/foo  
|-- include  
|   '-- foo.h  
`-- lib  
    '-- libfoo.so
```

Why?

- Library local to our project (not “installed”)
- We may be testing a different version of a library already installed in a system location
- Provided by package managers (Conan, vcpkg, Conda) ...
- We may not have root privileges to write in system locations
- Library is vendored-in
- ...

Linux example: library in non-default location

```
/opt/foo  
|-- include  
|   '-- foo.h  
`-- lib  
    '-- libfoo.so
```

```
● ● ●  
  
g++ my_program.cpp -lfoo -o my_program  
/usr/bin/ld: cannot find -lfoo: No such file or directory  
collect2: error: ld returned 1 exit status  
    return go(f, seed, [])  
}
```

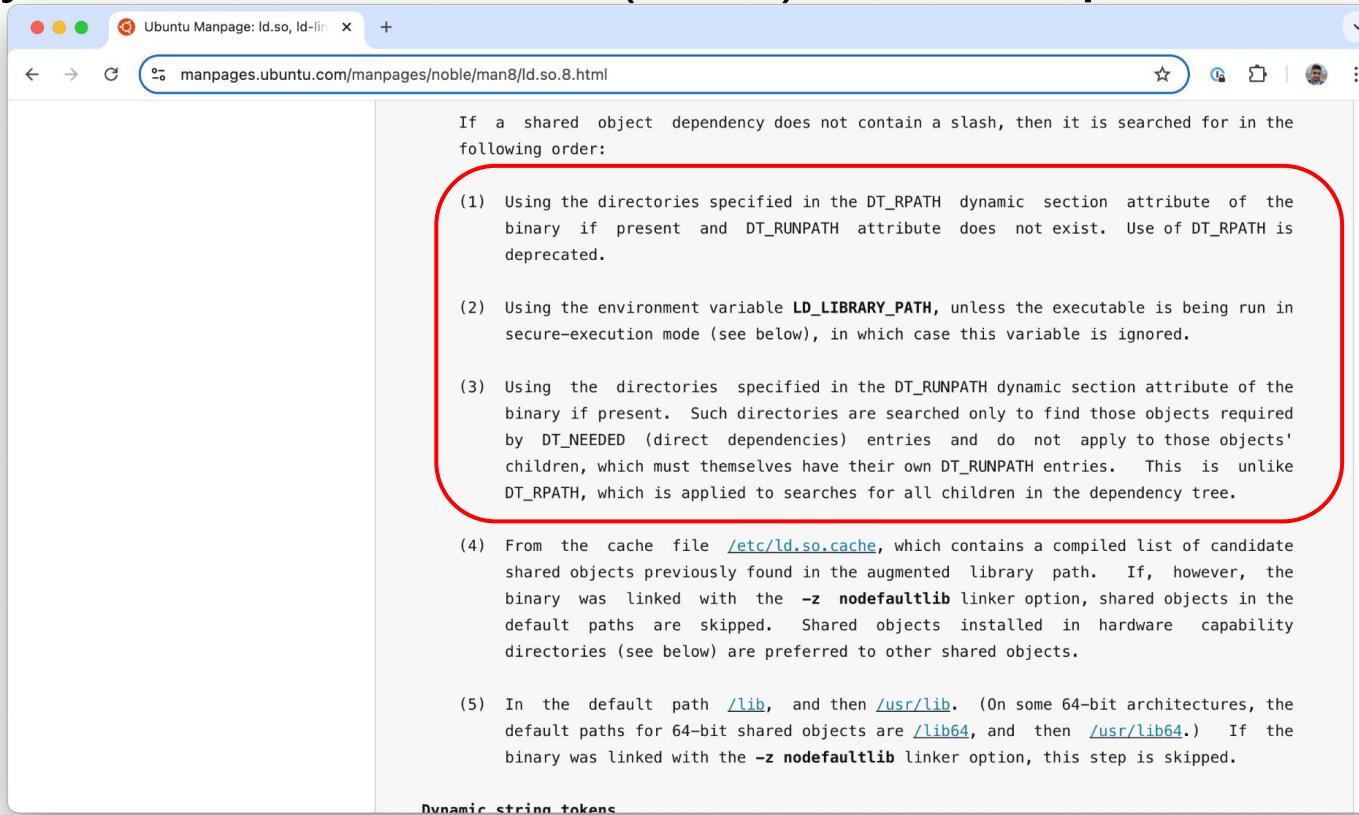
```
g++ my_program.cpp -L/opt/foo/lib -lfoo -o my_program
```

Linux example: library in non-default location

```
$ ./my_program
./my_program: error while loading shared libraries: libfoo.so: cannot open
shared object file: No such file or directory
```

The runtime linker/loader can not find the library now

The dynamic linker loader (ld.so) - search procedure



If a shared object dependency does not contain a slash, then it is searched for in the following order:

- (1) Using the directories specified in the DT_RPATH dynamic section attribute of the binary if present and DT_RUNPATH attribute does not exist. Use of DT_RPATH is deprecated.
- (2) Using the environment variable `LD_LIBRARY_PATH`, unless the executable is being run in secure-execution mode (see below), in which case this variable is ignored.
- (3) Using the directories specified in the DT_RUNPATH dynamic section attribute of the binary if present. Such directories are searched only to find those objects required by DT_NEEDED (direct dependencies) entries and do not apply to those objects' children, which must themselves have their own DT_RUNPATH entries. This is unlike DT_RPATH, which is applied to searches for all children in the dependency tree.
- (4) From the cache file `/etc/ld.so.cache`, which contains a compiled list of candidate shared objects previously found in the augmented library path. If, however, the binary was linked with the `-z nodefaultlib` linker option, shared objects in the default paths are skipped. Shared objects installed in hardware capability directories (see below) are preferred to other shared objects.
- (5) In the default path `/lib`, and then `/usr/lib`. (On some 64-bit architectures, the default paths for 64-bit shared objects are `/lib64`, and then `/usr/lib64`.) If the binary was linked with the `-z nodefaultlib` linker option, this step is skipped.

Dynamic string tokens

Linux example: library in non-default location

```
$ LD_LIBRARY_PATH=/opt/foo/lib LD_DEBUG=libs ./my_program
    225:      find library=libfoo.so [0]; searching
    225:          search path=/opt/foo/lib                  (LD_LIBRARY_PATH)
    225:              trying file=/opt/foo/lib/libfoo.so
    [...]
```



```
$ ld.so --library-path /opt/foo/lib ./my_program
```

Runtime linker/loader - rpath

- Runtime search path embedded in the ELF executable or library

```
g++ my_program.cpp -I/opt/foo/include -L/opt/foo/lib -lfoo -Wl,-rpath,/opt/foo/lib -o my_program
```

```
readelf -d my_program
```

```
Dynamic section at offset 0xfd80 contains 29 entries:
```

Tag	Type
0x0000000000000001	(NEEDED)
0x0000000000000001	(NEEDED)
0x0000000000000001d	(RUNPATH)

Name/Value
Shared library: [libfoo.so]
Shared library: [libc.so.6]
Library runpath: [/opt/foo/lib]

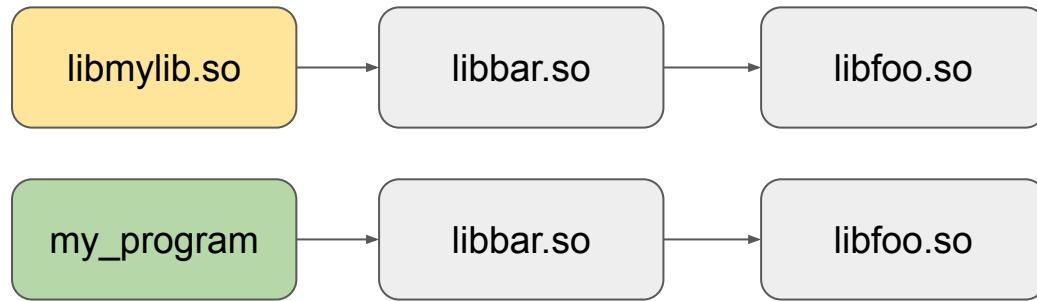
Runtime linker/loader - rpath

- Runtime search path embedded in the ELF executable or library
- The \$ORIGIN keyword: the directory containing the executable (allows for relocatability)

```
$ LD_DEBUG=libs ./my_program
    255:      find library=libfoo.so [0]; searching
    255:      search path=/opt/foo/lib          (RUNPATH from file ./my_program)
    255:      trying file=/opt/foo/lib/libfoo.so
```

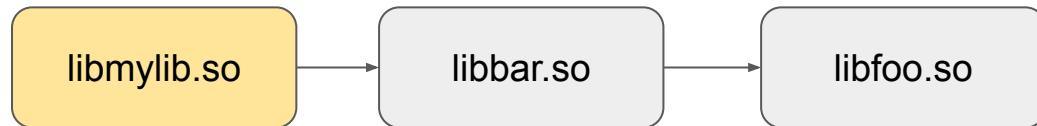
Linux example: transitive dependencies

```
/opt/foobar
|-- include
|   |-- bar.h
|   `-- foo.h
`-- lib
    |-- libbar.so
    '-- libfoo.so
```



Linux example: transitive dependencies

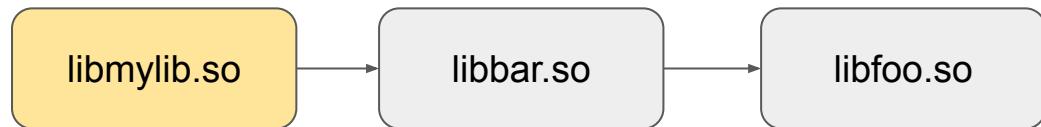
```
/opt/foobar
|-- include
|   |-- bar.h
|   `-- foo.h
`-- lib
    |-- libbar.so
    '-- libfoo.so
```



```
$ g++ my_lib.cpp -shared -o libmylib.so -I/opt/foobar/include
```

Linux example: transitive dependencies

```
/opt/foobar
|-- include
|   |-- bar.h
|   `-- foo.h
`-- lib
    |-- libbar.so
    '-- libfoo.so
```



```
$ readelf -d libmylib.so
```

```
Dynamic section at offset 0xfe00 contains 24 entries:
```

Tag	Type	Name/Value
-----	------	------------

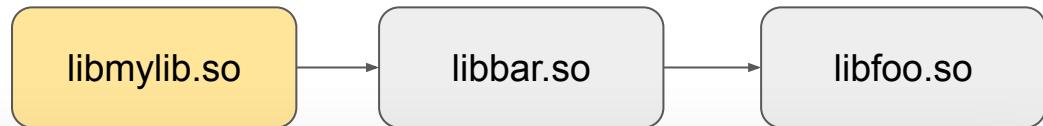
```
0x0000000000000001 (NEEDED)
```

```
[...]
```



No recorded dependency on
libfoo.so!

Linux example: transitive dependencies



```
$ g++ my_lib.cpp -shared -o libmylib.so -I/opt/foobar/include -L/opt/foobar/lib -lbar
```

Linux example: transitive dependencies



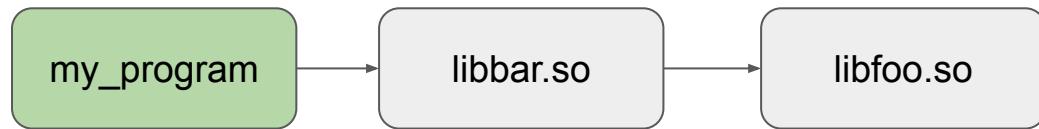
```
$ readelf -d libmylib.so
```

```
Dynamic section at offset 0xfd0 contains 25 entries:
```

Tag	Type	Name/Value
0x0000000000000001	(NEEDED)	Shared library: [libbar.so]
0x0000000000000001	(NEEDED)	Shared library: [libstdc++.so.6]
[...]		

Linux example: transitive dependencies

```
/opt/foobar
|-- include
|   |-- bar.h
|   `-- foo.h
`-- lib
    |-- libbar.so
    '-- libfoo.so
```



An executable!

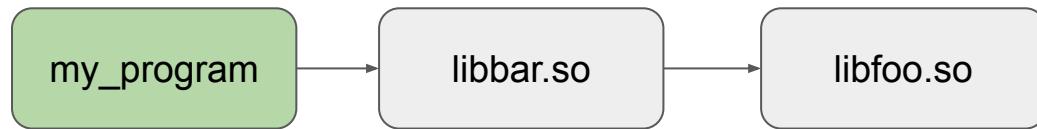


???

```
$ g++ my_program.cpp -o my_program -I/opt/foobar/include -L/opt/foobar/lib -lbar
/usr/bin/ld: warning: libfoo.so, needed by /opt/foobar/lib/libbar.so, not found (try using
-rpath or -rpath-link)
/usr/bin/ld: /opt/foobar/lib/libbar.so: undefined reference to `i_am_foo()'
collect2: error: ld returned 1 exit status
```

Linux example: transitive dependencies

```
/opt/foobar
|-- include
|   |-- bar.h
|   `-- foo.h
`-- lib
    |-- libbar.so
    '-- libfoo.so
```



```
$ g++ my_program.cpp -o my_program -I/opt/foobar/include -L/opt/foobar/lib -lbar -Wl,-rpath,/opt/foobar/lib
```

Linux example: transitive dependencies

```
--allow-shlib-undefined  
--no-allow-shlib-undefined
```

to
are

Allows or disallows undefined symbols in shared libraries. This switch is similar

--no-undefined except that it determines the behaviour when the undefined symbols

in a shared library rather than a regular object file. It does not affect how undefined symbols in regular object files are handled.

The default behaviour is to report errors for any undefined symbols referenced in shared libraries if the linker is being used to create an executable, but to allow them if the linker is being used to create a shared library.

Linux example: transitive dependencies

-L searchdir

--library-path=searchdir

Add path searchdir to the list of paths that ld will search for archive libraries

and

ld control scripts. You may use this option any number of times. The directories

are

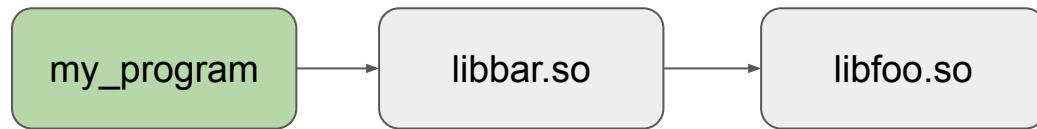
searched in the order in which they are specified on the command line. Directories specified on the command line are searched before the default directories. A **11 -L options apply to all -l options**, regardless of the order in which the options

appear.

-L options do not affect how ld searches for a linker script unless -T option is specified.

Linux example: transitive dependencies

```
/opt/foobar
|-- include
|   |-- bar.h
|   `-- foo.h
`-- lib
    |-- libbar.so
    '-- libfoo.so
```



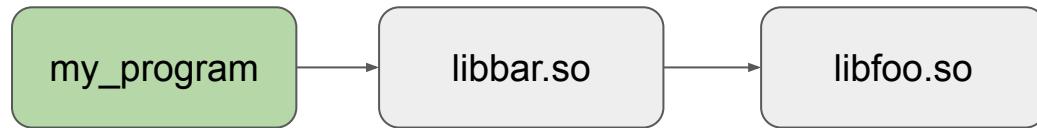
```
readelf -d my_program
```

```
Dynamic section at offset 0xfd70 contains 30 entries:
```

Tag	Type	Name/Value
0x0000000000000001	(NEEDED)	Shared library: [libbar.so]
0x0000000000000001	(NEEDED)	Shared library: [libstdc++.so.6]
0x0000000000000001	(NEEDED)	Shared library: [libc.so.6]
0x0000000000000001d	(RUNPATH)	Library runpath: [/opt/foobar/lib]

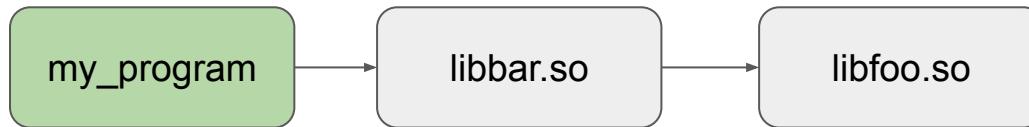
Linux example: transitive dependencies

```
/opt/foobar
|-- include
|   |-- bar.h
|   `-- foo.h
`-- lib
    |-- libbar.so
    '-- libfoo.so
```



```
$ ./my_program
./my_program: error while loading shared libraries: libfoo.so: cannot open
shared object file: No such file or directory
```

Linux example: RPATH vs RUNPATH



- RUNPATH only applies to direct dependencies.
- RPATH (legacy) can be used for indirect dependencies too
 - ! Problem: RPATH has priority over LD_LIBRARY_PATH

Possible solutions here:

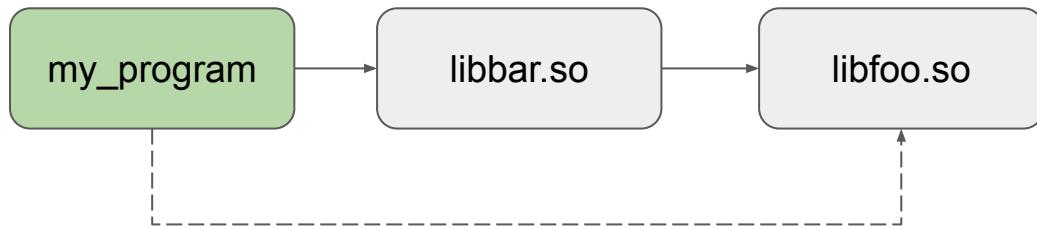
- Use the old ELF tag instead (pass `-Wl,--disable-new-dtags`)
- Embed a RUNPATH to `libbar.so`
 - When creating it:
 - After the fact: use `patchelf`
- Use `LD_LIBRARY_PATH` at runtime

Example 4: executable makes call to function in indirect library



```
$ g++ my_program.cpp -o my_program -I/opt/foobar/include -L/opt/foobar/lib -lbar -Wl,-rpath,/opt/foobar/lib
```

Example 4: call to function in (indirect) transitive library



```
$ g++ my_program.cpp -o my_program -I/opt/foobar/include -L/opt/foobar/lib -lbar -Wl,-rpath,/opt/foobar/lib
/usr/bin/ld: /tmp/ccG1fIWE.o: undefined reference to symbol '_Z15function_in_foo'
/usr/bin/ld: /opt/foobar/lib/libfoo.so: error adding symbols: DSO missing from command line
collect2: error: ld returned 1 exit status
```

Example 4: call to function in (indirect) transitive library



```
$ g++ my_program.cpp -o my_program -I/opt/foobar/include -L/opt/foobar/lib -lbar -lfoo
```

Apple OS platforms

macOS summary

- (static) linker: `ld`
 - `ld64`
 - `ld-prime` (available since Xcode 15)
- Dynamic linker/loader: `dyld`
- File extensions:
 - `dylib`, `so`, `framework`, (`tbd`)
- File inspection:
 - `otool`, `vtool`
- Post-build modifications:
 - `install_name_tool`

macOS static linker [ld] search path

- Similar to Linux: use -L for additional search paths
- Differences:
 - Only /usr/local/lib is searched by default in the local computer
 - System libraries are in the SDK
 - Frameworks have their own search paths

```
$ clang++ my_program.cpp -o my_program -lfoo -Xlinker -v
@(#)PROGRAM:ld PROJECT:ld-1053.12
BUILD 10:15:51 Mar 29 2024
[...]
Library search paths:
    /usr/local/lib
    /Applications/Xcode.app/Contents/Developer/Platforms/MacOSX.platform/Developer/SDKs/MacOSX.sdk/usr/lib
    /Applications/Xcode.app/Contents/Developer/Platforms/MacOSX.platform/Developer/SDKs/MacOSX.sdk/usr/lib/swift
Framework search paths:
    /Applications/Xcode.app/Contents/Developer/Platforms/MacOSX.platform/Developer/SDKs/MacOSX.sdk/System/Library/Frameworks
```

Mach-o binaries

otool: object file displaying tool

-L: display shared library dependencies

```
$ otool -L my_program
my_program:

libfoo.dylib (compatibility version 0.0.0, current version 0.0.0)
/usr/lib/libc++.1.dylib (compatibility version 1.0.0, current version 1700.255.5)
/usr/lib/libSystem.B.dylib compatibility version 1.0.0, current version 1345.120.2)
```

Dynamic library identification

Every dynamic library has an install name. This name identifies the library to the dynamic linker [dyld]. When you link to a dynamic library, the static linker [ld] records the library's install name in your Mach-O image [executable, dynamic library]

Dynamic Library Identification, Apple Developer Forums

<https://developer.apple.com/forums/thread/736719>

Dynamic library identification (cont'd)

Install name needs to be provided at link time.

Alternatively, after being built



```
clang++ libfoo.cpp -o libfoo.dylib -shared -Wl,-install_name,@rpath/libfoo.dylib
```



```
install_name_tool -id @rpath/libfoo.dylib libfoo.dylib
```

Dynamic loader [dyld] search procedure

Will depend on install name of dependency:

- `/usr/lib/libc++.1.dylib`
- `@rpath/libfoo.dylib`
 - Search in LC_RPATH entries in Mach-O header
 - RPATH is a stack of directories built from LC_RPATH entries of the chain that led to the current dyld load
 - Start at the executable
 - Similar to legacy DT_RPATH on Linux
- `libfoo.dylib` (leaf name)
 - For historical reasons, but no longer recommended on macOS

Dynamic loader [dyld] search procedure (cont'd)



```
DYLD_PRINT_SEARCHING=1 ./my_program
dyld[4556]: find path "@rpath/libfoo.dylib" @rpath/libfoo.dylib
dyld[4556]:   LC_RPATH '/path/to/example_1' from '/path/to/example_1/my_program'
dyld[4556]:   possible path(@path expansion): "/path/to/example_1/libfoo.dylib"
dyld[4556]:   found: dylib-from-disk: "/path/to/example_1/libfoo.dylib"
dyld[4556]: find path "/usr/lib/libc++.1.dylib"
dyld[4556]:   possible path(original path on disk): "/usr/lib/libc++.1.dylib"
dyld[4556]:   possible path(cryptex prefix): "/System/Volumes/Preboot/Cryptexes/OS/usr/lib/libc++.1.dylib"
dyld[4556]:   possible path(original path): "/usr/lib/libc++.1.dylib"
dyld[4556]:   found: dylib-from-cache: (0x000A) "/usr/lib/libc++.1.dylib"
dyld[4556]: find path "/usr/lib/libSystem.B.dylib"
dyld[4556]:   possible path(original path on disk): "/usr/lib/libSystem.B.dylib"
dyld[4556]:   possible path(cryptex prefix): "/System/Volumes/Preboot/Cryptexes/OS/usr/lib/libSystem.B.dylib"
dyld[4556]:   possible path(original path): "/usr/lib/libSystem.B.dylib"
dyld[4556]:   found: dylib-from-cache: (0x00AB) "/usr/lib/libSystem.B.dylib"
[...]
```

Dynamic loader [dyld] search procedure (cont'd)



```
$ clang++ my_program.cpp -o my_program -L/path/to/foo -lfoo -Wl,-rpath,/path/to/foo
```

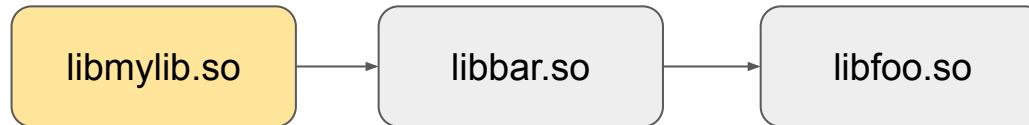


```
$ install_name_tool -add_rpath /path/to/lib my_program
```

Dynamic loader [dyld] search procedure (cont'd)

```
$ otool -l my_program
[...]
Load command 13
    cmd LC_LOAD_DYLIB
    cmdsize 48
        name @rpath/libfoo.dylib (offset 24)
    time stamp 2 Wed Dec 31 17:00:02 1969
        current version 0.0.0
compatibility version 0.0.0
Load command 14
    cmd LC_LOAD_DYLIB
    cmdsize 48
        name /usr/lib/libc++.1.dylib (offset 24)
    time stamp 2 Wed Dec 31 17:00:02 1969
        current version 1700.255.5
compatibility version 1.0.0
Load command 15
    cmd LC_LOAD_DYLIB
    cmdsize 56
        name /usr/lib/libSystem.B.dylib (offset 24)
    time stamp 2 Wed Dec 31 17:00:02 1969
        current version 1345.120.2
compatibility version 1.0.0
Load command 16
    cmd LC_RPATH
    cmdsize 80
        path /path/to/libfoo/ (offset 12)
[...]
```

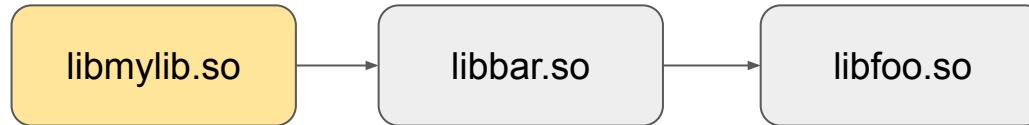
Indirect dependencies / missing symbols



```
$ g++ my_lib.cpp -shared -o libmylib.so -I/opt/foobar/include
```

Linux

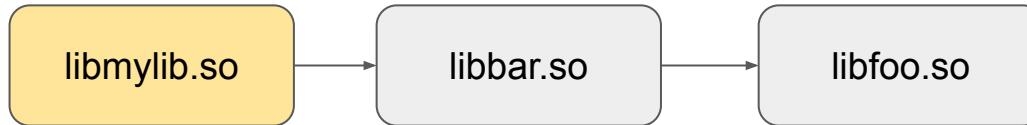
Indirect dependencies / missing symbols



```
clang++ libbar.cpp -I./libfoo/include -shared -o libbar.dylib
Undefined symbols for architecture arm64:
  "i_am_foo()", referenced from:
    i_am_bar() in libbar-f688bb.o
ld: symbol(s) not found for architecture arm64
```

macOS: error

Indirect dependencies / missing symbols



-undefined *treatment*

Specifies how undefined symbols are to be treated. Options are: error, warning, suppress, or dynamic_lookup. The default is error.

ld manpages on macOS

The two level namespace

Mach-O uses a two-level namespace. When a Mach-O image imports a symbol, it references the symbol name and the library where it expects to find that symbol.

An Apple Library Primer, Apple Developer Forums

<https://developer.apple.com/forums/thread/736719>



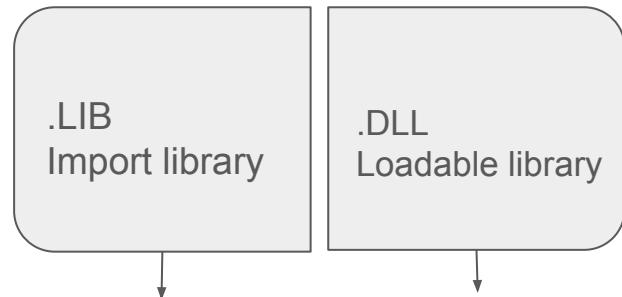
```
$ nm -m ./macos/libbar.dylib | grep foo | c++filt  
(undefined) external i_am_foo() (from libfoo)
```

Windows

Dynamically Loaded Libraries (DLL)

File format:

- PE: portable executable
- Shared libraries are “split” in two files:
 - **.lib** - import library, used at link time
 - Not to be confused with “.lib” static libraries
 - **.dll** - loaded at runtime
 - .pyd is used for loadable python modules



Tools:

- `cl.exe` - MSVC compiler
- `link.exe` - Linker
- `dumpbin.exe` - binary file dumper

Link.exe search path

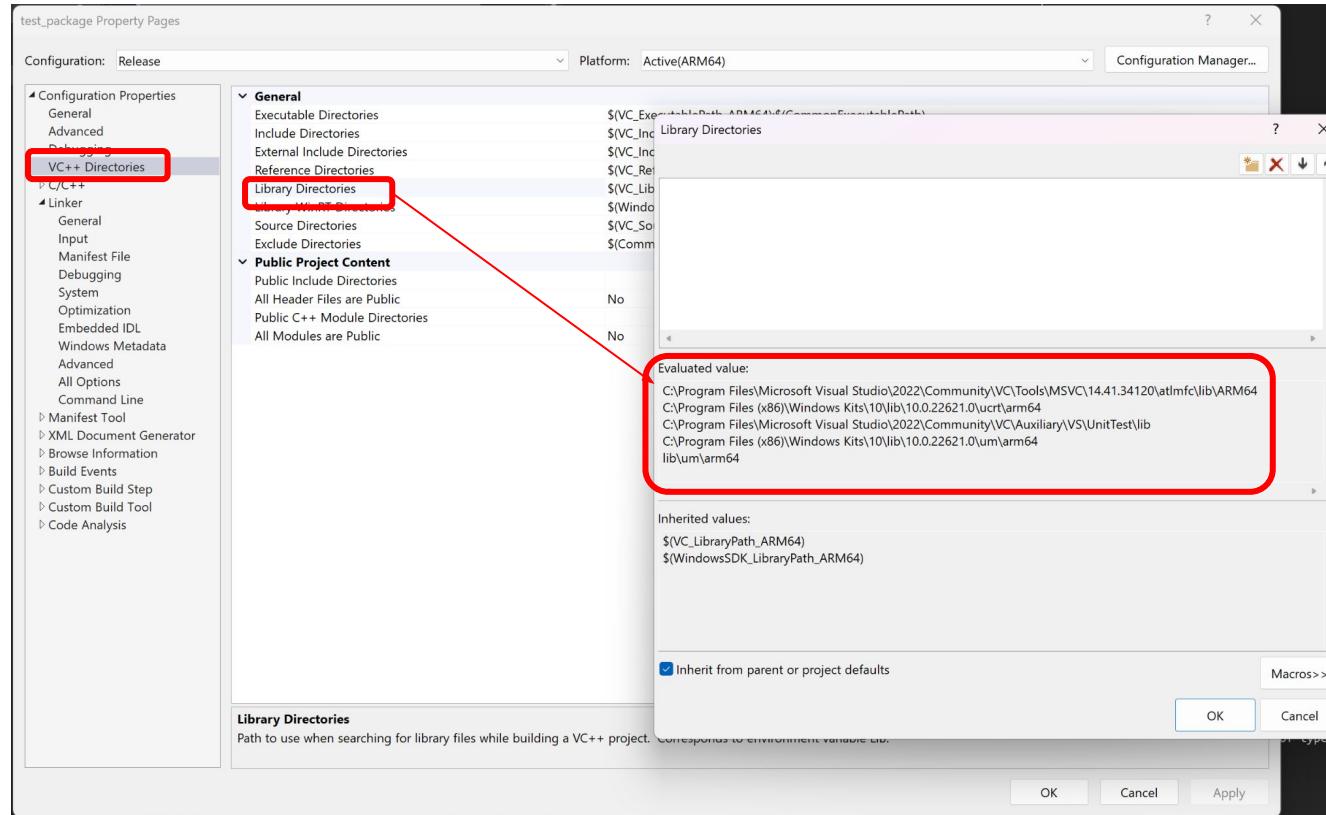
- Link.exe can be invoked by the compiler (cl.exe)
 - but typically build systems (CMake, MSBuild) will invoke link.exe directly
- Linker inputs:
 - Library names, e.g. foo.lib
 - File paths, e.g. C:/path/to/foo.lib
- Linker search path:
 - Content of **LIB** environment variable
 - **/LIBPATH (prepend paths)**

Link .exe default search path

```
> echo %LIB%
LIB=C:\Program Files\Microsoft Visual Studio\2022\Community\VC\Tools\MSVC\14.41.34120\ATLMFC\lib\ARM64;
C:\Program Files\Microsoft Visual Studio\2022\Community\VC\Tools\MSVC\14.41.34120\lib\ARM64;
C:\Program Files (x86)\Windows Kits\10\lib\10.0.22621.0\ucrt\arm64;
C:\Program Files (x86)\Windows Kits\10\\lib\10.0.22621.0\\um\arm64
```

- Set by Visual Studio Tools Command Prompt (“vcvars”)
- Depends on specific toolset and Windows SDK
- Pass **/VERBOSE:LIB** to print search procedure

Link .exe default search path



The DLL search path at runtime

The screenshot shows a Microsoft Learn article page. At the top, there's a navigation bar with tabs like 'Dynamic-link library search', 'Macrium Software | DLL Redir...', and a search bar. Below the navigation is a sidebar with a 'Filter by title' dropdown and a list of dynamic-link library topics. The main content area has a heading 'Standard search order for unpackaged apps'. It discusses the standard DLL search order based on safe DLL search mode. It includes a code snippet for disabling safe DLL search mode using the SetDllDirectory function. Below this, it lists 12 steps for the search order. To the right, there's a sidebar with 'Additional resources' sections for 'Events' (with a link to Nov 18, 4 PM - Nov 22, 4 PM) and 'Training' (with a link to 'Manage user files - Training'). At the bottom, there's a 'Documentation' section with links to 'Dynamic-link library redirection - Win32 apps' and 'Dynamic-Link Library Security - Win32 apps'. A note at the very bottom states: 'Safe DLL search mode is disabled when the search order is as follows.'

- Typically:
 - Next to the .exe file
 - System directory
 - Entries in PATH environment variable

Dynamic-link library search order

Microsoft Learn

<https://learn.microsoft.com/en-us/windows/win32/dlls/dynamic-link-library-search-order>

Listing dependencies

```
> dumpbin /dependents test_package.exe
Microsoft (R) COFF/PE Dumper Version 14.41.34120.0
Copyright (C) Microsoft Corporation. All rights reserved.

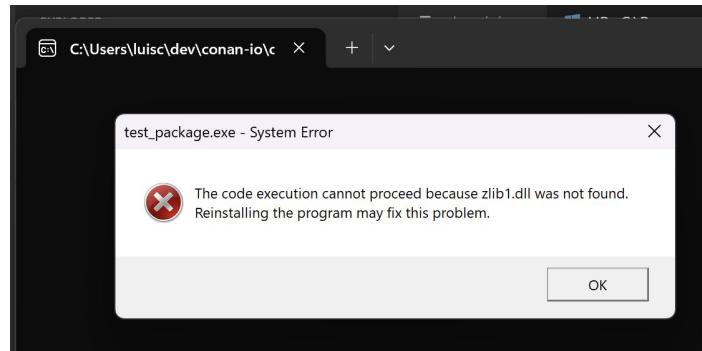
Dump of file test_package.exe

File Type: EXECUTABLE IMAGE

Image has the following dependencies:

zlib1.dll
VCRUNTIME140.dll
api-ms-win-crt-stdio-l1-1-0.dll
api-ms-win-crt-runtime-l1-1-0.dll
api-ms-win-crt-math-l1-1-0.dll
api-ms-win-crt-locale-l1-1-0.dll
api-ms-win-crt-heap-l1-1-0.dll
KERNEL32.dll
```

DLL not found



Launched from File explorer

Launched from command prompt A simple gray arrow pointing to the right, indicating the flow from the file explorer launch description to the command prompt launch description.

A screenshot of a macOS terminal window. The window title bar has three colored dots (red, yellow, green). The terminal output shows:

```
C:\Users\luisc\dev>test_package.exe  
C:\Users\luisc\dev>echo %errorlevel%  
-1073741515
```

DLL troubleshooting

Our options are limited:

- Dependency walker
 - “depends.exe”
- Windows debugger

The screenshot shows the Dependency Walker application interface. On the left, a tree view displays the dependencies of the executable 'TEST PACKAGE.EXE'. The root node is 'ZLIB1.DLL', which depends on 'VCRUNTIME140.DLL' and 'KERNEL32.DLL'. 'VCRUNTIME140.DLL' depends on several Microsoft CRT runtime DLLs: 'API-MS-WIN-CRT-RUNTIME-L1-1-0.DLL', 'API-MS-WIN-CRT-STRING-L1-1-0.DLL', 'API-MS-WIN-CRT-HEAP-L1-1-0.DLL', 'API-MS-WIN-CRT-CONVERT-L1-1-0.DLL', and 'API-MS-WIN-CRT-STDIO-L1-1-0.DLL'. The second 'KERNEL32.DLL' entry depends on 'API-MS-WIN-CRT-STDIO-L1-1-0.DLL', 'API-MS-WIN-CRT-RUNTIME-L1-1-0.DLL', 'API-MS-WIN-CRT-MATH-L1-1-0.DLL', 'API-MS-WIN-CRT-LOCALE-L1-1-0.DLL', and 'API-MS-WIN-CRT-HEAP-L1-1-0.DLL'. Below these, there is another 'KERNEL32.DLL' entry that depends on 'API-MS-WIN-CORE-RTLSUPPORT-L1-1-0.DLL' and 'API-MS-WIN-CORE-RTI SUPPORT-L1-2-0.DLL'.

On the right, two tables provide detailed information about the imports. The top table lists imports from 'ZLIB1.DLL' with the following data:

PI	Ordinal ^	Hint	Function	Entry Point
0x00000000	N/A	11 (0x000B)	deflate	Not Bound
0x00000000	N/A	14 (0x000E)	deflateEnd	Not Bound
0x00000000	N/A	17 (0x0011)	deflateInit_	Not Bound
0x00000000	N/A	85 (0x0055)	zlibVersion	Not Bound

The bottom table lists imports from 'TEST PACKAGE.EXE' with the following data:

E	Ordinal ^	Hint	Function	Entry Point
0x00000000				

At the bottom of the interface, two error messages are displayed in red:

Error: At least one required implicit or forwarded dependency was not found.
Warning: At least one delay-load dependency module was not found.

Build Systems

Build systems: Why locate libraries?

The linker already tells us if it can't find a library.

We could just do `-lopencv_core` and let the linker fail if the library can't be found.

But we may see things like the following in our build scripts:

- `find_package(OpenCV)`
- `PKG_CHECK_MODULES([OPENCV], [opencv >= 2.0])`

Why locate libraries? (cont'd)

- Error out early
 - Saves time!
- Additional checks
 - Is it the right version?
 - Is it the right architecture? Configured properly?
- Portability/cross platform for abstraction: filenames are different!
- Correct compiler and linker flags:
 - Compiler definitions
 - Paths to non-default locations: Include paths, linker paths (-L, -rpath-link)

Locating libraries - build system approaches

Try and link a “probe” executable during the configuration phase



→ OK for “system” libraries with well known names
👎 brings runtime loader into the “build” mix

“Find” the library file ahead of time

→ 👎 requires knowing library (file) names ahead of time
👎 need to replicate the linker search path (varies across systems/distros)

Package descriptor files to derive compiler/linker flags

→ Provides better abstraction, only need to know a generic “package” name

Requires a convention to search for these files... and their names!

Package descriptor files



pkg-config .pc files



CMake [package]-config.cmake:
imported targets



Common Package
Specification

Package descriptor files

Package descriptor files and where to find them

Coming to a C++ conference near you in 2025

CMake: Build time vs runtime

- CMake help on platforms that support RPATH:
 - macOS, Linux and other ELF platforms

“Build tree”

`cmake --build`

CMake passes -rpath to linker - for paths it knows about!

All linked binaries have absolute paths - this is why we can run binaries from local builds

“Install tree”

`cmake --install`

CMake patches the RPATHs, either removing them, or replacing them with the developer's choice (relative RPATH or path to install location)

CMake: Build time vs runtime - Windows

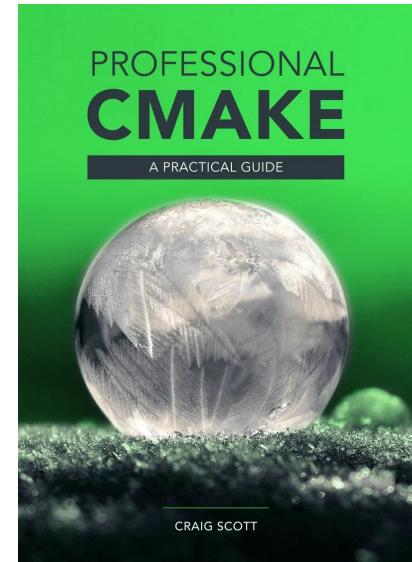
On Windows we're not so lucky.

- When launching executables from the build folder, we need DLLs
 - Either in the same folder as the executables
 - Or in the PATH environment variable
- Problems:
 - DLLs of dependencies in non-default locations
 - Projects that place DLLs and executables in different subfolders
 - Won't be able to run, debug, run tests

CMake: Build time vs runtime - Windows

```
# In this example, Algo is assumed to be a shared library defined elsewhere in
# the project and whose binary will be in a different directory to test_Algo
add_executable(test_Algo ...)
add_test(NAME CheckAlgo COMMAND test_Algo)

if(WIN32)
    # Ensure the required DLLs can be found at runtime
    set(algoDir "$<SHELL_PATH:$<TARGET_FILE_DIR:Algo>>")
    set(otherDllDir "C:\\path\\to\\another\\dll")
    set_property(TEST CheckAlgo APPEND PROPERTY
        ENVIRONMENT_MODIFICATION
        PATH=path_list_prepend:${algoDir}
        PATH=path_list_prepend:${otherDllDir}
    )
endif()
```



Snippet from “Professional CMake: A Practical Guide” by Craig Scott

CMake: Build time vs runtime - Windows

```
find_package(foo CONFIG REQUIRED) # package generated by install(EXPORT)

add_executable(exe main.c)
target_link_libraries(exe PRIVATE foo::foo foo::bar)
add_custom_command(TARGET exe POST_BUILD
  COMMAND ${CMAKE_COMMAND} -E copy -t ${TARGET_FILE_DIR:exe}
  $<TARGET_RUNTIME_DLLS:exe>
  COMMAND_EXPAND_LISTS
)
```

CMake documentation for
generator expressions

CMake: Build time vs runtime - Windows

Package managers to the rescue!

Conan CMake integrations:

- Command line: the `conanrun.bat` (PATH env var)
- Visual Studio generators (requires CMake 3.27 or higher)
 - Conan sets `CMAKE_VS_DEBUGGER_ENVIRONMENT`
- Optional:
 - Copy all DLLs to build folder - generate() method during conan install
 - `CONAN_RUNTIME_LIB_DIRS` - defined in toolchain

vcpkg:

- The `add_executable` built-in function is overridden to copy DLL dependencies right next to the executable

CMake: Build time vs runtime - Windows



```
C:\Users\dev\my_project>conan install . -o "*:shared=True"

C:\Users\dev\my_project>cmake --preset conan-default

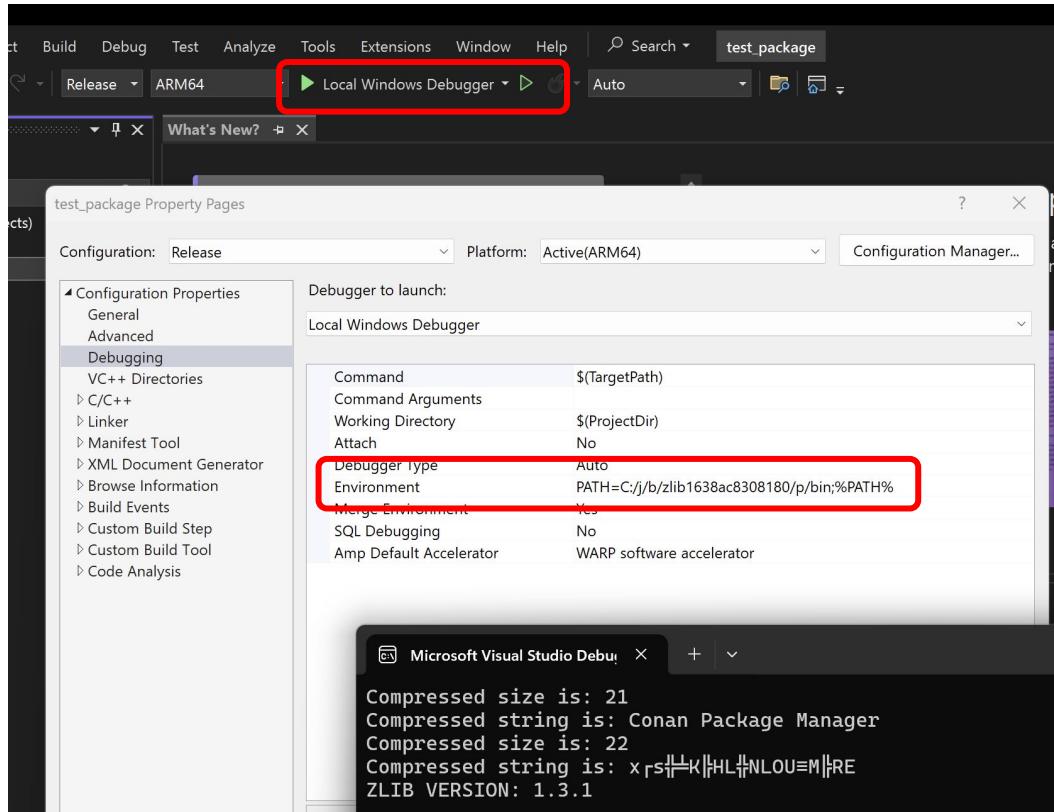
C:\Users\dev\my_project>cmake --build --preset conan-release

C:\Users\dev\my_project>build\generators\conanrun.bat

C:\Users\dev\my_project>build\Release\test_package.exe

Compressed size is: 21
Compressed string is: Conan Package Manager
Compressed size is: 22
Compressed string is: x0s00K0HL0NLOU0M0RE
ZLIB VERSION: 1.3.1
```

CMake: Build time vs runtime - Windows



Creating relocatable bundles

Installers, packages, bundles

- Goal: build on one machine, run in another
- Application must find all its shared library dependencies when launched

Who provides the shared
library dependencies

Where the application lives
once installed

System package manager

Fixed location

Bundled with the application

Arbitrary location (chosen at
install time)

Installers, packages, bundles

- Goal: build on one machine, run in another
- Application must find all its shared library dependencies when launched

Who provides the shared
library dependencies

Where the application lives
once installed

System package manager

Fixed location

Typical Linux

Bundled with the application

Arbitrary location (chosen at
install time)

macOS,
Windows...

Linux system package managers

- Linux users like these: JustWorks™
- External shared library dependencies are provided by other packages
 - We just need to package our executable files (and first-party libraries)
 - And record dependencies on other packages
- But has some downsides:
 - Dependencies available and versions are decided for us
 - Software vendors have to provide different packages for each distro (and each version!)
 - Some choose to host their own servers, users need to set them up first
 - Linux only!

Bundling dependencies

- Bundle the shared libraries needed by our executable(s)
- Give the runtime loader a helping hand:
 - ELF platforms, Apple platforms: Relative RPATH (\$ORIGIN, @executable_path)
 - Place DLLs next to executable on Windows

First we need to determine what to bundle!

System runtime libraries

First-party shared libraries

Third-party shared libraries



We typically don't bundle these (or can't)

Bundle formats

- Could be a folder with executable and libraries
- Or specific formats:
 - macOS App bundles (.app)
 - AppImage on Linux
 - Android APK

Gathering shared libraries for bundles

Linux: we can rely on the linker loader to give us a list (e.g. using `ldd`)

- Based on the output, determine if the library can/should be copied:
 - E.g. exclude system libraries from `/usr/lib`

Some examples:

- `linuxdeployqt`
 - (No longer maintained as of 2024)

Downsides:

- Mostly Linux and platforms that have `ldd`
- Uses the actual loader, won't work for cross-built binaries

Gathering shared libraries for bundles

```
① README.md  -eo appdirtool.go X
src > appimagetool > -eo appdirtool.go > ↗ getDirsFromSoConf

1114
1115     func findLibrary(filename string) (string, error) {
1116
1117         // Look for libraries in commonly used default locations
1118         locs := []string{"/usr/lib64", "/lib64", "/usr/lib", "/lib",
1119                         "/usr/lib/x86_64-linux-gnu/libfakeroot",
1120                         "/usr/local/lib",
1121                         "/usr/local/lib/x86_64-linux-gnu",
1122                         "/lib/x86_64-linux-gnu",
1123                         "/usr/lib/x86_64-linux-gnu",
1124                         "/lib32",
1125                         "/usr/lib32"}
1126
1127         for _, loc := range locs {
1128             libraryLocations = helpers.AppendIfMissing(libraryLocations, filepath.Clean(loc))
1129         }
1130
1131         // Additionally, look for libraries in the same locations in which glibc ld.so looks for libraries
1132         if helpers.Exists("/etc/ld.so.conf") {
1133             locs := getDirsFromSoConf("/etc/ld.so.conf")
1134             for _, loc := range locs {
1135                 libraryLocations = helpers.AppendIfMissing(libraryLocations, filepath.Clean(loc))
1136             }
1137
1138         // Also look for libraries in in LD_LIBRARY_PATH
1139         ldpstr := os.Getenv("LD_LIBRARY_PATH")
1140         ldps := strings.Split(ldpstr, ":")
1141         for _, ldp := range ldps {
1142             if ldp != "" {
1143                 libraryLocations = helpers.AppendIfMissing(libraryLocations, filepath.Clean(ldp))
1144             }
1145         }
1146 }
```

Alternatively: replicate the search logic of the runtime loader

Example from go-appimage:

probonopd/go-appimage



Go implementation of AppImage tools

A 26 Contributors 91 Issues ⭐ 710 Stars 71 Forks



Gathering shared libraries for bundles

```
install(TARGETS my_executable
        RUNTIME_DEPENDENCY_SET my_app_deps
)
install(RUNTIME_DEPENDENCY_SET my_app_deps
        PRE_EXCLUDE_REGEXES
        [[api-ms-win-.*]]
        [[ext-ms-.*]]
        [[kernel32\.dll]]
        [[libc\.so\..*]] [[libgcc_s\.so\..*]] [[libm\.so\..*]] [[libstdc\+\+\+.so\..*]]
        POST_EXCLUDE_REGEXES
        [[.*/system32/.*\.dll]]
        [[^/lib.*]]
        [[^/usr/lib.*]])
)
```

Docker

- A container is a process that is isolated from the rest of the system
- Which includes the filesystem
- When an application requests files via system calls:
 - The ELF interpreter, e.g. /lib/ld-linux-aarch64.so.1
 - Runtime dependencies
- The kernel creates an illusion around the process

If everything else fails...

... just link everything statically

Questions?

Shared Libraries and Where to Find Them



CONAN

C/C++ Package Manager



THANK YOU



Luis Caro Campos
R&D Team Lead, JFrog

