

+ 24

When Lock-Free Still Isn't Enough:

An Introduction to Wait-Free Programming and
Concurrency Techniques

DANIEL ANDERSON



20
24





Daniel Anderson

Assistant Teaching Professor @ CMU

danielanderson.net

**Carnegie
Mellon
University**

What we'll learn today

- Very quick review of concurrency and lock-free programming
- Review the “bread and butter” of lock-free design patterns
- Define **wait-free algorithms**, understand the definition and practical implications
- An example of an elegant wait-free algorithm and wait-free design
- Some simple benchmarks

Some assumed knowledge

- You know what `std::atomic` does and what it is used for
- You've heard of lock-free programming and know what a *compare_exchange* is

Our motivating problem

- Sometimes referred to as a “sticky counter” (it gets stuck at zero)

```
struct Counter {  
    // If the counter is greater than zero, add one and return true  
    // otherwise do nothing and return false  
    bool increment_if_not_zero();  
  
    // Decrement the counter. If the counter now equals zero,  
    // return true. Otherwise return false.  
    // Precondition: The counter is not zero  
    bool decrement();  
  
    uint64_t read(); // Return the current value of the counter  
};
```

- Required by `std::weak_ptr<T>::lock`
- Also, very useful for atomic memory management / concurrent data structures

Our first implementation

```
struct Counter {  
  
    bool increment_if_not_zero() {  
        if (counter > 0) {  
            counter++;  
            return true;  
        }  
        return false;  
    }  
  
    bool decrement() {  
        return (--counter == 0);  
    }  
  
    uint64_t read() { return counter; }  
  
    uint64_t counter{1};  
};
```

Thread 1:

```
increment_if_not_zero()  
    if (counter > 0) {  
  
        counter++;  
        return true;  
    }
```

Thread 2:

```
decrement() {  
  
    return (--counter == 0);  
}
```

Making it thread safe

```
struct Counter {  
  
    bool increment_if_not_zero() {  
        std::lock_guard g{m};  
        if (counter > 0) {  
            counter++;  
            return true;  
        }  
        return false;  
    }  
  
    bool decrement() {  
        std::lock_guard g{m};  
        return (--counter == 0);  
    }  
  
    std::mutex m;  
    uint64_t counter{1};  
};
```

Thread 1:

```
increment_if_not_zero()  
    std::lock_guard g{m};  
    if (counter > 0) {
```

Thread 2:

```
decrement() {  
    std::lock_guard g{m};
```



**One
Eternity
Later**

Making it thread safe

```
struct Counter {  
  
    bool increment_if_not_zero() {  
        std::lock_guard g{m};  
        if (counter > 0) {  
            counter++;  
            return true;  
        }  
        return false;  
    }  
  
    bool decrement() {  
        std::lock_guard g{m};  
        return (--counter == 0);  
    }  
  
    std::mutex m;  
    uint64_t counter{1};  
};
```

Thread 1:

```
increment_if_not_zero()  
    std::lock_guard g{m};  
    if (counter > 0) {  
  
        counter++;  
        return true;  
    }
```

Thread 2:

```
decrement() {  
    std::lock_guard g{m};
```


What have we learned so far?

Locks:

- Locks can eliminate **most** problems in concurrency!
- They do so by effectively eliminating concurrency...
- Locks often (but not always) have significant performance implications

Necessary disclaimer: Always measure performance. Never guess.

The rest of this talk:

- How to guess about performance
- We'll do some benchmarks too I promise

Progress guarantees

- Progress guarantees are a way to ***theoretically*** categorize concurrent algorithms:
- **Blocking**: No guarantee
- **Obstruction free (progress in isolation)**: A single thread executed in isolation will complete the operation in a bounded number of steps.
 - Obstruction-free algorithms are immune to deadlock
- **Lock free (at least one thread makes progress)**: At any given time, ***at least one*** thread is making progress on its operation
 - Guarantees system-wide throughput. Some operations are always completing, but individual operations are never guaranteed to ever complete
- **Wait free (all threads make progress)**: Every operation completes in a bounded number of steps regardless of other concurrent operations
 - Guaranteed bounded completion time for every individual operation.

A lock-free solution

```
struct Counter {  
  
    bool increment_if_not_zero() {  
        auto current = counter.load();  
        while (current > 0 && !counter.compare_exchange_weak(current, current + 1)) {}  
        return current > 0;  
    }  
  
    bool decrement() {  
        return counter.fetch_sub(1) == 1;  
    }  
  
    uint64_t read() { return counter.load(); }  
  
    std::atomic<uint64_t> counter{1};  
};
```

A lock-free solution

```
struct Counter {
```

```
    bool increment_if_not_zero() {  
        auto current = counter.load();  
        while (current > 0 && !counter.compare_exchange_weak(current, current + 1)) {}  
        return current > 0;  
    }
```

```
    bool decrement() {  
        return counter.fetch_sub(1) == 1;  
    }
```

```
    uint64_t read() { return counter.load(); }
```

```
    std::atomic<uint64_t> counter{1};  
};
```

```
compare_exchange(expected&, desired) {  
    if (current_value == expected) {  
        current_value = desired; return true; }  
    else { expected = current_value; return false; }  
}
```

A lock-free solution

```
struct Counter {
```

```
    bool increment_if_not_zero() {  
        auto current = counter.load();  
        while (current > 0 && !counter.compare_exchange_weak(current, current + 1)) {}  
        return current > 0;  
    }
```

```
    bool decrement() {  
        return counter.fetch_sub(1) == 1;  
    }
```

```
    uint64_t read() { return counter.load(); }
```

```
    std::atomic<uint64_t> counter{1};  
};
```

```
compare_exchange(expected&, desired) {  
    if (current_value == expected) {  
        current_value = desired; return true; }  
    else { expected = current_value; return false; }  
}
```

A lock-free solution

```
struct Counter {
```

```
    bool increment_if_not_zero() {  
        auto current = counter.load();  
        while (current > 0 && !counter.compare_exchange_weak(current, current + 1)) {}  
        return current > 0;  
    }
```

```
    bool decrement() {  
        return counter.fetch_sub(1) == 1;  
    }
```

```
    uint64_t read() { return counter.load(); }
```

```
    std::atomic<uint64_t> counter{1};  
};
```

```
compare_exchange(expected&, desired) {  
    if (current_value == expected) {  
        current_value = desired; return true; }  
    else { expected = current_value; return false; }  
}
```

The “CAS loop”

- The so-called “CAS loop” (compare-and-swap loop) is the bread and butter of lock-free algorithms and data structures
 - Read the current state of the data structure
 - Compute the new desired state from the current state
 - Commit the change only if no one else has already changed it (compare-exchange)
 - If someone else changed it, try again
- Progress is **lock free** because if an operation fails to make progress (the compare-exchange returns false) it can only be because a different operation made progress
- Progress is **not wait free** because a particular operation can fail the CAS loop forever because of competing operations succeeding

Tools for wait freedom

- A wait-free algorithm can not contain an unbounded CAS loop
 - This does not mean you can not use compare-exchange, just not in an unbounded loop!
- Most wait-free algorithms will make use of **atomic read-modify-write operations**:
 - **compare_exchange_weak/strong(expected, desired)**: Atomically replaces the current value with desired if current equals expected, otherwise loads the current value
 - **fetch_add(x) / fetch_sub(x)**: Atomically add/subtract x from the given variable and return the original value
 - **exchange(desired)**: Stores the value desired and returns the old value

Towards wait-free algorithms

- The “bread and butter” of lock-free algorithms was the **CAS loop**
- Threads **clobber** other threads that are trying to make progress, i.e., threads are competing to complete their operation first, at the expense of the others!
- To achieve wait freedom, threads should not be blocked by competing threads

Wait-free algorithm design

- Instead of being competitive, operations should be **collaborative**
- The “bread and butter” of wait-free algorithm design is called **helping**
- Operations that execute concurrently with another in-progress operation try to **help them make progress** instead of waiting on them or competing against them
- Requires a way for operations to detect others that are in progress, not always easy

The wait-free counter design

- Operations need a way detect that other operations are in progress
- We want a way to signal to other threads that we are planning to or already have set the counter to zero.
- **Big idea:** Steal some high-order bits of the counter to use as flags

0010101011101001110101010

flags counter

- An operation wants to **announce that the counter has been set to zero by setting the top flag**. (To any reading thread, top flag = 1 implies that the counter is zero).
- The second flag will be used for **helping**

Step 1: Wait-free counter without read

```
struct Counter {  
    static constexpr uint64_t is_zero = 1ull << 63;  
  
    bool increment_if_not_zero() {  
        return (counter.fetch_add(1) & is_zero) == 0;  
    }  
  
    bool decrement() {  
        if (counter.fetch_sub(1) == 1) {  
            uint64_t e = 0;  
            return counter.compare_exchange_strong(e, is_zero);  
        }  
        return false;  
    }  
};  
  
std::atomic<uint64_t> counter{1};  
};
```

← We must have counter == 0

Step 1: Wait-free counter without read

```
struct Counter {  
    static constexpr uint64_t is_zero = 1ull << 63;  
  
    bool increment_if_not_zero() {  
        return (counter.fetch_add(1) & is_zero) == 0;  
    }  
  
    bool decrement() {  
        if (counter.fetch_sub(1) == 1) {  
            uint64_t e = 0;  
            return counter.compare_exchange_strong(e, is_zero);  
        }  
        return false;  
    }  
};  
  
std::atomic<uint64_t> counter{1};  
};
```


Step 1: Wait-free counter without read

```
struct Counter {
    static constexpr uint64_t is_zero = 1ull << 63;

    bool increment_if_not_zero() {
        return (counter.fetch_add(1) & is_zero) == 0;
    }

    bool decrement() {
        if (counter.fetch_sub(1) == 1) {
            uint64_t e = 0;
            return counter.compare_exchange_strong(e, is_zero);
        }
        return false;
    }
}
```

What if this compare_exchange returns false??

The decrement *linearizes* after the increment.

```
std::atomic<uint64_t> counter{1};
};
```

Step 2: Adding a read operation. How hard can it be?

```

struct Counter {
    static constexpr uint64_t is_zero = 1ull << 63;

    bool increment_if_not_zero() {
        return (counter.fetch_add(1) & is_zero) == 0;
    }

    bool decrement() {
        if (counter.fetch_sub(1) == 1) {
            uint64_t e = 0;
            return counter.compare_exchange_strong(e, is_zero);
        }
        return false;
    }

    uint64_t read() {
        auto val = counter.load();
        return (val & is_zero) ? 0 : val;
    }

    std::atomic<uint64_t> counter{1};
};

```

counter = 0

Thread 1:

```

decrement()
    counter.fetch_sub(1)

```

Thread 2:

```

read()
    counter.load() // 0
    return 0

```

Thread 3:

```

increment_if_not_zero()
    counter.fetch_add(1)

```

If counter == 0, we can't affirm
that the counter is *actually* zero!
(It might go back to 1 because of
an increment)

Step 2: Wait-free counter with read?

```

struct Counter {
    static constexpr uint64_t is_zero = 1ull << 63;

    bool increment_if_not_zero() {
        return (counter.fetch_add(1) & is_zero) == 0;
    }

    bool decrement() {
        if (counter.fetch_sub(1) == 1) {
            uint64_t e = 0;
            return counter.compare_exchange_strong(e, is_zero);
        }
        return false;
    }

    uint64_t read() {
        auto val = counter.load();
        if (val == 0 then what?
    }

    std::atomic<uint64_t> counter{1};

```

Let's help them!!

There must be a thread
about to do this!

Step 2: Wait-free counter with read?

```

struct Counter {
    static constexpr uint64_t is_zero = 1ull << 63;

    bool increment_if_not_zero() {
        return (counter.fetch_add(1) & is_zero) == 0;
    }

    bool decrement() {
        if (counter.fetch_sub(1) == 1) {
            uint64_t e = 0;
            return counter.compare_exchange_strong(e, is_zero);
        }
        return false;
    }

    uint64_t read() {
        auto val = counter.load();
        if (val == 0 && counter.compare_exchange_strong(val, is_zero)) return 0; // helping!
        return (val & is_zero) ? 0 : val;
    }
};

std::atomic<uint64_t> counter{1};

```

Did we fix the bug?

Yes! But we added another one...

If a read helps to set the `is_zero` flag, none of the decrements return true...

Step 3: Almost there

```

struct Counter {
    static constexpr unsigned uint64_t is_zero = 1ull << 63;
    static constexpr unsigned uint64_t helped = 1ull << 62;

    bool increment_if_not_zero() {
        return (counter.fetch_add(1) & is_zero) == 0; }

    bool decrement() {
        if (counter.fetch_sub(1) == 1) {
            uint64_t e = 0;
            if (counter.compare_exchange_strong(e, is_zero)) return true;
            else if ((e & helped) && (counter.exchange(is_zero) & helped)) return true;
        }
        return false;
    }

    uint64_t read() {
        auto val = counter.load();
        if (val == 0 && counter.compare_exchange_strong(val, is_zero | helped)) return 0; // helping!
        return (val & is_zero) ? 0 : val; }

    std::atomic<uint64_t> counter{1};

```

Step 3: Almost there

```

struct Counter {
    static constexpr uint64_t is_zero = 1ull << 63;
    static constexpr uint64_t helped = 1ull << 62;

    bool increment_if_not_zero() {
        return (counter.fetch_add(1) & is_zero) == 0; }

    bool decrement() {
        if (counter.fetch_sub(1) == 1) {
            uint64_t e = 0;
            if (counter.compare_exchange_strong(e, is_zero)) return true;
            // else if ((e & helped) && (counter.exchange(is_zero) & helped)) return true;
        }
        return false;
    }

    uint64_t read() {
        auto val = counter.load();
        if (val == 0 && counter.compare_exchange_strong(val, is_zero | helped)) return 0; // helping!
        return (val & is_zero) ? 0 : val; }

    std::atomic<uint64_t> counter{1};

```

Step 4: Got it!

```

struct Counter {
    static constexpr uint64_t is_zero = 1ull << 63;
    static constexpr uint64_t helped = 1ull << 62;

    bool increment_if_not_zero() {
        return (counter.fetch_add(1) & is_zero) == 0; }

    bool decrement() {
        if (counter.fetch_sub(1) == 1) {
            uint64_t e = 0;
            if (counter.compare_exchange_strong(e, is_zero)) return true;
            else if ((e & helped) && (counter.exchange(is_zero) & helped)) return true;
        }
        return false;
    }

    uint64_t read() {
        auto val = counter.load();
        if (val == 0 && counter.compare_exchange_strong(val, is_zero | helped)) return 0; // helping!
        return (val & is_zero) ? 0 : val; }

    std::atomic<uint64_t> counter{1};

```

Careful
multiplication
after

Solution: One and *only one*
decrement must “take credit”
for zeroing the counter

Benchmarks

- My **atomic<shared_ptr>** implementation using the wait-free counter versus the lock-free counter. This affects the performance of the **load** operation.
- Benchmark #1:** p threads **loading** from an **atomic<shared_ptr>**

Load Latency

	1%	50%	99%
$p = 1$	21n	24.8n	32.7n
$p = 28$	217n	1.75u	13.1u
$p = 56$	535n	5.36u	31.1u

Lock Free Counter

Load Latency

	1%	50%	99%
$p = 1$	21.7n	24.8n	31.7n
$p = 28$	158n	1.31u	8.40u
$p = 56$	146n	2.43u	13.3u

Wait Free Counter

Benchmarks

- My **atomic<shared_ptr>** implementation using the wait-free counter versus the lock-free counter. This affects the performance of the **load** operation.
- Benchmark #2:** p threads, **50% loading** from an **atomic<shared_ptr>**, the other **50% storing**

Load Latency

	1%	50%	99%
$p = 28$	212n	1.19u	8.67u
$p = 56$	215n	2.41u	26.7u

Lock Free Counter

Load Latency

	1%	50%	99%
$p = 28$	164n	1.21u	7.65u
$p = 56$	153n	2.05u	26.4u

Wait Free Counter

Benchmarks

- My **atomic<shared_ptr>** implementation using the wait-free counter versus the lock-free counter. This affects the performance of the **load** operation.
- Benchmark #3:** p threads, **10% loading** from an **atomic<shared_ptr>**, the other **90% storing**

Load Latency

	1%	50%	99%
$p = 28$	217n	1.73u	14.1u
$p = 56$	240n	6.40u	82.6u

Lock Free Counter

Load Latency

	1%	50%	99%
$p = 28$	188n	1.41u	13.3u
$p = 56$	244n	7.46u	86.4u

Wait Free Counter

Benchmarks

Remember: Always measure performance. Never guess.

- Which algorithm is best often depends on the **workload**
 - How many reads vs how many writes
 - How many threads/cores
- Wait free was better for **read-mostly**, but lock free looks better for **write-mostly**

Take-home messages

Progress guarantees

- Useful theoretical classification of concurrent algorithms that can inform algorithm design
- Lock-free algorithms guarantee that one thread is making progress, while *wait-free algorithms guarantee that every thread is making progress*

Wait-free algorithm design

- The bread-and-butter technique is *helping*. Operations help concurrent operations rather than waiting for them (blocking) or compete with them (lock-free)

Performance Implications

- Never guess about performance
- **But do** hypothesize about performance by analyzing an algorithm's progress guarantees, and use these progress guarantees to guide the design of your algorithm
- **Then** benchmark it