



Why Is My Build So Slow?

Compilation Profiling and Visualization

SAMUEL PRIVETT



20
24



Disclaimer: The views and opinions expressed in this presentation are solely those of the presenters and do not necessarily represent the views, positions, or policies of Johnson & Johnson MedTech.

Table of Contents

1. Introduction
2. Assumptions
3. Visualizing Compilation
4. Single File
5. Project Level
6. Higher Order
7. Takeaways
8. Questions



Introduction

High Level Concepts

About Me

- Sam Privett (He/Him)
- Robotics Software Engineer @ JnJ
- Curious about compilers



What is this talk?


Visualizing Compilation with Ninja and Clang



Speed up compilers



Write code that compiles faster






Introduction


Assumptions

Visualizing Compilation


Lines of Code \approx Longer Build Times




Lines of Code




Build Times




Build Times



Build Times



Build Times



Boiling Frog



Developer Productivity



<https://xkcd.com/303/>



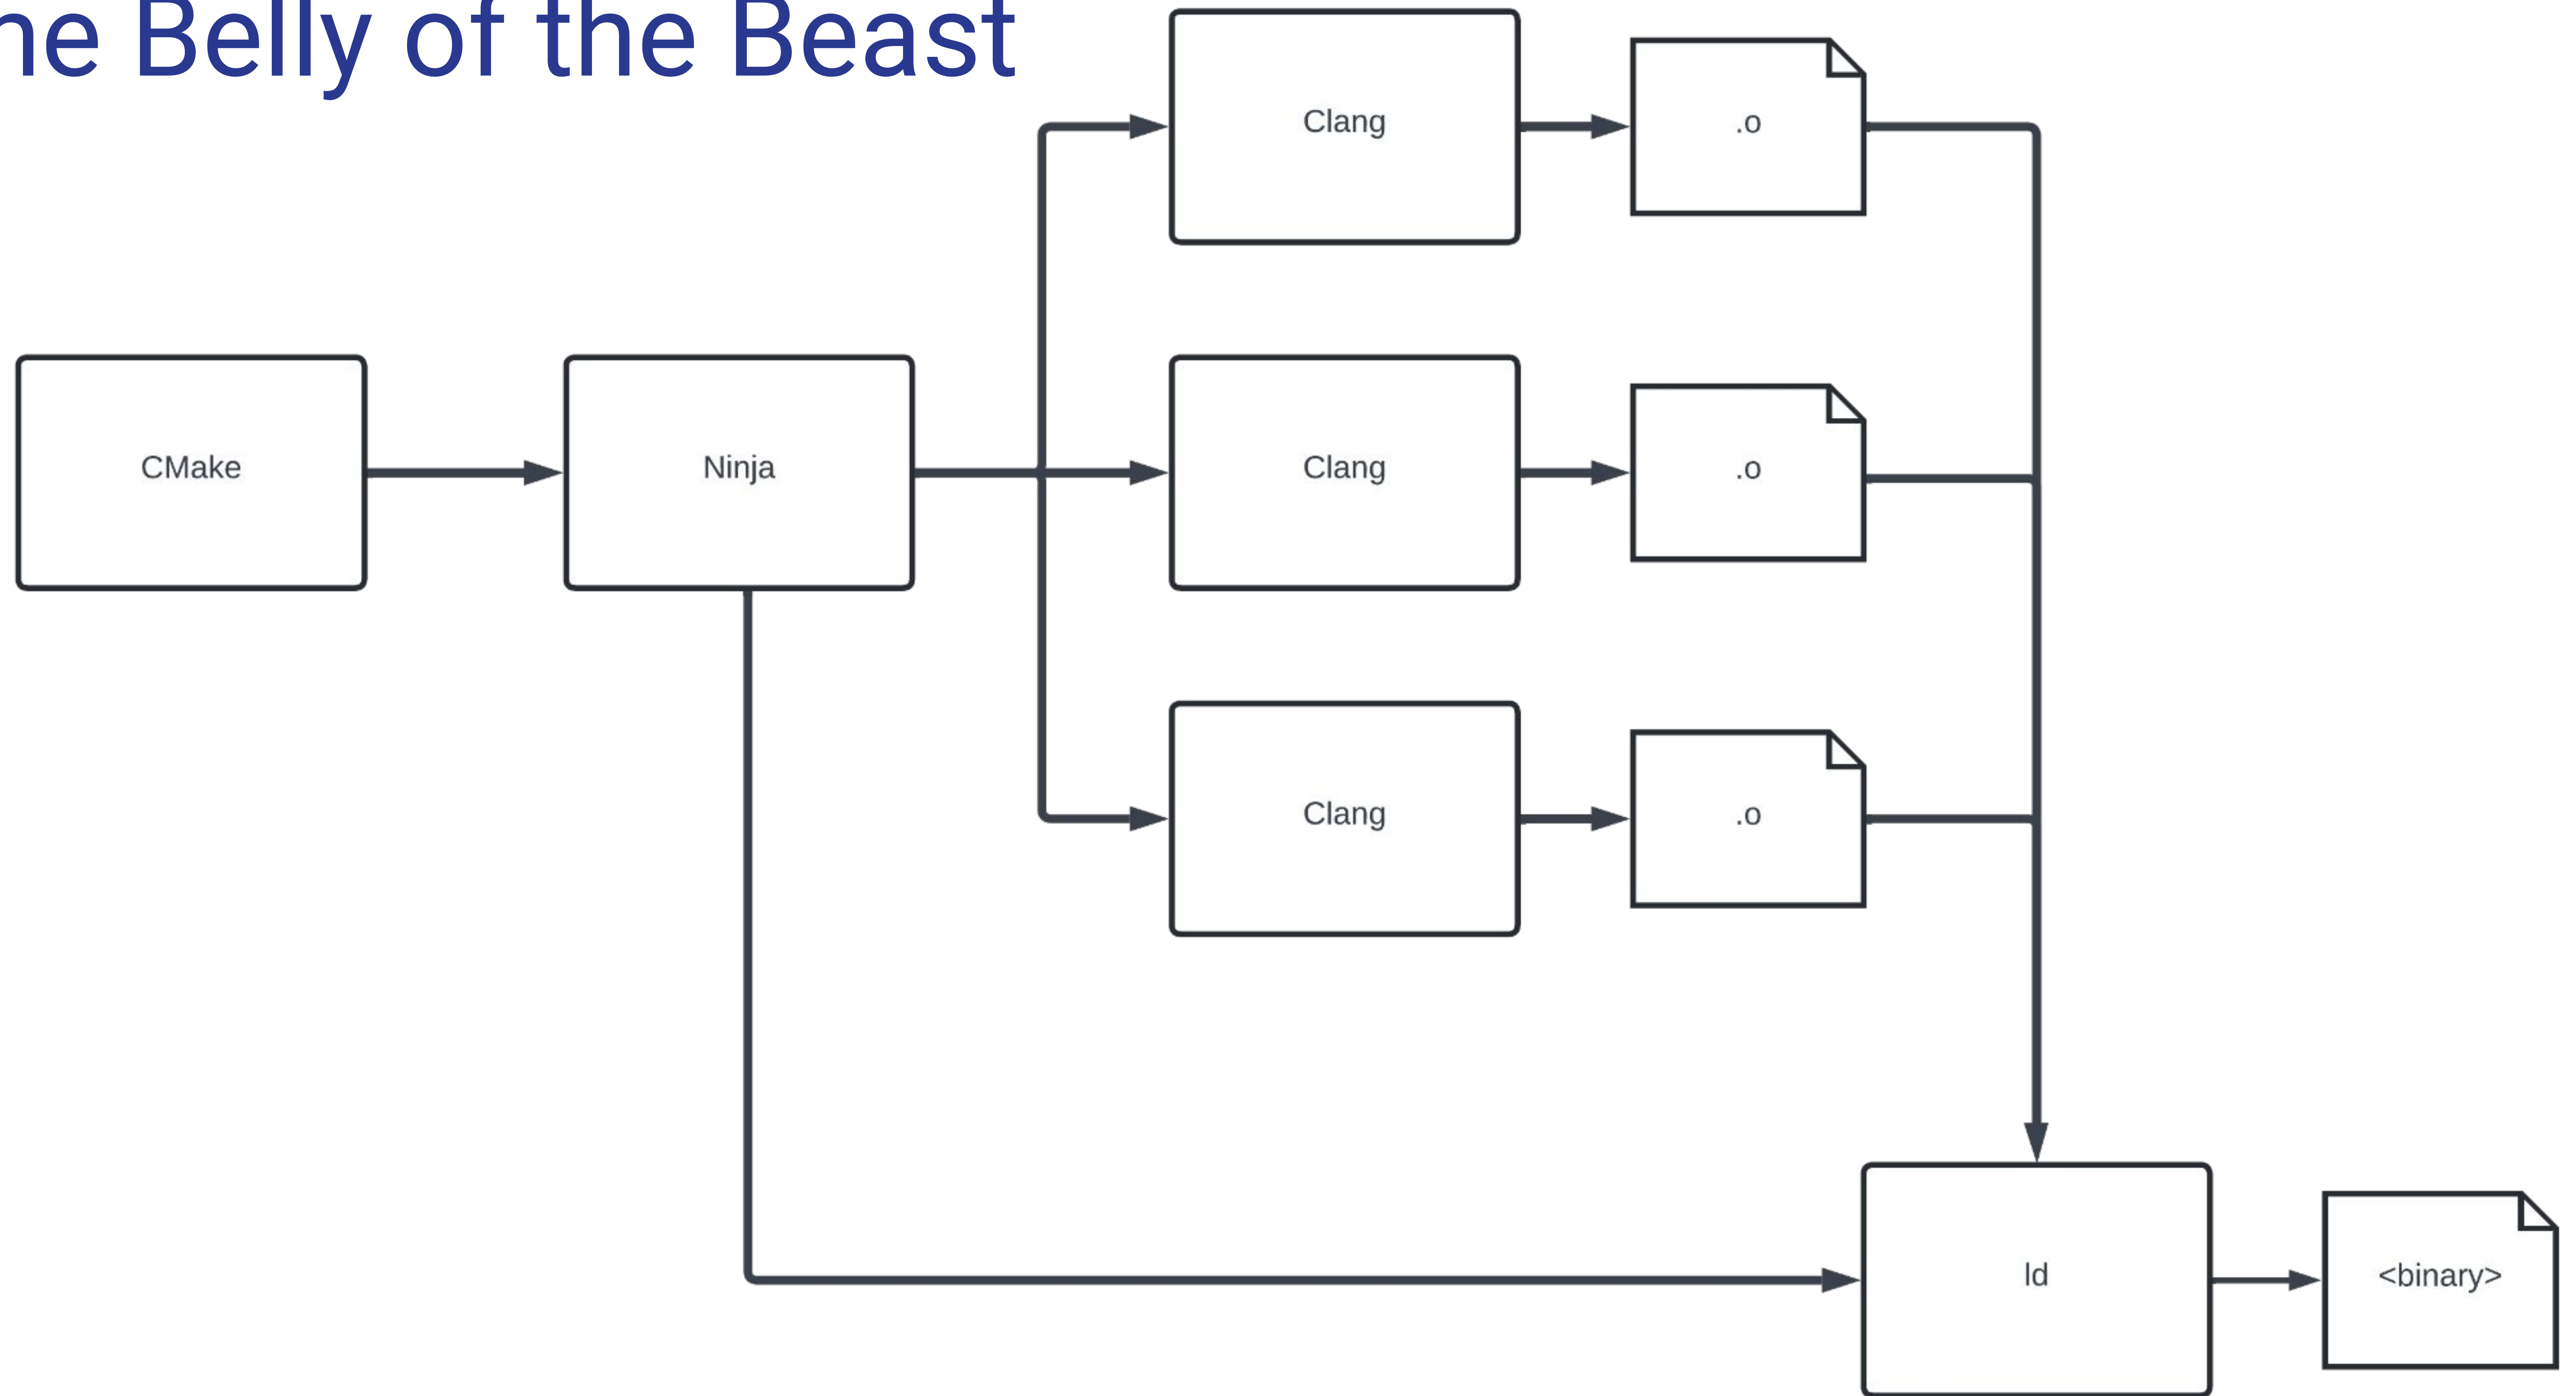


Assumptions

Visualizing Compilation

Single File Gotchas

Into the Belly of the Beast



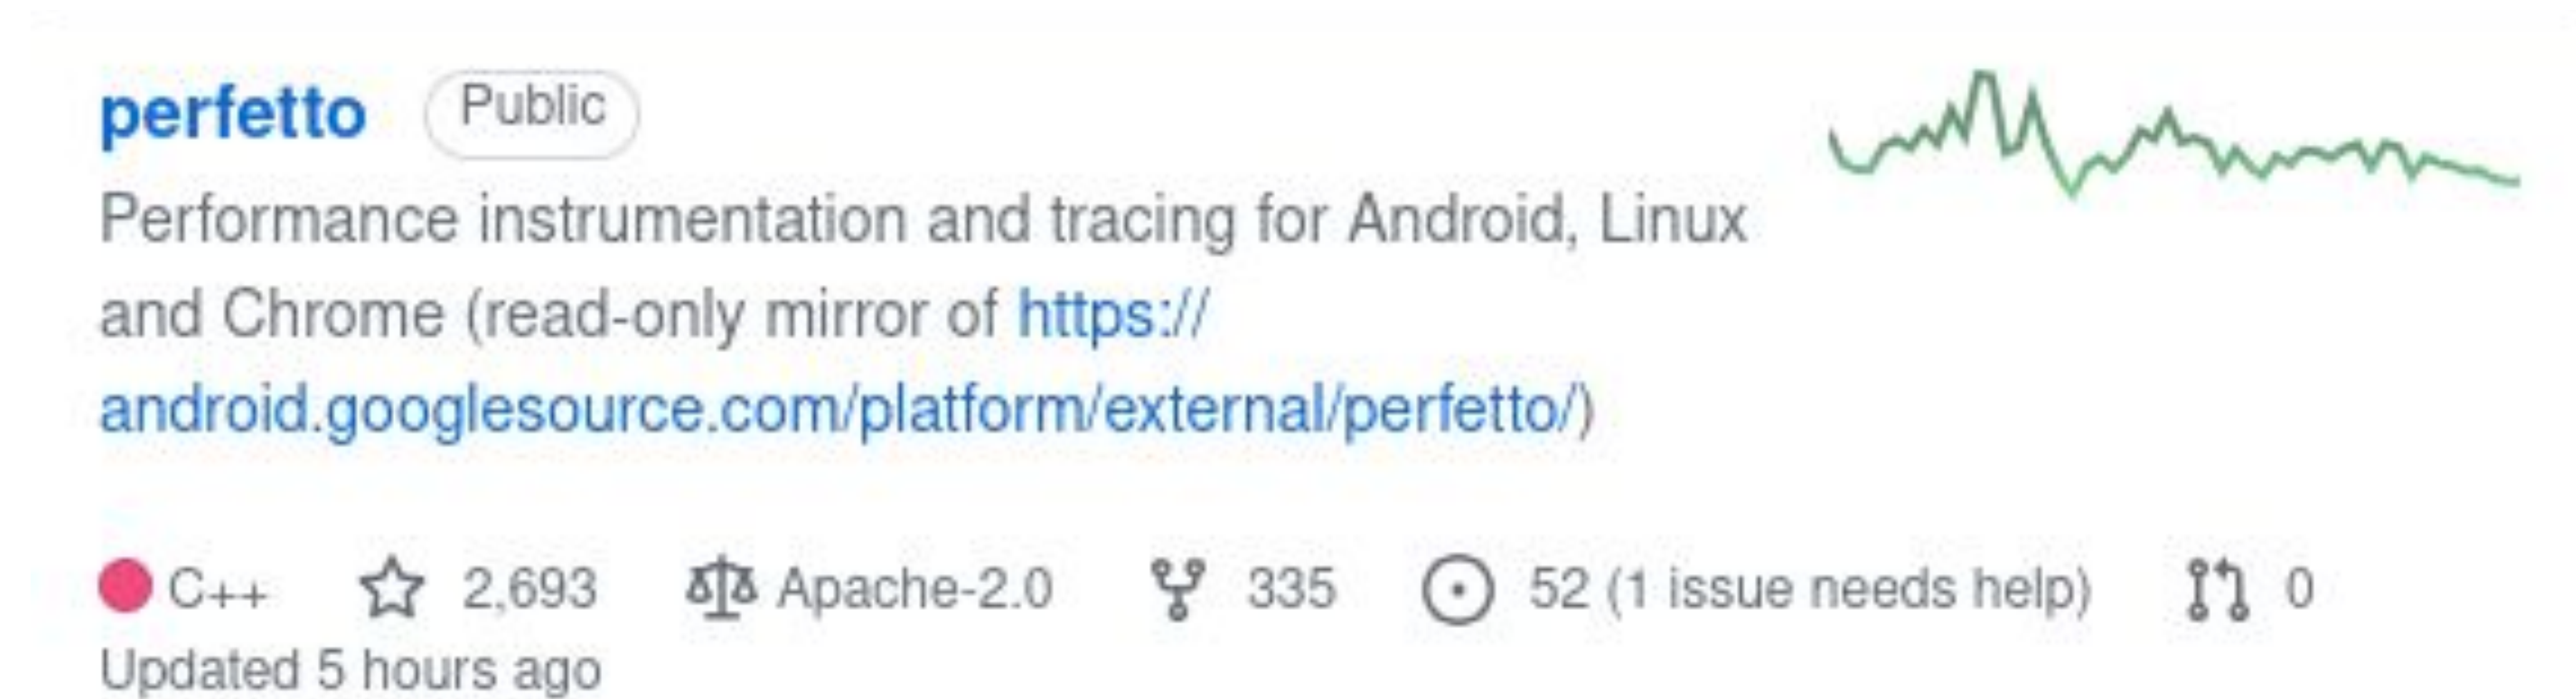
.ninja_log

- Tracks start/stop time in ms
- Name of the output file
- Hash of the command used (compiler, linker, etc)

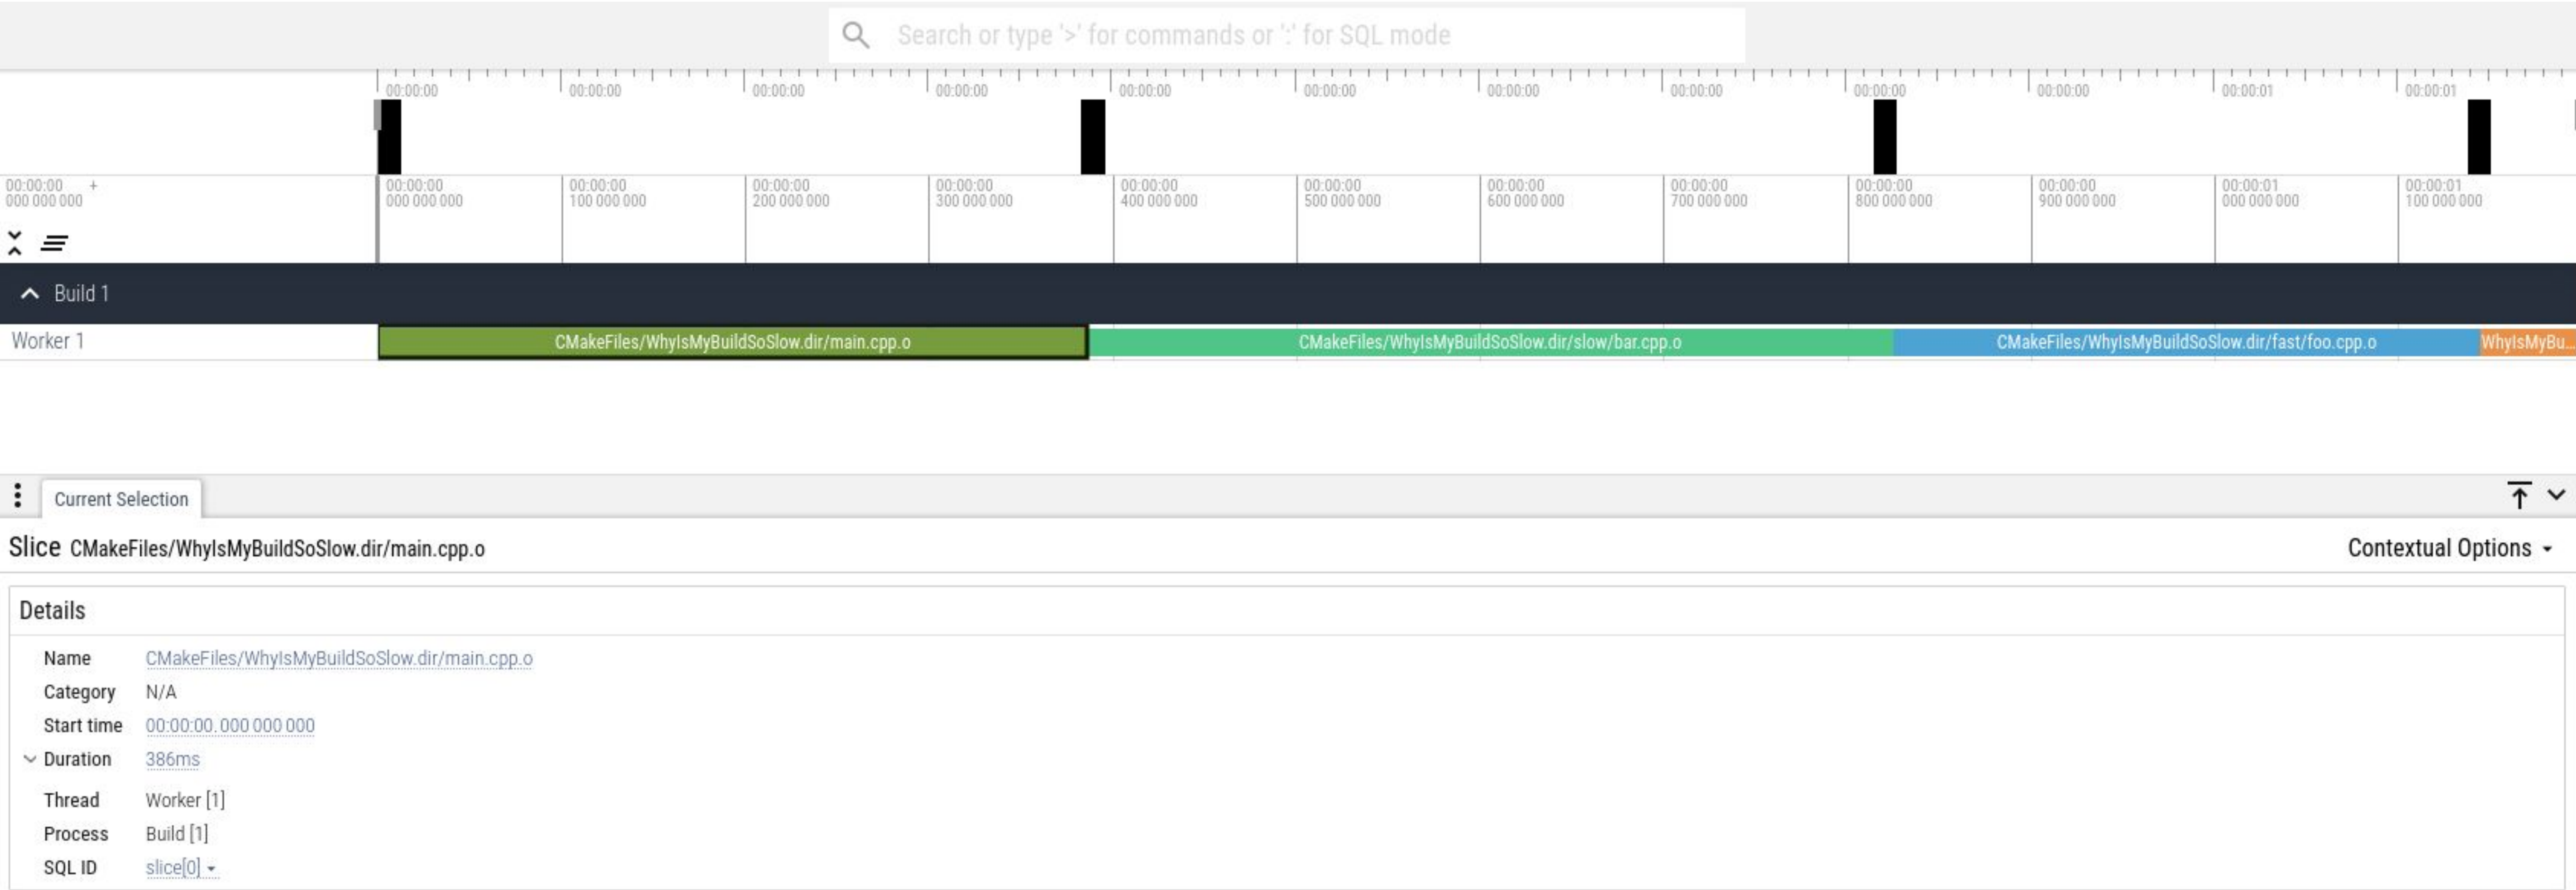
```
1 # ninja log v5
2 0      386      1724180262104460299      CMakeFiles/WhyIsMyBuildSoSlow.dir/main.cpp.o      696e8180c4325128
3 386    825      1724180262544460313      CMakeFiles/WhyIsMyBuildSoSlow.dir/slow/bar.cpp.o      f80e0b2e43465a47
4 825    1144     1724180262864460323      CMakeFiles/WhyIsMyBuildSoSlow.dir/fast/foo.cpp.o      7c134e49e6045f83
5 1144   1198     1724180262920460324      WhyIsMyBuildSoSlow      605a48889fdb150e
```

Visualization

- Interactive Trace Viewer
- <https://ui.perfetto.dev/>
 - Can build and run server locally as well
- Chrome Event Tracing Format JSON
- ... and more!

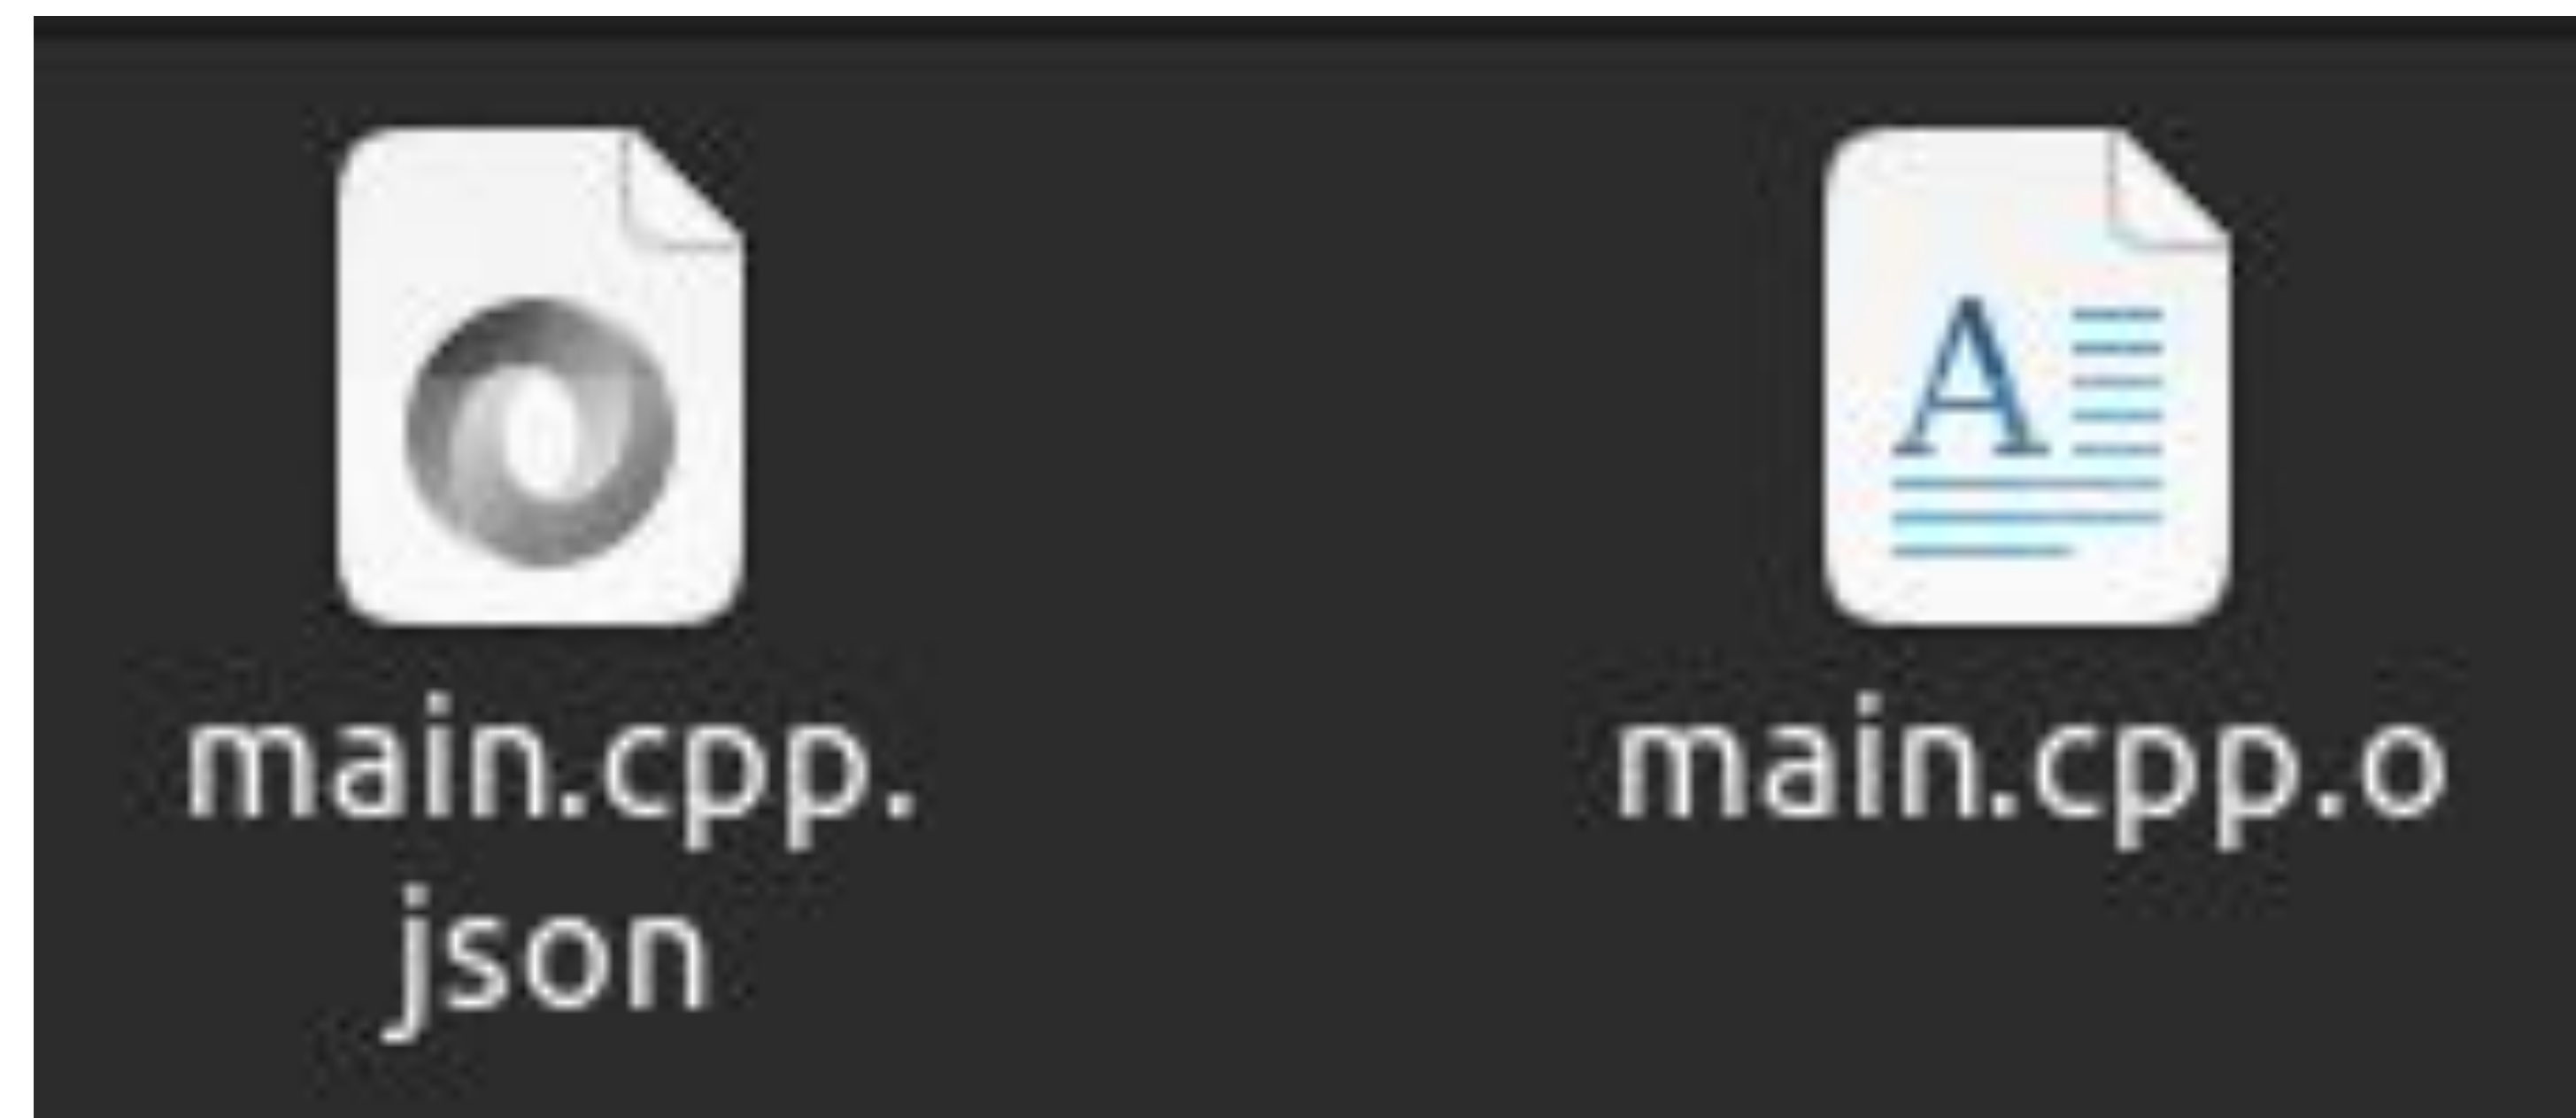


Visualization

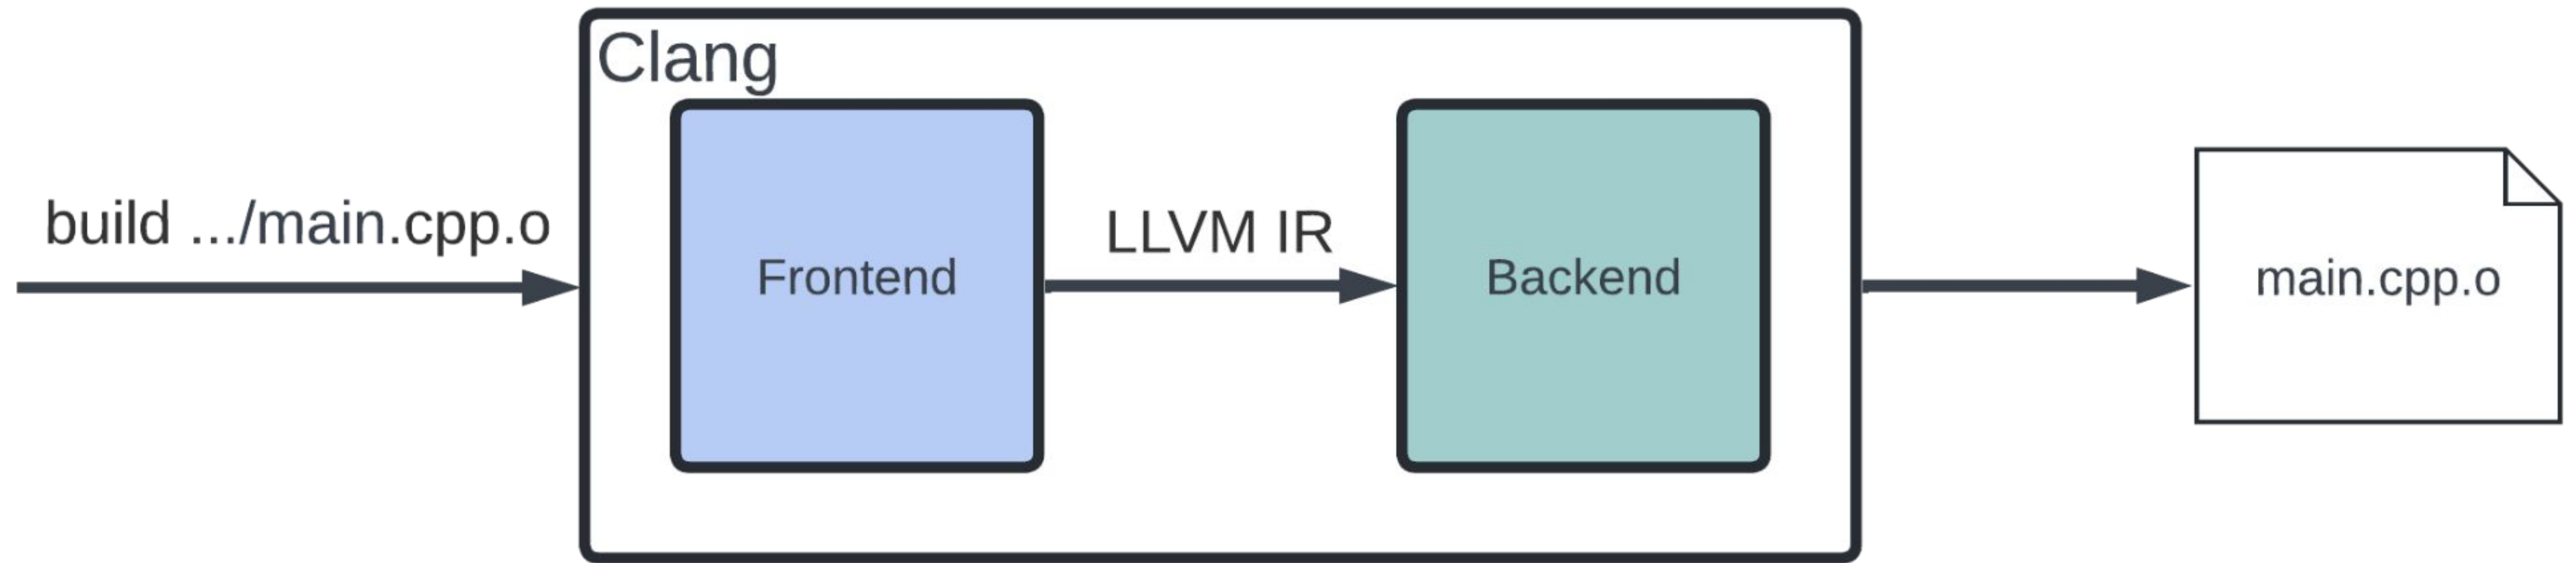


Clang -ftime-trace

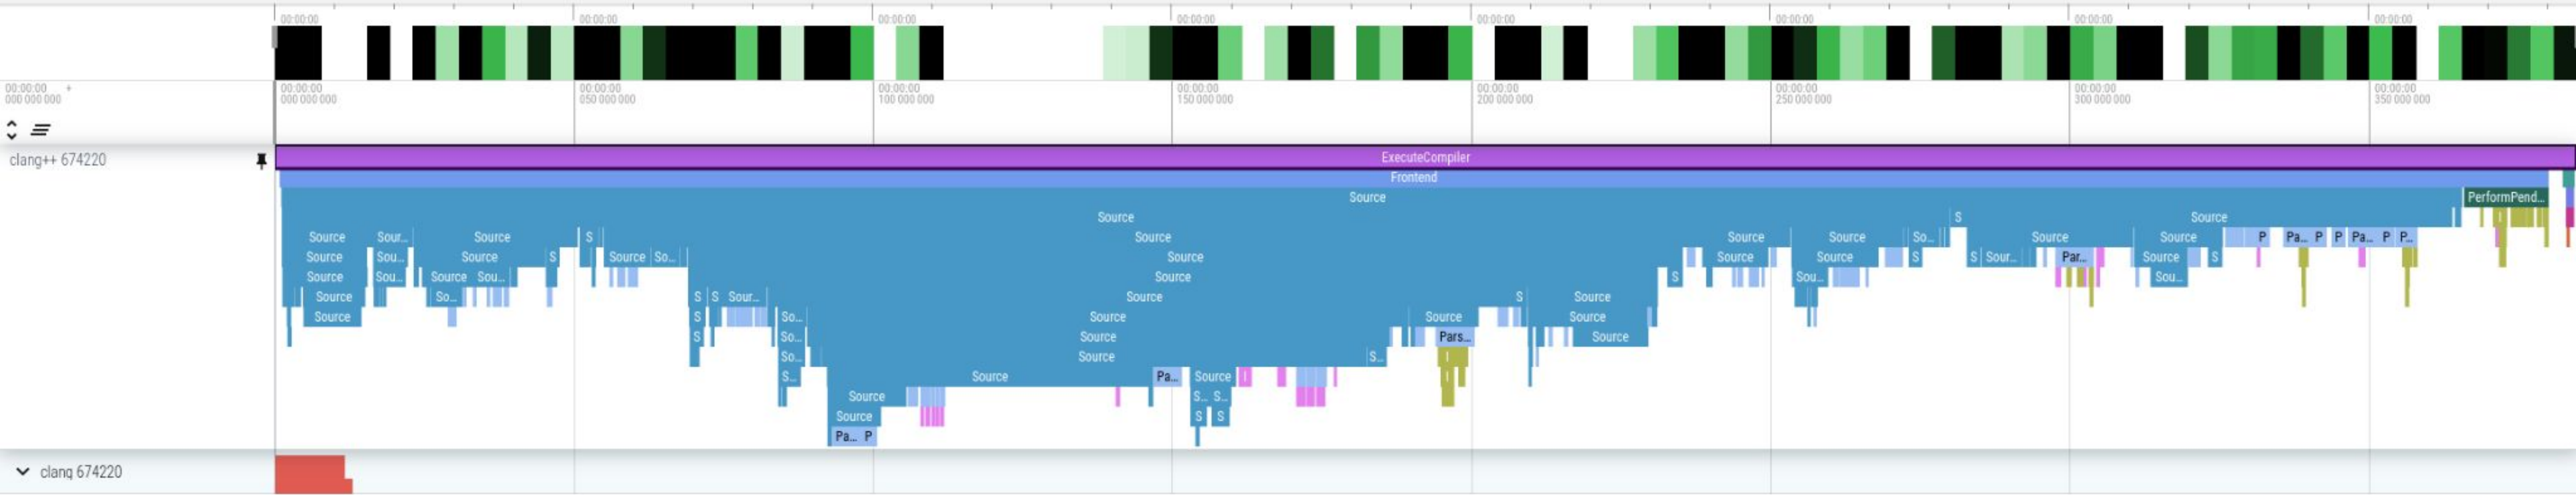
- Generates JSON file based on output filename
- Detailed information on where the compiler spent time



We Need to Go Deeper



Visualization



Current Selection

Slice ExecuteCompiler

Contextual Options

Details

Name

ExecuteCompiler

Category

N/A

Start time

00:00:00.000 011 000

Duration

384ms 718us

Thread

clang++ [674220]

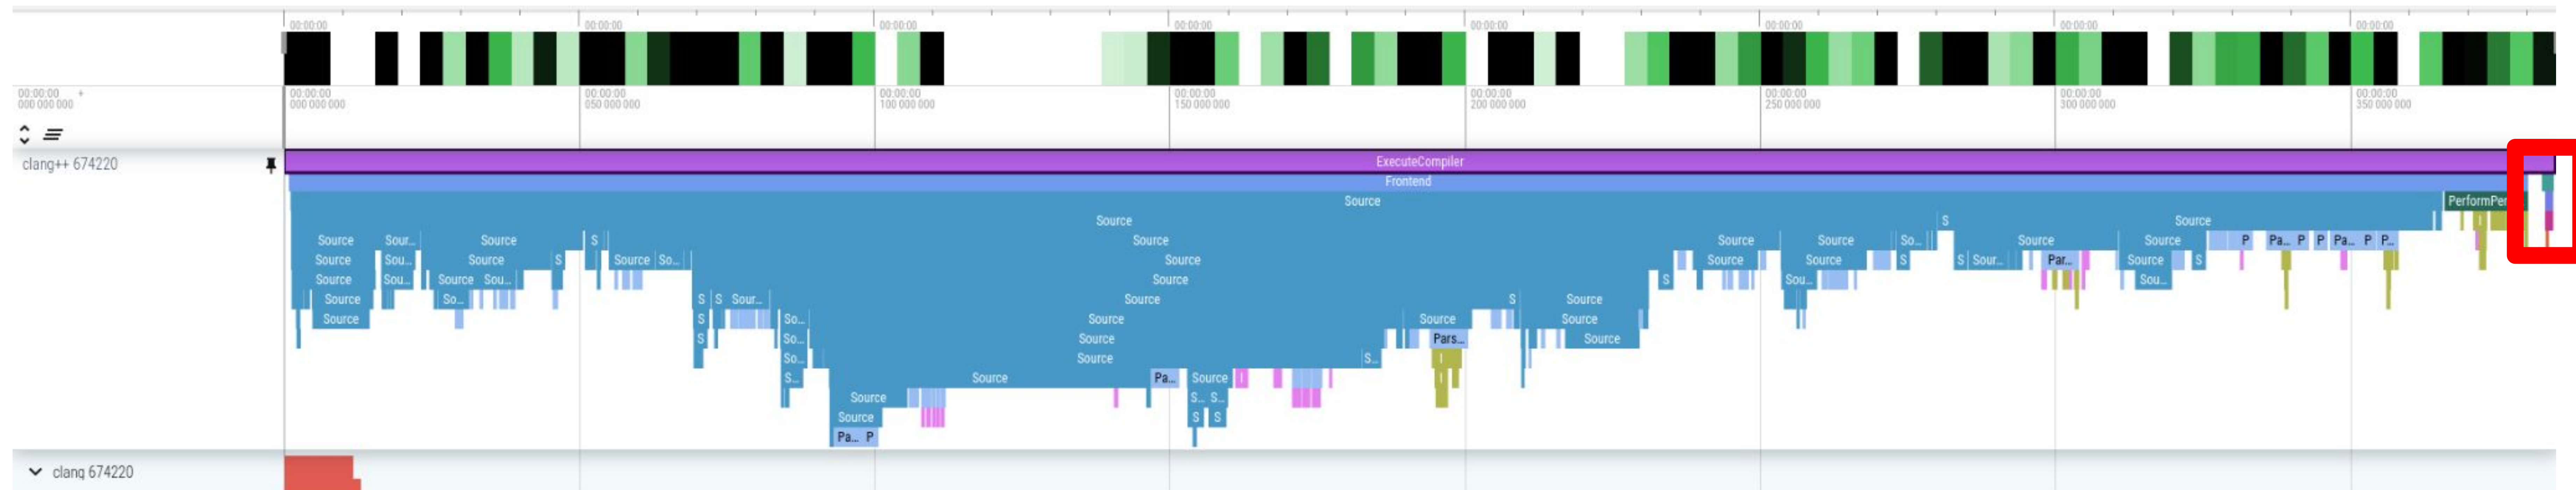
Process

clang [674220]

SQL ID

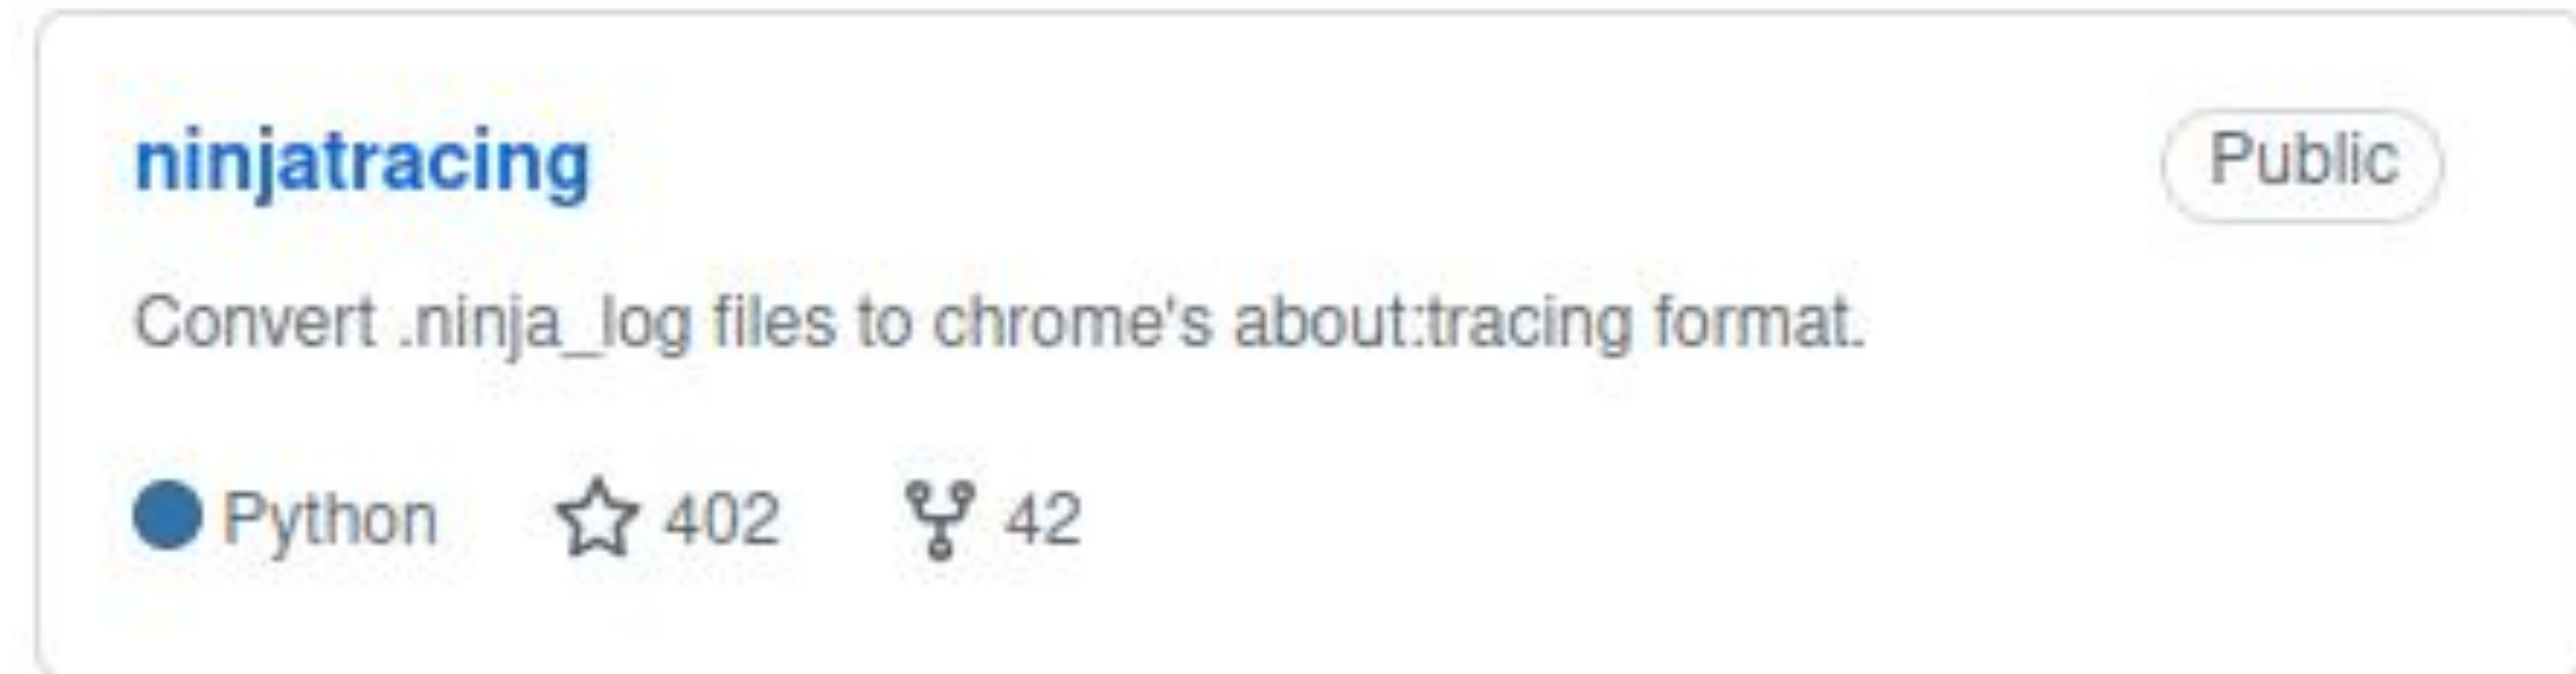
slice[31]

Visualization

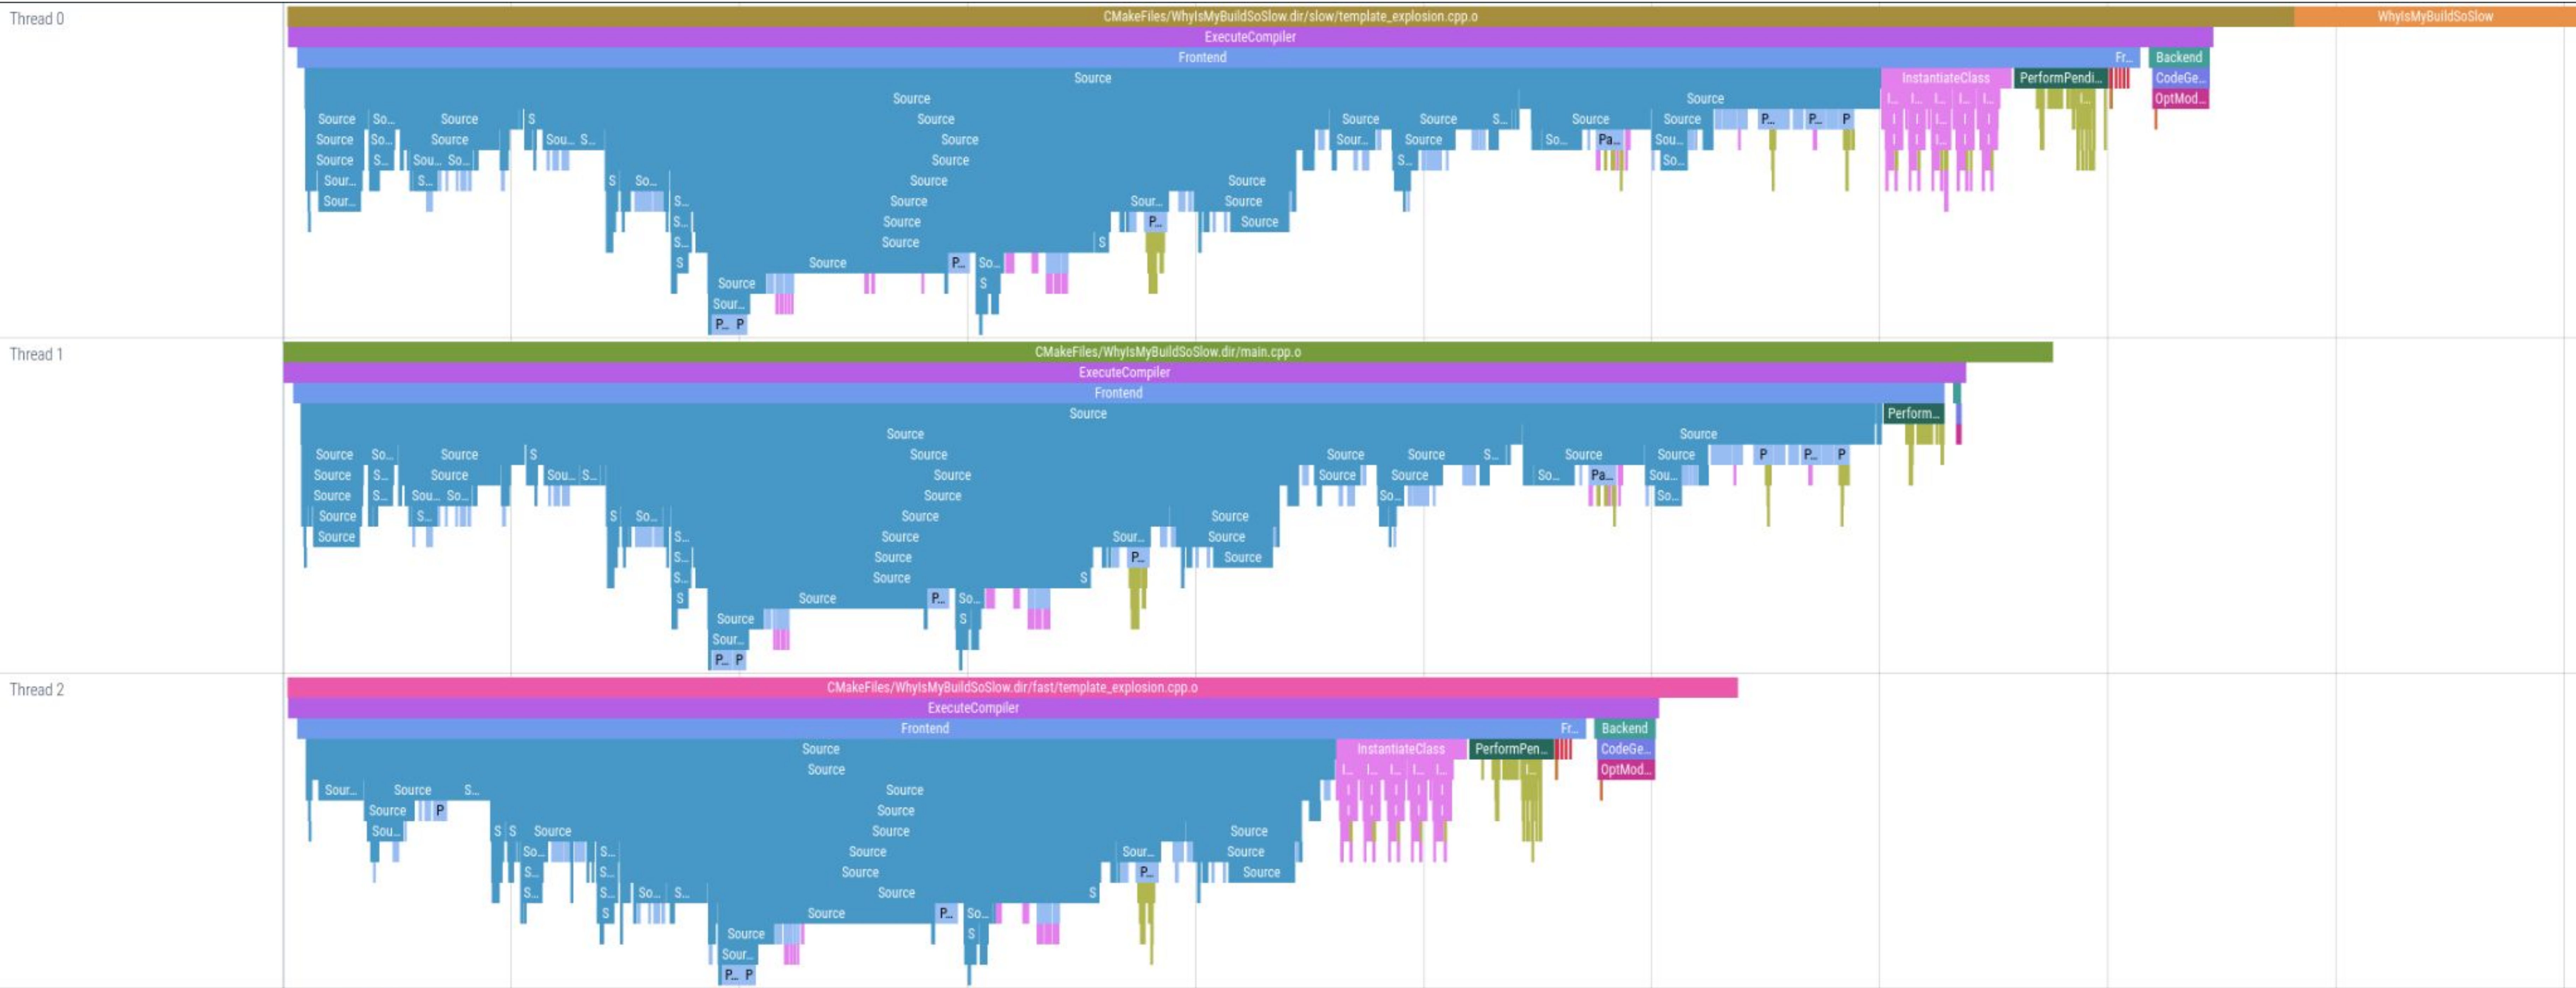


ninjatrace

- Convert .ninja_log files to chrome tracing format
- **Embed -ftime-trace .json files**



Combining Visualizations



Visualizing Compilation

Single File

Project Level

Includes

```
#include <iostream>
```

```
int main() {  
    std::cout << "Hello, world!" << std::endl;  
}
```


Includes

```
#include <iostream>
```

~69,000 Lines!

```
int main() {  
    std::cout << "Hello, world!" << std::endl;  
}
```

Includes

```
clang main.cpp -stdlib=libc++ -E &> prep_main.cpp
```

```
$ du -bh main.cpp
```

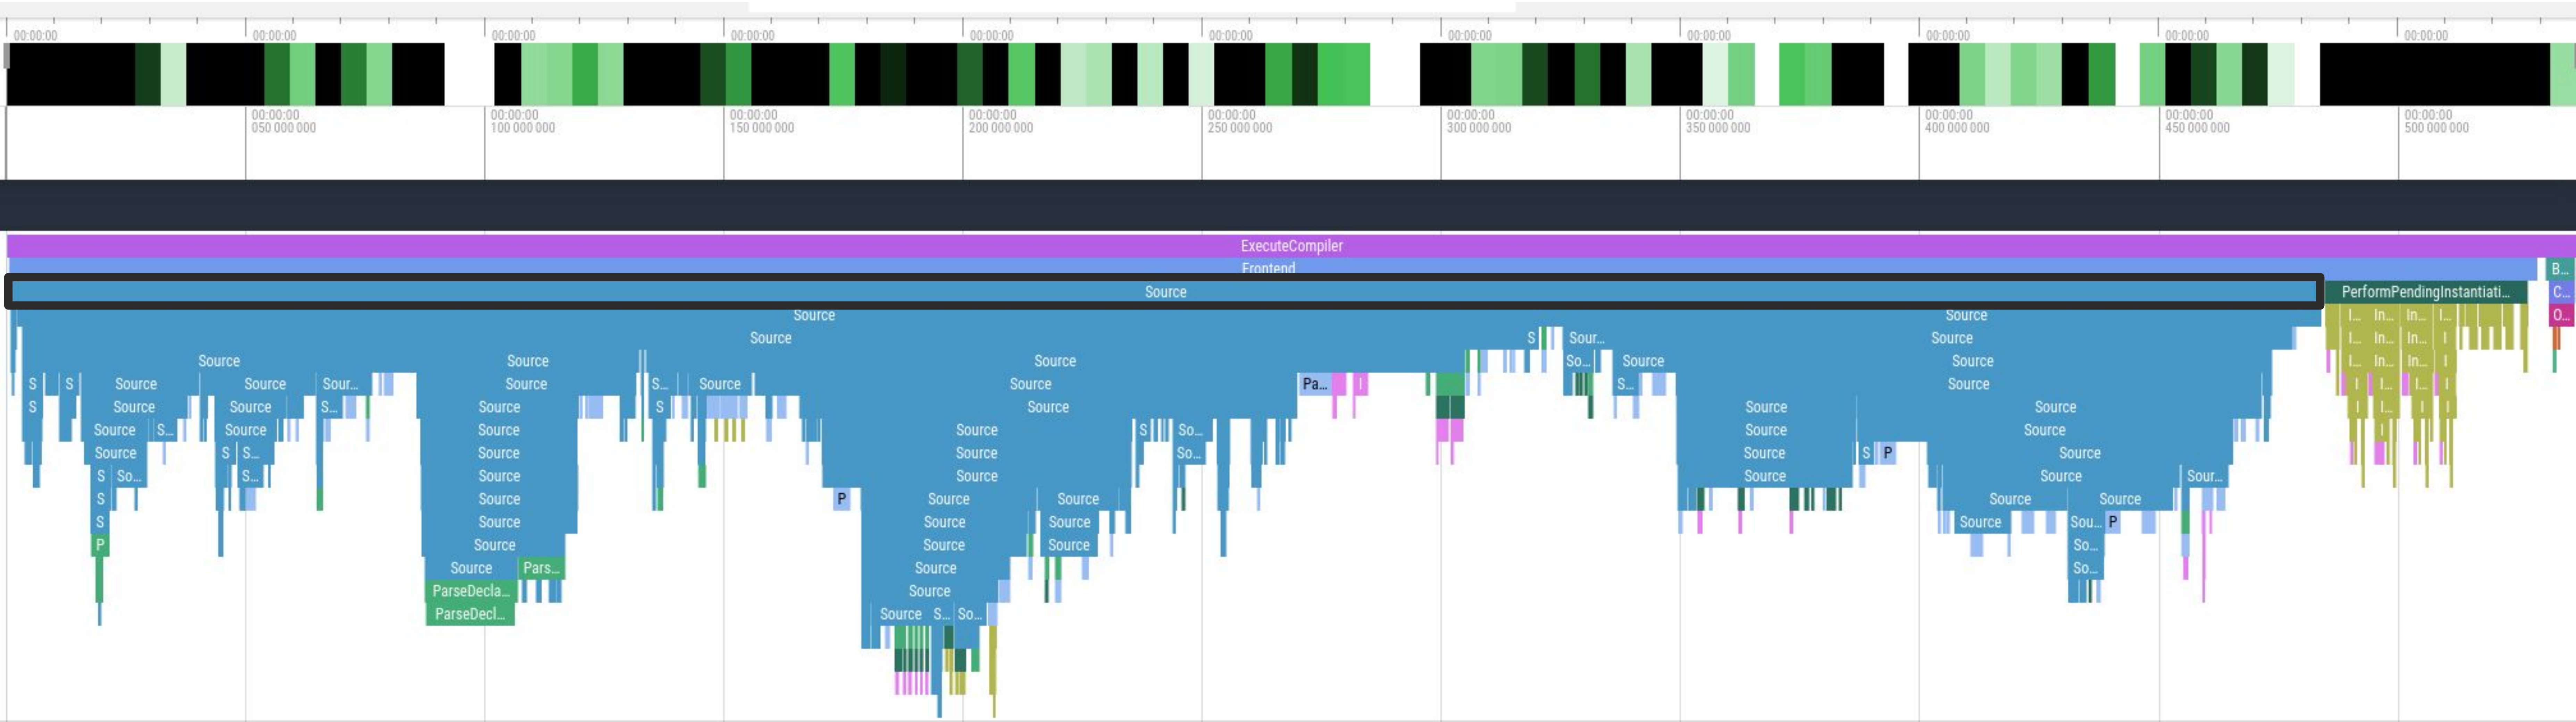
```
82 main.cpp
```

```
$ du -bh prep_main.cpp
```

```
2.9M prep_main.cpp
```


Includes

```
clang main.cpp -stdlib=libc++ -ftime-trace
```



Includes - Recommendations

- Refactor massive header files
 - Having smaller header files gives consumers a better chance at including only what's strictly necessary for them
- Forward declarations
 - Requires no external tooling
 - Frowned upon for entities defined in another project
 - Can obfuscate dependencies
- Include What You Use
 - IWYU Project

Templates

```
template <int N>  
struct Sum {  
    static const int value = N + Sum<N - 1>::value;  
};
```

```
template <>  
struct Sum<0> {  
    static const int value = 0;  
};
```

Templates

```
#include <iostream>
```

```
template <int N>  
struct Sum {  
    static const int value = N + Sum<N - 1>::value;  
};
```

```
template <>  
struct Sum<0> {  
    static const int value = 0;  
};
```

```
int main()  
{  
    std::cout << Sum<5>::value << std::endl;  
}
```


Templates

```
#include <iostream>
```

```
template <int N>
struct Sum {
    static const int value = N + Sum<N - 1>::value;
};
```

```
template <>
struct Sum<0> {
    static const int value = 0;
};
```

```
int main()
{
    std::cout << Sum<5>::value << std::endl;
}
```

cppinsights.io

```
#include <iostream>
```

```
template<int N>
struct Sum
{
    static const int value = N + Sum<N - 1>::value;
};
```

```
/* First instantiated from: insights.cpp:5 */
#ifdef INSIGHTS_USE_TEMPLATE
template<>
struct Sum<4>
{
    static const int value = 4 + Sum<3>::value;
};
```

```
#endif
/* First instantiated from: insights.cpp:5 */
#ifdef INSIGHTS_USE_TEMPLATE
template<>
struct Sum<3>
{
    static const int value = 3 + Sum<2>::value;
};
```

```
// Sum<2>, Sum<1>, Sum<0>
```

```
int main()
{
    std::cout.operator<<(Sum<5>::value).operator<<(std::endl);
    return 0;
}
```

Templates

Sum<**8192**>::value



~33,000 Lines!



Templates

- Template Instantiation is a type of monomorphization
- Trade build time (and binary size) for faster run times



Templates - Recommendations

- Consider whether you need to use Template Meta Programming
 - **constexpr** and **constexpr** can go a long way
 - Is there a way to more directly express intent to the compiler?



Template-less Meta-programming
Kris Jusiak

- Re-evaluate your API
 - Do you need to be generic over that extra type?
 - Can you eliminate recursion?



Single File

Project Level

Higher Order

ClangBuildAnalyzer

ClangBuildAnalyzer

Public

Clang build analysis tool using -ftime-trace

● C++ ☆ 982 🍴 61

- High level overview of hot spots
 - Files that took longest to parse
 - Templates that took longest to instantiate
 - Functions that took longest to compile
 - **Expensive headers (with include chains!)**
- Support for incremental builds

Translation Units

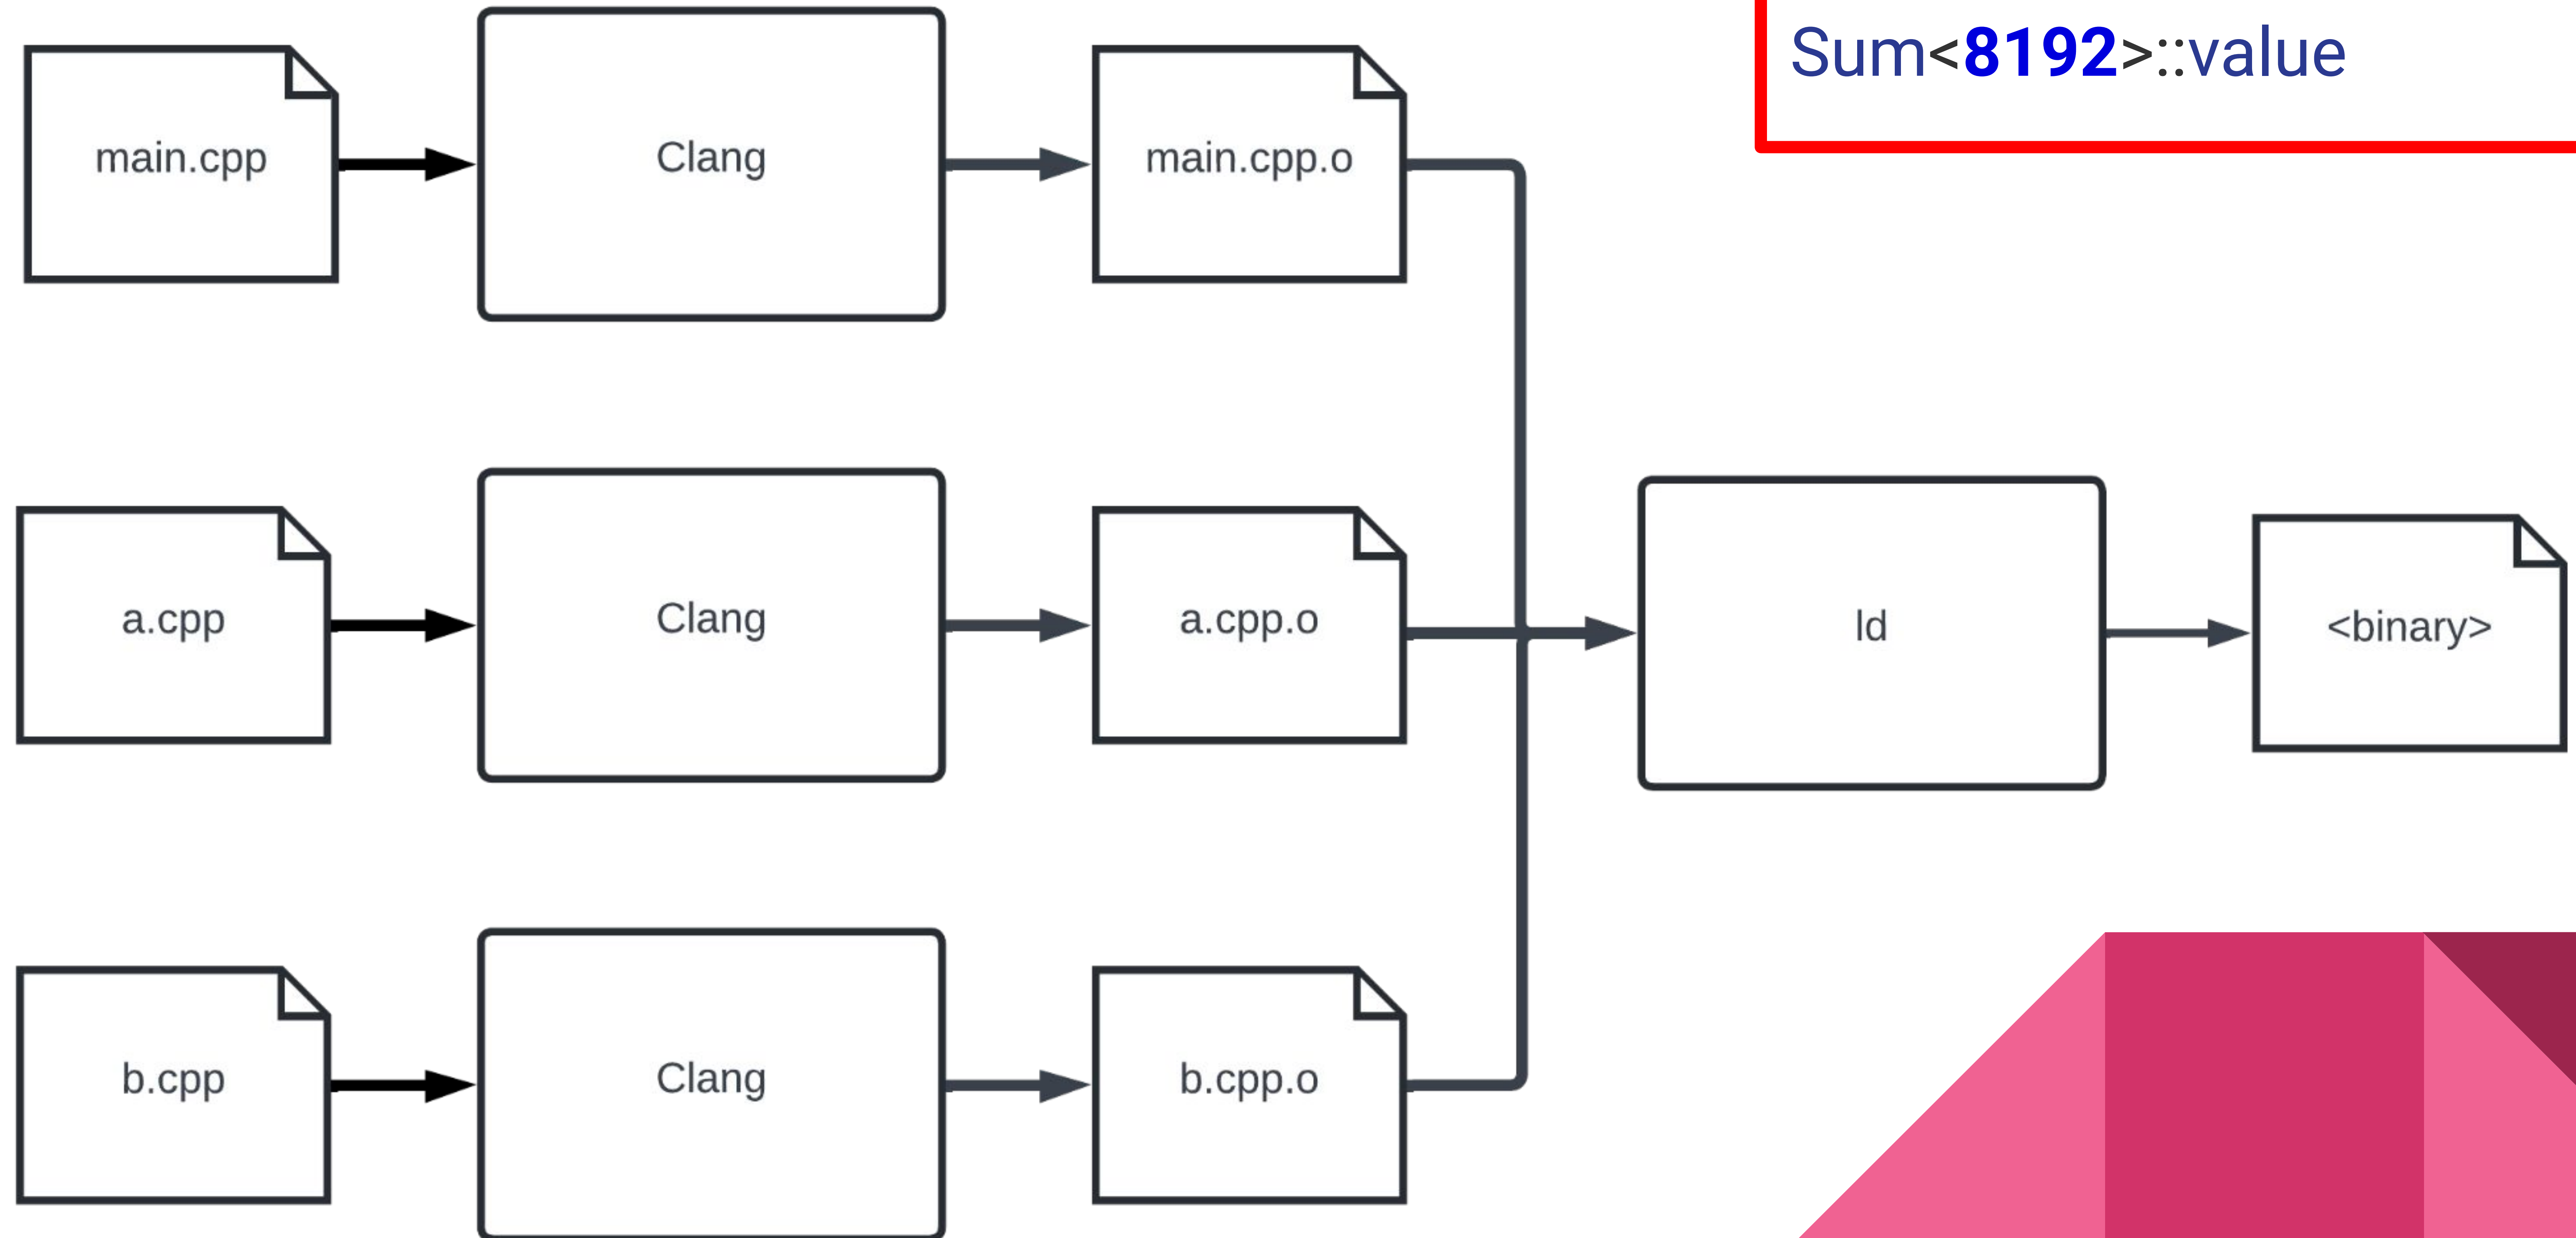
```
// templates.hpp
```

```
#include <iostream>
```

```
template <int N>  
struct Sum {  
    static const int value = N + Sum<N - 1>::value;  
};
```


```
template <>  
struct Sum<0> {  
    static const int value = 0;  
};
```

Translation Units



Translation Units

≈1.91 seconds



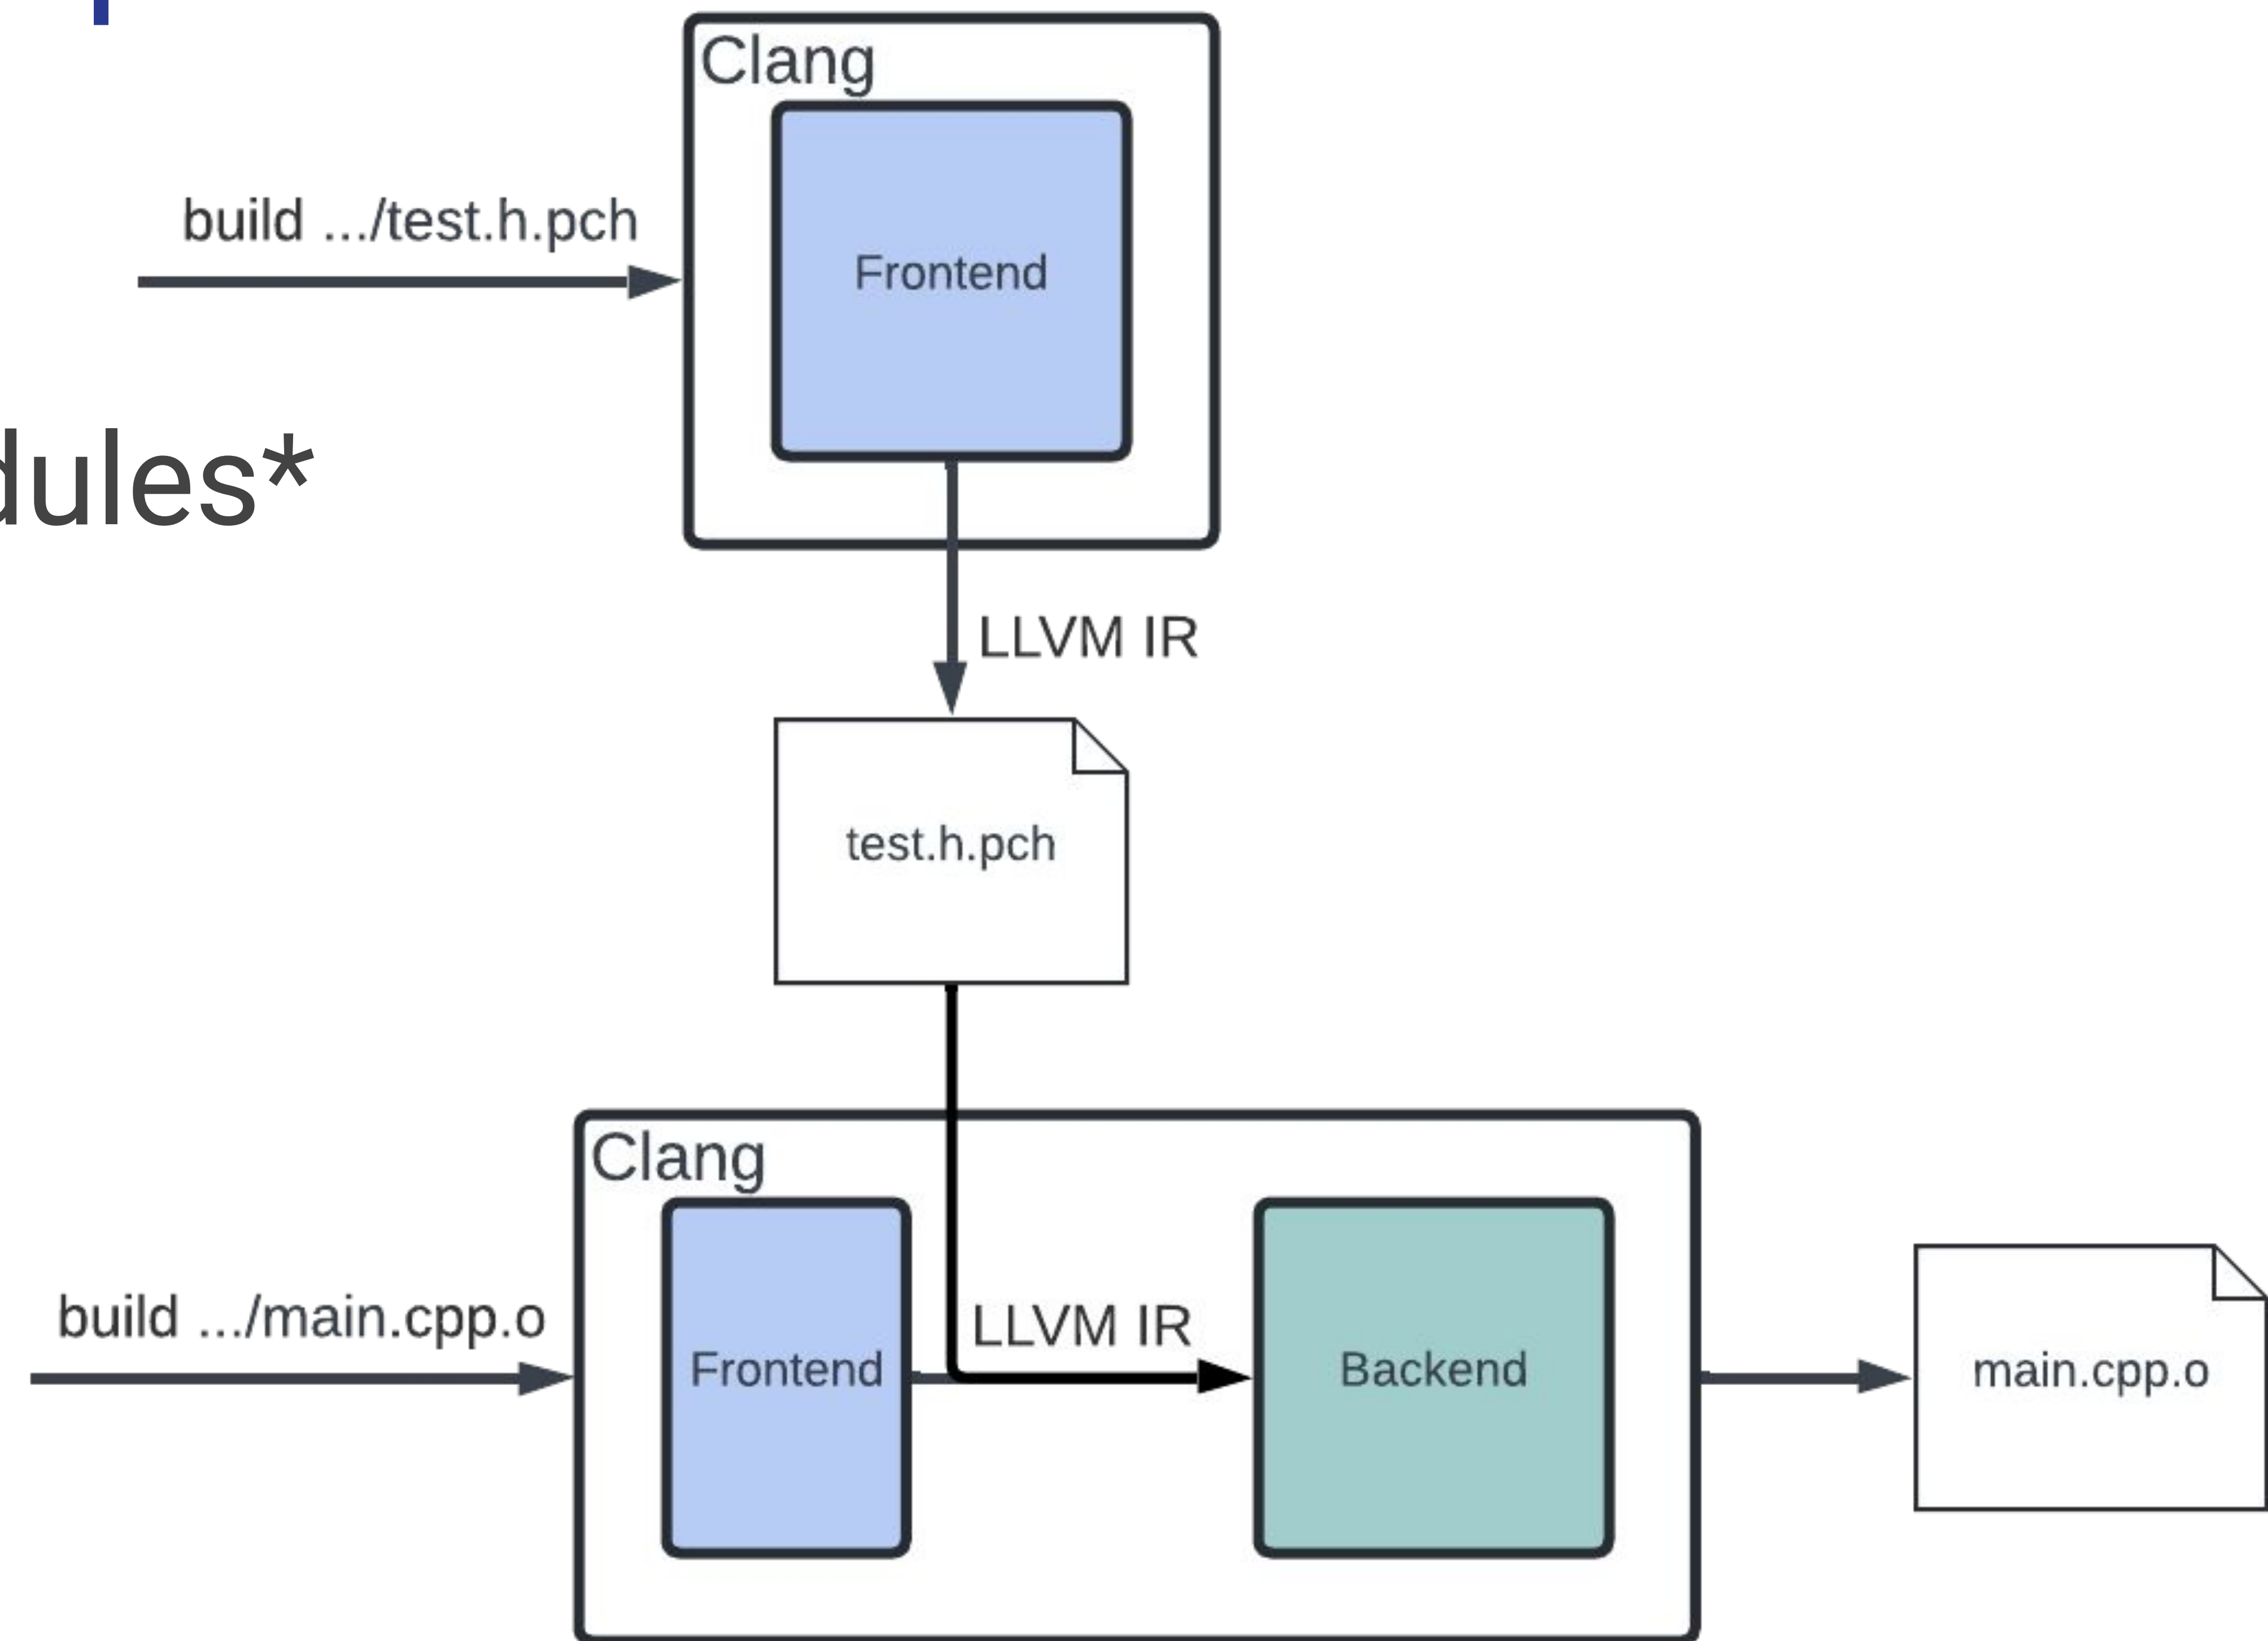
Translation Units

- `#include` and **template** costs are per translation unit
 - `#pragma once` and header guards won't save you



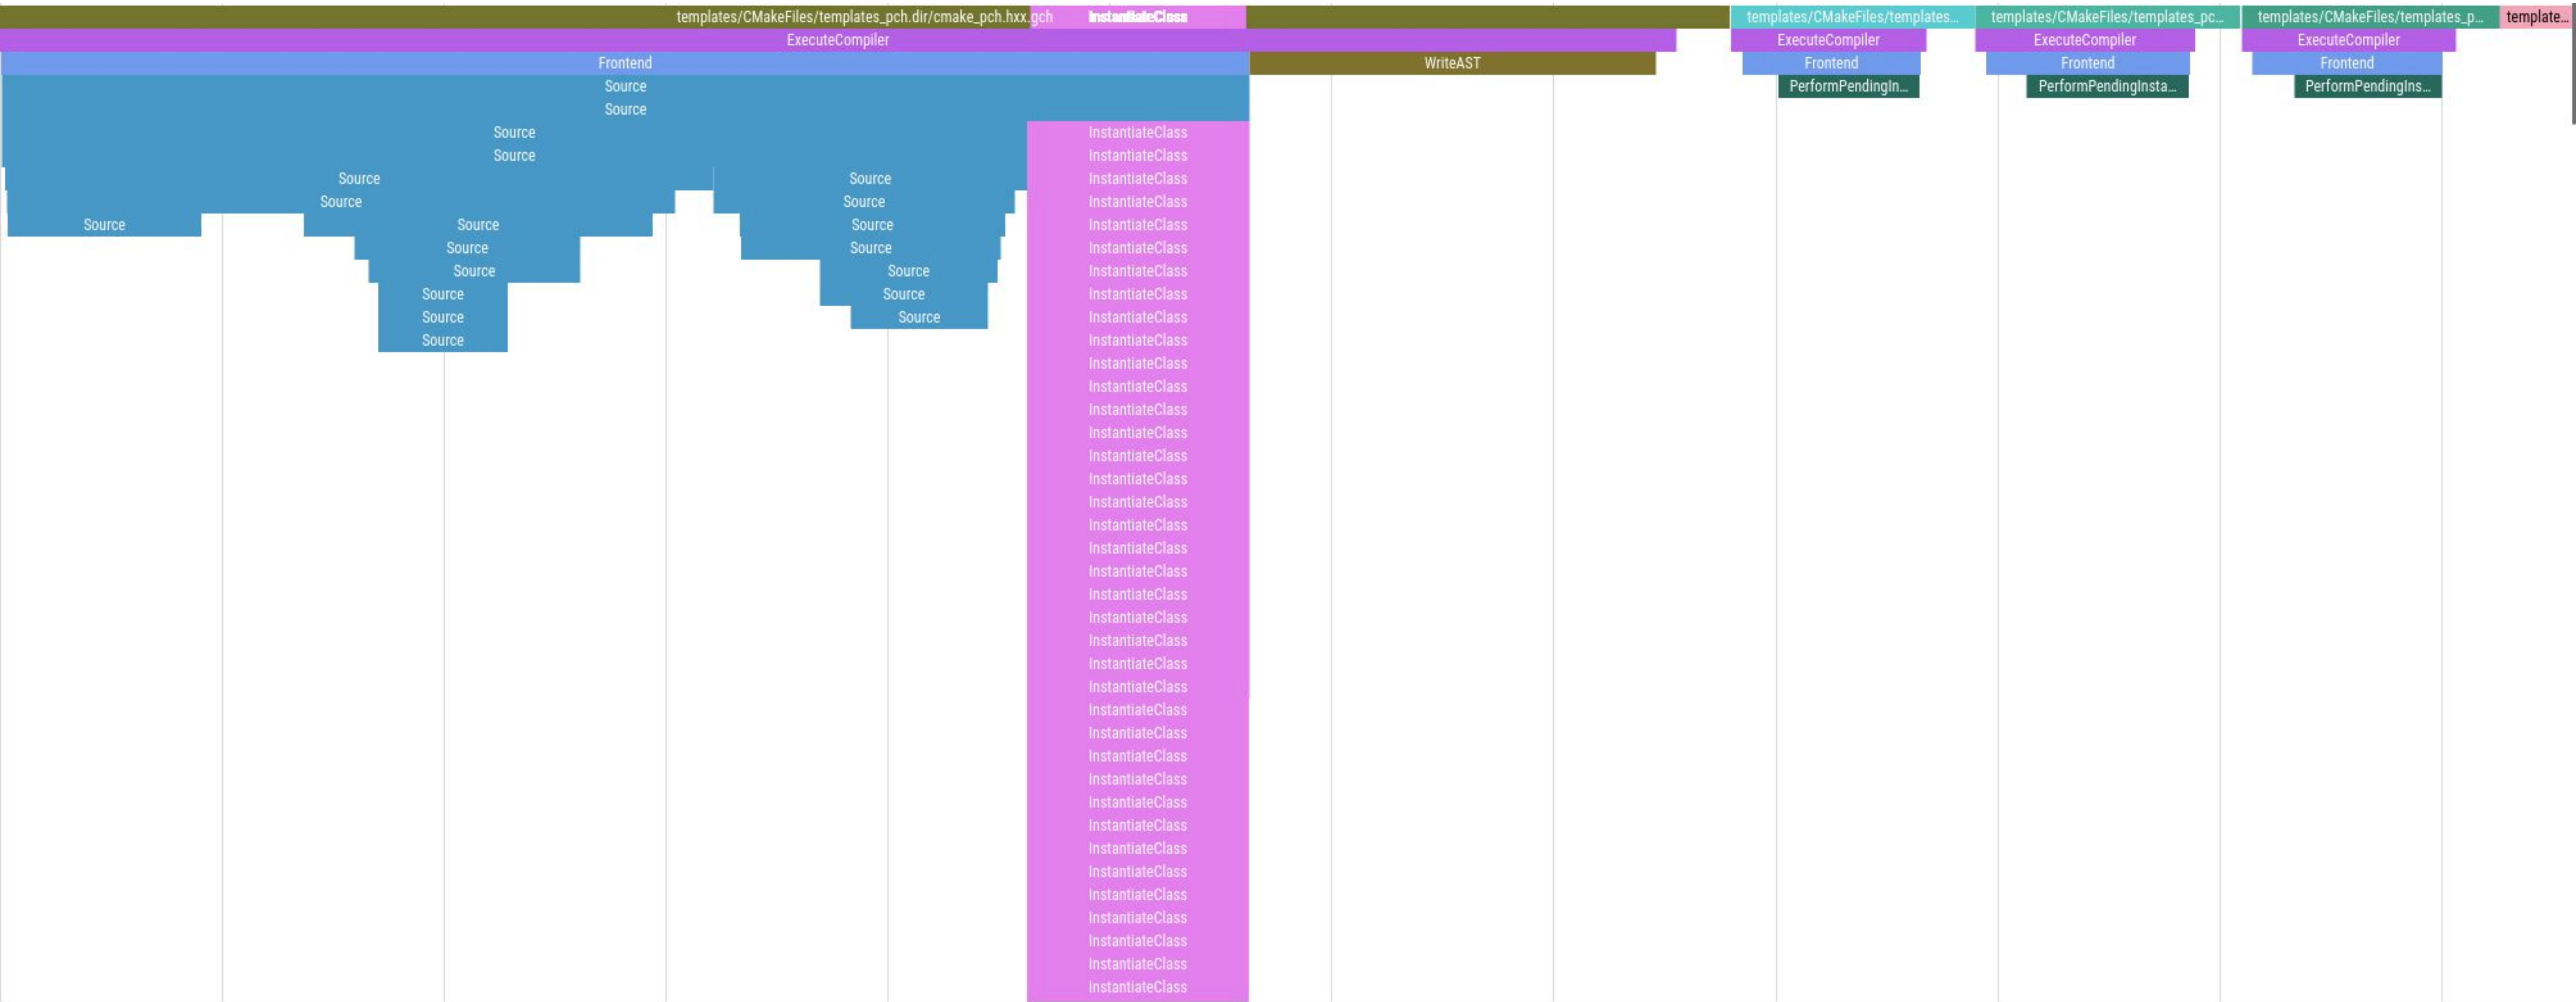
Wait, Why Do Modules Help?

- Eliminate redundant parsing
- Precompiled Headers \approx Modules*
 - Modules are more flexible



Using Precompiled Headers

≈1.16 seconds



Translation Units - Recommendations

- Precompiled headers
 - Not a standard feature
 - [target_precompile_headers](#)
- Modules
 - Standard feature (as of C++20)!
 - [Few compilers have full support](#)
 - [cmake-cxxmodules](#)

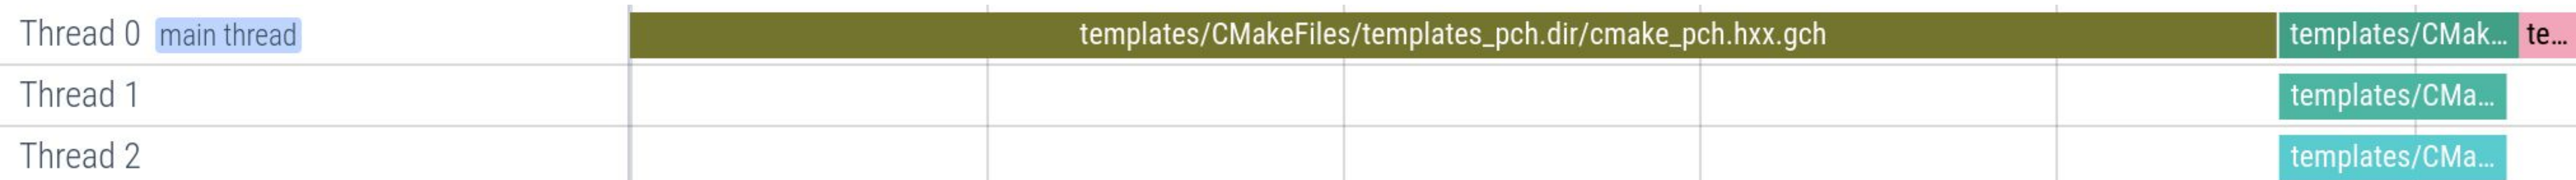


import CMake; // Mastering C++ Modules

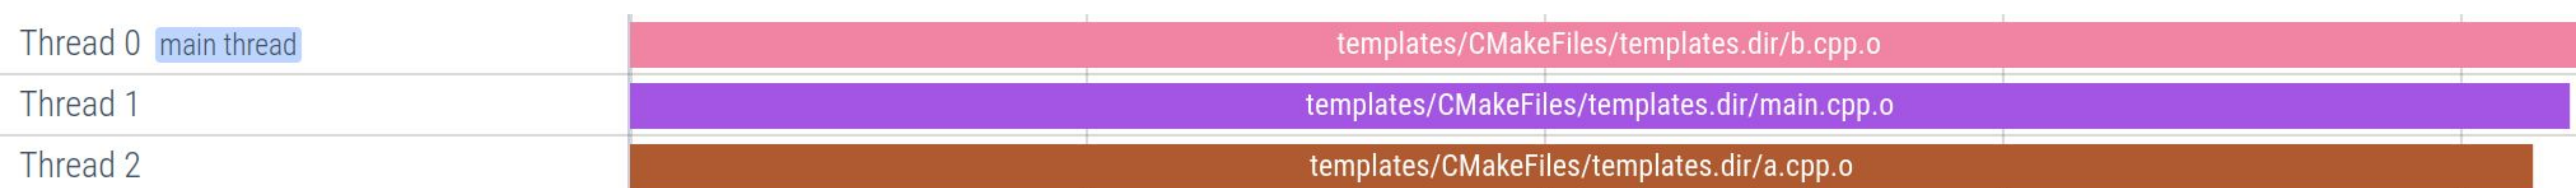
Bill Hoffman

Sometimes, Duplicating Work is Faster

1089ms

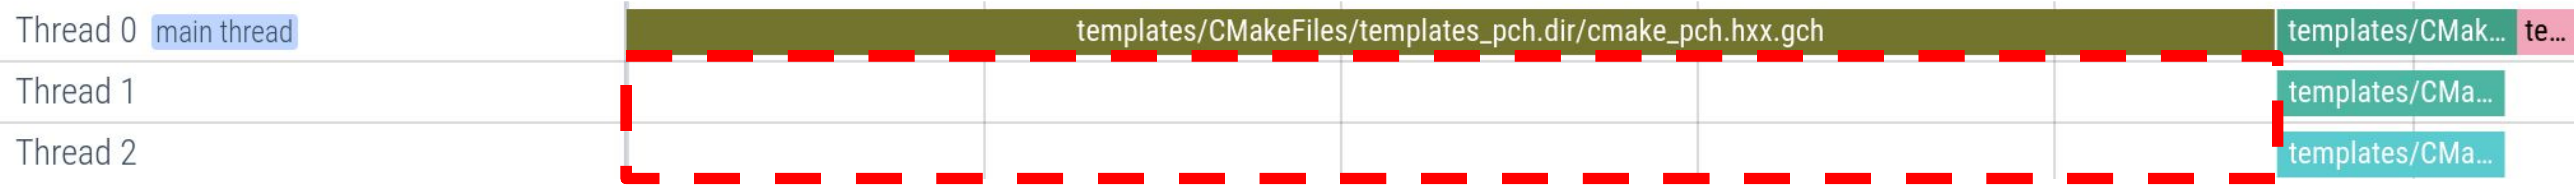


890ms



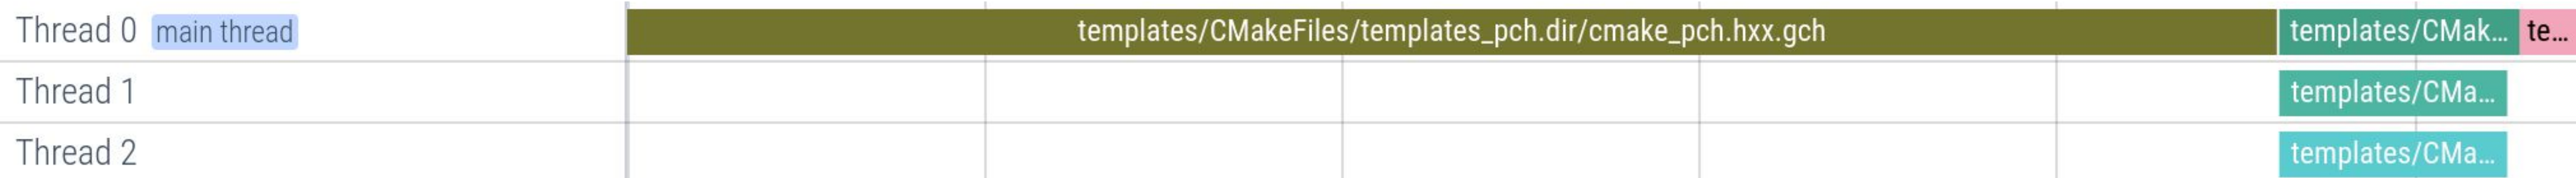
But...

1089ms



Free House for You, Jim!

1089ms



1137ms



Poor Dependency Management

- Builds should be purely functional
 - Single-core and parallel builds should just work
- Avoid large dependency bottlenecks
 - Can force the build to be synchronous
 - Especially important when generating code (i.e., Protobuf)
- Prefer smaller targets
- **Explicitly expressed dependencies enables efficient hardware utilization while maintaining build correctness**


Compiler Include Paths

```
#include <iostream>
```


```
int main() {  
    std::cout << "Hello, world!" << std::endl;  
}
```

```
clang -I... -stdlib=libc++ main.cpp
```


Include Paths Impact on Build Time



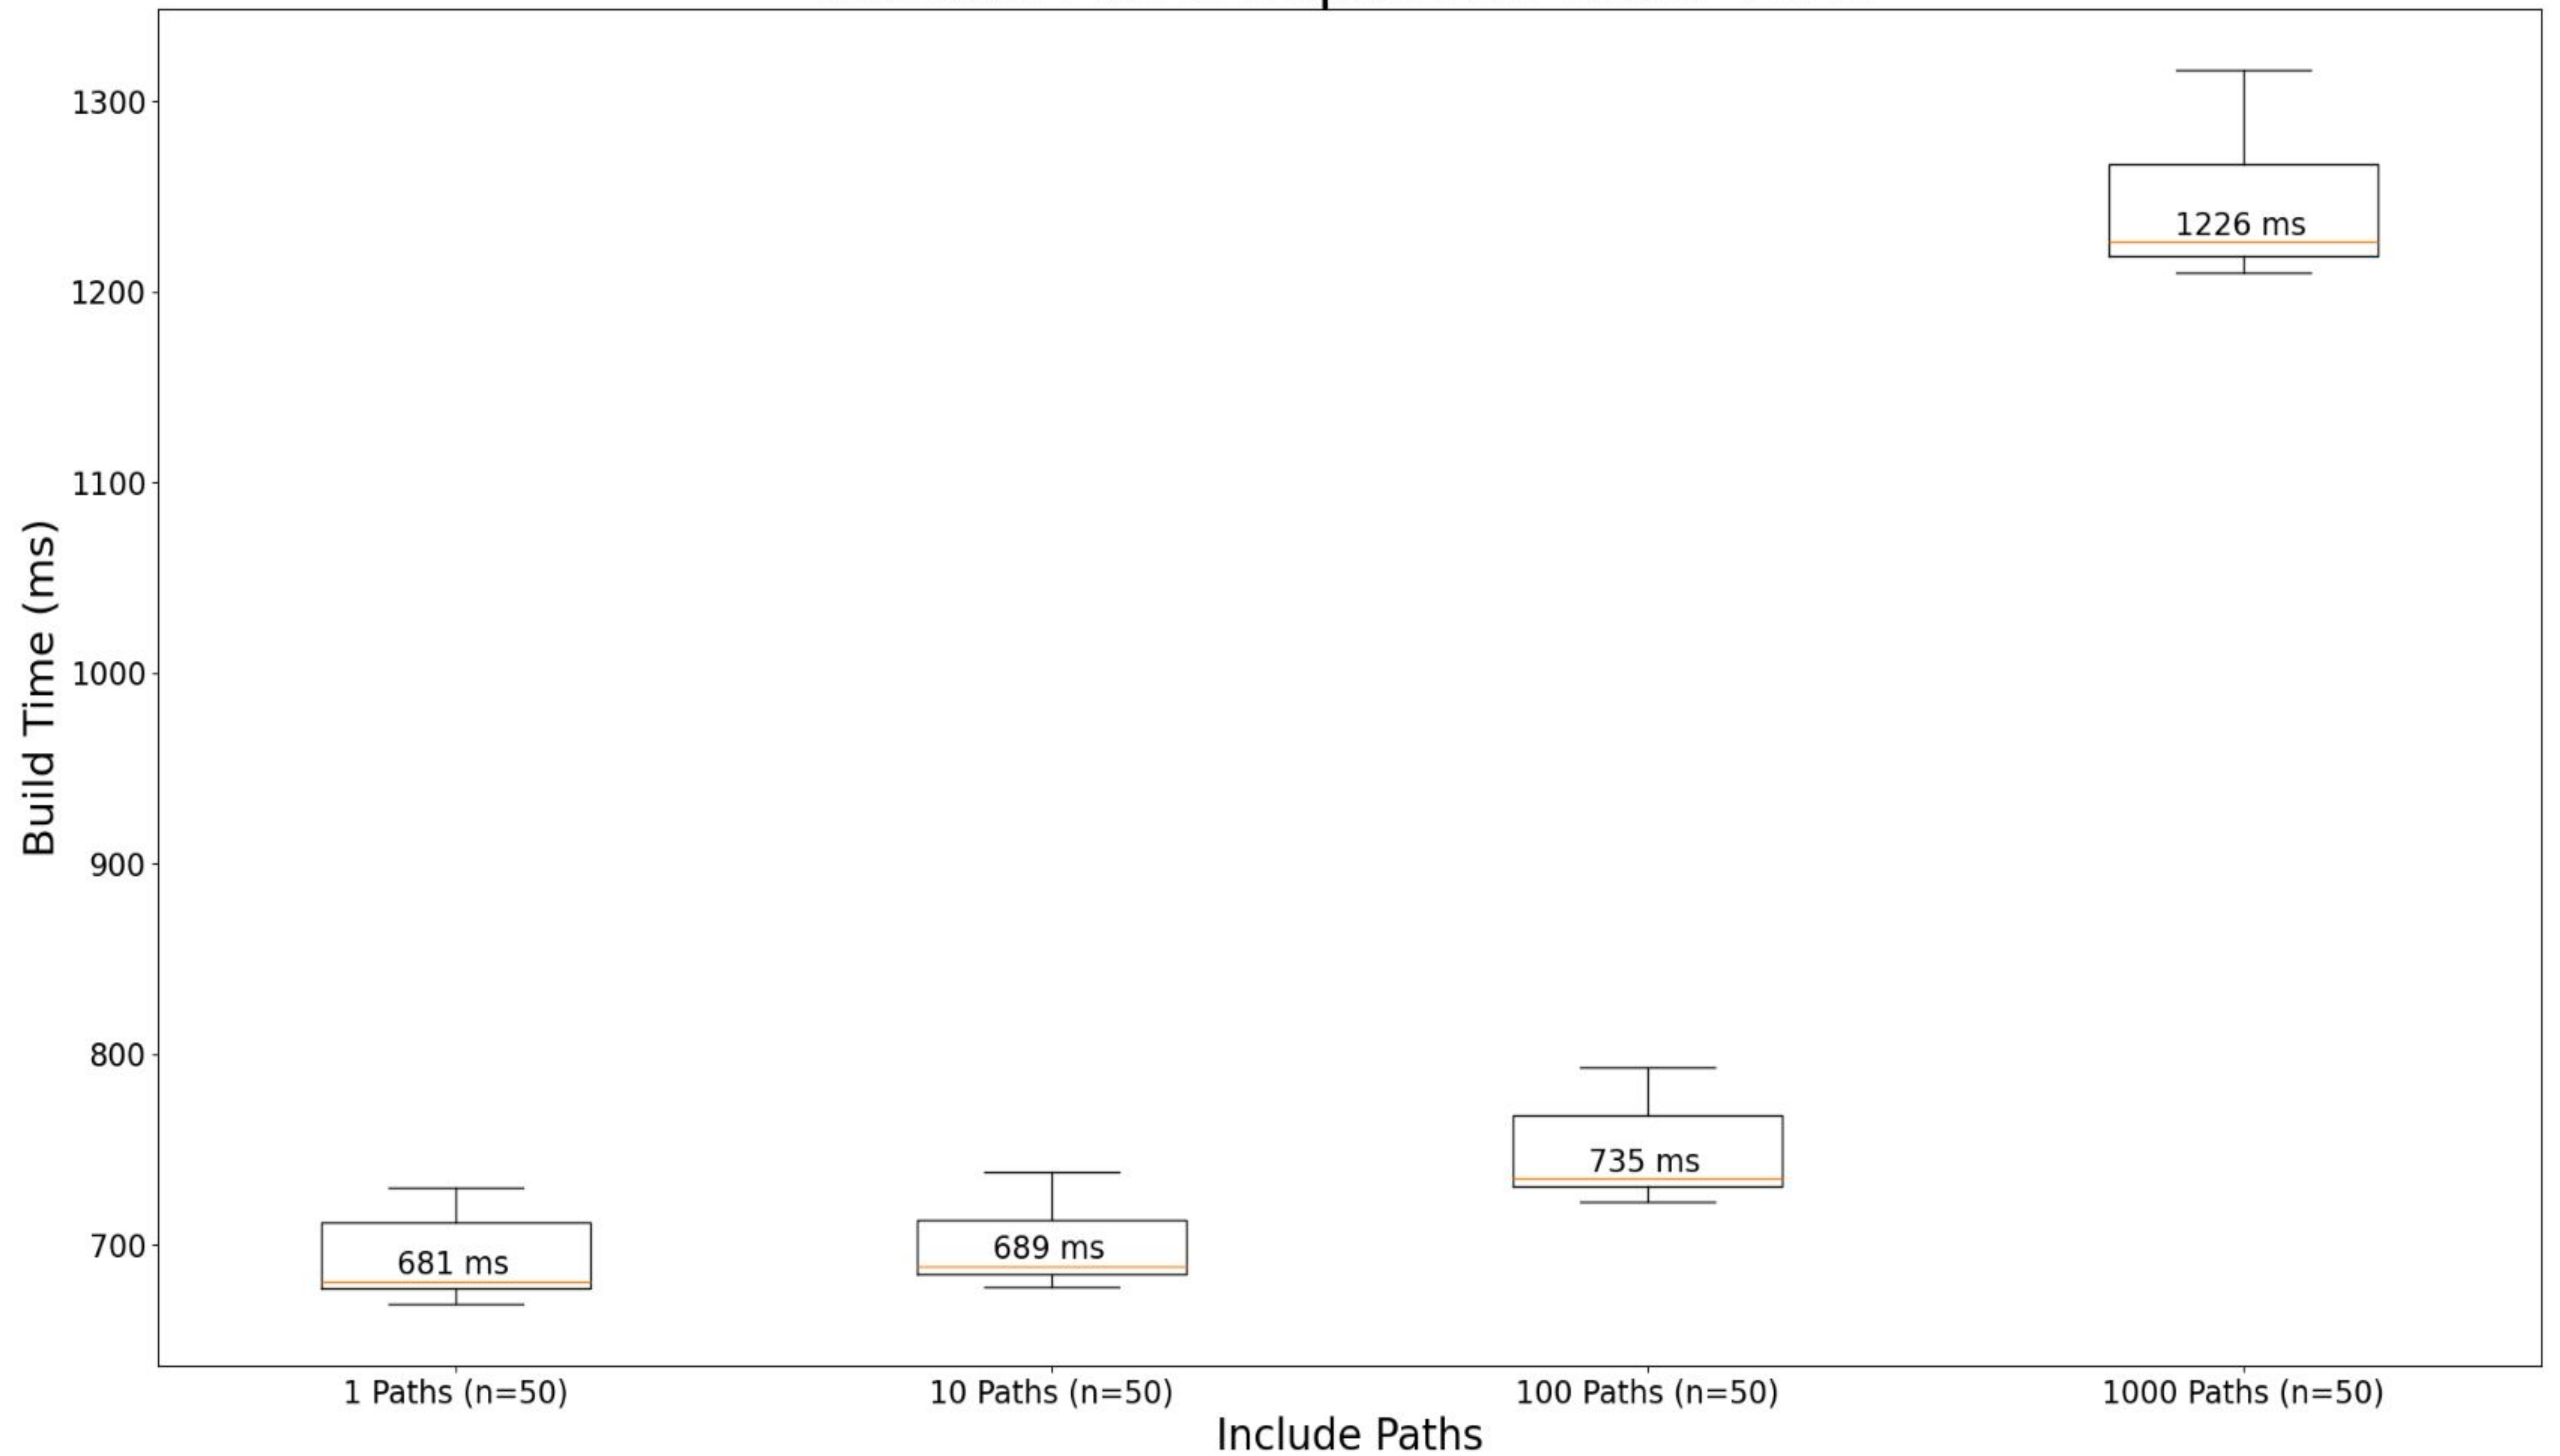
Include Paths Impact on Build Time



Include Paths Impact on Build Time



Include Paths Impact on Build Time




Compiler Include Paths

- Interacting with the filesystem is not free
- Caches abound but still expensive at scale
- Scales per translation unit



Querying Traces

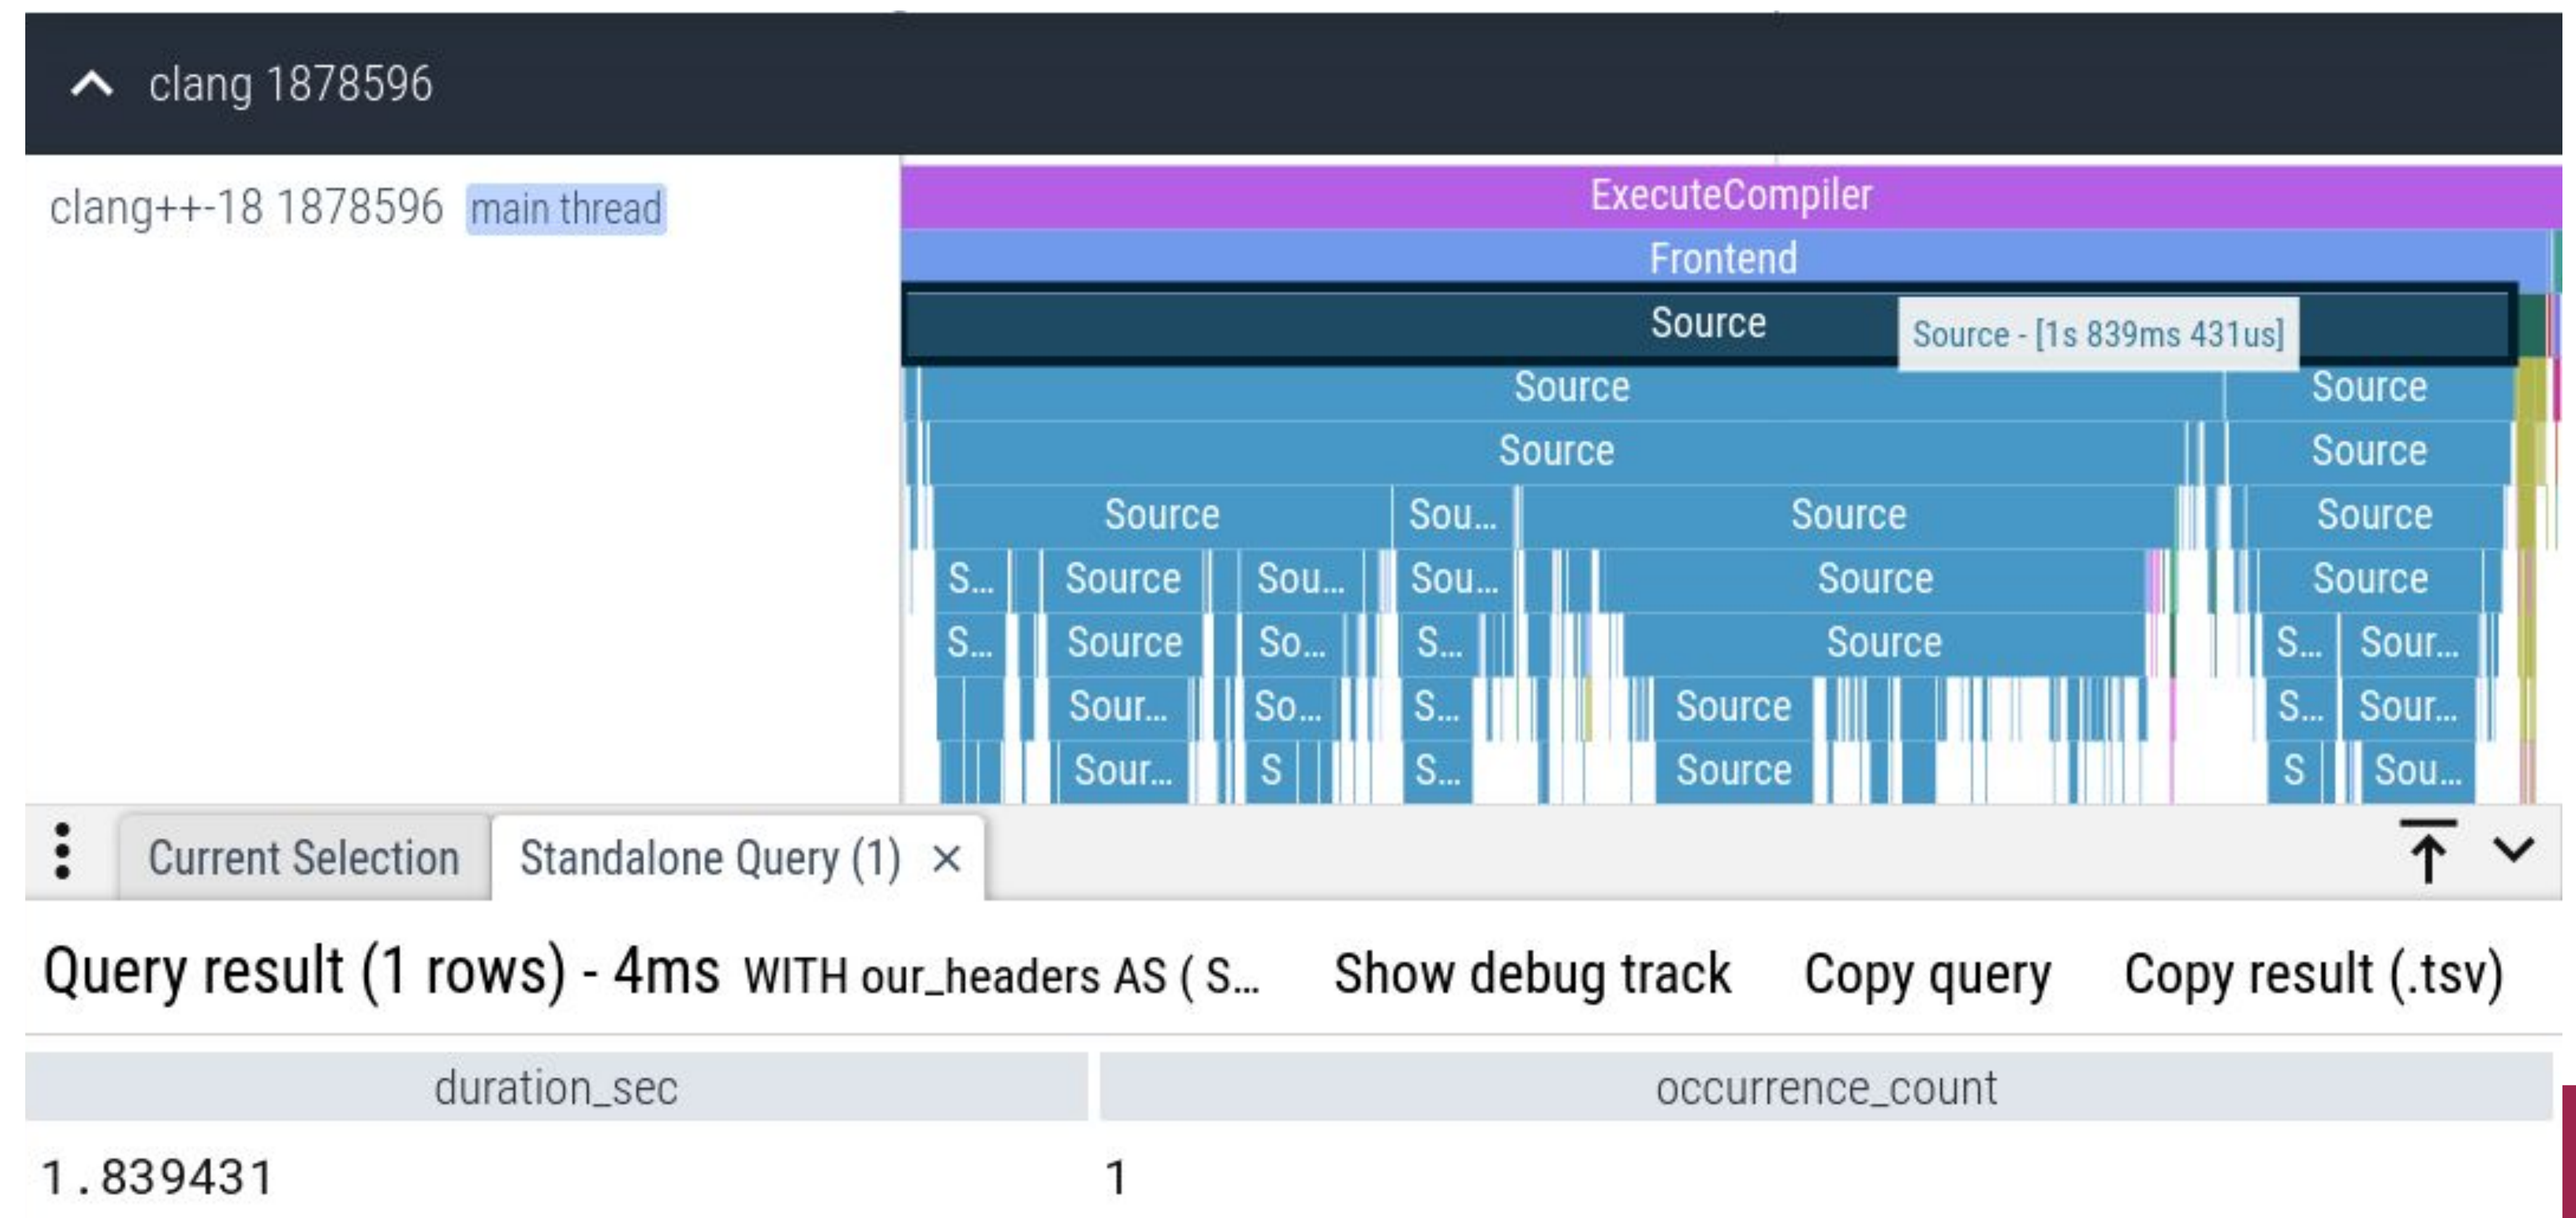
- Perfetto's superpower
- Answer complex questions like
 - “How much of the build is spent including this specific header?”
 - “What is the total impact on our build by using this external library?”



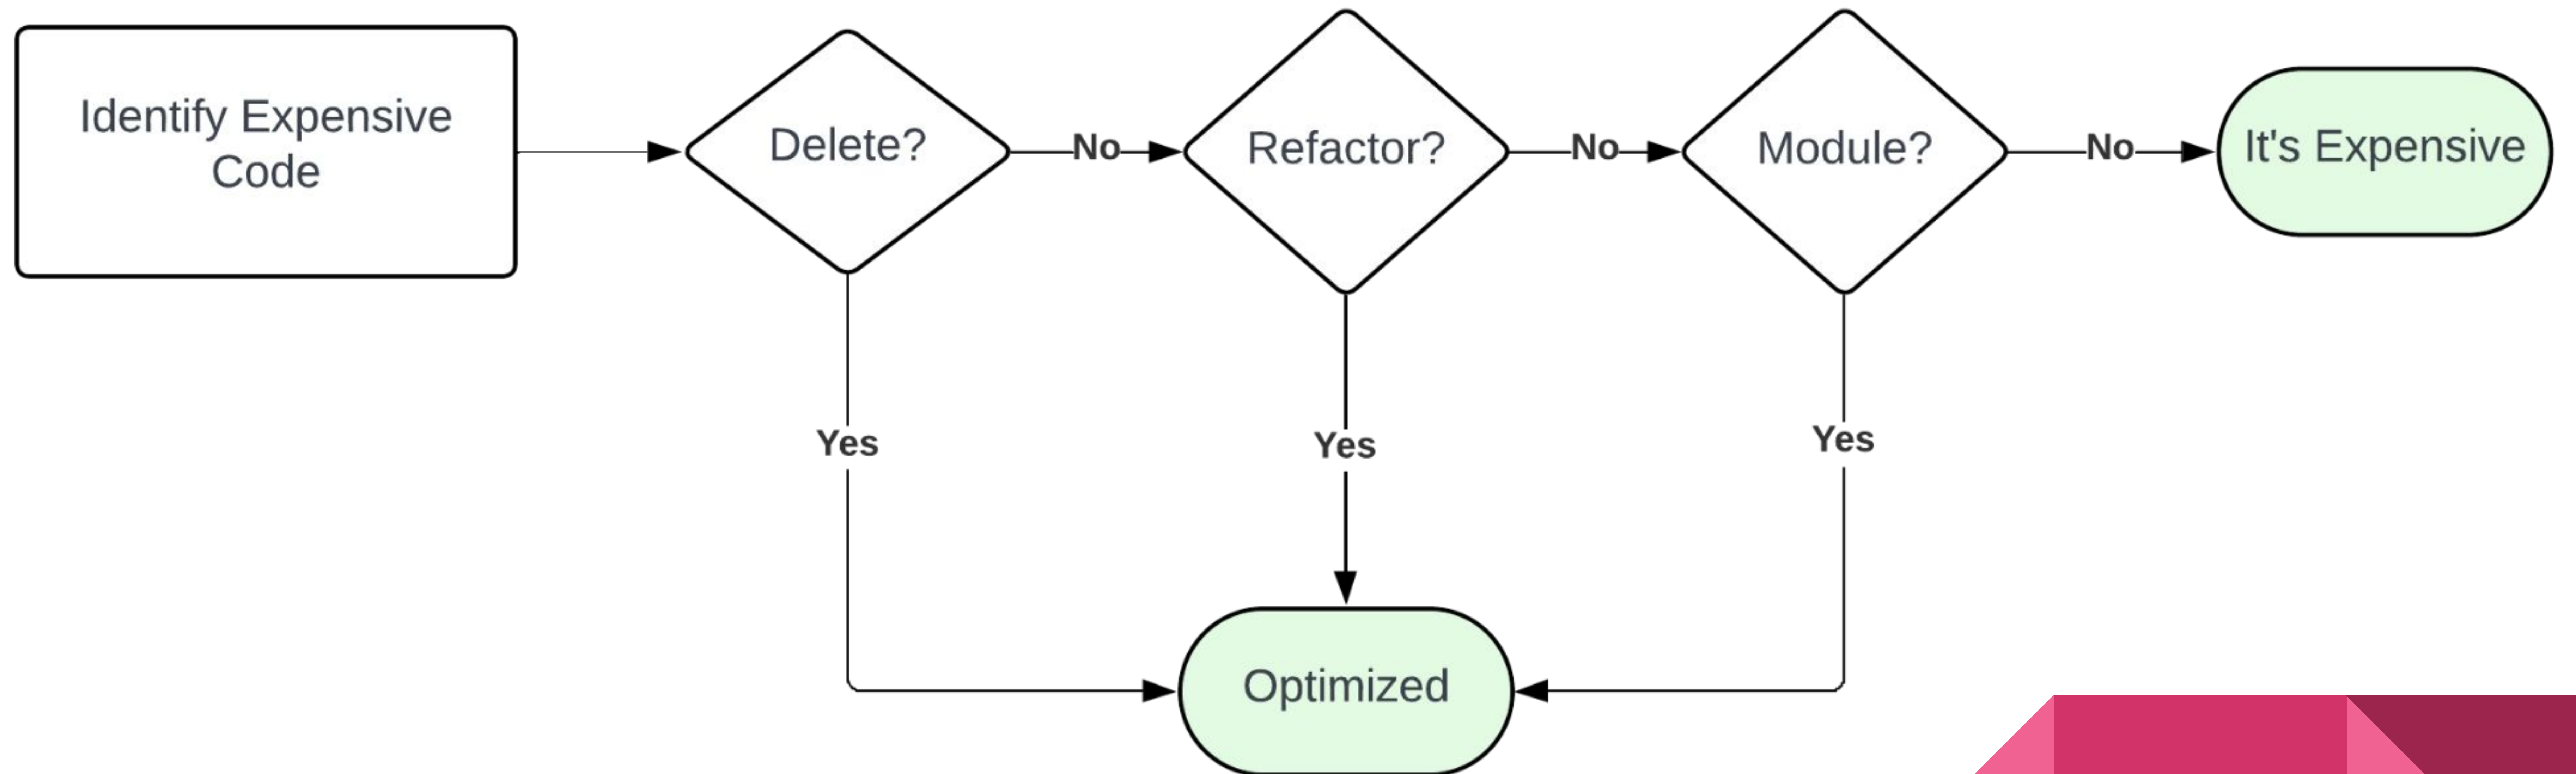
Querying Traces

“How much of the build is spent including this specific header?”

```
WITH our_headers AS (  
  SELECT  
    DISTINCT arg_set_id,  
    display_value  
  FROM  
    args  
  WHERE  
    KEY = 'args.detail'  
    AND (display_value LIKE 'my_header.h')  
) SELECT  
  SUM (slice.dur) / 1e+9 AS duration_sec,  
  COUNT (*) AS occurrence_count  
FROM  
  our_headers JOIN slice  
  ON our_headers.arg_set_id = slice.arg_set_id
```



Project Level Analysis



Is the Juice Worth the Squeeze?

- Optimizing build times is a lot of work
- Focus on low hanging fruit





Project Level

Higher Order

Takeaways


The Easy Stuff Is Gone, Now What?

- Do less work
 - Incremental build
 - Use a package manager
- Compiler caching
- Throw hardware at it
- Distributed builds



What Andy Giveth, Bill Taketh Away

- Andy and Bill's Law
 - New software will always consume any increase in computing power that new hardware can provide
- Wirth's law
 - Software is getting slower more rapidly than hardware is becoming faster





Higher Order

Takeaways

Questions?

Takeaways


- Profile your builds
- Compilation requires parsing which isn't free
- Better hardware won't always save us from ourselves
- We are all stewards of our build times
- Give a compiler engineer a hug




Thank You!



<https://github.com/maspe36/WhyIsMyBuildSoSlow>



<https://www.linkedin.com/in/sam-privett/>






careers.jnj.com



Takeaways

Questions?

GCC Profiling



- Externis, GCC Plugin “similar” to -ftime-trace
- Write a script to parse -ftime-report

Profile Memory Consumption

Really nice presentation on this from an LLVM Performance Workshop

<https://llvm.org/devmtg/2023-02-25/slides/basic-memory-profiler.pdf>

TL;DR: Nothing yet, but maybe soon